



**Oracle 19c Multi Tenant  
LABs**

## Multitenant Labs

- [Lab 1 Installing the HR Schema](#)
- [Lab 2 Reviewing the Multitenant configuration](#)
- [Lab 3 Exploring the PDB configuration](#)
- [Lab 4: Viewing Initialization Parameters by Using SQL\\*Plus](#)
- [Lab 5: Modifying Database parameters](#)
- [Lab 6: Starting and Stopping the container](#)
- [Lab 7: View Diagnostic information](#)
- [Lab 8: Creating a Local Users](#)
- [Lab 9: Granting Privilege](#)
- [Lab 10: Creating a Default Role for a User](#)
- [Lab 11 : Creating a PDB from Seed](#)
- [Lab 12: Cloning a PDB](#)
- [Lab 13: Unplugging and Plugging in a PDB](#)
- [Lab 14: Dropping a PDB](#)
- [Lab 15: Viewing Tablespace Information](#)
- [Lab 16: Creating a Tablespace](#)
- [Lab 17: Verifying that the Control File is Multiplexed](#)
- [Lab 18: Configuring the Size of the Fast Recovery Area](#)
- [Lab 19: Verifying that the Redo Log File is Multiplexed](#)
- [Lab 20: Verifying that ARCHIVELOG Mode is Configured](#)
- [Lab 21: Verifying Automatic Backups of the Control File and SPFILE](#)
- [Lab 22: Creating a Whole Database Backup](#)
- [Lab 23: Creating Partial Database Backups](#)
- [Lab 24: Recovering from the Loss of a System-Critical Data File](#)

## Lab 1: Installing HR

---

### Overview

In this practice, you will manually install the `HR` sample schema.

### Assumptions

You have a connection to the compute node through PuTTY or SSH and are logged in as the `oracle` user.

### Tasks

1. In your terminal window, navigate to the `$ORACLE_HOME/demo/schema/human_resources` directory.

```
[oracle@MYDBCS ~]$ cd $ORACLE_HOME/demo/schema/human_resources
[oracle@MYDBCS human_resources]$
```

2. Use the `ls` command to view the contents of the `human_resources` directory. In a later step, you will execute the `hr_main.sql` to create the `HR` user, objects and load data into the `HR` tables.

```
[oracle@MYDBCS human_resources]$ ls
hr_analz.sql  hr_comnt.sql  hr_drop_new.sql  hr_idx.sql
hr_main.sql
hr_code.sql   hr_cre.sql    hr_drop.sql      hr_main_new.sql
hr_popul.sql
[oracle@MYDBCS human_resources]$
```

3. Start SQL\*Plus and connect to the root container as the `SYS` user with the `SYSDBA` privilege.

```
[oracle@MYDBCS human_resources]$ sqlplus / as sysdba
...
SQL>
```

4. Switch to `PDB1`.

```
SQL> ALTER SESSION SET CONTAINER = PDB1;
Session altered.

SQL>
```

5. Execute the `hr_main.sql` script and respond to the prompts as follows.
  - a. Enter the password for the `HR` user as specified in the *Course Practice Environment: Security Credentials*.
  - b. Enter `USERS` as the default tablespace for the `HR` user.
  - c. Enter `TEMP` as the temporary tablespace for the `HR` user.
  - d. Enter `$ORACLE_HOME/demo/schema/log/` for the log directory.

```

SQL> @hr_main

specify password for HR as parameter 1:
Enter value for 1: password

specify default tablespace for HR as parameter 2:
Enter value for 2: USERS

specify temporary tablespace for HR as parameter 3:
Enter value for 3: TEMP

specify log path as parameter 4:
Enter value for 4: $ORACLE_HOME/demo/schema/log/

PL/SQL procedure successfully completed.

User created.

User altered.

Grant succeeded.

...

Comment created.

Commit complete.

PL/SQL procedure successfully completed.

SQL>

```

e. Exit from SQL \*Plus.

```

SQL> exit

...

[oracle@MYDBCS human_resources]$

```

6. Query the USER\_TABLES view as the HR user to verify that the user and tables were created.

a. Connect as the HR user. Be sure to provide the correct service name for your PDB as you did in Practice 5-4, step 3c.

```

[oracle@MYDBCS human_resources]$ sqlplus
hr/password@localhost:1521/PDB1.588436052.oraclecloud.internal

...

SQL>

```

- b. Query `USER_TABLES`.

```
SQL> SELECT table_name FROM user_tables;
```

```
TABLE_NAME
```

```
-----  
REGIONS  
COUNTRIES  
LOCATIONS  
DEPARTMENTS  
JOBS  
EMPLOYEES  
JOB_HISTORY
```

```
7 rows selected.
```

```
SQL>
```

7. Exit from SQL\*Plus and close the connection to the computenode.

```
SQL> exit
```

```
Disconnected from Oracle Database 19c Enterprise Edition Release  
18.0.0.0.0 - Production
```

```
Version 18.1.0.0.0
```

```
[oracle@MYDBCS human_resources]$ exit
```

## Lab 2: Reviewing the Multitenant configuration

---

### Overview

In this practice, you will learn how to do the following things:

- Set the Oracle environment variables
- Connect to the root container by using SQL\*Plus
- Query the data dictionary to view information about the containers, data files, users, instance, and services in a CDB
- List the services created automatically for each container

Some things to remember when you want to query the data dictionary for multiple PDBs or the whole CDB:

- Log in to the root container as a common user. A CDB common user is a database account created in the root container and is inherited by all PDBs in the CDB.
- Query container data objects, such as views whose names begin with V\$ and CDB\_. For more information, refer to the following sections in *Oracle Database Administrator's Guide*:
  - About Viewing Information When the Current Container is the CDB Root
  - Viewing Information About the Containers in a CDB

In some of the steps below, you will format columns by using the `COLUMN` command. For example, applying the format `A55` specifies an alphabetic format of 55 characters wide. Format `999` is an example of a numeric format. See *SQL\*Plus User's Guide and Reference* for additional information.

Commands in the practices are in uppercase and variables are in lower case. Any commands that you need to enter are bolded, for example:

```
SQL> SELECT regions FROM hr.departments;
```

### Assumptions

You are connected to the compute node as the `oracle` user. See Practice 5-2 for detail.

### Tasks

1. Set the Oracle environment variables. You need to set these each time you open a new terminal window.
  - a. In the terminal window, list the search path that holds the `oraenv` script.
  - b. Source the `oraenv` script. `oraenv` sets the required environment variables needed for you to connect to your database instance. The `oraenv` script sets the `ORACLE_SID` and `ORACLE_HOME` environment variables and includes the `$ORACLE_HOME/bin` directory in the `PATH` environment variable setting. Environment variables that this

script sets will persist in the terminal window until you close it. For the `ORACLE_SID` value, enter `ORCL`.

```
[oracle@MYDBCS ~]$ . oraenv
ORACLE_SID = [ORCL] ? ORCL
The Oracle base has been set to /u01/app/oracle
[oracle@MYDBCS ~]$
```

- c. View the environment variables set by the `oraenv` command.

```
[oracle@MYDBCS ~]$ set | grep ORACLE
OLD_ORACLE_BASE=/u01/app/oracle
ORACLE_BASE=/u01/app/oracle
ORACLE_HOME=/u01/app/oracle/product/18.0.0/dbhome_1
ORACLE_HOSTNAME=MYDBCS.compute-588436052.oraclecloud.internal
ORACLE_SID=ORCL
ORACLE_UNQNAME=ORCL
[oracle@MYDBCS ~]$
```

**Note:** Remember that from this point on, each time you open a terminal window you will need to source the `oraenv` script to set the environment variables for your CDB.

2. Connect to the root container by using SQL\*Plus.
  - a. Start SQL\*Plus and log in to the root container of your CDB as the `SYS` user with the `SYSDBA` privilege. You can connect to a database without a password when you have a local connection (on the same machine) and the current operating system user is a member of the privileged `OSDBA` group.
  - b. `Sqlplus / as sysdba`
  - c. `Sql> select * from v$containers;`
  - b. Verify that you are logged in to the root container as the `SYS` user by using the `SHOW USER` command.

```
SQL> SHOW user
USER is "SYS"
SQL>
```

3. View information about the containers in your CDB.

- a. Verify that you have a container database by querying the `V$DATABASE` view. The `NAME` column should contain `ORCL`, the `CDB` column should contain `YES`, and the `ID` should be 0 (zero). A value of zero is used for rows containing data that pertain to the entire CDB. This value is also used for rows in non-CDBs.

```
SQL> SELECT name, cdb, con_id FROM v$database;
```

| NAME | CDB | CON_ID |
|------|-----|--------|
| ORCL | YES | 0      |

```
SQL>
```

- b. Show the current container name. Because you're currently connected to the root container, the name should be `CDB$ROOT`.

```
SQL> SHOW con_name
```

```
CON_NAME
```

```
-----  
CDB$ROOT
```

```
SQL>
```

- c. Show the current container ID. Because you're currently connected to the root container, the ID should be 1.

```
SQL> SHOW con_id
```

```
CON_ID
```

```
-----  
1
```

```
SQL>
```

- d. Determine the version of Oracle Database by querying the `V$VERSION` view. This view displays version numbers of core library components in Oracle Database.

```
SQL> SELECT banner FROM v$version;
```

```
-----
```

- e. List all the containers in your CDB by querying the `V$CONTAINERS` view. The results should show three containers—the root container (`CDB$ROOT`), the seed PDB (`PDB$SEED`), and `PDB1`.



```
SQL> COLUMN name FORMAT A8
SQL> SELECT name, con_id FROM v$containers ORDER BY con_id;

NAME                CON_ID
-----
CDB$ROOT             1
PDB$SEED             2
PDB1                 3

SQL>
```

- f. List the PDBs in the CDB by using the `SHOW` command. The result should show two PDBs—the seed PDB (`PDB$SEED`) and `PDB1`. You can also list PDBs by querying the `V$PDBS` view. The `SHOW` command includes information about the open mode of each PDB and whether the PDB is restricted. The open mode for a PDB determines what type of activities a PDB will allow at that time. `PDB$SEED` is in `READ ONLY` mode and `PDB1` is in `READ WRITE` mode. The `RESTRICTED` column indicates whether only users possessing the `RESTRICTED SESSION` privilege can connect to the PDB.

```
SQL> SHOW pdbs

      CON_ID CON_NAME                                OPEN MODE  RESTRICTED
-----
          2 PDB$SEED                                READ ONLY   NO
          3 PDB1                                     READ WRITE NO

SQL>
```

- g. View the status of all PDBs in the CDB by querying the `CDB_PDBS` view. The status of a PDB describes the state of the PDB. For example, if the PDB is new, but never opened, the status is `NEW`. If it is available and ready for use, the status is `NORMAL`.

```
SQL> COLUMN pdb_name FORMAT A8
SQL> SELECT pdb_name, status FROM cdb_pdb$seeds ORDER BY 1;

PDB_NAME STATUS
-----
PDB1       NORMAL
PDB$SEED   NORMAL

SQL>
```

4. View information about the data files in your CDB.
- a. List all the data files in the CDB (for the root container and all PDBs) by querying the `CDB_DATA_FILES` view. The order of your results may vary.

```
SQL> COLUMN file_name FORMAT A50
SQL> COLUMN tablespace_name FORMAT A10
```

```
SQL> SELECT file_name, tablespace_name FROM cdb_data_files;
```

| FILE_NAME  | TABLESPACE |
|--|------------|
| /u02/app/oracle/oradata/ORCL/users01.dbf           | USERS      |
| /u02/app/oracle/oradata/ORCL/undotbs01.dbf         | UNDOTBS1   |
| /u02/app/oracle/oradata/ORCL/system01.dbf          | SYSTEM     |
| /u02/app/oracle/oradata/ORCL/sysaux01.dbf          | SYSAUX     |
| /u02/app/oracle/oradata/ORCL/PDB1/system01.dbf     | SYSTEM     |
| /u02/app/oracle/oradata/ORCL/PDB1/sysaux01.dbf     | SYSAUX     |
| /u02/app/oracle/oradata/ORCL/PDB1/undotbs01.dbf    | UNDOTBS1   |
| /u02/app/oracle/oradata/ORCL/PDB1/PDB1_users01.dbf | USERS      |

8 rows selected.

```
SQL>
```

- b. List all the tablespaces in the CDB (for both the root container and all the PDBs) by querying the V\$DATAFILE and V\$TABLESPACE views.

```
SQL> COL name FORMAT A12
```

```
SQL> SELECT d.file#, ts.name, ts.ts#, ts.con_id
  2  FROM v$datafile d, v$tablespace ts
  3  WHERE d.ts#=ts.ts# AND d.con_id=ts.con_id
  4  ORDER BY 4;
```

| FILE# | NAME     | TS# | CON_ID |
|-------|----------|-----|--------|
| 1     | SYSTEM   | 0   | 1      |
| 3     | SYSAUX   | 1   | 1      |
| 4     | UNDOTBS1 | 2   | 1      |
| 7     | USERS    | 4   | 1      |
| 6     | SYSAUX   | 1   | 2      |
| 13    | USERS    | 5   | 2      |
| 8     | UNDOTBS1 | 2   | 2      |
| 5     | SYSTEM   | 0   | 2      |
| 9     | SYSTEM   | 0   | 3      |
| 10    | SYSAUX   | 1   | 3      |
| 11    | UNDOTBS1 | 2   | 3      |
| 12    | USERS    | 5   | 3      |

12 rows selected.

```
SQL>
```

- c. List all temp files in the CDB (for the root container and all PDBs) by querying the CDB\_TEMP\_FILES view.

```
SQL> SELECT file_name, tablespace_name FROM cdb_temp_files;

FILE_NAME                                TABLESPACE
-----
/u04/app/oracle/oradata/temp/temp01.dbf    TEMP
/u02/app/oracle/oradata/ORCL/PDB1/pdbseed_temp0120 TEMP
18-02-19_18-48-12-642-PM.dbf

SQL>
```

- d. List all the redo log files in the CDB (for the root container and all PDBs) by querying the V\$LOGFILE view.

```
SQL> COLUMN member FORMAT A42
SQL> SELECT group#, member, con_id FROM v$logfile;

GROUP# MEMBER                                CON_ID
-----
3 /u04/app/oracle/redo/redo03.log            0
2 /u04/app/oracle/redo/redo02.log            0
1 /u04/app/oracle/redo/redo01.log            0

SQL>
```

- e. List the control files in the CDB by querying the V\$CONTROLFILE view. There should be two—control01.ctl and control02.ctl.

```
SQL> COLUMN name FORMAT A55
SQL> SELECT name, con_id FROM v$controlfile;

NAME                                CON_ID
-----
/u02/app/oracle/oradata/ORCL/control01.ctl    0
/u03/app/oracle/fast_recovery_area/ORCL/control02.ctl    0

SQL>
```

5. View information about the pre-created users in your CDB.

- a. List only the common users in the CDB by querying the CDB\_USERS view.

```
SQL> SELECT DISTINCT username FROM cdb_users
2 WHERE common = 'YES' ORDER BY 1;

USERNAME
```

```
ANONYMOUS
APPQOSSYS
AUDSYS
C##DBAAS_BACKUP
...
SYSTEM
WMSYS
XDB
XS$NULL

38 rows selected.

SQL>
```

- b. List all the users in every PDB in the CDB by querying the `CDB_USERS` view. In the results, notice that the `SYS`, `SYSTEM`, and `PDBADMIN` user accounts are listed for `PDB1`. The root container's id is 1 and `PDB1`'s id is 3.

```
SQL> COLUMN username FORMAT A25
SQL> SELECT con_id, username FROM cdb_users
2 ORDER BY username, con_id;
```

```
SQL>
```

6. View information about the database instance and the services.

- a. View the database instance name, its status, and which container database it is associated with by querying the `V$INSTANCE` view. The instance's status is `OPEN`, which means users can access the CDB and PDB.

```
SQL> SELECT instance_name, status, con_id FROM v$instance;
```

| INSTANCE_NAME | STATUS | CON_ID |
|---------------|--------|--------|
| ORCL          | OPEN   | 0      |

```
SQL>
```

- b. List the services for all the containers in the CDB by querying the `V$SERVICES` view. The query returns five services. The `PDB$SEED` service is not listed because no one should connect to it and no operation should be performed with it. It is reserved as a template to create other PDBs.

```
SQL> SELECT con_id, name FROM v$services ORDER BY 1;
```

| CON_ID | NAME                                   |
|--------|--|
| 1      | SYSS\$BACKGROUND                       |
| 1      | ORCL.588436052.oraclecloud.internal    |
| 1      | ORCL.588436052.oraclecloud.internalXDB |
| 1      | SYSS\$USERS                            |
| 3      | pdb1                                   |

```
SQL>
```

7. Exit SQL\*Plus.

```
SQL > exit
```

```
Disconnected from Oracle Database 19c EE High Perf Release  
18.0.0.0.0 - Production  
Version 18.1.0.0.0  
[oracle@MYDBCS ~]$
```

## Lab 3: Exploring the PDB configuration

---

### Tasks

1. Connect to `PDB1` indirectly through the root container.
  - a. Start SQL\*Plus and connect to the root container as the `SYS` user with the `SYSDBA` privilege. Oracle allows any DBA group user at the operating system level to log into SQL\*Plus without any authentication.

```
[oracle@MYDBCS ~]$ sqlplus / as sysdba
```

```
...
```

```
SQL>
```

- b. Verify that `PDB1` is open. After DBCA creates a PDB, it opens it automatically. The results below indicate that the open mode is `READ WRITE`, which means `PDB1` is open. PDB users with the `SYSDBA`, `SYSOPER`, `SYSBACKUP`, `SYSDG`, `SYSKM`, or `SYSRAC` privilege can connect to a closed PDB; however, all other PDB users can connect only when the PDB is open.

```
SQL> COLUMN con_id FORMAT 999
```

```
SQL> COLUMN name FORMAT A10
```

```
SQL> SELECT con_id, name, open_mode FROM v$pdb;
```

| CON_ID | NAME | OPEN_MODE |
|--------|------|-----------|
|--------|------|-----------|

```
-----
2 PDB$SEED    READ ONLY
3 PDB1        READ WRITE

SQL>
```

c. If PDB1 is closed for some reason and its open mode was MOUNTED in the previous step, open it by using the ALTER PLUGGABLE DATABASE command.

```
SQL> ALTER PLUGGABLE DATABASE PDB1 OPEN;

Pluggable database altered.

SQL>
```

d. Switch to PDB1. When logged in to a CDB as an appropriately privileged user, you can use the ALTER SESSION command to switch between containers within the CDB. From this point on, your queries against the data dictionary will retrieve information for PDB1 only.

```
SQL> ALTER SESSION SET CONTAINER = PDB1;

Session altered.

SQL>
```

e. Verify that the container name is PDB1.

```
SQL> SHOW con_name

CON_NAME
-----
PDB1
SQL>
```

2. Query the data dictionary to list the data files and temp files for PDB1.

a. List the data files for PDB1 and the tablespaces to which they belong by querying the DBA\_DATA\_FILES view.

```
SQL> col file_name format a60
SQL> col tablespace_name format a10
SQL> SELECT file_name, tablespace_name FROM dba_data_files;

FILE_NAME                                                                 TABLESPACE
-----
/u02/app/oracle/oradata/ORCL/PDB1/system01.dbf                        SYSTEM
/u02/app/oracle/oradata/ORCL/PDB1/sysaux01.dbf                       SYSAUX
/u02/app/oracle/oradata/ORCL/PDB1/undotbs01.dbf                      UNDOTBS1
```

```
/u02/app/oracle/oradata/ORCL/PDB1/PDB1_users01.dbf      USERS

SQL>
```

- b. List the temp files for PDB1 and the tablespaces to which they belong by querying the DBA\_TEMP\_FILES view.

```
SQL> SELECT file_name, tablespace_name FROM dba_temp_files;

FILE_NAME                                                    TABLESPACE
-----
/u02/app/oracle/oradata/ORCL/PDB1/pdbseed_temp012018-02-19_1 TEMP
8-48-12-642-PM.dbf

SQL>
```

- c. List the local users for PDB1 by querying the DBA\_USERS view.

```
SQL> SELECT DISTINCT username FROM dba_users WHERE common='NO';

USERNAME
-----
PDBADMIN
APEX_LISTENER
APEX_PUBLIC_USER
APEX_REST_PUBLIC_USER
FLOWS_FILES
APEX_050100
APEX_INSTANCE_ADMIN_USER
SCOTT

8 rows selected.

SQL>
```

3. Make a direct connection to PDB1 by using the Easy Connect syntax. The Easy Connect syntax enables you to connect to the PDB without 1) requiring a connection to the root container and 2) having to set up a net service name for the PDB.

- a. Disconnect from the PDB.

```
SQL > DISCONNECT

Disconnected from Oracle Database 19c EE High Perf Release
18.0.0.0.0 - Production
Version 18.1.0.0.0

SQL>
```

- b. Verify that you aren't connected as any user. The SHOW user command returns " " indicating that you are not connected.

```
SQL> SHOW user
```



```
USER is ""  
SQL>
```

- c. Connect to PDB1 directly as the SYSTEM user by using the Easy Connect syntax. See *Course Practice Environment: Security Credentials* for the SYSTEM user password. In Practice 5-3, step 6b, you queried V\$SERVICES. Append the value in the query results following ORCL to pdb1 to create the service name as shown in this example.

```
SQL> CONNECT  
system/password@localhost:1521/pdb  
  
Connected.  
SQL>
```

- d. Verify that you are now connected as the SYSTEM user by using the SHOW USER command again.

```
SQL> SHOW user  
SQL> USER is "SYSTEM"  
SQL>
```

4. Exit SQL\*Plus.

```
SQL> EXIT  
...  
[oracle@MYDBCS ~]$
```

## Lab 4: Viewing Initialization Parameters by Using SQL\*Plus

---

1. Start SQL\*Plus and connect to the root container as the `SYS` user with the `SYSDBA` privilege.

```
$ sqlplus / as sysdba
..
SQL>
```

2. View the values of the `DB_NAME` and `DB_DOMAIN` parameters. Together, these values create the global database name.
  - a. View the value of the `DB_NAME` parameter. This parameter specifies the current database identifier of up to eight characters. If you have multiple databases, the value of this parameter should match the Oracle instance identifier of each one to avoid confusion with other databases running on the system.

| SQL> SHOW PARAMETER db_name |        |       |
|-----------------------------|--------|-------|
| NAME                        | TYPE   | VALUE |
| db_name                     | string | ORCL  |
| SQL>                        |        |       |

- b. View the value of the `DB_DOMAIN` parameter. In a distributed database system, `DB_DOMAIN` specifies the logical location of the database within the network structure. You should set this parameter if this database is or ever will be part of a distributed system. There is no default value.

```
SQL> SHOW PARAMETER db_domain
NAME                                TYPE      VALUE
```

---

```
db_domain          string  588436052.oraclecloud.internal
SQL>
```

3. View the `DB_RECOVERY_FILE_DEST` and `DB_RECOVERY_FILE_DEST_SIZE` parameters. These parameters set the location of the fast recovery area and its size.

The `DB_RECOVERY_FILE_DEST` parameter specifies the default location for the fast recovery area. The fast recovery area contains multiplexed copies of current control files and online redo logs, as well as archived redo logs, flashback logs, and Recovery Manager (RMAN) backups. If you specify a value for `DB_RECOVERY_FILE_DEST`, you must also specify a value for the `DB_RECOVERY_FILE_DEST_SIZE` initialization parameter.

The `DB_RECOVERY_FILE_DEST_SIZE` parameter specifies (in bytes) the hard limit on the total space to be used by target database recovery files created in the fast recovery area.

```
SQL> SHOW PARAMETER db_recovery_file_dest
NAME                                TYPE      VALUE
-----                                -
db_recovery_file_dest               string    /u03/app/oracle/fast_recovery_area
db_recovery_file_dest_size          big integer 4G
SQL>
```

4. View the `SGA_TARGET` and `SGA_MAX_SIZE` parameters.

`SGA_TARGET` specifies the total amount of SGA memory available to a database instance and `SGA_MAX_SIZE` sets a maximum size for the SGA.

If you set the `SGA_TARGET` parameter, you enable the Automatic Shared Memory Management (ASMM) feature. The Oracle Database server will automatically distribute memory among the various SGA memory pools (buffer cache, shared pool, large pool, java pool, and streams pool), ensuring the most effective memory utilization. Note, the log buffer pool, other buffer caches (such as `KEEP` and `RECYCLE`), other block sizes, fixed SGA, and other internal allocations must be manually sized and are not affected by ASMM. The memory allocated to these pools is deducted from the total available memory for `SGA_TARGET` when ASMM is enabled.

The manageability monitor process (MMON) computes the values of the automatically tuned memory pools to support ASMM.

In addition to `SGA_TARGET` and `SGA_MAX_SIZE`, you can set minimum nonzero values for each memory pool if an application component needs a minimum amount of memory to function properly. ASMM will treat those values as minimum levels.

The range of values for `SGA_TARGET` can be from 64 MB to an operating system-dependent value. You can't modify this value in a PDB.

```
SQL> SHOW PARAMETER sga
NAME                                TYPE      VALUE
-----                                -
allow_group_access_to_sga          boolean    FALSE
lock_sga                           boolean    FALSE
pre_page_sga                        boolean    TRUE
```

|                              |             |         |
|------------------------------|-------------|---------|
| sga_max_size                 | big integer | 2640M   |
| sga_min_size                 | big integer | 0       |
| sga_target                   | big integer | 2640M   |
| unified_audit_sga_queue_size | integer     | 1048576 |
| SQL>                         |             |         |

5. View the `UNDO_TABLESPACE` parameter. This parameter specifies the undo tablespace to be used when an instance starts. Oracle Database creates and manages information that is used to roll back, or undo, changes to the database. Such information consists of records of the actions of transactions, primarily before they are committed. These records are collectively referred to as undo and are stored in the undo tablespace. The results below indicate that the undo tablespace in your environment is `UNDOTBS1`.

| SQL> <b>SHOW PARAMETER undo_tablespace</b> |        |          |
|--|--------|----------|
| NAME                                       | TYPE   | VALUE    |
| undo_tablespace                            | string | UNDOTBS1 |
| SQL>                                       |        |          |

6. View the `COMPATIBLE` parameter. This parameter specifies the release with which Oracle must maintain compatibility. It enables you to use a new release of Oracle, while at the same time guaranteeing backward compatibility with an earlier release. This is helpful if it becomes necessary to revert to the earlier release. By default, the value for the compatible entry for this parameter is equal to the version of the Oracle Database that you have installed.

| SQL> <b>SHOW PARAMETER compatible</b> |         |        |
|---------------------------------------|---------|--------|
| NAME                                  | TYPE    | VALUE  |
| compatible                            | string  | 18.0.0 |
| noncdb_compatible                     | boolean | FALSE  |
| SQL>                                  |         |        |

7. View the `CONTROL_FILES` initialization parameter. This parameter specifies one or more control files, separated by commas, and including paths. One to eight file names are listed. Oracle strongly recommends that you multiplex and mirror control files. The output has been formatted for legibility.

| SQL> <b>SHOW PARAMETER control_files</b>              |        |
|---|--------|
| NAME  | TYPE   |
| -----   | -----  |
| control_files   | string |
| VALUE   |        |
| -----   |        |
| /u02/app/oracle/oradata/ORCL/control01.ctl,           |        |
| /u03/app/oracle/fast_recovery_area/ORCL/control02.ctl |        |
| SQL>  |        |

8. View the PROCESSES, SESSIONS, and TRANSACTIONS initialization parameters.

- a. View the PROCESSES parameter. This parameter specifies the maximum number of operating system user processes that can simultaneously connect to an Oracle server. This value should allow for all background processes and user processes. The default values of the SESSIONS and TRANSACTIONS initialization parameters are derived from the PROCESSES parameter. Therefore, if you change the value of PROCESSES, you should evaluate whether to adjust the values of those derived parameters. The range of values is from six to an OS-dependent value. The default value is dynamic and dependent on the number of CPUs.

| SQL> SHOW PARAMETER processes |         |       |
|-------------------------------|---------|-------|
| NAME                          | TYPE    | VALUE |
| aq_tm_processes               | integer | 1     |
| db_writer_processes           | integer | 1     |
| gcs_server_processes          | integer | 0     |
| global_txn_processes          | integer | 1     |
| job_queue_processes           | integer | 4000  |
| log_archive_max_processes     | integer | 4     |
| processes                     | integer | 300   |
| SQL>                          |         |       |

- b. View the SESSIONS parameter. This parameter specifies the maximum number of sessions that can be created in the system. Because every login requires a session, this parameter effectively determines the maximum number of concurrent users in the system. Notice in the results that the session entry has a value of 472. You should always set this parameter explicitly to a value equivalent to your estimate of the maximum number of concurrent users, plus the number of background processes, plus approximately 10% for recursive sessions.

| SQL> SHOW PARAMETER sessions |         |       |
|------------------------------|---------|-------|
| NAME                         | TYPE    | VALUE |
| java_max_sessionspace_size   | integer | 0     |
| java_soft_sessionspace_limit | integer | 0     |
| license_max_sessions         | integer | 0     |
| license_sessions_warning     | integer | 0     |
| sessions                     | integer | 472   |
| shared_server_sessions       | integer |       |
| SQL>                         |         |       |

- c. View the TRANSACTIONS parameter. This parameter specifies how many rollback segments to bring online when the UNDO\_MANAGEMENT initialization parameter is equal to MANUAL. A transaction is assigned to a rollback segment when the transaction starts, and it can't change for the life of the transaction. A transaction table exists in the rollback segment header with limited space, limiting how many transactions a single

segment can support. Therefore, X number of concurrent transactions require at least X number of rollback segments. With Oracle Automatic Undo Management, the database creates rollback segments, brings them online, takes them offline, and drops them as needed.

|                                   |         |       |
|-----------------------------------|---------|-------|
| SQL> SHOW PARAMETER transactions  |         |       |
| NAME                              | TYPE    | VALUE |
| transactions                      | integer | 519   |
| transactions_per_rollback_segment | integer | 5     |
| SQL>                              |         |       |

9. View the configuration for the `DB_FILES` initialization parameter. This parameter specifies the maximum number of database files that can be opened for this database. The range of values is OS-dependent.

|                              |         |       |
|------------------------------|---------|-------|
| SQL> SHOW PARAMETER db_files |         |       |
| NAME                         | TYPE    | VALUE |
| db_files                     | integer | 500   |
| SQL>                         |         |       |

**View Advanced Parameters**

In this section, you use the `SHOW PARAMETER` command to view advanced parameters.

1. View the `COMMIT_LOGGING` parameter. This parameter is used to control how redo is batched by the Log Writer process. There is no default value, as shown below. You can modify this parameter in a PDB.

|                                    |        |       |
|------------------------------------|--------|-------|
| SQL> SHOW PARAMETER commit_logging |        |       |
| NAME                               | TYPE   | VALUE |
| -----                              | -----  | ----- |
| commit_logging                     | string |       |
| SQL>                               |        |       |

2. View the `COMMIT_WAIT` parameter. This parameter is used to control when the redo for a commit is flushed to the redo logs. There is no default value.

|                                 |        |       |
|---------------------------------|--------|-------|
| SQL> SHOW PARAMETER commit_wait |        |       |
| NAME                            | TYPE   | VALUE |
| commit_wait                     | string |       |
| SQL>                            |        |       |

3. View the `SHARED_POOL_SIZE` parameter. This parameter specifies the size of the shared pool in bytes. The shared pool contains objects such as shared cursors, stored procedures, control structures, and parallel execution message buffers. The range of values is OS-dependent. The default value is zero if the `SGA_TARGET` parameter is set. Otherwise, the value is 128 MB for a 64-bit platform or 48 MB for a 32-bit platform.

| SQL> <b>SHOW PARAMETER shared_pool_size</b> |             |       |
|---|-------------|-------|
| NAME  | TYPE        | VALUE |
| -----                                       |             |       |
| shared_pool_size                            | big integer | 0     |
| SQL>  |             |       |

4. View the `DB_BLOCK_SIZE` parameter. This parameter specifies the standard Oracle database block size (in bytes) and is used by all tablespaces by default. Its value is set during database creation and cannot be subsequently changed. The range of values is from 2048 to 32768 (OS-dependent). The default value is 8192.

| SQL> <b>SHOW PARAMETER db_block_size</b> |         |       |
|--|---------|-------|
| NAME                                     | TYPE    | VALUE |
| db_block_size                            | integer | 8192  |
| SQL>                                     |         |       |

5. View the `DB_CACHE_SIZE` initialization parameter. You configure this parameter to specify the size of the standard block buffer cache (default buffer pool). The range of values is at least 4 MB times the number of CPUs. Smaller values are automatically rounded up to this value. The default value is zero if the `SGA_TARGET` initialization parameter is set, otherwise the larger of 48 MB or (4 MB\*CPU\_COUNT).

| SQL> <b>SHOW PARAMETER db_cache_size</b> |             |       |
|--|-------------|-------|
| NAME                                     | TYPE        | VALUE |
| -----                                    |             |       |
| db_cache_size                            | big integer | 0     |
| SQL>                                     |             |       |

6. View the `UNDO_MANAGEMENT` parameter. This parameter specifies the undo space management mode that the system should use. When set to `AUTO`, the instance is started in automatic undo management mode. Otherwise, it is started in rollback undo mode. In rollback undo mode, undo space is allocated as rollback segments. In automatic undo mode, undo space is allocated as undo tablespaces. The value is `AUTO` or `MANUAL`. If the `UNDO_MANAGEMENT` parameter is omitted when the instance is started, the default value `AUTO` is used.

| SQL> <b>SHOW PARAMETER undo_management</b> |        |       |
|--|--------|-------|
| NAME                                       | TYPE   | VALUE |
| undo_management                            | string | AUTO  |
| SQL>                                       |        |       |

7. View the `MEMORY_TARGET` and `MEMORY_MAX_TARGET` parameters. `MEMORY_TARGET` specifies the Oracle system-wide usable memory. The database server tunes memory to the `MEMORY_TARGET` value, reducing or enlarging the SGA and PGA as needed. `MEMORY_MAX_TARGET` sets a maximum value for `MEMORY_TARGET`.

In a PFILE, if you omit MEMORY\_MAX\_TARGET and include a value for MEMORY\_TARGET, the database automatically sets MEMORY\_MAX\_TARGET to the value of MEMORY\_TARGET. If you omit the line for MEMORY\_TARGET and include a value for MEMORY\_MAX\_TARGET, the MEMORY\_TARGET parameter defaults to zero. After startup, you can dynamically change MEMORY\_TARGET to a nonzero value if it does not exceed the value of MEMORY\_MAX\_TARGET. For MEMORY\_TARGET, values range from 152 MB to MEMORY\_MAX\_TARGET.

- a. View the MEMORY\_TARGET parameter.

|                                   |             |       |
|-----------------------------------|-------------|-------|
| SQL> SHOW PARAMETER memory_target |             |       |
| NAME                              | TYPE        | VALUE |
| -----                             |             |       |
| memory_target                     | big integer | 0     |
| SQL>                              |             |       |

- b. View the MEMORY\_MAX\_TARGET parameter.

|                                       |             |       |
|---------------------------------------|-------------|-------|
| SQL> SHOW PARAMETER memory_max_target |             |       |
| NAME                                  | TYPE        | VALUE |
| -----                                 |             |       |
| memory_max_target                     | big integer | 0     |
| SQL>                                  |             |       |

8. View the PGAAggregateTarget parameter. This parameter specifies the amount of Program Global Area (PGA) memory available to all server processes attached to the database instance. This memory does not reside in the System Global Area (SGA). The database uses this parameter as a target amount of PGA memory to use. When setting this parameter, subtract the SGA from the total memory on the system that is available to the Oracle instance. The minimum value is 10 MB and the maximum value is 4096 GB minus 1. The default value is 10 MB or 20% of the size of the SGA, whichever is greater.

|  |             |            |
|--|-------------|------------|
| SQL> SHOW PARAMETER pga_aggregate_target |             |            |
| NAME                                     | TYPE        | VALUE      |
| -----                                    |             |            |
| pga_aggregate_target                     | big integer | 1837647360 |
| SQL>                                     |             |            |

Query Views for Parameter Values

In this section, you query views to learn about parameters.

1. Query the data dictionary to find views that contain the word "parameter." The query below returns 66 rows. Not all of these views contain information about initialization parameters. Among these rows are the V\$PARAMETER, V\$SPPARAMETER, V\$PARAMETER2, and V\$SYSTEM\_PARAMETER views, which you'll examine next.

|  |  |  |
|--|--|--|
| SQL> SET PAGES 100                                     |  |  |
| SQL> SELECT table_name FROM dict WHERE table_name LIKE |  |  |
| '%PARAMETER%';   |  |  |



```

TABLE_NAME

USER_ADVISOR_EXEC_PARAMETERS
USER_ADVISOR_PARAMETERS
USER_ADVISOR_SQLW_PARAMETERS
USER_XS_ACL_PARAMETERS
ALL_APPLY_PARAMETERS
ALL_CAPTURE_PARAMETERS
ALL_XS_ACL_PARAMETERS
...
V$PARAMETER_VALID_VALUES
V$SPPARAMETER
V$SYSTEM_PARAMETER
V$SYSTEM_PARAMETER2
V$SYSTEM_RESET_PARAMETER
V$SYSTEM_RESET_PARAMETER2

66 rows selected.

SQL>

```

2. Explore the `V$PARAMETER` view. This view displays the current parameter values in the current session.
  - a. View the columns in the `V$PARAMETER` view by using the `DESCRIBE` command. This command returns column names, whether null values are allowed (`NOT NULL` is displayed if the value cannot be null), and column data types.

The results below contain a column named `ISSYS_MODIFIABLE`. This column is important because it tells you whether a parameter is static or dynamic. If its value is `FALSE`, then the parameter is static; otherwise it's dynamic. To change a static parameter, you must shut down and restart the database; however, you can modify a dynamic parameter in real time while the database is online.

```

SQL> DESCRIBE v$parameter

```

| Name             | Null? | Type            |
|------------------|-------|-----------------|
| NUM              |       | NUMBER          |
| NAME             |       | VARCHAR2 (80)   |
| TYPE             |       | NUMBER          |
| VALUE            |       | VARCHAR2 (4000) |
| DISPLAY_VALUE    |       | VARCHAR2 (4000) |
| DEFAULT_VALUE    |       | VARCHAR2 (255)  |
| ISDEFAULT        |       | VARCHAR2 (9)    |
| ISSES_MODIFIABLE |       | VARCHAR2 (5)    |
| ISSYS_MODIFIABLE |       | VARCHAR2 (9)    |

|                       |                |
|-----------------------|----------------|
| ISPDB_MODIFIABLE      | VARCHAR2 (5)   |
| ISINSTANCE_MODIFIABLE | VARCHAR2 (5)   |
| ISMODIFIED            | VARCHAR2 (10)  |
| ISADJUSTED            | VARCHAR2 (5)   |
| ISDEPRECATED          | VARCHAR2 (5)   |
| ISBASIC               | VARCHAR2 (5)   |
| DESCRIPTION           | VARCHAR2 (255) |
| UPDATE_COMMENT        | VARCHAR2 (255) |
| HASH                  | NUMBER         |
| CON_ID                | NUMBER         |

SQL>

- b. Query NAME, ISSYS\_MODIFIABLE, and VALUE in the V\$PARAMETER view. The query returns many rows.

The TRANSACTIONS parameter is static as indicated by FALSE in the ISSYS\_MODIFIABLE column. The PLSQL\_WARNINGS parameter is dynamic as indicated by IMMEDIATE in the ISSYS\_MODIFIABLE column.

Optional: Before entering the following command, you can enter SET PAUSE ON to cause a pause after each page output. Press Enter to display each next page. After all pages have been displayed, you can issue the SET PAUSE OFF command to stop this feature.

|  |           |             |
|--|-----------|-------------|
| SQL> SELECT name, issys_modifiable, value FROM v\$parameter; |           |             |
| NAME   | ISSYS_MOD | VALUE       |
| lock_name_space  | FALSE     |             |
| processes  | FALSE     | 300         |
| ...  |           |             |
| multishard_query_data_consistency                            | IMMEDIATE | strong      |
| multishard_query_partial_results                             | IMMEDIATE | not allowed |
| 433 rows selected.   |           |             |
| SQL>   |           |             |

- c. Query the V\$PARAMETER view again, but this time be more specific. Include a WHERE clause to specify all parameters that contain the word "pool." The query returns eight parameters that contain the word "pool."

|   |                           |
|---|---------------------------|
| SQL> COLUMN name FORMAT A30               |                           |
| SQL> COLUMN value FORMAT A10              |                           |
| SQL> SELECT name, value FROM v\$parameter |                           |
| 2   | WHERE name LIKE '%pool%'; |

| NAME                      | VALUE    |
|---------------------------|----------|
| -----                     |          |
| shared_pool_size          | 0        |
| large_pool_size           | 0        |
| java_pool_size            | 0        |
| streams_pool_size         | 0        |
| shared_pool_reserved_size | 26843545 |
| memoptimize_pool_size     | 0        |
| buffer_pool_keep          |          |
| buffer_pool_recycle       |          |
| olap_page_pool_size       | 0        |
| 9 rows selected.          |          |
| SQL>                      |          |

3. Explore the `V$SPPARAMETER` view. This view contains information about the contents of the server parameter file. If a server parameter file was not used to start the instance, each row of the view will contain `FALSE` in the `ISSPECIFIED` column.

a. View the columns in the `V$SPPARAMETER` view by using the `DESCRIBE` command.

| SQL> DESCRIBE v\$spparameter |       |                |
|------------------------------|-------|----------------|
| Name                         | Null? | Type           |
| -----                        |       |                |
| FAMILY                       |       | VARCHAR2 (80)  |
| SID                          |       | VARCHAR2 (80)  |
| NAME                         |       | VARCHAR2 (80)  |
| TYPE                         |       | VARCHAR2 (11)  |
| VALUE                        |       | VARCHAR2 (255) |
| DISPLAY_VALUE                |       | VARCHAR2 (255) |
| ISSPECIFIED                  |       | VARCHAR2 (6)   |
| ORDINAL                      |       | NUMBER         |
| UPDATE_COMMENT               |       | VARCHAR2 (255) |
| CON_ID                       |       | NUMBER         |
| SQL>                         |       |                |

b. Query `NAME` and `VALUE` in the `V$SPPARAMETER` view. Browse the rows returned by the query. The results below have been formatted for easier viewing and show only a small portion of the results.

| SQL> SELECT name, value FROM v\$spparameter; |       |
|--|-------|
| NAME   | VALUE |
| -----  |       |

```

lock_name_space
processes                300
sessions
timed_statistics
timed_os_statistics
...
shrd_dupl_table_refresh_rate
multishard_query_data_consistency
multishard_query_partial_results

437 rows selected.

SQL>

```

4. Explore the `V$PARAMETER2` view. This view contains information about the initialization parameters that are currently in effect for the session, with each parameter value appearing as a row in the view. A new session inherits parameter values from the instance-wide values displayed in the `V$SYSTEM_PARAMETER2` view.
  - a. View the columns in the `V$PARAMETER2` view by using the `DESCRIBE` command.

```

SQL> DESCRIBE v$parameter2
Name                                         Null?     Type
-----
NUM                                         NUMBER
NAME                                       VARCHAR2(80)
TYPE                                       NUMBER
VALUE                                       VARCHAR2(4000)
DISPLAY_VALUE                             VARCHAR2(4000)
ISDEFAULT                                  VARCHAR2(6)
ISSES_MODIFIABLE                           VARCHAR2(5)
ISSYS_MODIFIABLE                           VARCHAR2(9)
ISPDB_MODIFIABLE                           VARCHAR2(5)
ISINSTANCE_MODIFIABLE                      VARCHAR2(5)
ISMODIFIED                                 VARCHAR2(10)
ISADJUSTED                                 VARCHAR2(5)
ISDEPRECATED                               VARCHAR2(5)
ISBASIC                                    VARCHAR2(5)
DESCRIPTION                                VARCHAR2(255)
ORDINAL                                    NUMBER
UPDATE_COMMENT                             VARCHAR2(255)
CON_ID                                     NUMBER

SQL>

```

- b. Query NAME and VALUE in the V\$PARAMETER2 view. Browse the rows returned by the query. The results below have been formatted for easier viewing and show only a very small portion of the results.

```
SQL> SELECT name, value FROM v$parameter2;
```

| NAME                              | VALUE       |
|-----------------------------------|-------------|
| lock_name_space                   |             |
| processes                         | 300         |
| sessions                          | 472         |
| timed_statistics                  | TRUE        |
| timed_os_statistics               | 0           |
| resource_limit                    | TRUE        |
| ...                               |             |
| shrd_dupl_table_refresh_rate      | 60          |
| multishard_query_data_consistency | strong      |
| multishard_query_partial_results  | not allowed |

438 rows selected.

```
SQL>
```

5. Explore the V\$SYSTEM\_PARAMETER view. This view contains information about the initialization parameters that are currently in effect for the instance.
- a. View the columns in the V\$SYSTEM\_PARAMETER view by using the DESCRIBE command.

```
SQL> DESCRIBE v$system_parameter
```

| Name                  | Null? | Type            |
|-----------------------|-------|-----------------|
| NUM                   |       | NUMBER          |
| NAME                  |       | VARCHAR2 (80)   |
| TYPE                  |       | NUMBER          |
| VALUE                 |       | VARCHAR2 (4000) |
| DISPLAY_VALUE         |       | VARCHAR2 (4000) |
| DEFAULT_VALUE         |       | VARCHAR2 (255)  |
| ISDEFAULT             |       | VARCHAR2 (9)    |
| ISSES_MODIFIABLE      |       | VARCHAR2 (5)    |
| ISSYS_MODIFIABLE      |       | VARCHAR2 (9)    |
| ISPDB_MODIFIABLE      |       | VARCHAR2 (5)    |
| ISINSTANCE_MODIFIABLE |       | VARCHAR2 (5)    |
| ISMODIFIED            |       | VARCHAR2 (8)    |
| ISADJUSTED            |       | VARCHAR2 (5)    |

|                |                |
|----------------|----------------|
| ISDEPRECATED   | VARCHAR2 (5)   |
| ISBASIC        | VARCHAR2 (5)   |
| DESCRIPTION    | VARCHAR2 (255) |
| UPDATE_COMMENT | VARCHAR2 (255) |
| HASH           | NUMBER         |
| CON_ID         | NUMBER         |

SQL>

- b. Query `NAME` and `VALUE` in the `V$SYSTEM_PARAMETER` view. Browse the rows returned by the query. The results below have been formatted for easier viewing and show only a very small portion of the results.

```
SQL> SELECT name, value FROM v$system_parameter;
```

| NAME                              | VALUE       |
|-----------------------------------|-------------|
| lock_name_space                   |             |
| processes                         | 300         |
| sessions                          | 472         |
| timed_statistics                  | TRUE        |
| timed_os_statistics               | 0           |
| resource_limit                    | TRUE        |
| ...                               |             |
| parallel_servers_target           | 16          |
| common_user_prefix                |             |
| multishard_query_data_consistency | strong      |
| multishard_query_partial_results  | not allowed |

457 rows selected.

SQL>

## 6. Exit SQL\*Plus.

```
SQL> EXIT
```

Disconnected from Oracle Database 19c Enterprise Edition Release  
18.0.0.0.0 - Production  
Version 18.1.0.0.0  
[oracle@MYDBCS ~]\$

## Lab 5: Modifying Database parameters

---

### Overview

In this practice, you modify the following kinds of initialization parameters (parameters) with SQL\*Plus:

- Session-level parameter
- Dynamic system-level parameter
- Static system-level parameter

### Assumptions

You are connected to the compute node as the `oracle` user.

### Tasks

#### Modify a Session-Level Parameter

In this section, you modify the `NLS_DATE_FORMAT` parameter. This parameter defines the default date format to use with the `TO_CHAR` and `TO_DATE` functions. The `NLS_TERRITORY` parameter determines the default value of `NLS_DATE_FORMAT`. `NLS_DATE_FORMAT` is one of the National Language Support (NLS) parameters that you can customize just for your session, therefore making it a session-level parameter. When your session ends, your modification expires, and the parameter is returned to its default value.

1. Start SQL\*Plus and log in to the database as the `SYS` user with the `SYSDBA` privilege.

```
$ sqlplus / as sysdba
..
SQL>
```

2. Learn about the `NLS_DATE_FORMAT` parameter by querying the `V$PARAMETER` view. Include a `WHERE` clause to narrow down the query to just the `NLS_DATE_FORMAT` parameter. Remember that in the `V$PARAMETER` view, the parameter names are in lowercase.

```
SQL> SELECT name, isses_modifiable, issys_modifiable,
2      ispdb_modifiable, value
3      FROM v$parameter
4      WHERE name = 'nls_date_format';
```

| NAME            | ISSES | ISSYS_MOD | ISPDB | VALUE |
|-----------------|-------|-----------|-------|-------|
| -----           |       |           |       |       |
| nls_date_format | TRUE  | FALSE     | TRUE  |       |

```
SQL>
```

3. Find out the default date format for the database by querying the `NLS_TERRITORY` parameter in the `V$PARAMETER` view. Include a `WHERE` clause to narrow down the query to just the `NLS_TERRITORY` parameter. Remember that in the `V$PARAMETER` view, the parameter names are in lowercase.

```
SQL> SELECT name, value FROM v$parameter
2 WHERE name = 'nls_territory';
```

| NAME          | VALUE   |
|---------------|---------|
| nls_territory | AMERICA |

```
SQL>
```

4. Connect to `PDB1`. Run a simple query against the sample data to view an example of the current default date format in use.

- a. Switch to `PDB1` by using the `ALTER SESSION` command.

```
SQL> ALTER SESSION SET container = PDB1;
Session altered.
```

```
SQL>
```

- b. Query the `LAST_NAME` and `HIRE_DATE` columns in the `HR.EMPLOYEES` table. Notice the date format is `dd-mon-rr`.

```
SQL> SELECT last_name, hire_date FROM hr.employees;
```

| LAST_NAME | HIRE_DATE |
|-----------|-----------|
| King      | 17-JUN-03 |
| Kochhar   | 21-SEP-05 |
| De Haan   | 13-JAN-01 |
| Hunold    | 03-JAN-06 |
| Ernst     | 21-MAY-07 |
| ...       |           |
| Mavris    | 07-JUN-02 |
| Baer      | 07-JUN-02 |
| Higgins   | 07-JUN-02 |
| Gietz     | 07-JUN-02 |

```
107 rows selected.
```

```
SQL>
```

5. Modify the `NLS_DATE_FORMAT` parameter to use the format `mon dd yyyy` by using the `ALTER SESSION` command.

```
SQL> ALTER SESSION SET nls_date_format = 'mon dd yyyy';
```



6. Rerun the query against the `HR.EMPLOYEES` table. Notice that the date format has changed from `dd-mon-rr` to `mon dd yyyy`.

```
SQL> SELECT last_name, hire_date FROM hr.employees;
LAST_NAME                HIRE_DATE
-----
King                      jun 17 2003
Kochhar                   sep 21 2005
De Haan                   jan 13 2001
Hunold                    jan 03 2006
Ernst                      may 21 2007
...
Mavris                    jun 07 2002
Baer                      jun 07 2002
Higgins                   jun 07 2002
Gietz                     jun 07 2002

107 rows selected.

SQL>
```

7. Query the `NLS_DATE_FORMAT` parameter again by using the `SHOW PARAMETER` command. The value column now reflects the custom date format.

```
SQL> SHOW PARAMETER nls_date_format
NAME                                TYPE      VALUE
-----
nls_date_format                     string    mon dd yyyy
SQL>
```

8. Disconnect from `PDB1` to end your session.

```
SQL> DISCONNECT
Disconnected from Oracle Database 19c Enterprise Edition Release
18.0.0.0.0 - Production
Version 18.1.0.0.0
SQL>
```

9. Connect to `PDB1` again as the `SYSTEM` user by using the Easy Connect syntax. See *Course Practice Environment: Security Credentials* for the `SYSTEM` user password. In Practice 5-3, step 6b, you queried `V$SERVICES`. Append the value in the query results following `ORCL` to `pdb1` to create the service name as shown in this example.

```
SQL> connect
system/password@localhost:1521/PDB1.588436052.oraclecloud.intern
al
Connected.
SQL>
```

10. Rerun the query against the `HR.EMPLOYEES` table. The date format has reverted back to the default format `dd-mon-rr`. A session-level parameter change only lasts for the duration of the session.

```
SQL> SELECT last_name, hire_date FROM hr.employees;
LAST_NAME                HIRE_DATE
-----
King                      17-JUN-03
Kochhar                   21-SEP-05
De Haan                   13-JAN-01
Hunold                    03-JAN-06
Ernst                     21-MAY-07
...
Mavris                    07-JUN-02
Baer                      07-JUN-02
Higgins                   07-JUN-02
Gietz                     07-JUN-02

107 rows selected.

SQL>
```

11. Query the `NLS_DATE_FORMAT` parameter again by using the `SHOW PARAMETER` command. The `VALUE` column no longer has the custom date format.

```
SQL> SHOW PARAMETER nls_date_format
NAME                                TYPE        VALUE
-----
nls_date_format                     string
```

### Modify a Dynamic System-Level Parameter

In this section, you modify the `JOB_QUEUE_PROCESSES` parameter. This parameter specifies the maximum number of job slaves per database instance that can be created for the execution of `DBMS_JOB` jobs and Oracle Scheduler (`DBMS_SCHEDULER`) jobs.

1. Exit SQL\*Plus, and connect to the root container with the `SYSDBA` privilege. If you try to update the `JOB_QUEUE_PROCESSES` parameter from `PDB1`, you'll get an error. Also, you'll need the `SYSDBA` privilege to restart the database instance later on.

```
SQL> exit
```

```
...
[oracle@MYDBCS ~]$ sqlplus / as sysdba
...
SQL>
```

2. Learn about the `JOB_QUEUE_PROCESSES` parameter by querying the `V$PARAMETER` view. Include a `WHERE` clause to narrow down the query to just the `JOB_QUEUE_PROCESSES` parameter. Remember that in the `V$PARAMETER` view, the parameter names are in lowercase.

```
SQL> SELECT name, isses_modifiable, issys_modifiable, value
FROM v$parameter WHERE name = 'job_queue_processes';
```

| NAME                | ISSES | ISSYS_MOD | VALUE |
|---------------------|-------|-----------|-------|
| job_queue_processes | FALSE | IMMEDIATE | 4000  |

```
SQL>
```

3. Change the `JOB_QUEUE_PROCESSES` parameter value to 15 by using the `ALTER SYSTEM` command. Set `SCOPE` equal to `BOTH` so that the change happens in both the database instance memory (which makes the change immediate) and in the `SPFILE` (which makes the change permanent).

```
SQL> ALTER SYSTEM SET job_queue_processes=15 SCOPE=BOTH;
```

```
System altered.
```

```
SQL>
```

4. Use the `SHOW PARAMETER` command to verify that the `JOB_QUEUE_PROCESSES` parameter value is now equal to 15. Notice that only `job` was entered with the `SHOW PARAMETER` command instead of the full name, `job_queue_processes`. Remember, when you use the `SHOW PARAMETER` command, you don't have to enter the full name. The database server will find all parameters that contain the letters `job`. In this example, the database server found two parameters that contain the letters `job`: `job_queue_processes` and `max_datapump_jobs_per_pdb`. The query result indicates that the `job_queue_processes` value in memory is now 15.

```
SQL> SHOW PARAMETER job
```

| NAME                      | TYPE    | VALUE |
|---------------------------|---------|-------|
| job_queue_processes       | integer | 15    |
| max_datapump_jobs_per_pdb | integer | 100   |

```
SQL>
```

5. Verify that the new value for the `JOB_QUEUE_PROCESSES` parameter persists after the database instance is restarted.

- a. Shut down the database instance with the `IMMEDIATE` mode.

```
SQL> SHUTDOWN IMMEDIATE
Database closed.
Database dismounted.
ORACLE instance shut down.
SQL>
```

- b. Start the database instance by using the `STARTUP` command.

```
SQL> STARTUP
ORACLE instance started.

Total System Global Area 2768239832 bytes
Fixed Size                  8899800 bytes
Variable Size              704643072 bytes
Database Buffers 1979711488 bytes  Redo
Buffers                    74985472 bytes
Database mounted.
Database opened.
SQL>
```

- c. View the configuration for the `JOB_QUEUE_PROCESSES` parameter again by using the `SHOW PARAMETER` command. The value is 15, which proves that your change to the parameter persisted after the database instance was restarted.

| SQL> SHOW PARAMETER job   |         |       |
|---------------------------|---------|-------|
| NAME                      | TYPE    | VALUE |
| job_queue_processes       | integer | 15    |
| max_datapump_jobs_per_pdb | integer | 100   |
| SQL>                      |         |       |

### Modify a Static System-Level Parameter

In this section, you modify the `SEC_MAX_FAILED_LOGIN_ATTEMPTS` parameter. This parameter specifies the number of authentication attempts that can be made by a client on a connection to the server process. These login attempts can be for multiple user accounts in the same connection. After the specified number of failure attempts, the connection will be automatically dropped by the server process.

1. Learn about the `SEC_MAX_FAILED_LOGIN_ATTEMPTS` parameter by querying the `V$PARAMETER` view. Include a `WHERE` clause to narrow down the query to just the `SEC_MAX_FAILED_LOGIN_ATTEMPTS` parameter. Remember that in the `V$PARAMETER` view, the parameter names are in lowercase. The query results below have been formatted for easier viewing.

| SQL> <b>SELECT name, isses_modifiable, issys_modifiable, value</b><br><b>FROM v\$parameter WHERE name = 'sec_max_failed_login_attempts';</b> |       |           |       |
|--|-------|-----------|-------|
| NAME   | ISSES | ISSYS_MOD | VALUE |
| sec_max_failed_login_attempts  | FALSE | FALSE     | 3     |
| SQL>   |       |           |       |

2. Change the SEC\_MAX\_FAILED\_LOGIN\_ATTEMPTS parameter value to 2 by using the ALTER SYSTEM command. Include the comment 'Reduce for tighter security' and set the scope equal to SPFILE so that the change is made only in the SPFILE. When you specify SCOPE as SPFILE or as BOTH, an optional COMMENT clause lets you associate a text string with the parameter update. The comment is written to the SPFILE.

|   |  |  |  |
|---|--|--|--|
| SQL> <b>ALTER SYSTEM SET sec_max_failed_login_attempts = 2</b><br><b>COMMENT='Reduce for tighter security.' SCOPE=SPFILE;</b> |  |  |  |
| System altered.   |  |  |  |
| SQL>  |  |  |  |

3. View the SEC\_MAX\_FAILED\_LOGIN\_ATTEMPTS parameter value by using the SHOW PARAMETER command. The query result indicates that the value hasn't been updated yet. It's still equal to 3 because you need to restart the database instance for the change to take effect, which is required for static parameters.

| SQL> <b>SHOW PARAMETER sec_max</b> |         |       |  |
|------------------------------------|---------|-------|--|
| NAME                               | TYPE    | VALUE |  |
| sec_max_failed_login_attempts      | integer | 3     |  |
| SQL>                               |         |       |  |

4. Restart the database and then verify that the new value for the SEC\_MAX\_FAILED\_LOGIN\_ATTEMPTS parameter is updated.
  - a. Shut down the database instance with the IMMEDIATE mode.

|                                |  |  |  |
|--------------------------------|--|--|--|
| SQL> <b>SHUTDOWN immediate</b> |  |  |  |
| Database closed.               |  |  |  |
| Database dismounted.           |  |  |  |
| ORACLE instance shut down.     |  |  |  |
| SQL>                           |  |  |  |

- b. Start the database instance by using the STARTUP command.

|                          |  |  |  |
|--------------------------|--|--|--|
| SQL> <b>STARTUP</b>      |  |  |  |
| ORACLE instance started. |  |  |  |

```
Total System Global Area 2768239832 bytes
Fixed Size                  8899800 bytes
Variable Size               704643072 bytes
Database Buffers 1979711488 bytes   Redo
Buffers                     74985472 bytes
Database mounted.
Database opened.
SQL>
```

- c. View the SEC\_MAX\_FAILED\_LOGIN\_ATTEMPTS parameter value again by using the SHOW PARAMETER command. The query result indicates that the parameter's value was successfully changed to 2.

```
SQL> SHOW PARAMETER sec_max
```

| NAME                          | TYPE    | VALUE |
|-------------------------------|---------|-------|
| sec_max_failed_login_attempts | integer | 2     |

```
SQL>
```

- d. View the NAME and UPDATE\_COMMENT columns in the V\$PARAMETER view for the SEC\_MAX\_FAILED\_LOGIN\_ATTEMPTS parameter. Notice that the comment you added is stored in this view. The results below are formatted for easier reading.

```
SQL> SELECT name, update_comment
FROM v$parameter WHERE name='sec_max_failed_login_attempts';
```

| NAME                          | UPDATE_COMMENT               |
|-------------------------------|------------------------------|
| sec_max_failed_login_attempts | Reduce for tighter security. |

```
SQL>
```

5. Exit SQL\*Plus and close the terminal window.

```
SQL> EXIT
```

## Lab 6: Starting and Stopping the container

---

1. Source the `oraenv` script.

```
[oracle@MYDBCS ~]$ . oraenv
ORACLE_SID = [ORCL] ?
The Oracle base remains unchanged with value /u01/app/oracle
[oracle@MYDBCS ~]$
```

2. Start SQL\*Plus and log in to the database as the `SYS` user with the `SYSDBA` privilege.

```
$ sqlplus / as sysdba
...
SQL>
```

3. Shut down the database instance in `NORMAL` mode. Normal is the default shutdown mode if no mode is specified. During this mode of shutdown, the database instance closes the database—all data files and online redo log files are closed. Next, the database instance dismounts the database—all control files associated with the database instance are closed. Lastly, the Oracle software shuts down the database instance—background processes are terminated, and the System Global Area (SGA) is removed from memory. When a database instance shuts down in normal mode, the database instance waits for all users to disconnect before completing the shutdown, and no new connections are allowed. Control is not returned to the session that initiates a database shutdown until shutdown is complete.

```
SQL> SHUTDOWN
Database closed.
Database dismounted.
ORACLE instance shut down.
SQL>
```

4. Show the current user. Note that SQL\*Plus is still running and the current user is `SYS`.

```
SQL> SHOW USER
USER is "SYS"
SQL>
```

5. Show the current container name. This step returns an error because the database is shut down.

```
SQL> SHOW con_name
ERROR:
ORA-01034: ORACLE not available
Process ID: 0
Session ID: 0 Serial number: 0

SP2-1545: This feature requires Database availability.
SQL>
```

6. Start up the database instance in `NOMOUNT` mode. During this step, the Oracle software locates the parameter file (SPFILE or PFILE), allocates memory to the System Global Area (SGA), starts the background processes, and opens the alert log and trace files. At this stage, the database instance is started; however, users cannot access it yet. You would usually start in `NOMOUNT` mode if you were creating a database, re-creating control files, or performing certain backup and recovery tasks.

```
SQL> STARTUP NOMOUNT
ORACLE instance started.

Total System Global Area 2768239832 bytes
Fixed Size                  8899800 bytes
Variable Size               704643072 bytes
Database Buffers 1979711488 bytes  Redo
Buffers                    74985472 bytes
SQL>
```

7. Mount the database by using the `ALTER DATABASE MOUNT` command. During this step, the database instance mounts the database. This means that the database instance locates and opens all the control files specified in the initialization parameter file and reads the control files to obtain the names and statuses of the data files and online redo log files. The database instance does not, however, verify the existence of the data files and online redo log files at this time. You must mount the database, but not open it when you want to rename data files, enable/disable online redo log file archiving options, or perform a full database recovery.

```
SQL> ALTER DATABASE MOUNT;
Database altered.

SQL>
```

8. Open the database by using the `ALTER DATABASE` command. During this step, the database instance opens the data files for the CDB and online redo log files and checks the consistency of the database. When the database is open, all users can access the database instance.



```
SQL> ALTER DATABASE OPEN;  
Database altered.  
  
SQL>
```

9. Show the current container name.

```
SQL> SHOW con_name  
  
CON_NAME  
-----  
CDB$ROOT  
  
SQL>
```

10. Show the current user.

```
SQL> SHOW user  
USER is "SYS"  
  
SQL>
```

11. Determine the state of the database: `select open_mode from v$database;`

12. Check whether PDB1 is open by querying the OPEN\_MODE column in the V\$PDBS view.

```
SQL> COLUMN con_id FORMAT 999  
SQL> COLUMN name FORMAT A10  
SQL> SELECT con_id, name, open_mode FROM v$pdbs;
```

## Lab 7: View Diagnostic information

---

1. Start SQL\*Plus and log in to the database as the `SYS` user with the `SYSDBA` privilege.

```
[oracle@MYDBCS ~]$ sqlplus / as sysdba
...
SQL>
```

2. View the locations of the various diagnostics directories in the ADR. The results below have been formatted for easier reading.
  - The path that corresponds to the `Diag Alert` entry in the `NAME` column is for the XML version. This path is `/u01/app/oracle/diag/rdbms/orcl/ORCL/alert`.
  - The path that corresponds to the `Diag Trace` entry is for the text-only version. This path is `/u01/app/oracle/diag/rdbms/orcl/ORCL/trace`.

```
SQL> SELECT name, value FROM v$diag_info;
```

| NAME                  | VALUE  |
|-----------------------|--|
| Diag Enabled          | TRUE   |
| ADR Base              | /u01/app/oracle  |
| ADR Home              | /u01/app/oracle/diag/rdbms/orcl/ORCL                         |
| Diag Trace            | /u01/app/oracle/diag/rdbms/orcl/ORCL/trace                   |
| Diag Alert            | /u01/app/oracle/diag/rdbms/orcl/ORCL/alert                   |
| Diag Incident         | /u01/app/oracle/diag/rdbms/orcl/ORCL/incident                |
| Diag Cdump            | /u01/app/oracle/diag/rdbms/orcl/ORCL/cdump                   |
| Health Monitor        | /u01/app/oracle/diag/rdbms/orcl/ORCL/hm                      |
| Default Trace File    | /u01/app/oracle/diag/rdbms/orcl/ORCL/trace/ORCL_ora_2600.trc |
| Active Problem Count  | 1  |
| Active Incident Count | 6  |

11 rows selected.

SQL>

### 3. Exit SQL\*Plus.

```
SQL> EXIT
```

## Use an Editor to View the Alert Log

### 1. View the XML version of the alert log. The log.xml file is the XML version of the alert log.

#### a. Browse to the /u01/app/oracle/diag/rdbms/orcl/ORCL/alert directory.

```
[oracle@MYDBCS ~]$ cd /u01/app/oracle/diag/rdbms/orcl/ORCL/alert
[oracle@MYDBCS alert]$
```

#### b. List the contents of the directory. Notice that there is a log.xml file in this directory.

```
[oracle@MYDBCS alert]$ ls
log.xml
[oracle@MYDBCS alert]$
```

#### c. Use cat or more to scroll through the file. Notice that it is a chronological log of messages about non-default initialization parameters used at startup, errors, SQL statements, and so on. Oracle Database uses the alert log to keep a record of these events as an alternative to displaying the information on an operator's console.

```
[oracle@MYDBCS alert]$ more log.xml
<msg time='2018-03-07T22:19:08.858+00:00' org_id='oracle'
comp_id='rdbms'
msg_id='opistr_real:1244:2538814769' type='NOTIFICATION'
group='startup'
level='16' host_id='MYDBCS' host_addr='10.18.24.38'
pid='8090' version='1' con_uid='1'
```

```
con_id='1' con_name='CDB$ROOT'>
<txt>Starting ORACLE instance (normal) (OS id: 8090)
...
```

2. View the text-only version of the alert log.

- a. Change to the `/u01/app/oracle/diag/rdbms/orcl/ORCL/trace` directory.

```
[oracle@MYDBCS alert]$ cd
/u01/app/oracle/diag/rdbms/orcl/ORCL/trace
[oracle@MYDBCS trace]$
```

- b. The `alert_ORCL.log` (format is `alert_SID.log`) file is the text-only version. In this directory, you also have server process trace files (TRC files) and trace map files (TRM files). Each server and background process can write to an associated trace file. When a process detects an internal error, it dumps information about the error to its trace file. Trace map files contain structural information about trace files and are used for searching and navigation.

```
[oracle@MYDBCS trace]$ ls
alert_ORCL.log                ORCL_ora_12783.trc
ORCL_ora_27673.trc
alert_ORCL_v20180310.log.gz   ORCL_ora_12783.trm
ORCL_ora_27673.trm
ORCL_aqpc_10857.trc           ORCL_ora_12785.trc
ORCL_ora_27674.trc
...
[oracle@MYDBCS trace]$
```

- c. Open the file with an editor or use a command such as `tail` to view the contents of the alert log.

```
[oracle@MYDBCS trace]$ tail -500 alert_ORCL.log
2018-03-15T20:41:22.507272+00:00
db_recovery_file_dest_size of 6144 MB is 33.63% used. This is a
user-specified limit on the amount of space that will be used by
this
database for recovery-related files, and does not reflect the
amount of
space available in the underlying filesystem or ASM diskgroup.
...
Pluggable database PDB1 opened read write
2018-03-16T15:10:48.837547+00:00
Completed: ALTER PLUGGABLE DATABASE ALL OPEN
Starting background process CJQ0
Completed: ALTER DATABASE OPEN
2018-03-16T15:10:49.521562+00:00
CJQ0 started with pid=65, OS id=4909
[oracle@MYDBCS trace]$
```

## Use ADRCI to View the Alert Log

1. Start the ADRCI tool. Recall that you set the Oracle environment variables at the beginning of this practice; however, only the `ORACLE_HOME` environment variable needs to be set prior to starting ADRCI. If you ever need to set just that one variable, you can do so by entering the following at the command prompt: `export`

`PATH=$PATH:$ORACLE_HOME/bin.`

```
[oracle@MYDBCS trace]$ adrci
```

2. View the alert log by using the `SHOW ALERT` command. The `SHOW ALERT` command opens the alert log file in the `vi` editor, by default.

```
adrci> SHOW ALERT
ADR Home = /u01/app/oracle/diag/rdbms/orcl/ORCL:
*****
*****
Output the results to file: /tmp/alert_11237_1404_ORCL_1.ado
2018-03-07 22:19:08.858000 +00:00
Starting ORACLE instance (normal) (OS id: 8090)
*****
*****
Dump of system resources acquired for SHARED GLOBAL AREA (SGA)
Per process system memlock (soft) limit = 128G
Expected per process system memlock (soft) limit to lock
SHARED GLOBAL AREA (SGA) into memory: 2642M
Available system pagesizes:
4K, 2048K
...
```

3. Enter **G** to move to bottom of the alert file.

```
2018-03-16 15:10:45.273000 +00:00
Opening pdb with no Resource Manager plan active
2018-03-16 15:10:47.226000 +00:00
Pluggable database PDB1 opened read write
2018-03-16 15:10:48.837000 +00:00
Completed: ALTER PLUGGABLE DATABASE ALL OPEN
```

```
Starting background process CJQ0
Completed: ALTER DATABASE OPEN
CJQ0 started with pid=65, OS id=4909
```

4. Enter **?Starting ORACLE instance?** and press return. Press **N** to search from the bottom of the file to find the last time the instance was started. The following will be similar to your alert log. Note: Here lowercase and uppercase are important because `vi` distinguishes them, unless you ignore them by setting `:set ic`.

```
Starting ORACLE instance (normal) (OS id: 8090)
*****
*****
Dump of system resources acquired for SHARED GLOBAL AREA (SGA)
Per process system memlock (soft) limit = 128G
Expected per process system memlock (soft) limit to lock
SHARED GLOBAL AREA (SGA) into memory: 2642M
Available system pagesizes:
  4K, 2048K
Supported system pagesize(s):
  PAGESIZE  AVAILABLE_PAGES  EXPECTED_PAGES  ALLOCATED_PAGES
ERROR(s)
          4K          Configured              4          675844
NONE
          2048K              0              1321              0
NONE
RECOMMENDATION:
  1. For optimal performance, configure system with expected
number
  of pages for every supported system pagesize prior to the next
instance restart operation.
*****
*****
LICENSE_MAX_SESSION = 0
LICENSE_SESSIONS_WARNING = 0
Initial number of CPU is 2
Number of processor cores in the system is 2
Number of processor sockets in the system is 1
search hit BOTTOM, continuing at TOP
```

5. Search forward by entering **/ ALTER** to find the line that starts with `ALTER DATABASE MOUNT`. Here lowercase and uppercase are important because `vi` distinguishes them.

```
ALTER DATABASE  MOUNT
2018-03-07 22:21:04.104000 +00:00
Using default pga_aggregate_limit of 3505 MB
2018-03-07 22:21:06.471000 +00:00
```

```
.... (PID:9128): Redo network throttle feature is disabled at
mount time
Successful mount of redo thread 1, with mount id 2299076813
Database mounted in Exclusive Mode
Lost write protection disabled
.... (PID:9128): Using STANDBY_ARCHIVE_DEST parameter default
value as USE_DB_RECOVERY_FILE_DEST [krsd.c:17695]
Completed: ALTER DATABASE MOUNT
```

6. Search forward again by entering / **ALTER** to find the line that starts with ALTER DATABASE OPEN. Notice that the stages that the database goes through during startup are MOUNT and OPEN.

```
ALTER DATABASE OPEN
Ping without log force is disabled:
  instance mounted in exclusive mode.
Buffer Cache Full DB Caching mode changing from FULL CACHING
DISABLED to FULL CACHING ENABLED
Crash Recovery excluding pdb 2 which was cleanly closed.
Crash Recovery excluding pdb 3 which was cleanly closed.
2018-03-07 22:21:08.617000 +00:00
Endian type of dictionary set to little
LGWR (PID:9095): STARTING ARCH PROCESSES
Starting background process ARC0
```

7. Exit the vi editor by entering :q and pressing Enter.
8. Exit adrci and close the terminal window.

```
adrci > exit
[oracle@MYDBCS trace]$
```

## Log DDL Statements in the DDL Log File

1. Determine if DDL logging is enabled in PDB1. If not, enable it by setting the value for the ENABLE\_DDL\_LOGGING initialization parameter to TRUE.
  - a. Start SQL\*Plus and log in to the database as the SYS user with the SYSDBA privilege.

```
[oracle@MYDBCS trace]$ sqlplus / as sysdba
...
SQL>
```

- b. Switch to PDB1.

```
SQL> ALTER SESSION SET CONTAINER = PDB1;
Session altered.

SQL>
```

- c. Issue the `SHOW PARAMETER` command to view the value for `ENABLE_DDL_LOGGING`. In Oracle Database Cloud Service, `ENABLE_DDL_LOGGING` is set to `TRUE` by default. The default value for `ENABLE_DDL_LOGGING` is `FALSE` in non-Cloud installations.

```
SQL> SHOW PARAMETER enable_ddl_logging
```

| NAME               | TYPE    | VALUE |
|--------------------|---------|-------|
| enable_ddl_logging | boolean | TRUE  |

```
SQL>
```

- d. If DDL logging was not enabled, you could enable it for just this session by using the `ALTER SESSION` command.

```
SQL> ALTER SESSION SET enable_ddl_logging = TRUE;
Session altered.

SQL>
```

2. Create and drop a table to generate statements that will logged.

```
SQL> CREATE TABLE TEST (name varchar2(15));
Table created.
SQL> DROP TABLE TEST;
Table dropped.
SQL>
```

3. Exit SQL\*Plus.

```
SQL> EXIT
...
[oracle@MYDBCS trace]$
```

4. Change to the directory where the text version of the DDL log file resides.

```
[oracle@MYDBCS trace]$ cd
/u01/app/oracle/diag/rdbms/orcl/ORCL/log
[oracle@MYDBCS log]$
```

5. List the contents of the log directory.

```
[oracle@MYDBCS log]$ ls
ddl  ddl_ORCL.log  debug  debug.log  hcs  imdb  test
[oracle@MYDBCS log]$
```

6. View the `ddl_ORCL.log` file by using the `cat` command. Your output will be different from the output shown below.

```
[oracle@MYDBCS log]$ cat ddl_ORCL.log
2018-03-16T19:31:30.795903+00:00
diag_adl:CREATE TABLE TEST (name varchar2(15) )
2018-03-16T19:31:57.762139+00:00
diag_adl:DROP TABLE TEST
[oracle@MYDBCS log]$
```



## Lab 8: Creating a Local Users

---

### Overview

In this practice, you log in to PDB1 as the local administrator (PDB1\_ADMIN1) and create a local user account called INVENTORY, which will own the new Inventory software application.

INVENTORY is an example of a user account that does not represent a person.

### Assumptions

You are logged in to the compute node as the `oracle` user.

### Tasks

#### Create the INVENTORY User Account

1. Start SQL\*Plus and connect to PDB1 as the PDB1\_ADMIN1 user.

```
[oracle@MYDBCS ~]$ sqlplus PDB1_ADMIN1/password@PDB1
...
SQL>
```

2. Create a local user account named INVENTORY. Set the default tablespace to the USERS tablespace and grant unlimited quota on that tablespace. Refer to *Course Practice Environment: Security Credentials* for the password value.

```
SQL> CREATE USER INVENTORY IDENTIFIED BY password DEFAULT
TABLESPACE users QUOTA UNLIMITED ON users;

User created.

SQL>
```

3. Grant the CREATE SESSION privilege to INVENTORY.

```
SQL> GRANT CREATE SESSION TO INVENTORY;

Grant succeeded.

SQL>
```

4. List the local user accounts for PDB1 by querying the DBA\_USERS view. The INVENTORY account is included in the list.

```
SQL> SELECT DISTINCT username FROM dba_users WHERE common='NO'
ORDER BY username;

USERNAME
-----
APEX_050100
APEX_INSTANCE_ADMIN_USER
```

```
APEX_LISTENER
APEX_PUBLIC_USER
APEX_REST_PUBLIC_USER
FLOWS_FILES
HR
INVENTORY
PDB1_ADMIN1
PDBADMIN
SCOTT

11 rows selected.

SQL>
```

### Connect as INVENTORY and Verify Privileges

1. Disconnect PDB1\_ADMIN1 from PDB1.

```
SQL> DISCONNECT
...
SQL>
```

2. Verify that the INVENTORY user account can connect to PDB1.

```
SQL> CONNECT INVENTORY/password@PDB1
Connected.
SQL>
```

3. List the privileges for INVENTORY by querying the SESSION\_PRIVS view. The results show that INVENTORY has the CREATE SESSION privilege.

```
SQL> SELECT * FROM session_privs ORDER BY privilege;

PRIVILEGE
-----
CREATE SESSION

SQL>
```

4. Exit SQL\*Plus.

```
SQL> EXIT
...
[oracle@MYDBCS ~]$
```

## Lab 9: Granting Privilege

---

### Explore the Privileges and Roles Granted to PDBADMIN

1. Start SQL\*Plus and connect as the SYS user with the SYSDBA privilege.

Note: PDBADMIN does not have the required privileges to view data from the DBA\_SYS\_PRIVS view in PDB1, which you will do in the next step.

```
$ sqlplus / AS SYSDBA
...
SQL>
```

2. List the system privileges granted to the PDBADMIN user by querying the DBA\_SYS\_PRIVS view. This view describes system privileges granted to users and roles. The results show that no system privileges are explicitly granted to PDBADMIN. However, there may be privileges granted through roles.

```
SQL> SELECT * FROM dba_sys_privs WHERE grantee='PDBADMIN';
no rows selected

SQL>
```

3. List the roles granted to the PDBADMIN user by querying the CDB\_ROLE\_PRIVS view. This view describes the roles granted to all users and roles in the database. The results show that PDBADMIN is granted the PDB\_DBA role. Also, the ADMIN OPTION is enabled (ADM=YES), which means that PDBADMIN can grant the PDB\_DBA role to other users.

```
SQL> col granted_role format a10
SQL> SELECT granted_role, admin_option FROM cdb_role_privs WHERE
grantee='PDBADMIN';

GRANTED_RO  ADM
-----
PDB_DBA      YES
```

```
DBA          NO
```

```
SQL>
```

4. List the system privileges granted to the `PDB_DBA` role by querying the `ROLE_SYS_PRIVS` view.

- a. Switch to `PDB1`. You must be connected to `PDB1` to retrieve data, and you must be connected as the `SYS` user.

```
SQL> ALTER SESSION SET CONTAINER = PDB1;
```

```
Session altered.
```

```
SQL>
```

- b. Query the `ROLE_SYS_PRIVS` view. This view describes system privileges granted to roles. Information is provided only about roles to which the user has access. Because you're connected to `PDB1` as the `SYS` user, you have access to all role information. The results show that the `PDB_DBA` role consists of three system privileges: `CREATE SESSION`, `SET CONTAINER`, and `CREATE PLUGGABLE DATABASE`.

```
SQL> SELECT privilege FROM role_sys_privs WHERE role='PDB_DBA'
ORDER BY privilege;
```

```
PRIVILEGE
```

```
-----
CREATE PLUGGABLE DATABASE
```

```
CREATE SESSION
```

```
SET CONTAINER
```

```
SQL>
```

5. List the roles that are granted to the `PDB_DBA` role by querying the `DBA_ROLE_PRIVS` view. The results show that the `PDB_DBA` role is granted the `CONNECT` role.

```
SQL> SELECT granted_role FROM dba_role_privs WHERE grantee =
'PDB_DBA';
```

```
GRANTED_RO
```

```
-----
CONNECT
```

```
SQL>
```

6. List the privileges granted to the `CONNECT` role by querying the `ROLE_SYS_PRIVS` view. The results show that the `CONNECT` role consists of the `SET CONTAINER` and `CREATE SESSION` privileges.

```
SQL> SELECT privilege FROM role_sys_privs WHERE role='CONNECT'
ORDER BY privilege;
```

```
PRIVILEGE

CREATE SESSION
SET CONTAINER

SQL>
```

7. Let's summarize our findings: From these queries, you learned that the `PDBADMIN` user is granted the `PDB_DBA` role by default, and that role consists of the `CONNECT` role and the `CREATE PLUGGABLE DATABASE` system privilege. The `CONNECT` role contains the `SET CONTAINER` and `CREATE SESSION` system privileges.

### Grant the DBARole to PDBADMIN

1. Grant the DBA role locally to PDBADMIN.

```
SQL> GRANT dba TO pdbadmin;

Grant succeeded.

SQL>
```

2. List the roles that are granted to PDBADMIN by querying the `DBA_ROLE_PRIVS` view. The results show that PDBADMIN is now granted the DBA and PDB\_DBA roles.

```
SQL> SELECT granted_role FROM dba_role_privs WHERE grantee =
'PDBADMIN' ORDER BY granted_role;

GRANTED_ROLE
-----
DBA
PDB_DBA

SQL>
```

3. Exit SQL\*Plus and close the terminal window.

```
SQL> EXIT

...

[oracle@MYDBCS ~]$ exit
```

## Lab 10: Creating a Default Role for a User

---

### Tasks

#### Configure a Default Role for JGOODMAN

1. Start SQL\*Plus and connect to PDB1 as the PDBADMIN user.

```
$ sqlplus PDBADMIN/password@PDB1
...
SQL>
```

2. View the current roles for JGOODMAN by querying the DBA\_ROLE\_PRIVS view. Also, show whether the roles are default roles. The results show that JGOODMAN is granted two roles, HRMANAGER and HRCLERK, and both are default roles (the DEF column = YES).

```
SQL> COLUMN granted_role FORMAT A20
SQL> SELECT granted_role, default_role FROM dba_role_privs WHERE
grantee='JGOODMAN';
```

| GRANTED_ROLE | DEF |
|--------------|-----|
|--------------|-----|

```
-----
HRCLERK                YES
HRMANAGER              YES

SQL>
```

```
SQL> ALTER USER JGOODMAN DEFAULT ROLE HRCLERK;
```

```
User altered.
```

```
SQL>
```

3. Set the default role for JGOODMAN to be HRCLERK only by using the ALTER USER command and DEFAULT ROLE clause.
4. View the current roles and default role settings for JGOODMAN again by querying the DBA\_ROLE\_PRIVS view. The results show that the default role is HRCLERK and the HRMANAGER role is no longer a default role. Jenny still has this role; however, she'll need to enable it to exercise its privileges.

```
SQL> SELECT granted_role, default_role FROM dba_role_privs WHERE  
grantee='JGOODMAN';
```

| GRANTED_ROLE | DEF  |
|--------------|------|
| -----        | ---- |
| HRCLERK      | YES  |
| HRMANAGER    | NO   |

```
SQL>
```

5. Disconnect PDBADMIN from PDB1.

```
SQL> DISCONNECT
```

```
...
```

```
SQL>
```

## Enable a Non-Default Role

1. Connect to PDB1 as JGOODMAN.

```
SQL> CONNECT JGOODMAN/password@PDB1
```

```
Connected.
```

```
SQL>
```

2. View the roles for the current session. Notice that the default role, HRCLERK, is in effect.

```
SQL> SELECT * FROM session_roles;
```

```
ROLE
```

```
-----  
HRCLERK
```

```
SQL>
```

```
SQL> SET ROLE HRMANAGER;
```

```
Role set.
```

```
SQL>
```

3. Suppose JGOODMAN needs to operate as an HR Manager, and not an HR Clerk. Change the enabled role to HRMANAGER. Caution: If you use the SET ROLE command, any roles not included in the command will be disabled.
4. View the roles for the current session again. The HRMANAGER role is now enabled.

```
SQL> SELECT * FROM session_roles;
```

```
ROLE
```

```
-----  
HRMANAGER
```

```
SQL>
```

5. Suppose JGOODMAN needs both roles. Use the SET ROLE command to enable them both.

```
SQL> SET ROLE HRMANAGER, HRCLERK;
```

```
Role set.
```

```
SQL>
```

6. View the roles for the current session again. The HRMANAGER and HRCLERK roles are now in effect.

```
SQL> SELECT * FROM session_roles;
```

```
ROLE
```

```
-----  
HRCLERK  
HRMANAGER
```

```
SQL>
```

7. Exit SQL\*Plus.

```
SQL> EXIT
```

```
...
```

```
[oracle@MYDBCS ~]$
```



## Lab 11 : Creating a PDB from Seed

---

### Overview

In this practice, you create an empty PDB named `PDB2` in your CDB by using the seed PDB.

Note: You can use Database Configuration Assistant, SQL Developer, or SQL commands to create a PDB from seed. This practice shows you how to do it by using SQL commands in SQL\*Plus.

### Assumptions

You are logged in as the `oracleuser`.

### Tasks

1. Open a new terminal window and connect to the compute node as the `oracleuser`.

```
$ sudo su - oracle
```

2. Source the `oraenv` script.

```
[oracle@MYDBCS ~]$ . oraenv
ORACLE_SID = [ORCL] ?
The Oracle base remains unchanged with value /u01/app/oracle
[oracle@MYDBCS ~]$
```

3. Start SQL\*Plus and log in to your CDB with the `SYSDBA` privilege.

```
[oracle@MYDBCS ~]$ sqlplus / as sysdba
...
SQL>
```

4. Create `PDB2` by using the `CREATE PLUGGABLE DATABASE` command. Specify an admin user named `PDB2ADMIN` and grant this user the `DBA` role. Refer to *Course Practice Environment: Security Credentials* for the password value.

```
SQL> CREATE PLUGGABLE DATABASE PDB2
      2 ADMIN USER PDB2ADMIN IDENTIFIED BY password
      3 ROLES= (dba) ;

Pluggable database created.

SQL>
```

In a non-DBCS installation of Oracle Database, the seed PDB does not have a `USERS` tablespace. You can include the `DEFAULT TABLESPACE USERS` clause to create a default permanent tablespace for any non-administrative users for which you do not specify a different permanent tablespace as shown in this example:

```
CREATE PLUGGABLE DATABASE PDB2
...
  DEFAULT TABLESPACE USERS
  DATAFILE '/u02/app/oracle/oradata/ORCL/PDB2/users01.dbf'
  SIZE 250M AUTOEXTEND ON
...
```

In DBCS, OMF is enabled by default, so the datafiles for PDB2 will be created in the location set by the `DB_CREATE_FILE_DEST` initialization parameter. In a database that is not OMF-enabled, you can specify the target location of the data files by using the `FILE_NAME_CONVERT` clause. This clause enables you to specify the target locations of the files based on the names of the source files. The first parameter in the clause is the source directory of the seed data files. The second is the destination directory for the new PDB data files. Here is an example using the `FILE_NAME_CONVERT` clause:

```
CREATE PLUGGABLE DATABASE PDB2
...
  FILE_NAME_CONVERT=(
    '/u02/app/oracle/oradata/ORCL/pdbseed/',
    '/u02/app/oracle/oradata/ORCL/PDB2/',
    '/u04/app/oracle/oradata/temp/',
    '/u04/app/oracle/oradata/temp/PDB2/')
...
```

## 5. Open PDB1.

- a. View the open mode for PDB1. After a PDB is created, its open mode is `MOUNTED`. When a PDB is in mounted mode, it behaves like a CDB in mounted mode. It does not allow changes to any objects, and it is accessible only to database administrators connected as `SYSDBA`. Information about the PDB is removed from memory caches. Cold backups of the PDB are possible.

```
SQL> COLUMN con_id FORMAT 999
SQL> COLUMN name FORMAT A10
SQL> SELECT con_id, name, open_mode FROM v$pdb;
```

| CON_ID | NAME      | OPEN_MODE  |
|--------|-----------|------------|
| 2      | PDB\$SEED | READ ONLY  |
| 3      | PDB1      | READ WRITE |
| 4      | PDB2      | MOUNTED    |

- b. Open PDB2 by using the `ALTER PLUGGABLE DATABASE` command.

```
SQL> ALTER PLUGGABLE DATABASE PDB2 OPEN;
```

Pluggable database altered.

SQL>

- c. Verify that the open mode for PDB2 is now READ WRITE.

SQL> **SELECT con\_id, name, open\_mode FROM v\$pdb;**

| CON_ID | NAME      | OPEN_MODE  |
|--------|-----------|------------|
| 2      | PDB\$SEED | READ ONLY  |
| 3      | PDB1      | READ WRITE |
| 4      | PDB2      | READ WRITE |

SQL>

6. View the list of services registered with the listener. When you create a PDB, a service is created and started. The name of the service is the same name as the PDB. You will connect to this service in the next step.

SQL> **!lsnrctl status**

LSNRCTL for Linux: Version 18.0.0.0.0 - Production on 22-MAR-2018 16:00:59

Copyright (c) 1991, 2017, Oracle. All rights reserved.

Connecting to

(DESCRIPTION=(ADDRESS=(PROTOCOL=TCP) (HOST=MYDBCS.compute-588436052.oraclecloud.internal) (PORT=1521)))

STATUS of the LISTENER

|              |                                       |
|--------------|---------------------------------------|
| Alias        | LISTENER                              |
| Version      | TNSLSNR for Linux: Version 18.0.0.0.0 |
| - Production |                                       |
| Start Date   | 19-MAR-2018 15:23:07                  |
| Uptime       | 3 days 0 hr. 37 min. 52 sec           |
| Trace Level  | off                                   |
| Security     | ON: Local OS Authentication           |
| SNMP         | OFF                                   |

Listener Parameter File

/u01/app/oracle/product/18.0.0/dbhome\_1/network/admin/listener.ora

Listener Log File

/u01/app/oracle/diag/tnslsnr/MYDBCS/listener/alert/log.xml

Listening Endpoints Summary...

(DESCRIPTION=(ADDRESS=(PROTOCOL=tcp) (HOST=MYDBCS.compute-588436052.oraclecloud.internal) (PORT=1521)))

(DESCRIPTION=(ADDRESS=(PROTOCOL=ipc) (KEY=EXTPROC1521)))

```

      (DESCRIPTION=(ADDRESS=(PROTOCOL=tcps) (HOST=MYDBCS.compute-
588436052.oraclecloud.internal) (PORT=5500)) (Security=(my_wallet_
directory=/u01/app/oracle/admin/ORCL/xdw_wallet)) (Presentation=H
TTP) (Session=RAW))
Services Summary...
Service
"66db11a937912ac9e0532618120ab4b9.588436052.oraclecloud.internal
" has 1 instance(s).
    Instance "ORCL", status READY, has 1 handler(s) for this
service...
Service
"68035adc452a1241e0532618120ad62d.588436052.oraclecloud.internal
" has 1 instance(s).
    Instance "ORCL", status READY, has 1 handler(s) for this
service...
Service "ORCL.588436052.oraclecloud.internal" has 1 instance(s).
    Instance "ORCL", status READY, has 1 handler(s) for this
service...
Service "ORCL.588436052.oraclecloud.internalXDB" has 1
instance(s).
    Instance "ORCL", status READY, has 1 handler(s) for this
service...
Service "PDB1.588436052.oraclecloud.internal" has 1 instance(s).
    Instance "ORCL", status READY, has 1 handler(s) for this
service...
Service "PDB2.588436052.oraclecloud.internal" has 1 instance(s).
    Instance "ORCL", status READY, has 1 handler(s) for this
service...
The command completed successfully

SQL>

```

7. Connect to PDB2 as the PDB2ADMIN user by using the Easy Connect method.

```

SQL> CONNECT
PDB2ADMIN/password@localhost:1521/PDB2.588436052.oraclecloud.int
ernal
Connected.
SQL>

```

Note: Alternatively, you could have switched to PDB2 by using the ALTER SESSION command.

```
ALTER SESSION SET container = PDB2;
```

8. Explore PDB2.

- a. Show the current container.

```
SQL> SHOW con_name
```

```
CON_NAME
```

```
-----
```

```
PDB2
```

```
SQL>
```

- b. Show the current container ID.

```
SQL> SHOW con_id
```

```
CON_ID
```

```
-----
```

```
4
```

```
SQL>
```

- c. List the service for PDB2 by querying the V\$SERVICES view.

```
SQL> COLUMN name FORMAT A20
```

```
SQL> SELECT name FROM v$services;
```

```
NAME
```

```
-----
```

```
PDB2
```

```
SQL>
```

- d. List the data files for PDB2 and their respective tablespaces by querying the DBA\_DATA\_FILES view. Recall that DBCS uses OMF by default, so the files were created with the OMF file naming format.

```
SQL> col file_name format a55
```

```
SQL> col tablespace_name format a10
```

```
SQL> SELECT file_name, tablespace_name FROM dba_data_files;
```

```
FILE_NAME
```

```
TABLESPACE
```

```
-----
```

```
/u02/app/oracle/oradata/ORCL/6D975E8B80B85F14E0537A051D SYSTEM  
0A3C0D/datafile/o1_mf_system_fk2t2tmq_.dbf
```

```
/u02/app/oracle/oradata/ORCL/6D975E8B80B85F14E0537A051D SYSAUX  
0A3C0D/datafile/o1_mf_sysaux_fk2t2tmv_.dbf
```

```
/u02/app/oracle/oradata/ORCL/6D975E8B80B85F14E0537A051D UNDOTBS1  
0A3C0D/datafile/o1_mf_undotbs1_fk2t2tmy_.dbf
```

```
/u02/app/oracle/oradata/ORCL/6D975E8B80B85F14E0537A051D USERS  
0A3C0D/datafile/o1_mf_users_fk2t2tn4_.dbf
```

```
SQL>
```

- e. List the temp files for PDB2 by querying the DBA\_TEMP\_FILES view. The query returns one temp file. Your temp file name will be different from the one shown below.

```
SQL> SELECT file_name, tablespace_name FROM dba_temp_files;

FILE_NAME                                TABLESPACE
-----
/u02/app/oracle/oradata/ORCL/6D975E8B80B85F14E0537A051D TEMP
0A3C0D/datafile/o1_mf_temp_fk2t2tn2_.dbf

SQL>
```

- f. List the local users for PDB2 by querying the DBA\_USERS view.

```
SQL> SELECT DISTINCT username FROM dba_users WHERE common='NO'
ORDER BY username;

USERNAME
-----
APEX_050100
APEX_INSTANCE_ADMIN_USER
APEX_LISTENER
APEX_PUBLIC_USER
APEX_REST_PUBLIC_USER
FLOWS_FILES
PDB2ADMIN

7 rows selected.

SQL>
```

- g. List the common users for PDB2 by querying the DBA\_USERS view.

```
SQL> SELECT DISTINCT username FROM dba_users WHERE common='YES'
ORDER BY username;

USERNAME
-----
ANONYMOUS
APPQOSSYS
AUDSYS
C##CDB_ADMIN1
...
SYSTEM
WMSYS
XDB
```

```
XS$NULL

39 rows selected.

SQL>
```

h. Exit SQL\*Plus.

```
SQL > EXIT

...

[oracle@MYDBCS ~]$
```

9. Add a service name entry to the `tnsnames.ora` file for PDB2.

a. Change the directory to `$ORACLE_HOME/network/admin`.

```
[oracle@MYDBCS ~]$ cd $ORACLE_HOME/network/admin
[oracle@MYDBCS admin]$
```

b. View the `tnsnames.ora` file by using the `cat` command.

```
[oracle@MYDBCS admin]$ cat tnsnames.ora

ORCL =
  (DESCRIPTION =
    (ADDRESS = (PROTOCOL = TCP) (HOST = MYDBCS.compute-
588436052.oraclecloud.internal) (PORT = 1521))
    (CONNECT_DATA =
      (SERVER = DEDICATED)
      (SERVICE_NAME = ORCL.588436052.oraclecloud.internal)
    )
  )

PDB1 =
  (DESCRIPTION =
    (ADDRESS = (PROTOCOL = TCP) (HOST = MYDBCS.compute-
588436052.oraclecloud.internal) (PORT = 1521))
    (CONNECT_DATA =
      (SERVER = DEDICATED)
      (SERVICE_NAME = pdb1.588436052.oraclecloud.internal)
    )
  )

[oracle@MYDBCS admin]$
```

c. Use an editor such as `vi` to add an entry for PDB2 to the `tnsnames.ora` file.

```
PDB2 =
  (DESCRIPTION =
```

```
(ADDRESS = (PROTOCOL = TCP) (HOST = MYDBCS.compute-588436052.oraclecloud.internal) (PORT = 1521))
  (CONNECT_DATA =
    (SERVER = DEDICATED)
    (SERVICE_NAME = pdb2.588436052.oraclecloud.internal)
  )
)
```

- d. Use the `cat` command to view the `tnsnames.ora` file and ensure that your new entry is formatted correctly.

```
[oracle@MYDBCS admin]$ cat tnsnames.ora

ORCL =
  (DESCRIPTION =
    (ADDRESS = (PROTOCOL = TCP) (HOST = MYDBCS.compute-588436052.oraclecloud.internal) (PORT = 1521))
    (CONNECT_DATA =
      (SERVER = DEDICATED)
      (SERVICE_NAME = ORCL.588436052.oraclecloud.internal)
    )
  )

PDB1 =
  (DESCRIPTION =
    (ADDRESS = (PROTOCOL = TCP) (HOST = MYDBCS.compute-588436052.oraclecloud.internal) (PORT = 1521))
    (CONNECT_DATA =
      (SERVER = DEDICATED)
      (SERVICE_NAME = pdb1.588436052.oraclecloud.internal)
    )
  )

PDB2 =
  (DESCRIPTION =
    (ADDRESS = (PROTOCOL = TCP) (HOST = MYDBCS.compute-588436052.oraclecloud.internal) (PORT = 1521))
    (CONNECT_DATA =
      (SERVER = DEDICATED)
      (SERVICE_NAME = pdb2.588436052.oraclecloud.internal)
    )
  )

[oracle@MYDBCS admin]$
```



10. Connect to PDB2 by using the new service name and verify the current container.

- a. Start SQL\*Plus and connect to PDB2 as the SYSTEM user by using the PDB2 net service name.

```
[oracle@MYDBCS admin]$ sqlplus system/password@PDB2
...
SQL>
```

- b. Verify that the current container name is PDB2.

```
SQL> SHOW con_name
```

```
CON_NAME
```

```
-----
```

```
PDB2
```

```
SQL>
```

- c. Exit SQL\*Plus.

```
SQL> exit
```

```
...
```

```
[oracle@MYDBCS admin]$
```

## Lab 12: Cloning a PDB

---

### Overview

In this practice, you use SQL\*Plus to hot clone PDB1 as PDB3 in the CDB.

### Tip

Because you use several windows at the same time in this practice, you may find it helpful to change the name of each of them in their banner at the top.

To set a title for a terminal window:

1. In the terminal window's menu, select **Terminal** and then **Set Title**. A Set Title dialog box is displayed.
2. In the Title box, enter the window number.
3. Click **OK**.

### Assumptions

You are connected to the compute node as the `oracle` user.

### Tasks

#### Window 1: Create a Directory for PDB3

1. Open a new terminal window and connect to the compute node as the `oracle` user. This terminal window will be called Window 1 throughout the practice.

```
$sudo su - oracle
```

2. Source the `oraenv` script.

```
[oracle@MYDBCS ~]$ . oraenv
ORACLE_SID = [ORCL] ?
The Oracle base remains unchanged with value /u01/app/oracle
[oracle@MYDBCS ~]$
```

3. Create a directory for the new PDB, named PDB3, under the CDB file location. You'll first determine the correct location and then you'll create the directory.
  - a. Log in to SQL\*Plus as the `SYS` user with the `SYSDBA` privilege.

```
[oracle@MYDBCS ~]$ sqlplus / as sysdba
...
SQL>
```

- b. Query `V$DATAFILE` to determine the location of the root container (CDB) datafiles.

```
SQL> select name from v$datafile;
```

NAME

```
-----  
/u02/app/oracle/oradata/ORCL/system01.dbf  
/u02/app/oracle/oradata/ORCL/sysaux01.dbf  
/u02/app/oracle/oradata/ORCL/undotbs01.dbf  
/u02/app/oracle/oradata/ORCL/pdbseed/system01.dbf  
/u02/app/oracle/oradata/ORCL/pdbseed/sysaux01.dbf  
/u02/app/oracle/oradata/ORCL/users01.dbf  
/u02/app/oracle/oradata/ORCL/pdbseed/undotbs01.dbf  
/u02/app/oracle/oradata/ORCL/PDB1/system01.dbf  
/u02/app/oracle/oradata/ORCL/PDB1/sysaux01.dbf  
/u02/app/oracle/oradata/ORCL/PDB1/undotbs01.dbf  
/u02/app/oracle/oradata/ORCL/PDB1/PDB1_users01.dbf  
/u02/app/oracle/oradata/ORCL/659622D851BF1AE2E0533620C40AD6D5/da  
tafile/o1_mf_users_fjv8c7l1_.dbf  
/u02/app/oracle/oradata/ORCL/6D975E8B80B85F14E0537A051D0A3C0D/da  
tafile/o1_mf_system_fk2t2tmq_.dbf  
/u02/app/oracle/oradata/ORCL/6D975E8B80B85F14E0537A051D0A3C0D/da  
tafile/o1_mf_sysaux_fk2t2tmv_.dbf  
/u02/app/oracle/oradata/ORCL/6D975E8B80B85F14E0537A051D0A3C0D/da  
tafile/o1_mf_undotbs1_fk2t2tmy_.dbf  
  
16 rows selected.  
  
SQL>
```

- c. Use the `host` command to exit to the operating system.

```
SQL> host  
[oracle@MYDBCS ~]$
```

- d. Create a new directory named `PDB3` in the location you determined in the previous step.

```
$ mkdir /u02/app/oracle/oradata/ORCL/PDB3  
[oracle@MYDBCS ~]$
```

- e. Enter `exit` to return to SQL\*Plus.

```
[oracle@MYDBCS ~]$ exit  
exit  
  
SQL>
```

## Window 1: Verify that the HR Account in PDB1 is Unlocked

1. Switch to PDB1.

```
SQL> ALTER SESSION SET CONTAINER = PDB1;
```

```
Session altered.
```

```
SQL>
```

2. Verify that the `HR` user account is unlocked, by checking for a status of `OPEN`.

```
SQL> col username format a15
```

```
SQL> SELECT username, account_status FROM dba_users where  
username = 'HR';
```

| USERNAME | ACCOUNT_STATUS |
|----------|----------------|
| HR       | OPEN           |

```
SQL>
```

3. Switch back to the root container (`CDB$ROOT`).

```
SQL> ALTER SESSION SET CONTAINER = CDB$ROOT;
```

```
Session altered.
```

```
SQL>
```

## Window 2: Start a Transaction in `PDB1`

Start a transaction in `PDB1` to determine what happens during the cloning operation when there is an uncommitted transaction.

1. Open a new terminal window and connect to the compute node as the `oracle` user. This window will be referred to as Window 2 throughout the practice.

```
[oracle@edvm ~]$ cd ~/.ssh  
[oracle@edvm .ssh]$ ssh -i your_private_key_file  
oracle@your_compute_node_IP_Address  
[oracle@MYDBCS ~]$
```

2. Source the `oraenv` script.

```
[oracle@MYDBCS ~]$ . oraenv  
ORACLE_SID = [ORCL] ?  
The Oracle base remains unchanged with value /u01/app/oracle  
[oracle@MYDBCS ~]$
```

3. Start `SQL*Plus` and connect to `PDB1` as the `HR` user. Refer to *Course Practice Environment: Security Credentials* for the password value.

```
[oracle@MYDBCS ~]$ sqlplus hr/password@PDB1  
...  
SQL>
```

4. Issue a query against the `EMPLOYEES` table to display the salary for employee ID 100.

```
SQL> SELECT salary FROM hr.employees WHERE employee_id = 100;

      SALARY
-----
      24000

SQL>
```

5. Update the `EMPLOYEES` table so that the employee salaries are increased by 10%. You will commit this transaction after you clone `PDB1`.

```
SQL> UPDATE employees SET salary=salary * 1.1;

107 rows updated.

SQL>
```

6. Display the salary for employee ID 100 again. The salary changed from 24000 to 26400. Do not commit this transaction at this time.

```
SQL> SELECT salary FROM hr.employees WHERE employee_id = 100;

      SALARY
-----
      26400

SQL>
```

### Window 1: Clone `PDB1` as `PDB3`

In this section, you clone `PDB1` as `PDB3`. `PDB1` is currently open and in `READ WRITE` mode. There is also a pending transaction in `PDB1`. Cloning `PDB1` while it is open and has a pending transaction is referred to as *hot cloning*.

1. In Window 1, create a clone named `PDB3` from `PDB1` by using the `CREATE PLUGGABLE DATABASE` statement.

In Oracle Database Cloud Service, PDBs are encrypted, so you must include the `KEYSTORE IDENTIFIED BY` clause. The keystore password is the administrative password you entered when you created the database deployment.

```
$ mkdir /u02/oradata
```

```
Sql> alter system set db_create_file_dest = '/u02/oradata';
```

```
SQL> CREATE PLUGGABLE DATABASE PDB3 FROM PDB1
      2  '

Pluggable database created.

SQL>
```

2. Verify that the open mode for PDB1 is READ WRITE and the open mode for PDB3 is MOUNTED by querying the V\$PDBS view.

```
SQL> COLUMN con_id FORMAT 999
SQL> COLUMN name FORMAT A10
SQL> SELECT con_id, name, open_mode FROM v$pdb;
```

| CON_ID | NAME      | OPEN_MODE  |
|--------|-----------|------------|
| 2      | PDB\$SEED | READ ONLY  |
| 4      | PDB2      | READ WRITE |
| 3      | PDB1      | READ WRITE |
| 5      | PDB3      | MOUNTED    |

```
SQL>
```

3. Open PDB3 so that its open mode is READ WRITE.

```
SQL> ALTER PLUGGABLE DATABASE PDB3 OPEN;
```

```
Pluggable database altered.
```

```
SQL>
```

4. Verify that the open mode for both PDB1 and PDB3 is READ WRITE.

```
SQL> SHOW PDBS
```

| CON_ID | CON_NAME  | OPEN MODE  | RESTRICTED |
|--------|-----------|------------|------------|
| 2      | PDB\$SEED | READ ONLY  | NO         |
| 4      | PDB2      | READ WRITE | NO         |
| 3      | PDB1      | READ WRITE | NO         |
| 5      | PDB3      | READ WRITE | NO         |

```
SQL>
```

## Window 2: Commit the Transaction

1. In Window 2, commit the pending transaction in PDB1.

```
SQL> COMMIT;
```

```
Commit complete.
```

```
SQL>
```

2. Display the new salary for employee ID 100. The salary is 26400.

```
SQL> SELECT salary FROM hr.employees WHERE employee_id = 100;
```

```
      SALARY
-----
      26400

SQL>
```

3. Question: Do you think the salaries are updated in the clone (PDB3)?

Answer: Continue to the next section to find out.

### Window 1: Explore PDB3

1. In Window 1, switch to PDB3 by using the `ALTER SESSION` command. This command connects you to PDB3 as the SYS user.

```
SQL> ALTER SESSION SET container = PDB3;

Session altered.

SQL>
```

2. What is the salary of employee ID 100 in PDB3?

```
SQL> SELECT salary FROM hr.employees WHERE employee_id = 100;

      SALARY
-----
      24000

SQL>
```

3. Question: The original salary was 24000. Earlier in Window 2, you updated the salary to 26400 in PDB1. Why isn't the salary showing as 26400 in PDB3?

Answer: The salary was not increased because you entered the `COMMIT` statement after the clone operation had completed.

4. Display the service name for PDB3 by querying the `V$SERVICES` view.

```
SQL> COLUMN name FORMAT A20
SQL> SELECT name FROM v$services;

NAME
-----
PDB3

SQL>
```

5. List the data files for PDB3 and their respective tablespaces by querying the DBA\_DATA\_FILES view. The results are formatted for easier viewing.

```
SQL> SELECT file_name, tablespace_name FROM dba_data_files;

FILE_NAME
-----
TABLESPACE_NAME
-----
/u02/app/oracle/oradata/ORCL/PDB3/ORCL/68196836353470ABE053FA5E8
60AAA15/datafile/o1_mf_system_fcbk0d_.dbf
SYSTEM
/u02/app/oracle/oradata/ORCL/PDB3/ORCL/68196836353470ABE053FA5E8
60AAA15/datafile/o1_mf_sysaux_fcbk0l_.dbf
SYSaux
/u02/app/oracle/oradata/ORCL/PDB3/ORCL/68196836353470ABE053FA5E8
60AAA15/datafile/o1_mf_undotbs1_fcbk1s_.dbf
UNDOTBS1
/u02/app/oracle/oradata/ORCL/PDB3/ORCL/68196836353470ABE053FA5E8
60AAA15/datafile/o1_mf_users_fcbk1z_.dbf
USERS

SQL>
```

6. Question: Do you notice a difference between the data file names in the previous step compared with the names when you created a PDB from seed?

Answer: In this case, Oracle Managed Files (OMF) names the data files for you because you used the CREATE\_FILE\_DEST clause, which only defines the directory for the data files. This clause comes from the initialization parameter DB\_CREATE\_FILE\_DEST. If you use this parameter, then all your PDB data files will end up in the same directory; whereas using the CREATE\_FILE\_DEST clause enables you to specify distinct directories for each PDB.

7. List the temp file(s) for PDB3 by querying the DBA\_TEMP\_FILES view. The query returns one temp file. The name of your temp file will be different from the one shown below.

```
SQL> SELECT file_name, tablespace_name FROM dba_temp_files;

FILE_NAME
-----
TABLESPACE_NAME
-----
/u02/app/oracle/oradata/ORCL/PDB3/ORCL/68196836353470ABE053FA5E8
60AAA15/datafile/o1_mf_temp_fcbk1x_.dbf
TEMP
```



```
SQL>
```

8. List the local users for PDB3 by querying the DBA\_USERS view.

```
SQL> col username format a30
SQL> SELECT DISTINCT username FROM dba_users WHERE common='NO'
ORDER BY username;

USERNAME
-----
APEX_050100
APEX_INSTANCE_ADMIN_USER
APEX_LISTENER
APEX_PUBLIC_USER
APEX_REST_PUBLIC_USER
DHAMBY
FLOWS_FILES
HR
INVENTORY
JGOODMAN
PDB1_ADMIN1
PDBADMIN
RPANDYA
SCOTT

14 rows selected.

SQL>
```

9. List the common users for PDB3 by querying the DBA\_USERS view.

```
SQL> SELECT DISTINCT username FROM dba_users WHERE common='YES'
ORDER BY username;

USERNAME
-----
ANONYMOUS
APPQOSSYS
AUDSYS
C##CDB_ADMIN1
...
SYSRAC
SYSTEM
WMSYS
XDB
XS$NULL
```

```
39 rows selected.
```

```
SQL>
```

10. Exit SQL\*Plus.

```
SQL> EXIT
```

```
...
```

```
[oracle@MYDBCS ~]$
```

## Window 2: Return Salary Values to Their Original Values

1. Return to Window 2. You should be logged in to PDB1 as the HR user.
2. Return the SALARY column values in the EMPLOYEES table back to their original values.

```
SQL> UPDATE employees SET salary=salary / 1.1;
```

```
107 rows updated.
```

```
SQL>
```

3. Commit the transaction.

```
SQL> COMMIT;
```

```
Commit complete.
```

```
SQL>
```

4. Display the salary for employee ID 100. The result is 24000.

```
SQL> SELECT salary FROM hr.employees WHERE employee_id = 100;
```

```
      SALARY
```

```
-----  
      24000
```

```
SQL>
```

5. Exit SQL\*Plus, close the connection to the compute node, and close the terminal window.

```
SQL> EXIT
```

```
...
```

```
[oracle@MYDBCS ~]$
```

## Lab 13: Unplugging and Plugging in a PDB

---

### Overview

In this practice, you unplug PDB3 from the ORCL CDB and plug it back into ORCL CDB. You give the PDB a new name (HRPDB) when you plug it back in.

### Assumptions

You are logged in to the compute node as the `oracle` user.

You completed Practice 10-2 Hot Cloning a PDB.

### Tasks

1. Unplug PDB3 from the ORCL CDB.
  - a. Start SQL\*Plus and log in to ORCL with the SYSDBA privilege.

```
$ sqlplus / as sysdba
...
SQL>
```

- b. Close PDB3.

PDBs must be closed before you can unplug them and drop them. If PDB3 is already closed, you will receive an error message.

```
SQL> ALTER PLUGGABLE DATABASE PDB3 CLOSE IMMEDIATE;

Pluggable database altered.

SQL>
```

- c. Unplug PDB3 into an XML file named `/u02/app/oracle/oradata/PDB3.xml`.

The unplugging operation makes changes in the PDB data files to record that the PDB was properly and successfully unplugged.

Because the PDB is encrypted, you must include the `ENCRYPT USING transport_secret` clause. If you do not include the clause, you will receive an `ORA-46680: master keys of the container database must be exported` error. Supply a value of `TransPDB3` for *transport\_secret* for the course practice.

Because the PDB is still part of the CDB, you can back it up in Oracle Recovery Manager (Oracle RMAN). This backup provides a convenient way to archive the unplugged PDB. After backing it up, you can then remove it from the CDB catalog. However, you must preserve the data files for any subsequent plugging operations.

```
SQL> ALTER PLUGGABLE DATABASE PDB3
  2  UNPLUG INTO '/u02/app/oracle/oradata/PDB3.xml';

Pluggable database altered.
```

```
SQL>
```

- d. Check the status of PDB3 by querying CDB\_PDBS.

```
SQL> col PDB_NAME format a10
```

```
SQL> SELECT pdb_name, status FROM cdb_pdb;
```

| PDB_NAME  | STATUS    |
|-----------|-----------|
| PDB2      | NORMAL    |
| PDB\$SEED | NORMAL    |
| PDB1      | NORMAL    |
| PDB3      | UNPLUGGED |

```
SQL>
```

- e. Drop PDB3 while it is closed, but keep its datafiles so you can plug the PDB back in.

```
SQL> DROP PLUGGABLE DATABASE PDB3 KEEP DATAFILES;
```

```
Pluggable database dropped.
```

```
SQL>
```

- f. Verify the status of the unplugged PDB3 by querying the CDB\_PDBS view. Note that PDB3 is not included.

```
SQL> SELECT pdb_name, status FROM cdb_pdb;
```

| PDB_NAME  | STATUS |
|-----------|--------|
| PDB2      | NORMAL |
| PDB\$SEED | NORMAL |
| PDB1      | NORMAL |

```
SQL>
```

2. Plug PDB3 back into the ORCL CDB. The method would be similar if you were to plug the PDB into a different CDB.

- a. Make sure that PDB3 is compatible with the ORCL CDB. Execution of the following PL/SQL block raises an error if it is not compatible.

Tip: Enter each line, followed by a return, and the whole procedure will run after you close with a slash.

```
SQL> set serveroutput on
```

```
SQL> DECLARE
```

```
2 compatible BOOLEAN := FALSE;
```

```

4      3  BEGIN
        compatible := DBMS_PDB.CHECK_PLUG_COMPATIBILITY(
            pdb_descr_file =>
'/u02/app/oracle/oradata/PDB3.xml');
        5  if compatible then
        6  DBMS_OUTPUT.PUT_LINE('PDB3 is compatible');
        7  else DBMS_OUTPUT.PUT_LINE('PDB3 is not compatible');
        8  end if;
        9  END;
       10 /

```

PDB3 is compatible

PL/SQL procedure successfully completed.

SQL>

- b. Plug PDB3 back into the ORCL CDB by using the NOCOPY method. Rename the plugged-in PDB as HRPDB.

Because the PDB is encrypted, you must include the `IDENTIFIED BY password` clause and include the keystore password. In Oracle Database Cloud Service, the keystore password is the value you supplied when you created the database deployment.

You must also include the `DECRYPT USING transport_secret` clause. The value for `transport_secret` is the same value you specified when you unplugged the PDB. For the course practice, the value is `TransPDB3`.

The original data files of the unplugged PDB now belong to the new plugged-in PDB.

```

SQL> CREATE PLUGGABLE DATABASE HRPDB
      2  USING '/u02/app/oracle/oradata/PDB3.xml'
      3  NOCOPY TEMPFILE REUSE
      4  KEYSTORE IDENTIFIED BY password
      4  DECRYPT USING TransPDB3;

```

Pluggable database created.

SQL>

3. Examine the plugged-in PDB.
  - a. List all the containers in your CDB by querying the `V$CONTAINERS` view. The results list five containers—the root container (`CDB$ROOT`), the seed PDB (`PDB$SEED`), `PDB1`, `PDB2`, and `HRPDB`.

```

SQL> COLUMN name FORMAT A8
SQL> SELECT name, con_id FROM v$containers ORDER BY con_id;

```

| NAME      | CON_ID |
|-----------|--------|
| CDB\$ROOT | 1      |
| PDB\$SEED | 2      |
| PDB1      | 3      |
| PDB2      | 4      |
| HRPDB     | 5      |
| SQL>      |        |

- b. Show the status of HRPDB by querying the CDB\_PDBS view.

|  |        |
|--|--------|
| SQL> SELECT pdb_name, status FROM cdb_pdb\$ WHERE<br>pdb_name='HRPDB'; |        |
| PDB_NAME   | STATUS |
| -----  | -----  |
| HRPDB  | NEW    |
| SQL>   |        |

- c. Show the open mode of HRPDB by querying the V\$PDBS view.

|   |  |
|---|--|
| SQL> SELECT open_mode FROM v\$pdb\$ WHERE name='HRPDB'; |  |
| OPEN_MODE   |  |
| -----   |  |
| MOUNTED   |  |
| SQL>  |  |

- d. List the data files of HRPDB by querying the V\$DATAFILE view. Recall that the HRPDB container's ID is 5. Your paths and data file names will differ from those shown below.

|   |  |
|---|--|
| SQL> COLUMN name FORMAT A50   |  |
| SQL> SELECT name FROM v\$datafile WHERE con_id=5;   |  |
| NAME  |  |
| -----   |  |
| /u02/app/oracle/oradata/ORCL/PDB3/ORCL/68196836353470ABE053FA5E8<br>60AAA15/datafile/o1_mf_system_fcbk0d_.dbf   |  |
| /u02/app/oracle/oradata/ORCL/PDB3/ORCL/68196836353470ABE053FA5E8<br>60AAA15/datafile/o1_mf_sysaux_fcbk0l_.dbf   |  |
| /u02/app/oracle/oradata/ORCL/PDB3/ORCL/68196836353470ABE053FA5E8<br>60AAA15/datafile/o1_mf_undotbs1_fcbk0m_.dbf |  |
| /u02/app/oracle/oradata/ORCL/PDB3/ORCL/68196836353470ABE053FA5E8<br>60AAA15/datafile/o1_mf_users_fcbk0n_.dbf    |  |

```
SQL>
```

4. Open and connect to HRPDB.

- a. Open HRPDB.

```
SQL> ALTER PLUGGABLE DATABASE HRPDB open;
```

```
Pluggable database altered.
```

```
SQL>
```

- b. Query V\$SERVICES.

```
SQL> SELECT name FROM v$services ORDER BY name;
```

```
NAME
```

```
-----
```

```
ORCL.588436052.oraclecloud.internal
```

```
ORCL.588436052.oraclecloud.internalXDB
```

```
SYS$BACKGROUND
```

```
SYS$USERS
```

```
PDB1
```

```
PDB2
```

```
hrpdb
```

```
7 rows selected.
```

```
SQL>
```

- c. Connect to HRPDB as the SYS user with the SYSDBA privilege. Recall that you need to append the values following ORCL as shown in this example.

```
SQL> CONNECT
```

```
SYS/password@localhost:1521/hrpdb.588436052.oraclecloud.internal  
AS SYSDBA
```

```
Connected.
```

```
SQL>
```

5. Verify the current container name is HRPDB.

```
SQL> SHOW con_name
```

```
CON_NAME
```

```
-----
```

```
HRPDB
```

```
SQL>
```

6. Exit from SQL\*Plus.

```
SQL> exit
```

```
...
```

```
[oracle@MYDBCS ~]$
```

## Lab 14: Dropping a PDB

---

### Overview

In this practice, you drop the `HRPDB` PDB.

### Assumptions

You are logged in to the compute node as the `oracle` user.

You completed the following practices in this lesson:

- Creating a PDB from Seed
- Hot Cloning a PDB

### Tasks

1. Start SQL\*Plus and connect to the root container as the `SYS` user with the `SYSDBA` privilege.

```
[oracle@MYDBCS ~]$ sqlplus / as sysdba
..
SQL>
```

2. List the PDBs in `ORCL`. The results show four PDBs: `PDB$SEED`, `PDB1`, `PDB2`, and `HRPDB`.

```
SQL> SHOW PDBS

  CON_ID CON_NAME              OPEN  MODE  RESTRICTED
-----
      2 PDB$SEED                READ  ONLY   NO
      3 PDB1                  READ  WRITE  NO
      4 PDB2                  READ  WRITE  NO
      5 HRPDB                 READ  WRITE  NO

SQL>
```

3. Close `HRPDB`.

```
SQL> ALTER PLUGGABLE DATABASE HRPDB CLOSE;

Pluggable database altered.

SQL>
```

4. Drop `HRPDB`, including its data files, by using the `DROP PLUGGABLE DATABASE` command.

```
SQL> DROP PLUGGABLE DATABASE HRPDB INCLUDING DATAFILES;

Pluggable database dropped.

SQL>
```



5. List the PDBs in ORCL. The results show three PDBs: PDB\$SEED, PDB1, and PDB2.

```
SQL> SHOW PDBS
```

| CON_ID | CON_NAME  | OPEN | MODE  | RESTRICTED |
|--------|-----------|------|-------|------------|
| 2      | PDB\$SEED | READ | ONLY  | NO         |
| 3      | PDB1      | READ | WRITE | NO         |
| 4      | PDB2      | READ | WRITE | NO         |

```
SQL>
```

6. Exit SQL\*Plus and close the terminal window.

```
SQL> EXIT
```

...

```
[oracle@MYDBCS ~]$ exit
```

# Lab 15: Viewing Tablespace Information

## Overview

In this practice, you use SQL\*Plus to query various views to learn about tablespace content in PDB1. You also view tablespace information with Enterprise Manager Database Express (EM Express).

## Assumptions

You are logged in as the `oracle` user.

## Tasks

- 1. Open a new terminal window and connect to the compute node as the `oracle` user.

```
$ sudo su - oracle
```

- 2. Source the `oraenv` script.

```
[oracle@MYDBCS ~]$ . oraenv
ORACLE_SID = [ORCL] ?
The Oracle base remains unchanged with value /u01/app/oracle
[oracle@MYDBCS ~]$
```

- 3. Start SQL\*Plus and connect to PDB1 as the PDBADMIN user. Refer to *Course Practice Environment: Security Credentials* for the password value.

```
[oracle@MYDBCS ~]$ sqlplus PDBADMIN/password@PDB1
...
SQL>
```

- 4. List the columns in the `DBA_TABLESPACES` view by using the `DESCRIBE` command.

```
SQL> DESCRIBE dba_tablespaces
Name                                         Null?    Type
-----
TABLESPACE_NAME                             NOT NULL VARCHAR2(30)
BLOCK_SIZE                                  NOT NULL NUMBER
INITIAL_EXTENT                               NUMBER
NEXT_EXTENT                                NUMBER
MIN_EXTENTS                                 NOT NULL NUMBER
...
DEF_CELLMEMORY                               VARCHAR2(14)
DEF_INMEMORY_SERVICE                        VARCHAR2(12)
DEF_INMEMORY_SERVICE_NAME                    VARCHAR2(1000)
LOST_WRITE_PROTECT                           VARCHAR2(7)
```

```
CHUNK_TABLESPACE                                VARCHAR2 (1)

SQL>
```

5. List the tablespaces in PDB1.

```
SQL> SELECT DISTINCT tablespace_name FROM dba_tablespaces ORDER
BY tablespace_name;

TABLESPACE_NAME
-----
SYSAUX
SYSTEM
TEMP
UNDOTBS1
USERS

SQL>
```

6. Find out which tablespace contains the HR schema by querying the ALL\_TABLES view.

```
SQL> SELECT DISTINCT tablespace_name FROM all_tables WHERE
owner='HR';

TABLESPACE_NAME
-----

USERS

SQL>
```

7. Query the STATUS, CONTENTS, LOGGING, PLUGGED\_IN, BIGFILE, EXTENT\_MANAGEMENT, and ALLOCATION\_TYPE columns in the DBA\_TABLESPACES view for the SYSAUX tablespace.

```
SQL> SELECT status, contents, logging, plugged_in, bigfile,
extent_management, allocation_type FROM dba_tablespaces where
tablespace_name='SYSAUX';

STATUS      CONTENTS          LOGGING PLU  BIG  EXTENT_MAN  ALLOCA
-----
ONLINE PERMANENT LOGGING NO NO LOCAL SYSTEM  SQL>
```

- STATUS shows the value ONLINE, indicating the tablespace is available to users.
- CONTENTS indicates the PERMANENT tablespace type.

- LOGGING shows the value LOGGING, indicating that certain DML operations are logged in the redo log file.
- PLUGGED\_IN shows the value NO, indicating that the tablespace is not plugged in.
- BIGFILE shows the value NO, indicating that the tablespace is a smallfile tablespace.
- EXTENT\_MANAGEMENT shows the value LOCAL, indicating that the tablespace is locally managed (not dictionary managed).
- ALLOCATION\_TYPE shows the value SYSTEM, indicating that the extents of the tablespace are managed by the system, and you cannot specify an extent size.

8. List the columns in the V\$TABLESPACE view by using the DESCRIBE command. This view displays tablespace information from the control file.

| SQL> DESCRIBE v\$tablespace |       |               |
|-----------------------------|-------|---------------|
| Name                        | Null? | Type          |
| TS#                         |       | NUMBER        |
| NAME                        |       | VARCHAR2 (30) |
| INCLUDED_IN_DATABASE_BACKUP |       | VARCHAR2 (3)  |
| BIGFILE                     |       | VARCHAR2 (3)  |
| FLASHBACK_ON                |       | VARCHAR2 (3)  |
| ENCRYPT_IN_BACKUP           |       | VARCHAR2 (3)  |
| CON_ID                      |       | NUMBER        |
| SQL>                        |       |               |

9. Query the V\$TABLESPACE view for the SYSAUX tablespace.

| SQL> SELECT * FROM v\$tablespace WHERE name='SYSAUX'; |        |     |     |     |     |        |
|---|--------|-----|-----|-----|-----|--------|
| TS#   | NAME   | INC | BIG | FLA | ENC | CON_ID |
| 1   | SYSAUX | YES | NO  | YES |     | 3      |
| SQL>  |        |     |     |     |     |        |

- INCLUDED\_IN\_DATABASE\_BACKUP contains the value YES, indicating that the tablespace is included in full database backups by using the BACKUP DATABASE RMAN command.
- BIGFILE contains the value NO, indicating that the tablespace is a smallfile tablespace.
- FLASHBACK\_ON contains the value YES, indicating that the tablespace participates in FLASHBACK DATABASE operations.
- ENCRYPT\_IN\_BACKUP contains the value null, indicating that encryption is neither explicitly turned on nor off at the tablespace level.

- CON\_ID indicates the container to which the data pertains. In this case, PDB1 is container ID 3.

10. List all the tables in the USERS tablespace owned by the HR account.

```
SQL> SELECT table_name FROM all_tables WHERE  
tablespace_name='USERS' and owner='HR';
```

```
TABLE_NAME  
-----
```

```
REGIONS  
LOCATIONS  
DEPARTMENTS  
JOBS  
EMPLOYEES  
JOB_HISTORY
```

```
6 rows selected.
```

```
SQL>
```

11. List all the indexes in the USERS tablespace owned by the HR account.

```
SQL> SELECT index_name FROM all_indexes WHERE  
tablespace_name='USERS' AND owner='HR' ORDER BY index_name;
```

```
INDEX_NAME  
-----
```

```
COUNTRY_C_ID_PK  
DEPT_ID_PK  
DEPT_LOCATION_IX  
EMP_DEPARTMENT_IX  
EMP_EMAIL_UK  
EMP_EMP_ID_PK  
EMP_JOB_IX  
EMP_MANAGER_IX  
EMP_NAME_IX  
JHIST_DEPARTMENT_IX  
JHIST_EMPLOYEE_IX  
JHIST_EMP_ID_ST_DATE_PK  
JHIST_JOB_IX  
JOB_ID_PK  
LOC_CITY_IX  
LOC_COUNTRY_IX  
LOC_ID_PK  
LOC_STATE_PROVINCE_IX
```

```

REG_ID_PK

19 rows selected.

SQL>

```

12. List the columns in the `DBA_DATA_FILES` view by using the `DESCRIBE` command. You can query this view to learn about the data files contained in a tablespace.

```

SQL> DESCRIBE dba_data_files
Name                                         Null?    Type
-----
FILE_NAME                                   VARCHAR2(513)
FILE_ID                                     NUMBER
TABLESPACE_NAME                           VARCHAR2(30)
BYTES                                       NUMBER
BLOCKS                                     NUMBER
STATUS                                     VARCHAR2(9)
RELATIVE_FNO                               NUMBER
AUTOEXTENSIBLE                            VARCHAR2(3)
MAXBYTES                                   NUMBER
MAXBLOCKS                                 NUMBER
INCREMENT_BY                              NUMBER
USER_BYTES                                NUMBER
USER_BLOCKS                              NUMBER
ONLINE_STATUS                             VARCHAR2(7)
LOST_WRITE_PROTECT                        VARCHAR2(7)

SQL>

```

13. List data file information for the `SYSAUX` tablespace by querying various columns in the `DBA_DATA_FILES` view.

```

SQL> COLUMN file_name FORMAT A50
SQL> SELECT file_name, autoextensible, bytes, maxbytes,
user_bytes FROM dba_data_files WHERE tablespace_name='SYSAUX';

FILE_NAME                                     AUT
BYTES      MAXBYTES    USER_BYTES
-----
/u02/app/oracle/oradata/ORCL/PDB1/sysaux01.dbf YES  828375040
3.4360E+10  827326464

SQL>

```

The results show thefollowing:

- `AUTOEXTENSIBLE` contains the value YES, indicating that the auto extend feature is enabled for a data file. The tablespace size can increase without you having to take any action.
- `BYTES` is the size of the file in bytes.
- `MAXBYTES` is the maximum file size allowed.
- `USER_BYTES` is the size of the file available for user data.

14. Find out how many segments are there in the `SYSAUX` tablespace by querying the `DBA_SEGMENTS` view.

```
SQL> SELECT count(segment_name) FROM dba_segments WHERE
tablespace_name='SYSAUX';

COUNT (SEGMENT_NAME)
-----
                        2356

SQL>
```

15. Find out which index in the `SYSAUX` tablespace takes up the most space by querying the `DBA_SEGMENTS` view. The results indicate that the `I_WRI$_OPTSTAT_H_OBJ#_ICOL#_ST` index takes up the most space.

```
SQL> col segment_name format a35
SQL> SELECT *
  2  FROM (SELECT segment_name, segment_type, bytes
  3          FROM dba_segments
  4          WHERE segment_type = 'INDEX' AND
  5          tablespace_name = 'SYSAUX'
  6          ORDER BY bytes desc)
  7  WHERE rownum < 2;

SEGMENT_NAME                                SEGMENT_TYPE          BYTES
-----
I_WRI$_OPTSTAT_H_OBJ#_ICOL#_ST             INDEX                  42991616

SQL>
```

16. Exit SQL\*Plus.

```
SQL> exit
..
[oracle@MYDBCS ~]$
```

## Lab 16: Creating a Tablespace

---

### Overview

In this practice, you create and populate a tablespace named `INVENTORY`.

### Assumptions

You are logged in to the compute node as the `oracle` user.

### Tasks

#### Use SQL\*Plus to Create the `INVENTORY` Tablespace and Table `x`

As the `PDBADMIN` user in SQL\*Plus, execute the `CreateINVENTORYTablespace.sql` script to create the `INVENTORY` tablespace. Next, execute a script named `CreateTableX.sql` to create and populate a table called `x` in the `INVENTORY` tablespace. At first, you will get an error trying to populate the table. In the next section, you correct the problem.

1. Start SQL\*Plus and connect to `PDB1` as the `PDBADMIN` user. Refer to *Course Practice Environment: Security Credentials* for the password value.

```
[oracle@MYDBCS ~]$ sqlplus PDBADMIN/password@PDB1
...
SQL>
```

2. Execute the `CreateINVENTORYTablespace.sql` script.

```
SQL> set echo on
SQL> @/home/oracle/labs/CreateINVENTORYTablespace.sql

SQL> CREATE SMALLFILE TABLESPACE INVENTORY
 2
          DATAFIL
E  3
'/u02/app/oracle/oradata/ORCL/PDB1/INVENTORY01.DBF' SIZE 5242880
 4          DEFAULT NOCOMPRESS
 5          ONLINE
 6          SEGMENT SPACE MANAGEMENT AUTO
 7          EXTENT MANAGEMENT LOCAL AUTOALLOCATE;

Tablespace

created.  SQL>
```



## Lab 17: Verifying that the Control File is Multiplexed

---

### Overview

In this practice, you verify that the control file is multiplexed.

A control file is a small binary file that describes the structure of the database. It must be available for writing by the Oracle server whenever the database is mounted or opened. Without this file, the database cannot be mounted, and recovery or re-creation of the control file is required. Your database should have a minimum of two control files on different storage devices to minimize the impact of a loss of one control file. The loss of a single control file causes the instance to fail because all control files must be available at all times. However, recovery can be a simple matter of copying one of the other control files. The loss of all control files is slightly more difficult to recover from, but is not usually catastrophic.

### Assumptions

You are logged in as the `oracleuser`.

### Tasks

1. Open a new terminal window and connect to the compute node as the `oracleuser`.

```
$sudo su - oracle
```

2. Source the `oraenv` script.

```
[oracle@MYDBCS ~]$ . oraenv
ORACLE_SID = [ORCL] ?
The Oracle base remains unchanged with value /u01/app/oracle
[oracle@MYDBCS ~]$
```

3. Start SQL\*Plus and connect to the CDB root as the `SYS` user with the `SYSDBA` privilege.

```
$ sqlplus / AS SYSDBA
...
SQL>
```

4. Find out how many control files exist in the database. The query returns the names of two control files (`control01.ctl` and `control02.ctl`), which verifies that the control files are multiplexed.

```
SQL> SELECT name FROM v$controlfile;

NAME
-----
/u02/app/oracle/oradata/ORCL/control01.ctl
/u03/app/oracle/fast_recovery_area/ORCL/control02.ctl
```

```
SQL>
```

When the CDB was created, DBCA created two control files. When you use the `CREATE DATABASE` command in SQL\*Plus to create a database, you configure the `CONTROL_FILES` parameter to generate two control files and set their names.

5. View the `CONTROL_FILES` parameter. Notice that the paths to the control files are stored in this parameter. The results below are formatted for easier viewing.

```
SQL> SHOW PARAMETER control_files

NAME                                TYPE
-----                                -----
VALUE
-----
control_files                       string
/u02/app/oracle/oradata/ORCL/control01.ctl,
/u03/app/oracle/fast_recovery_area/ORCL/control02.ctl
SQL>
```

6. Create a parameter file (PFILE) from the server parameter file (SPFILE).

```
SQL> CREATE PFILE FROM SPFILE;

File created.

SQL>
```

7. Shut down the database instance in `IMMEDIATE` mode.

```
SQL> SHUTDOWN IMMEDIATE
Database closed.
Database dismounted.
ORACLE instance shut down.
SQL>
```

8. Exit SQL\*Plus.

```
SQL> EXIT
```

9. Create a directory for the new control file.

```
[oracle@MYDBCS ~]$ mkdir -p
/u01/app/oracle/controlfiles_dir/ORCL
[oracle@MYDBCS ~]$
```

10. Before you edit your PFILE, make a backup copy of it.

```
[oracle@MYDBCS ~]$ cp $ORACLE_HOME/dbs/initORCL.ora
$ORACLE_HOME/dbs/backup_initORCL.ora
[oracle@MYDBCS ~]$
```

11. Copy one of the control files to the directory you created in a previous step (/u01/app/oracle/controlfiles\_dir/ORCL) and name the file as `control03.ctl`.

```
[oracle@MYDBCS ~]$ cp /u02/app/oracle/oradata/ORCL/control01.ctl
/u01/app/oracle/controlfiles_dir/ORCL/control03.ctl
[oracle@MYDBCS ~]$
```

12. Open the PFILE (initORCL.ora) in the vi editor and add the name of the new control file to the end of the list of control files. Include the path. Be certain not to enter spaces between the single quotes and commas in the control\_files= line. Be certain that this line is one continuous line, without line breaks. Save and close the file (:wq).

```
$ vi $ORACLE_HOME/dbs/initORCL.ora
...
*.control_files='/u02/app/oracle/oradata/ORCL/control01.ctl',
'/u03/app/oracle/fast_recovery_area/ORCL/control02.ctl',
'/u01/app/oracle/controlfiles_dir/ORCL/control03.ctl'
...
$
```

13. Start SQL\*Plus and connect to the root container as the SYS user with the SYSDBA privilege. You are connected to an idle instance.

```
$ sqlplus / AS SYSDBA
...
SQL>
```

14. Start the database instance.

```
SQL> STARTUP
ORACLE instance started.

Total System Global Area 2768239832 bytes
Fixed Size                  8899800 bytes
Variable Size              704643072 bytes
Database Buffers 1979711488 bytes  Redo
Buffers                    74985472 bytes
Database mounted.
Database opened.
SQL>
```

15. View the CONTROL\_FILES parameter again.

```
SQL> SHOW PARAMETER control_files

NAME                                TYPE
-----
VALUE
-----
control_files                      string
/u02/app/oracle/oradata/ORCL/control01.ctl,
/u03/app/oracle/fast_recovery_area/ORCL/control02.ctl
```

```
SQL>
```

16. Question: Why does the `CONTROL_FILES` parameter still show only two control files?

Answer: By default, the database instance starts up with the SPFILE. If an SPFILE does not exist, then the instance starts up with a PFILE. In this case, both an SPFILE and PFILE are present, so the SPFILE takes precedence. You configured the PFILE, not the SPFILE. The SPFILE still contains only two references.

17. Re-create the third control file because the current version is no longer an exact copy of the others.

- a. Shut down the database instance with the `IMMEDIATE` option.

```
SQL> SHUTDOWN IMMEDIATE
Database closed.
Database dismounted.
ORACLE instance shut down.
SQL>
```

- b. Exit SQL\*Plus.

```
SQL> EXIT
```

- c. Use the `cp` command to re-create `control03.ctl`.

```
[oracle@MYDBCS ~]$ cp /u02/app/oracle/oradata/ORCL/control01.ctl
/u01/app/oracle/controlfiles_dir/ORCL/control03.ctl
[oracle@MYDBCS ~]$
```

18. Re-create the SPFILE from the updated PFILE.

- a. Start SQL\*Plus and connect to the CDB root as the `SYS` user with the `SYSDBA` privilege. You are connected to an idle instance.

```
[oracle@MYDBCS ~]$ sqlplus / AS SYSDBA
...
SQL>
```

- b. Create the SPFILE.

```
SQL> CREATE SPFILE FROM PFILE;

File created.

SQL>
```

19. Start the database instance.

```
SQL> STARTUP
ORACLE instance started.

Total System Global Area 2768239832 bytes
Fixed Size                  8899800 bytes
Variable Size              704643072 bytes
Database Buffers          1979711488 bytes
```

```
Redo Buffers              74985472 bytes
Database mounted.
Database opened.
SQL>
```

20. View the `CONTROL_FILES` parameter again. The third control file is now included in the list, which indicates that the SPFILE is configured properly. The results below are formatted for easier viewing.

```
SQL> SHOW PARAMETER control_files

NAME                                TYPE
-----
VALUE
-----
control_files                      string
/u02/app/oracle/oradata/ORCL/control01.ctl,
/u03/app/oracle/fast_recovery_area/ORCL/control02.ctl,
/u01/app/oracle/controlfiles_dir/ORCL/control03.ctl
SQL>
```

21. Query the `V$CONTROLFILE` view to confirm the number of control files. The result indicates that three control files are defined.

```
SQL> SELECT name FROM v$controlfile;

NAME
-----
/u02/app/oracle/oradata/ORCL/control01.ctl
/u03/app/oracle/fast_recovery_area/ORCL/control02.ctl
/u01/app/oracle/controlfiles_dir/ORCL/control03.ctl

SQL>
```

## Lab 18: Configuring the Size of the Fast Recovery Area

---

### Overview

In this practice, you review the fast recovery area (FRA) configuration and change its size to 12GB.

### Assumptions

You are logged in to SQL\*Plus from the previous practice.

### Tasks

1. Question: How can you evaluate the space needed for the FRA?

Answer: The amount of disk space to allocate for the FRA depends on the size and activity levels of your database. As a general rule, the larger the FRA, the more useful it is. Ideally, the FRA should be large enough for copies of your data and control files, as well as for flashback, online redo, and archived logs needed to recover the database with the backups kept based on the retention policy (covered in one of the next practices). In short, the FRA should be at least twice the size of the database so that it can hold one backup and several archived logs.

2. View the values of the `DB_RECOVERY_FILE_DEST` and `DB_RECOVERY_FILE_DEST_SIZE` initialization parameters.

```
SQL> SHOW PARAMETER db_recovery_file_dest

NAME                                TYPE
-----
VALUE
-----
db_recovery_file_dest               string
/u03/app/oracle/fast_recovery_area
db_recovery_file_dest_size          big integer
4G
SQL>
```

3. Question: Is the fast recovery area enabled?

Answer: Yes. The `DB_RECOVERY_FILE_DEST` and `DB_RECOVERY_FILE_DEST_SIZE` parameters values are not null, indicating that the fast recovery area is enabled.

4. Question: Which changes can you make to the fast recovery area?

Answer: You can change the location and size for the fast recovery area.

5. Question: Does changing the size of the fast recovery area require the database to be restarted?

Answer: No, a restart is not required for this change because the `DB_RECOVERY_FILE_DEST_SIZE` parameter is dynamic.

```
SQL> ALTER SYSTEM SET db_recovery_file_dest_size = 12G
SCOPE=both;

System altered.

SQL>
```

- 6. Change the size of the fast recovery area to 12GB and set the scope to BOTH.

Note: If the archived redo log file destination fills up or cannot be written to, the database will halt. You would then need to remove archived redo log files from the archived redo log file destination so that the database could resume operations. This activity is covered in one of the next practices.

- 7. View the DB\_RECOVERY\_FILE\_DEST\_SIZE initialization parameter again. The result verifies that the size has been set to 12GB.

| SQL> SHOW PARAMETER db_recovery_file_dest_size |             |       |
|--|-------------|-------|
| NAME   | TYPE        | VALUE |
| db_recovery_file_dest_size                     | big integer | 12G   |
| SQL>   |             |       |

## Lab 19: Verifying that the Redo Log File is Multiplexed

---

### Overview

Ensure that there are at least two redo log members in each group. If you are using file system storage, then each member should be distributed on separate disks or controllers so that no single equipment failure impacts an entire log group. The loss of an entire current log group is one of the most serious media failures because it can result in data loss. The loss of a single member of a multi-member log group is trivial and does not affect database operation (other than causing an alert to be published in the alert log). One set of members should be stored in the FRA.

### Assumptions

You are logged in to SQL\*Plus from the previous practice.

### Tasks

1. Query **V\$LOGFILE** to determine the configuration (number of members) for each redo log group. The result shows that there are currently three log groups (1, 2, and 3) and only one member in each group.

```
SQL> SELECT group#, status, member FROM v$logfile;
```

| GROUP# | STATUS | MEMBER                          |
|--------|--------|---------------------------------|
| 3      |        | /u04/app/oracle/redo/redo03.log |
| 2      |        | /u04/app/oracle/redo/redo02.log |
| 1      |        | /u04/app/oracle/redo/redo01.log |

```
SQL>
```

2. Question: Why is it recommended to have three groups when two would be sufficient?  
Answer: The Oracle Database server treats the online redo log groups as a circular buffer in which to store transaction information, filling one group and then moving on to the next. After all groups have been written to, the Oracle Database server begins overwriting information in the first log group. If the database is configured in **ARCHIVELOG** mode, the LGWR cannot overwrite data in the first log group if it has not been archived.
3. Question: Can multiplexing redo logs impact database performance?  
Answer: Multiplexing redo logs may heavily influence database performance because a commit cannot complete until the transaction information has been written to the logs by LGWR. You must place your redo log files on your fastest disks served by your fastest controllers. If possible, do not place any other database files on the same disks as your redo log files. Because only one group is written to at a given time, there is no performance impact in having members from several groups on the same disk.



4. Add another member to each redo log group. Name each member as `redonnb.log` where `nn` represents the group number.

```
SQL> ALTER DATABASE ADD LOGFILE MEMBER
'/u03/app/oracle/fast_recovery_area/ORCL/redo01b.log' TO GROUP
1;

Database altered.

SQL> ALTER DATABASE ADD LOGFILE MEMBER
'/u03/app/oracle/fast_recovery_area/ORCL/redo02b.log' TO GROUP
2;

Database altered.

SQL> ALTER DATABASE ADD LOGFILE MEMBER
'/u03/app/oracle/fast_recovery_area/ORCL/redo03b.log' TO GROUP
3;

Database altered.

SQL>
```

5. Verify that the redo log files are now multiplexed. The query result shows that each group has two members, and therefore, the redo log files are multiplexed. Observe the `INVALID` status of the newly added redo log members. This status is expected because the new members have not yet been written to by LGWR. When a log switch occurs and the group containing the new member becomes `CURRENT`, the new member's status will change to `NULL`.

```
SQL> SELECT group#, status, member FROM v$logfile ORDER BY 1, 3;

GROUP# STATUS  MEMBER
-----
      1 INVALID
/u03/app/oracle/fast_recovery_area/ORCL/redo01b.log
      1      /u04/app/oracle/redo/redo01.log
      2 INVALID
/u03/app/oracle/fast_recovery_area/ORCL/redo02b.log
      2      /u04/app/oracle/redo/redo02.log
      3 INVALID
/u03/app/oracle/fast_recovery_area/ORCL/redo03b.log
      3      /u04/app/oracle/redo/redo03.log

6 rows selected.

SQL>
```

6. Switch the log files and observe the changes.
- a. Find out which log group is the current log group. In this example, the query result shows that group 3 is the current group. Your current group may be different.

```
SQL> SELECT group#, members, archived, status FROM v$log;
```

| GROUP# | MEMBERS | ARC | STATUS   |
|--------|---------|-----|----------|
| 1      | 2       | YES | INACTIVE |
| 2      | 2       | YES | INACTIVE |
| 3      | 2       | NO  | CURRENT  |

```
SQL>
```

- b. Switch the log files three times.

```
SQL> ALTER SYSTEM SWITCH LOGFILE;
```

```
System altered.
```

```
SQL> ALTER SYSTEM SWITCH LOGFILE;
```

```
System altered.
```

```
SQL> ALTER SYSTEM SWITCH LOGFILE;
```

```
System altered.
```

```
SQL>
```

- c. Query the V\$LOGFILE view again. Notice that the switch caused the new members' statuses to change to null.

```
SQL> SELECT group#, status, member FROM v$logfile ORDER BY 1, 3;
```

| GROUP#           | STATUS | MEMBER  |
|------------------|--------|---|
| 1                |        | /u03/app/oracle/fast_recovery_area/ORCL/redo01b.1<br>og |
| 1                |        | /u04/app/oracle/redo/redo01.log                         |
| 2                |        | /u03/app/oracle/fast_recovery_area/ORCL/redo02b.1<br>og |
| 2                |        | /u04/app/oracle/redo/redo02.log                         |
| 3                |        | /u03/app/oracle/fast_recovery_area/ORCL/redo03b.1<br>og |
| 6 rows selected. |        | /u04/app/oracle/redo/redo03.log                         |

```
SQL>
```

- d. Query the `V$LOG` view again to learn which log group is now the current group. In this example, the results show that the LGWR is writing to group 3. Your group may be different. Your statuses may be different too. An `INACTIVE` status means the log group is no longer needed for database instance recovery.

```
SQL> SELECT group#, members, archived, status FROM v$log;
```

| GROUP# | MEMBERS | ARC | STATUS   |
|--------|---------|-----|----------|
| 1      | 2       | YES | INACTIVE |
| 2      | 2       | YES | INACTIVE |
| 3      | 2       | NO  | CURRENT  |

```
SQL>
```

- e. Switch the log file.

```
SQL> ALTER SYSTEM SWITCH LOGFILE;
```

```
System altered.
```

```
SQL>
```

- f. Query the `V$LOG` view again. The current group has changed to group 1, and the former current group's status is now `ACTIVE`. Your current group may be different. An `ACTIVE` status means that the log group is active, but it's not the current log group. It is needed for crash recovery. It may be in use for block recovery. It may or may not be archived.

```
SQL> SELECT group#, members, archived, status FROM v$log;
```

| GROUP# | MEMBERS | ARC | STATUS   |
|--------|---------|-----|----------|
| 1      | 2       | NO  | CURRENT  |
| 2      | 2       | YES | INACTIVE |
| 3      | 2       | YES | ACTIVE   |

```
SQL>
```

- g. Switch the log file again.

```
SQL> ALTER SYSTEM SWITCH LOGFILE;
```

```
System altered.
```

```
SQL>
```

- h. Query the `V$LOG` view again. The current group has changed again to group 2, and the status of both the other groups is now `ACTIVE`. Your current group may be different.

```
SQL> SELECT group#, members, archived, status FROM v$log;
```

| GROUP# | MEMBERS | ARC | STATUS  |
|--------|---------|-----|---------|
| 1      | 2       | YES | ACTIVE  |
| 2      | 2       | NO  | CURRENT |
| 3      | 2       | YES | ACTIVE  |

```
SQL>
```

- i. Question: Can the LGWR background process write to only one member of the `CURRENT` group in case the other members are missing or damaged?

Answer: Yes, it can. As long as there is one member left in the `CURRENT` group, LGWR can work.

7. To save space in your course practice environment, drop the redo log file members you created in step 4.

- a. Determine which redo log group is current. You cannot drop a member of the current group.

```
SQL> SELECT group#, status FROM v$log;
```

| GROUP# | STATUS   |
|--------|----------|
| 1      | INACTIVE |
| 2      | CURRENT  |
| 3      | INACTIVE |

```
SQL>
```

- b. Drop the member in the previous group and then perform a log switch. In this example, group 2 is current, so the command drops a member in group 1.

```
SQL> ALTER DATABASE DROP LOGFILE MEMBER  
'/u03/app/oracle/fast_recovery_area/ORCL/redo01b.log';
```

Database altered.

```
SQL> alter system switch logfile;
```

System altered.

```
SQL>
```

- c. Drop the member in the next group and then perform a log switch.

```
SQL> ALTER DATABASE DROP LOGFILE MEMBER  
'/u03/app/oracle/fast_recovery_area/ORCL/redo02b.log';
```

Database altered.

```
SQL> alter system switch logfile;
```

```
System altered.
```

```
SQL>
```

- d. Drop the member in the final group and then perform a log switch.

```
SQL> ALTER DATABASE DROP LOGFILE MEMBER
```

```
 '/u03/app/oracle/fast_recovery_area/ORCL/redo03b.log' ;
```

```
Database altered.
```

```
SQL> alter system switch logfile;
```

```
System altered.
```

```
SQL>
```

- e. Verify that each group now has only one member.

```
SQL> SELECT group#, members, archived, status FROM v$log;
```

| GROUP# | MEMBERS | ARC | STATUS  |
|--------|---------|-----|---------|
| 1      | 1       | YES | ACTIVE  |
| 2      | 1       | NO  | CURRENT |
| 3      | 1       | YES | ACTIVE  |

- f. Exit from SQL\*Plus.

```
SQL> exit
```

```
...
```

```
[oracle@MYDBCS ~]$
```

- g. Remove the physical files from the operating system.

```
[oracle@MYDBCS ~]$ rm  
/u03/app/oracle/fast_recovery_area/ORCL/redo*.log
```

- h. Verify that the redo log files have been removed.

```
[oracle@MYDBCS ~]$ ls /u03/app/oracle/fast_recovery_area/ORCL  
archivelog autobackup control02.ctl flashback onlinelog
```

## Lab 20: Verifying that ARCHIVELOG Mode is Configured

### Overview

In this practice, you verify that your database is in `ARCHIVELOG` mode so that redo logs are archived.

### Assumptions

### Tasks

1. Log in to SQL\*Plus as the `SYS` user with the `SYSDBA` privilege.

---

```
[oracle@MYDBCS ~]$ sqlplus / AS SYSDBA
...
SQL
```

2. Issue the `ARCHIVE LOG LIST` command to verify that the database is in `ARCHIVELOG` mode. The result confirms that the database log mode is set to `ARCHIVELOG` mode.

```
SQL> archive log list
Database log mode           Archive Mode
Automatic archival          Enabled
Archive destination         USE_DB_RECOVERY_FILE_DEST
Oldest online log sequence  27
Next log sequence to archive 29
Current log sequence        29
SQL>
```

In Oracle Database Cloud Service, the database is set to `ARCHIVELOG` mode by default. If you are creating your own Oracle database, you will need to place the database in `ARCHIVELOG` mode as described in this lesson.

3. Exit from SQL\*Plus.

```
SQL> EXIT
...
[oracle@MYDBCS ~]$
```

4) Place the database in archivelog mode

5) \$ mkdir /u01/fra

6) Sqlplus / as sysdba

```
sql> alter system set db_recovery_file_dest_size=10g;
```

```
sql> alter system set db_recovery_file_dest='/u01/fra'
```

```
sql> shutdown immediate
```

```
sql> startup mount
```

```
sql> alter database archivelog;
```

```
sql> alter database open;
```

```
sql> alter system switch logfile;
```

## Lab 21: Verifying Automatic Backups of the Control File and SPFILE

---

### Overview

In this practice, you use Recovery Manager (RMAN) to configure automatic backups of the control file and server parameter file (SPFILE) when a backup of the database is made and when there is a structural change to the database.

### Assumptions

You are logged in to the compute node as the `oracle` user.

You completed the practices in Lesson 17.

### Tasks

1. Start Recovery Manager and connect to the CDB root (target database) as the `SYS` user.

```
[oracle@MYDBCS ~]$ rman target /

Recovery Manager: Release 18.0.0.0.0 - Production on Wed Apr 4
20:42:53 2018
Version 18.1.0.0.0

Copyright (c) 1982, 2018, Oracle and/or its affiliates. All
rights reserved.

connected to target database: ORCL (DBID=2299035716)

RMAN>
```

2. Show all RMAN settings. Notice the `CONFIGURE CONTROLFILE AUTOBACKUP ON;` setting.

```
RMAN> SHOW ALL;

using target database control file instead of recovery catalog
RMAN configuration parameters for database with db_unique_name
ORCL are:
CONFIGURE RETENTION POLICY TO REDUNDANCY 1; # default
CONFIGURE BACKUP OPTIMIZATION ON;
CONFIGURE DEFAULT DEVICE TYPE TO DISK; # default
CONFIGURE CONTROLFILE AUTOBACKUP ON;
CONFIGURE CONTROLFILE AUTOBACKUP FORMAT FOR DEVICE TYPE DISK TO
'%F'; # default
CONFIGURE DEVICE TYPE DISK PARALLELISM 1 BACKUP TYPE TO
BACKUPSET; # default
```



```
CONFIGURE DATAFILE BACKUP COPIES FOR DEVICE TYPE DISK TO 1; #
default
CONFIGURE ARCHIVELOG BACKUP COPIES FOR DEVICE TYPE DISK TO 1; #
default
CONFIGURE MAXSETSIZE TO UNLIMITED; # default
CONFIGURE ENCRYPTION FOR DATABASE OFF; # default
CONFIGURE ENCRYPTION ALGORITHM 'AES128'; # default
CONFIGURE COMPRESSION ALGORITHM 'BASIC' AS OF RELEASE 'DEFAULT'
OPTIMIZE FOR LOAD TRUE ; # default
CONFIGURE RMAN OUTPUT TO KEEP FOR 7 DAYS; # default
CONFIGURE ARCHIVELOG DELETION POLICY TO NONE; # default
CONFIGURE SNAPSHOT CONTROLFILE NAME TO
'/u01/app/oracle/product/18.0.0/dbhome_1/dbs/snapcf_ORCL.f'; #
default

RMAN>
```

3. Question: In your configuration, does RMAN automatically back up the control file and server parameter file (SPFILE) with every backup and database structural change?

Answer: Yes, it does because the `CONTROLFILE AUTOBACKUP` attribute is set to `ON`.

4. Question: Will a backup operation back up all control files or only one of the multiplexed control files?

Answer: It will back up only one of the multiplexed control files because all control files in a database are identical.

5. Exit RMAN.

```
RMAN> EXIT
```

```
Recovery Manager complete.
```

```
[oracle@MYDBCS ~]$
```

## Lab 22: Creating a Whole Database Backup

---

### Overview

In this practice, you use Recovery Manager to back up your entire database, including the archived redo log files, the SPFILE, and the control files. The backup should be the base for an incremental backup strategy.

### Assumptions

You are logged in to the compute node as the `oracle` user.

You completed the following practices:

- Practice 17-2 Configure the Size of the Fast Recovery Area
- Practice 18-2 Verifying Automatic Backups of the Control File and SPFILE
- Practice 18-5: Checking Storage Availability

### Tasks

1. Start Oracle Recovery Manager (RMAN) and connect to the CDB root as the `SYS` user.

```
[oracle@MYDBCS ~]$ rman target /

Recovery Manager: Release 18.0.0.0.0 - Production on Thu Apr 5
13:43:04 2018
Version 18.1.0.0.0

Copyright (c) 1982, 2018, Oracle and/or its affiliates. All
rights reserved.

connected to target database: ORCL (DBID=2299035716)

RMAN>
```

2. View the structure of the CDB in terms of PDBs, tablespaces, and data files (permanent and temporary). Your file numbers will differ from those shown below.

```
RMAN> REPORT schema;

using target database control file instead of recovery catalog
Report of database schema for database with db_unique_name ORCL

List of Permanent Datafiles
=====
File Size(MB) Tablespace          RB segs Datafile Name
-----
1      870      SYSTEM              YES
/u02/app/oracle/oradata/ORCL/system01.dbf
```

```

3      1480      SYSAUX              NO
/u02/app/oracle/oradata/ORCL/sysaux01.dbf
4       60      UNDOTBS1             YES
/u02/app/oracle/oradata/ORCL/undotbs01.dbf
5      340      PDB$SEED:SYSTEM      NO
/u02/app/oracle/oradata/ORCL/pdbseed/system01.dbf
6      620      PDB$SEED:SYSAUX      NO
/u02/app/oracle/oradata/ORCL/pdbseed/sysaux01.dbf
7       5       USERS                NO
/u02/app/oracle/oradata/ORCL/users01.dbf
8      200      PDB$SEED:UNDOTBS1     NO
/u02/app/oracle/oradata/ORCL/pdbseed/undotbs01.dbf
12     350      PDB1:SYSTEM           YES
/u02/app/oracle/oradata/ORCL/PDB1/system01.dbf
13     790      PDB1:SYSAUX          NO
/u02/app/oracle/oradata/ORCL/PDB1/sysaux01.dbf
14     200      PDB1:UNDOTBS1        YES
/u02/app/oracle/oradata/ORCL/PDB1/undotbs01.dbf
15     50       PDB1:USERS            NO
/u02/app/oracle/oradata/ORCL/PDB1/PDB1_users01.dbf
19     350      PDB2:SYSTEM           YES
/u02/app/oracle/oradata/ORCL/PDB2/system01.dbf
20     670      PDB2:SYSAUX          NO
/u02/app/oracle/oradata/ORCL/PDB2/sysaux01.dbf
21     200      PDB2:UNDOTBS1        YES
/u02/app/oracle/oradata/ORCL/PDB2/undotbs01.dbf
38     7        PDB1:INVENTORY       NO
/u02/app/oracle/oradata/ORCL/PDB1/INVENTORY01.DBF

```

# List of Temporary Files

=====

| File | Size (MB) | Tablespace | Maxsize (MB) | Tempfile Name |
|------|-----------|------------|--------------|---------------|
|------|-----------|------------|--------------|---------------|

|   |     |                |       |  |
|---|-----|----------------|-------|--|
| 1 | 130 | TEMP           | 32767 | /u04/app/oracle/oradata/temp/temp01.dbf  |
| 2 | 62  | PDB\$SEED:TEMP | 32767 | /u04/app/oracle/oradata/temp/pdbseed_temp012018-02-08_13-49-27-256-PM.dbf      |
| 3 | 62  | PDB2:TEMP      | 32767 | /u04/app/oracle/oradata/temp/PDB2/pdbseed_temp012018-02-08_13-49-27-256-PM.dbf |
| 4 | 62  | PDB1:TEMP      | 32767 | /u04/app/oracle/oradata/temp/temp012018-02-08_13-49-27-256-PM.dbf              |

RMAN>

3. Back up the whole database. Your results will be different from the results shown below; for example, the piece handle names will be different.

```
RMAN> BACKUP DATABASE;

Starting backup at 09-APR-18
using target database control file instead of recovery catalog
allocated channel: ORA_DISK_1
channel ORA_DISK_1: SID=46 device type=DISK
channel ORA_DISK_1: starting full datafile backup set
channel ORA_DISK_1: specifying datafile(s) in backup set
input datafile file number=00001
name=/u02/app/oracle/oradata/ORCL/system01.dbf
input datafile file number=00003
name=/u02/app/oracle/oradata/ORCL/sysaux01.dbf
input datafile file number=00004
name=/u02/app/oracle/oradata/ORCL/undotbs01.dbf
input datafile file number=00007
name=/u02/app/oracle/oradata/ORCL/users01.dbf
channel ORA_DISK_1: starting piece 1 at 09-APR-18
channel ORA_DISK_1: finished piece 1 at 09-APR-18
piece
handle=/u03/app/oracle/fast_recovery_area/ORCL/backupset/2018_04_09/o1_mf_nnndf_TAG20180409T160501_fdqkvh9t_.bkp
tag=TAG20180409T160501 comment=NONE
channel ORA_DISK_1: backup set complete, elapsed time: 00:01:25
channel ORA_DISK_1: starting full datafile backup set
channel ORA_DISK_1: specifying datafile(s) in backup set
input datafile file number=00013
name=/u02/app/oracle/oradata/ORCL/PDB1/sysaux01
...
channel ORA_DISK_1: finished piece 1 at 09-APR-18
piece
handle=/u03/app/oracle/fast_recovery_area/ORCL/69350B8874FA03C8E053A23F160AC9F7/backupset/2018_04_09/o1_mf_nnndf_TAG20180409T160501_fdql1xjr_.bkp tag=TAG20180409T160501 comment=NONE
channel ORA_DISK_1: backup set complete, elapsed time: 00:00:55
Finished backup at 09-APR-18

Starting Control File and SPFILE Autobackup at 09-APR-18
piece
handle=/u03/app/oracle/fast_recovery_area/ORCL/autobackup/2018_04_09/o1_mf_s_973008564_fdql3o2s_.bkp comment=NONE
Finished Control File and SPFILE Autobackup at 09-APR-18

RMAN>
```

4. Question: Do you have to shut down the database to back it up?

Answer: No, as long as the database is in ARCHIVELOG mode, the backup can take place while the database is opened. This is a hot backup (or online backup). A cold backup (or offline backup) is a backup completed while the database is closed and is required if the database is in NOARCHIVELOG mode.

5. Question: Are hot backups consistent?

Answer: Online backups are inconsistent because with the database opened, there is no guarantee that the data files are synchronized with the control files. However, offline backups taken while the database is not opened are consistent because, at the time of the backup, the system change number (SCN) in data file headers matches the SCN in the control files.

6. Question: What would allow hot backups (inconsistent backups) to perform a complete database recovery?

Answer: During a complete recovery, restored online backups are recovered until the current SCN is matched, with the use of the archive log files and online redo log files.

7. Question: Did the backup include the SPFILE and control files?

Answer: Yes. This is the last operation completed at the end of the backup command.

```
...
Starting Control File and SPFILE Autobackup at 09-APR-18
piece
handle=/u03/app/oracle/fast_recovery_area/ORCL/autobackup/2018_0
4_09/o1_mf_s_973008564_fdql3o2s_.bkp comment=NONE
Finished Control File and SPFILE Autobackup at 09-APR-18
```

8. Question: Does the complete operation create a single backupset?

Answer: No. The operation creates multiple backupsets.

- Four backup sets including data files (one for each of the containers): CDB root, PDB seed, PDB1, PDB2
- One backup set for the SPFILE and control files.

9. List the backup sets. Look for Piece Name in the results for each backup set.

```
RMAN> LIST BACKUP;

List of Backup Sets
=====

BS Key    Type LV Size       Device Type Elapsed Time Completion
Time
-----
---
2         Full  17.95M    DISK          00:00:01      05-APR-18
        BP Key: 2   Status: AVAILABLE Compressed: NO  Tag:
TAG20180405T170539
```

```

      Piece Name:
/u03/app/oracle/fast_recovery_area/ORCL/autobackup/2018_04_05/o1
_mf_s_972666339_fdf3x4bf_.bkp
      SPFILE Included: Modification time: 05-APR-18
      SPFILE db_unique_name: ORCL
      Control File Included: Ckp SCN: 2845204          Ckp time: 05-APR-
18
...

BS Key   Type LV Size          Device Type Elapsed Time Completion
Time
-----
---
10       Full    18.23M      DISK          00:00:01      09-APR-18
      BP Key: 10      Status: AVAILABLE  Compressed: NO  Tag:
TAG20180409T160924
      Piece Name:
/u03/app/oracle/fast_recovery_area/ORCL/autobackup/2018_04_09/o1
_mf_s_973008564_fdql3o2s_.bkp
      SPFILE Included: Modification time: 09-APR-18
      SPFILE db_unique_name: ORCL
      Control File Included: Ckp SCN: 3108607          Ckp time: 09-APR-
18

RMAN>

```

#### 10. Exit RMAN.

```
RMAN> EXIT
```

#### 11. Verify that the files are stored on disk in the FRA.

```

$ cd /u03/app/oracle/fast_recovery_area/ORCL
$ ls -ltR
.:
total 18608
-rwxr-x--- 1 oracle oinstall 19021824 Apr  9 16:19 control02.ctl
drwxr-x--- 3 oracle oinstall    4096 Apr  9 16:08
69350B8874FA03C8E053A23F160AC9F7
drwxr-x--- 3 oracle oinstall    4096 Apr  9 16:07
64B4AE270F061AFDE0536604C40A9006
drwxr-x--- 3 oracle oinstall    4096 Apr  9 16:06
692121551ED82717E053E20D130AEB6C
drwxr-x--- 3 oracle oinstall    4096 Apr  9 16:05 backupset
drwxr-x--- 5 oracle oinstall    4096 Apr  9 15:01 autobackup
drwxr-x--- 5 oracle oinstall    4096 Apr  9 14:33 archivelog
drwxr-x--- 2 oracle oinstall    4096 Apr  5 16:51 flashback
drwxr-x--- 2 oracle oinstall    4096 Apr  5 16:49 onlinelog

```

```

./69350B8874FA03C8E053A23F160AC9F7:
total 4
drwxr-x--- 3 oracle oinstall 4096 Apr  9 16:08 backupset

...
./archivelog/2018_04_06:
total 0

./archivelog/2018_04_05:
total 0

./flashback:
total 2097176
-rwxr-x--- 1 oracle oinstall 1073750016 Apr  9 16:15
ol_mf_fdf30qbz_.flb
-rwxr-x--- 1 oracle oinstall 1073750016 Apr  5 16:52
ol_mf_fdf32pm6_.flb

./onlinelog:
total 0
[oracle@MYDBCS ORCL]$

```

- a. Question: Where are the backups of control files and SPFILE located?

Answer: They are created in the `autobackup` subdirectory.

- b. Question: How are backups deleted?

Answer: Space management in the FRA is governed by a backup retention policy. A retention policy determines when files are obsolete, which means that they are no longer needed to meet your data recovery objectives. The Oracle Database server automatically manages this storage by deleting files that are no longer needed.

## 12. View the backup retention policy.

- a. Start RMAN and connect to the CDB root as the `SYS` user.

```
$ rman target /
```

- b. Issue the `SHOW RETENTION POLICY` command. The policy is `REDUNDANCY 1`.

```

RMAN> SHOW RETENTION POLICY;

using target database control file instead of recovery catalog
RMAN configuration parameters for database with db_unique_name
ORCL are:
CONFIGURE RETENTION POLICY TO REDUNDANCY 1; # default

RMAN>

```

13. Question: How does Oracle determine when files are obsolete?

Answer: There are two retention policy parameters that are mutually exclusive:

- If a retention policy is enabled with `RECOVERY WINDOW OF 5 DAYS`, the window stretches from the current time (`SYSDATE`) to the point of recoverability, which is the earliest date to which you want to recover. The point of recoverability is `SYSDATE - integer days` in the past.
- If a retention policy is enabled with `REDUNDANCY r`, then RMAN skips backups only if at least  $n$  backups of an identical file exist on the specified device, where  $n=r+1$  (default is 1).

RMAN automatically deletes obsolete backup sets and copies in the FRA when space is needed.

14. Manually delete obsolete files by issuing the `DELETE OBSOLETE` command.

```
RMAN> delete obsolete;

using target database control file instead of recovery catalog
RMAN retention policy will be applied to the command
RMAN retention policy is set to redundancy 1
allocated channel: ORA_DISK_1
channel ORA_DISK_1: SID=4 device type=DISK
Deleting the following obsolete backups and copies:
Type                               Key      Completion Time      Filename/Handle
-----
---
Backup Set                         2        05-APR-18
  Backup Piece                     2        05-APR-18
/u03/app/oracle/fast_recovery_area/ORCL/autobackup/2018_04_05/o1
_mf_s_972666339_fdf3x4bf_.bkp
Backup Set                         3        06-APR-18
  Backup Piece                     3        06-APR-18
/u03/app/oracle/fast_recovery_area/ORCL/autobackup/2018_04_06/o1
_mf_s_972751516_fdhq2y5d_.bkp
Backup Set                         4        09-APR-18
  Backup Piece                     4        09-APR-18
/u03/app/oracle/fast_recovery_area/ORCL/autobackup/2018_04_09/o1
_mf_s_973004494_fdqg4gqp_.bkp
Backup Set                         5        09-APR-18
  Backup Piece                     5        09-APR-18
/u03/app/oracle/fast_recovery_area/ORCL/autobackup/2018_04_09/o1
_mf_s_973005697_fdqhb2cd_.bkp

Do you really want to delete the above objects (enter YES or
NO)? yes
deleted backup piece
```



```

backup piece
handle=/u03/app/oracle/fast_recovery_area/ORCL/autobackup/2018_0
4_05/o1_mf_s_972666339_fdf3x4bf_.bkp RECID=2 STAMP=972666340
deleted backup piece
backup piece
handle=/u03/app/oracle/fast_recovery_area/ORCL/autobackup/2018_0
4_06/o1_mf_s_972751516_fdhq2y5d_.bkp RECID=3 STAMP=972751517
deleted backup piece
backup piece
handle=/u03/app/oracle/fast_recovery_area/ORCL/autobackup/2018_0
4_09/o1_mf_s_973004494_fdqg4gqp_.bkp RECID=4 STAMP=973004494
deleted backup piece
backup piece
handle=/u03/app/oracle/fast_recovery_area/ORCL/autobackup/2018_0
4_09/o1_mf_s_973005697_fdqhb2cd_.bkp RECID=5 STAMP=973005698
Deleted 4 objects

RMAN>

```

15. Back up the database and archive logs as image copies. At the same time, free space in the FRA by deleting the archive log files once they are backed up.

- a. Perform the backup.

```

RMAN> BACKUP AS COPY DATABASE PLUS ARCHIVELOG DELETE INPUT;

Starting backup at 06-JUN-18
current log archived
using target database control file instead of recovery catalog
allocated channel: ORA_DISK_1
channel ORA_DISK_1: SID=284 device type=DISK
channel ORA_DISK_1: starting archived log copy
input archived log thread=1 sequence=15 RECID=15 STAMP=978029619
output file
name=/u03/app/oracle/fast_recovery_area/ORCL/archivelog/2018_06_
06/o1_mf_1_15_fkjlfpws_.arc RECID=21 STAMP=978105797
...
Finished backup at 06-JUN-18

Starting backup at 06-JUN-18
using channel ORA_DISK_1
channel ORA_DISK_1: starting datafile copy
input datafile file number=00003
name=/u02/app/oracle/oradata/ORCL/sysaux01.dbf
...
Starting Control File and SPFILE Autobackup at 06-JUN-18

```

```
piece
handle=/u03/app/oracle/fast_recovery_area/ORCL/autobackup/2018_0
6_06/o1_mf_s_978106455_fkj22r3p_.bkp comment=NONE
Finished Control File and SPFILE Autobackup at 06-JUN-18

RMAN>
```

- b. Question: What would you do if an error such as the following occurs?

```
RMAN-00571:
=====
RMAN-00569: ===== ERROR MESSAGE STACK FOLLOWS
=====
RMAN-00571:
=====
RMAN-03002: failure of backup plus archivelog command at
01/23/2017 11:05:08
ORA-19809: limit exceeded for recovery files
ORA-19804: cannot reclaim 67108864 bytes disk space from
19327352832 bytes limit
```

Answer: Increase the DB\_RECOVERY\_FILE\_DEST\_SIZE parameter value to 30G by issuing the following command:

```
RMAN> ALTER SYSTEM SET db_recovery_file_dest_size = 30G
SCOPE=both;
```

- c. Question: What is the advantage of creating backups as image copies?

Answer: The advantage of creating a backup as an image copy is improved granularity of the restore operation. With an image copy, only the file or files need to be retrieved from your backup location. With backup sets, the entire backup set must be retrieved from your backup location before you extract the file or files that are needed.

- d. Question: What is the advantage of creating backups as backup sets?

Answer: The advantage of creating backups as backup sets is better space usage. In most databases, 20% or more of the data blocks are empty blocks. Image copies back up every data block, even if the data block is empty. Backup sets significantly reduce the space required by the backup. In most systems, the advantages of backup sets outweigh the advantages of image copies.

- e. Question: How many image copies of the data files are created?

Answer: There are 15 image copies, one image copy for each data file in the CDB, PDBs included.

#### 16. Exit RMAN.

```
RMAN> EXIT
```

## Lab 23: Creating Partial Database Backups

---

### Overview

In this practice, you use Recovery Manager to back up PDB1, including the archived redo log files. This is a partial database backup.

### Assumptions

You are logged in to the compute node as the `oracle` user.

### Tasks

1. Start Recovery Manager (RMAN) and connect to the CDB root as the SYS user.

```
[oracle@MYDBCS ORCL]$ rman target /

Recovery Manager: Release 18.0.0.0.0 - Production on Wed Jun 6
16:48:45 2018
Version 18.1.0.0.0

Copyright (c) 1982, 2018, Oracle and/or its affiliates. All
rights reserved.

connected to target database: ORCL (DBID=1505229725)

RMAN>
```

2. Back up PDB1, including the archived redo log files.

```
RMAN> BACKUP PLUGGABLE DATABASE PDB1 PLUS ARCHIVELOG;

Starting backup at 06-JUN-18
current log archived
using target database control file instead of recovery catalog
allocated channel: ORA_DISK_1
channel ORA_DISK_1: SID=41 device type=DISK
channel ORA_DISK_1: starting archived log backup set
channel ORA_DISK_1: specifying archived log(s) in backup set
input archived log thread=1 sequence=14 RECID=22 STAMP=978105812
...
Starting Control File and SPFILE Autobackup at 06-JUN-18
piece
handle=/u03/app/oracle/fast_recovery_area/ORCL/autobackup/2018_0
6_06/o1_mf_s_978108946_fkj4jm29_.bkp comment=NONE
Finished Control File and SPFILE Autobackup at 06-JUN-18

RMAN>
```

3. Exit RMAN.

```
RMAN> EXIT  
Recovery Manager complete.  
[oracle@MYDBCS ORCL]$
```

4. Question: Did the partial backup automatically include the SPFILE and control files?

Answer: Yes. The setting verified in Practice 18-2 Verifying Automatic Backups of the Control File and SPFILE is also valid for partial backups.

5. Question: How many backup sets are created?

Answer: Four backup sets: one for the PDB data files, one for the SPFILE and control file, one for the archived log files before the data file backup set, and one for the archived log files after the data file backup set.

6. Question: Can you connect in RMAN directly to the PDB to perform the same backup?

Answer: Yes. In this case, you do not have to specify that you want to back up a PDB. Instead, you can use the simple BACKUP DATABASE command.

7. Perform a partial database backup in PDB1 directly.

a. Start RMAN and connect to PDB1 as the SYS user.

```
[oracle@MYDBCS ORCL]$ rman target SYS/password@PDB1  
...  
connected to target database: ORCL:PDB1 (DBID=2133829969)  
  
RMAN>
```

b. Execute the BACKUP DATABASE command. Notice that the SPFILE and control file are not backed up.

```
RMAN> BACKUP DATABASE;  
Starting backup at 06-JUN-18  
using target database control file instead of recovery catalog  
allocated channel: ORA_DISK_1  
channel ORA_DISK_1: SID=41 device type=DISK  
...  
channel ORA_DISK_1: starting piece 1 at 06-JUN-18  
channel ORA_DISK_1: finished piece 1 at 06-JUN-18  
piece  
handle=/u03/app/oracle/fast_recovery_area/ORCL/6D5DA61701C1285EE  
0533E89810ACAA5/backupset/2018_06_06/o1_mf_nnndf_TAG20180606T165  
917_fkj4q7w6_.bkp tag=TAG20180606T165917 comment=NONE  
channel ORA_DISK_1: backup set complete, elapsed time: 00:01:25  
Finished backup at 06-JUN-18  
  
RMAN>
```

8. Try to configure the recovery setting for the PDB so that the SPFILE and control file are backed up too. You get an error message because you must be connected to the CDB root to configure any recovery settings.

```
RMAN> CONFIGURE CONTROLFILE AUTOBACKUP ON;
RMAN-00571:
=====
RMAN-00569: ===== ERROR MESSAGE STACK FOLLOWS
=====
RMAN-00571:
=====
RMAN-03002: failure of configure command at 06/06/2018 17:02:27
RMAN-07536: command not allowed when connected to a Pluggable
Database

RMAN>
```

9. Exit RMAN.

```
RMAN> EXIT
```

10. Back up the TBS\_APP tablespace in PDB2.

- a. Connect to PDB2 as the SYS user.

```
$ rman target SYS/password@PDB2
...
connected to target database: ORCL:PDB2 (DBID=3237529478)

RMAN>
```

- b. Back up the tablespace.

```
RMAN> BACKUP TABLESPACE tbs_app;
Starting backup at 06-JUN-18
using target database control file instead of recovery catalog
allocated channel: ORA_DISK_1
channel ORA_DISK_1: SID=55 device type=DISK
channel ORA_DISK_1: starting full datafile backup set
channel ORA_DISK_1: specifying datafile(s) in backup set
input datafile file number=00059
name=/u02/app/oracle/oradata/ORCL/PDB2/tbs_app01.dbf
channel ORA_DISK_1: starting piece 1 at 06-JUN-18
channel ORA_DISK_1: finished piece 1 at 06-JUN-18
piece
handle=/u03/app/oracle/fast_recovery_area/ORCL/6D975E8B80B85F14E
0537A051D0A3C0D/backupset/2018_06_06/o1_mf_nnndf_TAG20180606T170
407_fkj5091o_.bkp tag=TAG20180606T170407 comment=NONE
channel ORA_DISK_1: backup set complete, elapsed time: 00:00:04
Finished backup at 06-JUN-18
```

```
RMAN>
```

- c. Exit RMAN.

```
RMAN> EXIT
```

```
Recovery Manager complete.
```

```
[oracle@DKKDBCS ORCL]$
```

11. Can you connect to the CDB root and perform the same operation?

- a. Start RMAN and connect to the CDB root as the SYS user.

```
$ rman target /
```

```
...
```

```
connected to target database: ORCL (DBID=1505229725)
```

```
RMAN>
```

- b. Back up the TBS\_APP tablespace in PDB2. You must specify the PDB in which the tablespace exists.

```
RMAN> BACKUP TABLESPACE PDB2:tbs_app;
```

```
Starting backup at 06-JUN-18
```

```
using target database control file instead of recovery catalog
```

```
allocated channel: ORA_DISK_1
```

```
channel ORA_DISK_1: SID=271 device type=DISK
```

```
channel ORA_DISK_1: starting full datafile backup set
```

```
channel ORA_DISK_1: specifying datafile(s) in backup set
```

```
input datafile file number=00059
```

```
name=/u02/app/oracle/oradata/ORCL/PDB2/tbs_app01.dbf
```

```
channel ORA_DISK_1: starting piece 1 at 06-JUN-18
```

```
channel ORA_DISK_1: finished piece 1 at 06-JUN-18
```

```
piece
```

```
handle=/u03/app/oracle/fast_recovery_area/ORCL/6D975E8B80B85F14E
```

```
0537A051D0A3C0D/backupset/2018_06_06/o1_mf_nnndf_TAG20180606T170
```

```
547_fkj53f1h_.bkp tag=TAG20180606T170547 comment=NONE
```

```
channel ORA_DISK_1: backup set complete, elapsed time: 00:00:03
```

```
Finished backup at 06-JUN-18
```

```
Starting Control File and SPFILE Autobackup at 06-JUN-18
```

```
piece
```

```
handle=/u03/app/oracle/fast_recovery_area/ORCL/autobackup/2018_0
```

```
6_06/o1_mf_s_978109551_fkj53l2p_.bkp comment=NONE
```

```
Finished Control File and SPFILE Autobackup at 06-JUN-18
```

```
RMAN>
```

- c. Question: Did the operation back up only the tablespace data files?

## Lab 24: Recovering from the Loss of a System-Critical Data File

---

### Overview

In this practice, you recover your CDB after the data file for the `SYSTEM` tablespace (in the CDB root) has been inadvertently removed.

### Tip

Because you use several windows at the same time in this practice, you may find it helpful to change the name of each of them in their banner at the top.

To set a title for a terminal window:

1. In the terminal window's menu, select **Terminal** and then **Set Title**. A Set Title dialog box is displayed.
2. In the Title box, enter the window number.
3. Click **OK**.

### Assumptions

You are logged in as the `oracle` user.

You completed the following practices:

- Practice 18-2 Verifying Automatic Backups of the Control File and SPFILE
- Practice 18-4 Creating a Whole Database Backup

### Tasks

#### Create a Loss of a System-Critical Data File

##### Window 1

1. Open a new terminal window and connect to the compute node. This window will be referred to as Window 1.
- 2) Identify the system datafile under the root container
- 3) Select \* from dba\_data\_files where tablespace\_name = 'SYSTEM' and con\_id = 1
- 4) From the OS rm the datafile associated to that tablespace

## Recover the Database by Using the `RESTORE` and `RECOVER` Commands

### Window 1

1. Start Recovery Manager (RMAN) and connect to the target database (CDB root).

```
[oracle@MYDBCS ~]$ rman target /  
...  
connected to target database: ORCL (DBID=1500451933)  
  
RMAN>
```

2. Issue the `RESTORE` command. You must provide the number for the missing data file.

```
RMAN> RESTORE DATAFILE 1;  
  
Starting restore at 13-JUN-18  
using target database control file instead of recovery catalog  
allocated channel: ORA_DISK_1  
channel ORA_DISK_1: SID=46 device type=DISK  
  
channel ORA_DISK_1: restoring datafile 00001  
input datafile copy RECID=2 STAMP=978678712 file  
name=/u03/app/oracle/fast_recovery_area/ORCL/datafile/ol_mf_syst  
em_fl1jwb73_.dbf  
destination for restore of datafile 00001:  
/u02/app/oracle/oradata/ORCL/system01.dbf  
RMAN-00571:  
=====
```

```
RMAN-00569: ===== ERROR MESSAGE STACK FOLLOWS  
=====
```

```
RMAN-00571:  
=====
```

```
RMAN-03002: failure of restore command at 06/13/2018 08:08:12  
ORA-19573: cannot obtain exclusive enqueue for datafile 1
```



```
ORA-45909: restore, recover or block media recovery may be in progress
ORA-19600: input file is datafile-copy 2
(/u03/app/oracle/fast_recovery_area/ORCL/datafile/o1_mf_system_f
11jwb73_.dbf)
ORA-19601: output file is datafile 1
(/u02/app/oracle/oradata/ORCL/system01.dbf)
RMAN>
```

3. Question: What does the error message "cannot obtain exclusive enqueue for datafile 1" mean?

Answer: The restore operation requires an exclusive enqueue lock on the data file 1 that is not obtainable. In this case, you have to open the database instance in `MOUNT` mode. This means that you have to shut down the database instance if it did not already abort.

4. Question: What does a database instance shut down do?

Answer: This closes all PDBs and therefore prevents all users from working during the recovery operation.

5. Try to shut down the database instance in `IMMEDIATE` mode. You get an error.

```
RMAN> SHUTDOWN IMMEDIATE

RMAN-00571:
=====
RMAN-00569: ===== ERROR MESSAGE STACK FOLLOWS
=====
RMAN-00571:
=====
RMAN-03002: failure of shutdown command at 06/13/2018 08:10:09
ORA-01116: error in opening database file 1
ORA-01110: data file 1:
'/u02/app/oracle/oradata/ORCL/system01.dbf'
ORA-27041: unable to open file
Linux-x86_64 Error: 2: No such file or directory
Additional information: 3

RMAN>
```

6. Question: Why does the `SHUTDOWN IMMEDIATE` command fail?

Answer: RMAN needs to close all data files cleanly by writing the current SCN to all data file headers. This cannot be done because data file 1 is missing.

7. Reconnect to the database.

8. Try to shut down the database instance in `ABORT` mode.

```
RMAN> SHUTDOWN ABORT

Oracle instance shut down
```

```
RMAN>
```

9. Start the database instance in MOUNT mode.

```
RMAN> STARTUP MOUNT
```

```
Oracle instance started  
database mounted
```

```
Total System Global Area      2768239832 bytes
```

```
Fixed Size                      8899800 bytes
```

```
Variable Size                   704643072 bytes
```

```
Database Buffers                1979711488 bytes
```

```
Redo Buffers                     74985472 bytes
```

```
RMAN>
```

10. Restore the missing data file.

```
RMAN> RESTORE DATAFILE 1;
```

```
Starting restore at 10-APR-18
```

```
allocated channel: ORA_DISK_1
```

```
channel ORA_DISK_1: SID=25 device type=DISK
```

```
channel ORA_DISK_1: starting datafile backup set restore
```

```
channel ORA_DISK_1: specifying datafile(s) to restore from  
backup set
```

```
channel ORA_DISK_1: restoring datafile 00001 to  
/u02/app/oracle/oradata/ORCL/system01.dbf
```

```
channel ORA_DISK_1: reading from backup piece
```

```
/u03/app/oracle/fast_recovery_area/ORCL/backupset/2018_04_09/o1_  
mf_nnndf_TAG20180409T160501_fdqkvh9t_.bkp
```

```
channel ORA_DISK_1: piece
```

```
handle=/u03/app/oracle/fast_recovery_area/ORCL/backupset/2018_04_  
_09/o1_mf_nnndf_TAG20180409T160501_fdqkvh9t_.bkp  
tag=TAG20180409T160501
```

```
channel ORA_DISK_1: restored backup piece 1
```

```
channel ORA_DISK_1: restore complete, elapsed time: 00:00:15
```

```
Finished restore at 10-APR-18
```

```
RMAN>
```

11. Recover the missing data file.

```
RMAN> RECOVER DATAFILE 1;
```

```
Starting recover at 10-APR-18
using channel ORA_DISK_1

starting media recovery

archived log for thread 1 with sequence 14 is already on disk as
file
/u03/app/oracle/fast_recovery_area/ORCL/archivelog/2018_04_10/o1
_mf_1_14_fdskgq34_.arc
archived log for thread 1 with sequence 15 is already on disk as
file
/u03/app/oracle/fast_recovery_area/ORCL/archivelog/2018_04_10/o1
_mf_1_15_fdsmhyob_.arc
archived log for thread 1 with sequence 16 is already on disk as
file
/u03/app/oracle/fast_recovery_area/ORCL/archivelog/2018_04_10/o1
_mf_1_16_fdsmhyqz_.arc
archived log file
name=/u03/app/oracle/fast_recovery_area/ORCL/archivelog/2018_04_
10/o1_mf_1_14_fdskgq34_.arc thread=1 sequence=14
media recovery complete, elapsed time: 00:00:01
Finished recover at 10-APR-18

RMAN>
```

12. Open the CDB root.

```
RMAN> ALTER DATABASE OPEN;

Statement processed

RMAN>
```

13. Open all PDBs.

```
RMAN> ALTER PLUGGABLE DATABASE ALL OPEN;

Statement processed

RMAN>
```

14. Exit RMAN.

```
RMAN> EXIT

Recovery Manager complete.
[oracle@MYDBCS ~]$
```

15. Start SQL\*Plus and connect to the CDB root as the `SYSTEM` user. Refer to *Course Practice Environment: Security Credentials* for the password value.

```
SQL> CREATE USER c##test IDENTIFIED BY DBAdmin_1;
```

```
User created.
```

```
SQL>
```

```
$ sqlplus SYSTEM/password
```

```
...
```

```
SQL>
```

16. Try creating the `c##test` user again. This time the user is created.

17. Keep Window 1 open for the next section.

### Use the Data Recovery Advisor to Recover the Database

1. **Window 2:** Open a new terminal window and execute the `RMAN_crash.sh` script to create a failure. This window will be referred to as Window 2.

Answer: Remember that the DBWR background process does not necessarily write immediately into the data files.

4. **Window 1:** Attempt to resize datafile 1. If it completes, execute the `ALTER SYSTEM SWITCH LOGFILE` command. You should receive an error message about the missing data file.

```
SQL> ALTER DATABASE DATAFILE 1 RESIZE 1G;
ALTER DATABASE DATAFILE 1 RESIZE 1G
*
ERROR at line 1:
ORA-01565: error in identifying file
'/u02/app/oracle/oradata/ORCL/system01.dbf'
ORA-27037: unable to obtain file status
Linux-x86_64 Error: 2: No such file or directory
Additional information: 7

SQL>
```

5. **Window 1:** Exit SQL\*Plus.

```
SQL> EXIT
...
[oracle@MYDBCS ~]$
```

6. **Window 1:** Start RMAN and connect to the target database.

```
$ rman target /
...
connected to target database (not started)

RMAN>
```

7. **Window 1:** Start the database instance in MOUNT mode.

```
RMAN> STARTUP MOUNT;

Oracle instance started
database mounted

Total System Global Area      2768239832 bytes

Fixed Size                     8899800 bytes
Variable Size                  704643072 bytes
Database Buffers               1979711488 bytes
Redo Buffers                    74985472 bytes

RMAN>
```

8. **Window 1:** Use the `LIST FAILURE` command to determine the error. The value in the Summary column tells you that `system01.dbf` is missing.

```
RMAN> LIST FAILURE;

using target database control file instead of recovery catalog
Database Role: PRIMARY

List of Database Failures
=====

Failure ID Priority Status      Time Detected Summary
-----
262          CRITICAL OPEN      10-APR-18    System datafile 1:
'/u02/app/oracle/oradata/ORCL/system01.dbf' is missing

RMAN>
```

9. **Window 1:** Display repair options. At the very end of the results, a repair script is listed.

```
RMAN> ADVISE FAILURE;

Database Role: PRIMARY

List of Database Failures
=====

Failure ID Priority Status      Time Detected Summary
-----
262          CRITICAL OPEN      10-APR-18    System datafile 1:
'/u02/app/oracle/oradata/ORCL/system01.dbf' is missing

analyzing automatic repair options; this may take some time
allocated channel: ORA_DISK_1
channel ORA_DISK_1: SID=25 device type=DISK
analyzing automatic repair options complete

Mandatory Manual Actions
=====
no manual actions available

Optional Manual Actions
=====
1. If file /u02/app/oracle/oradata/ORCL/system01.dbf was
unintentionally renamed or moved, restore it
```

```
Automated Repair Options
=====
Option Repair Description
-----
1      Restore and recover datafile 1
      Strategy: The repair includes complete media recovery with no
data loss
      Repair script:
/u01/app/oracle/diag/rdbms/orcl/ORCL/hm/reco_2771153169.hm

RMAN>
```

10. **Window 1:** Use the REPAIR FAILURE PREVIEW command to generate a script with all repair actions and comments.

```
RMAN> REPAIR FAILURE PREVIEW;

Strategy: The repair includes complete media recovery with no
data loss
Repair script:
/u01/app/oracle/diag/rdbms/orcl/ORCL/hm/reco_2771153169.hm

contents of repair script:
# restore and recover datafile
restore ( datafile 1 );
recover datafile 1;
sql 'alter database datafile 1 online';

RMAN>
```

11. **Window 1:** Use the REPAIR FAILURE command to repair database failures identified by the Data Recovery Advisor. When prompted, enter YES to execute the repair. When prompted to open the database, enter YES.

```
RMAN> REPAIR FAILURE;

Strategy: The repair includes complete media recovery with no
data loss
Repair script:
/u01/app/oracle/diag/rdbms/orcl/ORCL/hm/reco_2771153169.hm

contents of repair script:
# restore and recover datafile
restore ( datafile 1 );
recover datafile 1;
sql 'alter database datafile 1 online';
```

Do you really want to execute the above repair (enter YES or NO)? **YES**

executing repair script

Starting restore at 10-APR-18

using channel ORA\_DISK\_1

channel ORA\_DISK\_1: starting datafile backup set restore

channel ORA\_DISK\_1: specifying datafile(s) to restore from backup set

channel ORA\_DISK\_1: restoring datafile 00001 to  
/u02/app/oracle/oradata/ORCL/system01.dbf

channel ORA\_DISK\_1: reading from backup piece

/u03/app/oracle/fast\_recovery\_area/ORCL/backupset/2018\_04\_09/o1\_  
mf\_nnndf\_TAG20180409T160501\_fdqkvh9t\_.bkp

channel ORA\_DISK\_1: piece

handle=/u03/app/oracle/fast\_recovery\_area/ORCL/backupset/2018\_04\_  
\_09/o1\_mf\_nnndf\_TAG20180409T160501\_fdqkvh9t\_.bkp  
tag=TAG20180409T160501

channel ORA\_DISK\_1: restored backup piece 1

channel ORA\_DISK\_1: restore complete, elapsed time: 00:00:07

Finished restore at 10-APR-18

Starting recover at 10-APR-18

using channel ORA\_DISK\_1

starting media recovery

archived log for thread 1 with sequence 14 is already on disk as file

/u03/app/oracle/fast\_recovery\_area/ORCL/archivelog/2018\_04\_10/o1\_  
\_mf\_1\_14\_fdskgq34\_.arc

archived log for thread 1 with sequence 15 is already on disk as file

...

archived log file

name=/u03/app/oracle/fast\_recovery\_area/ORCL/archivelog/2018\_04\_  
10/o1\_mf\_1\_16\_fdsmhyqz\_.arc thread=1 sequence=16

archived log file

name=/u03/app/oracle/fast\_recovery\_area/ORCL/archivelog/2018\_04\_  
10/o1\_mf\_1\_17\_fdsn3sf0\_.arc thread=1 sequence=17

media recovery complete, elapsed time: 00:00:01

Finished recover at 10-APR-18

~~sql statement: alter database datafile 1 online~~



```
repair failure complete
```

```
Do you want to open the database (enter YES or NO)? YES  
database opened
```

```
RMAN>
```

12. **Window 1:** Open all the PDBs.

```
RMAN> ALTER PLUGGABLE DATABASE ALL OPEN;
```

```
Statement processed
```

```
RMAN>
```

13. **Window 1:** Exit RMAN.

```
RMAN> EXIT
```

```
Recovery Manager complete.
```

```
[oracle@MYDBCS ~]$
```

14. **Window 2:** Close the terminal window.