

Oracle 19c New Features Lab Guide

Table of Contents

Course Practice Environment: Security Credentials.....	5
Course Practice Environment: Security Credentials.....	6
Practices for Lesson 1: Introduction	7
Practices for Lesson 1: Overview	8
Practice 1-1: Discovering the Practice Environment	9
Practices for Lesson 2: Leveraging Multitenant Enhancements.....	13
Practices for Lesson 2: Overview	14
Practice 2-1: Managing a CDB Fleet	15
Practice 2-2: Managing and Using PDB Snapshots	16
Practice 2-3: Using a Dynamic Container Map	40
Practice 2-4: Using Static and Dynamic Lockdown Profiles.....	47
Practice 2-5: Switching Over Refreshable Cloned PDBs	67
Practices for Lesson 3: Managing Security.....	75
Practices for Lesson 3: Overview	76
Practice 3-1: Creating Schema-Only Accounts	77
Practice 3-2: Managing PDB Keystores	80
Practice 3-3: Creating User-Defined TDE Master Encryption Keys	92
Practice 3-4: Exporting and Importing Fixed-User Database Links.....	100
Practice 3-5: Encrypting Sensitive Data in Database Replay Files	109
Practices for Lesson 4: Using RMAN Enhancements	115
Practices for Lesson 4: Overview	116
Practice 4-1: Recovering a Plugged Non-CDB Using Preplugin Backups.....	117
Practice 4-2: Recovering a Plugged PDB Using Preplugin Backups.....	118
Practice 4-3: Duplicating a PDB into an Existing CDB	132
Practice 4-4: Duplicating an On-Premises CDB for Cloud.....	133
Practices for Lesson 5: Using General Database Enhancements	135
Practices for Lesson 5: Overview	136
Practice 5-1: Managing Private Temporary Tables	137
Practice 5-2: Using Data Pump Import with <code>CONTINUE_LOAD_ON_FORMAT_ERROR</code>	147
Practice 5-3: Converting <code>HASH</code> Partitioned Tables to <code>RANGE</code> Partitioned Tables Online	154
Practice 5-4: Converting <code>LIST</code> Partitioned Tables to <code>LIST AUTOMATIC</code> Partitioned Tables Online	161
Practice 5-5: Converting <code>LIST AUTOMATIC</code> Partitioned Tables to Subpartitioned Tables Online	166
Practice 5-6: Merging Partitions of a Partitioned Table Online	171
Practice 5-7: Using Batched DDL.....	178

Practices for Lesson 6: Improving Performance 197

 Practices for Lesson 6: Overview..... 198

 Practice 6-1: Configuring and Using AIM..... 199

 Practice 6-2: Tracking IM Expressions Within a Window Capture 206

Practices for Lesson 7: Handling Enhancements in Big Data and Data Warehousing..... 215

 Practices for Lesson 7: Overview..... 216

 Practice 7-1: Querying Inlined External Tables..... 217

 Practice 7-2: Populating External Tables in the In-Memory Column Store..... 221

 Practice 7-3: Using Hierarchy-Based Predicates and Calculated Measures on Analytic Views..... 227

 Practice 7-4: Using Polymorphic Table Functions..... 237

Practices for Lesson 8: Describing Sharding Enhancements..... 249

 Practices for Lesson 8 250

Course Practice Environment: Security Credentials

For product-specific credentials used in this course, see the following table:

Product-Specific Credentials		
Product/Application	Username	Password
VM (for on-premise DBs)		
ORCL and PDB1	Any user	Welcome_1
Other PDBs created in ORCL	Any user	Welcome_1
Practice 3-2: CDB keystore		Welcome_1
Practice 3-2: PDB keystore		<i>Your_defined_password</i> You can set Welcome_1 or any other value. Then distinguish <i>password_CDB</i> and <i>password_PDB</i> .
Practice 3-4: steps i and j, Data Pump export and import encryption password		<i>Your_defined_password</i> You can set any value for <i>enc_password</i> . Then use the same value set at export (step i), at import (step j).
SSO wallet in practice 3-5 step 4 (d)		Welcome1

Lab Environment

In this practice, you will start oracle database which has been already created as follows:


ORACLE_SID=fenagodb

ORACLE_PDB=fenagodb1

ORACLE_PWD=fenago

Tasks

1. Log in to your lab environment and double click **“Start Oracle Database”** shortcut to start database server.




Overview


In this practice, you will switch to oracle user from terminal

Tasks

- a. Open terminal and run “xhost +” command as root user:



- b. Run and run “su - oracle” command in the terminal to switch to **oracle** user:



Practices for Lesson 1: Introduction

Practices for Lesson 1: Overview

Overview

In this practice, you will discover the system environment for the practices of this course.

Notes

- All practices are independent from one lesson to another.
 - There are two solutions to clean up your CDB.
1. At the beginning of the practices of each lesson, you can execute the cleanup shell script from previous practices, then the first shell script of the current practice.
 2. Alternatively, you can execute the `/home/oracle/labs/admin/cleanup_PDBs.sh` shell script. The shell script drops all PDBs that may have been created by any of the practices, and finally recreates `PDB1`.

Practice 1-1: Discovering the Practice Environment

Your system has one compute node (VM). The VM is dedicated to Oracle Database 19c, installed with a pre-created CDB and one pre-created PDB.

- The ORCL CDB with PDB1 PDB
- Net service names for all PDBs created in practices are already logged in the `$ORACLE_HOME/network/admin/tnsnames.ora` file.

The configuration for the ORCL database on the VM corresponds to the preconfigured directories on an Oracle Cloud compute node associated with the pre-created ORCL database of a database deployment (or Database Cloud Service instance).

- CDB data files in `/u02/app/oracle/oradata/`
- CDB root data files in `/u02/app/oracle/oradata/<dbname>`
- PDB data files in `/u02/app/oracle/oradata/<dbname>/PDB1`
- Control files in `/u02/app/oracle/oradata/<dbname>` and `/u03/app/oracle/fast_recovery_area/<dbname>`
- All redo log files in `/u04/app/oracle/redo` (The preconfigured directory for redo logs of the unique CDB on Oracle Cloud is `/u04/app/oracle/redo`.)
- All backup files in `/u03/app/oracle/fast_recovery_area/<dbname>`
- Password and initialization parameter files in `$ORACLE_HOME/dbs`
- Diagnostics files in `/u01/app/oracle/diag/rdbms/<dbname_lower_case>/<dbname>/...`
- TDE keystore in `/u01/app/oracle/admin/<dbname>/tde_wallet`
- Net files in `$ORACLE_HOME/network/admin`

1. On the Welcome page of the Secure Global Desktop, in the Cloud Login Instructions section, you find important information:
 - The IP address of the VM
 - The text of the SSH private key that will allow you to connect to the VM. Be ready to copy this text into an SSH private key file that you are going to create. The text starts with `-----BEGIN RSA PRIVATE KEY-----` and ends with `-----END RSA PRIVATE KEY-----`. Both lines will be included in the SSH private key file.
2. In the left navigator of the Secure Global Desktop, click My Training Environments.
3. Select the Linux Desktop and click Start.
4. Enter the `oracle` password and change the password as requested. Then click OK to log in to the Linux Desktop.
5. Right-click Terminal and select Open to open a Linux session.
6. Edit the `/home/oracle/.ssh/config` file and update the first line `Host <IP address of DB instance from OPC>` to `Host IP_Address` with the IP address that you obtained in step 1. Save the file.

7. Open a text editor to create the file, `mykey` and copy the SSH private key text into the file as defined in step 1. Save the file.
8. Change the permissions of the `mykey` file.

```
$ chmod 600 mykey
$
```

9. Use the SSH private key file, `mykey`, to log in to the VM whose IP address has been retrieved in step 1.

```
$ ssh -i mykey oracle@IP_Address
...
[oracle@DBCS18-99699623-DB1 ~]
```

You are now connected to the VM. Observe that the UNIX prompt contains the Unix username and the database deployment name. The prompt is different for every student. Therefore, it is now displayed as `$` in all the following instructions.

10. Check the connections to ORCL CDB and verify the PDB1 PDB open status.

```
$ env | grep ORA
ORACLE_UNQNAME=ORCL
ORACLE_SID=ORCL
ORACLE_HOSTNAME=my_vm
ORACLE_BASE=/u01/app/oracle
ORACLE_HOME=/u01/app/oracle/product/18.0.0/dbhome_1
$ sqlplus / AS SYSDBA

SQL> SELECT name, cdb FROM V$DATABASE;

NAME                                CDB
-----
ORCL                                YES

SQL> SHOW PDBS

      CON_ID  CON_NAME                                OPEN MODE  RESTRICTED
-----
          2  PDB$SEED                                READ ONLY  NO
          3  PDB1                                READ WRITE NO

SQL> EXIT
$
```

Note: The `ORACLE_HOME` value is for the time being 18.0.0 and may change to 18.1.0 over time.

11. Verify that the listener is running and listening to `ORCL` and `pdb1` services. The value `xxxxxxxxxx` is assigned during the automated Database Cloud Service instance deployment before the course is started.

```
$ lsnrctl status
...
Service "ORCL.xxxxxxxxxx.oraclecloud.internal" has 1 instance(s).
  Instance "ORCL", status READY, has 1 handler(s) for this
service
...
Service "pdb1.xxxxxxxxxx.oraclecloud.internal" has 1 instance(s).
  Instance "ORCL", status READY, has 1 handler(s) for this
service...
The command completed successfully
$
```


Practices for Lesson 2:
Leveraging Multitenant Enhancements

Practices for Lesson 2: Overview

Overview

In these practices, you will discover how to monitor PDBs of different CDBs centrally from one CDB lead; create, manage, and restore carousel PDB snapshots; use the dynamic container map; handle lockdown profiles at different levels of the CDB; use static and dynamic lockdown profiles; and finally, switch over refreshable cloned PDBs.

Practice 2-1: Managing a CDB Fleet

Overview

In this practice, you learn how to create a CDB fleet with a CDB lead and its CDB members. You also learn how to monitor PDBs across all CDBs from the CDB lead.

Tasks

1. You use the Oracle By Example “Managing CDB Fleets” to see how to proceed. Launch a browser and click the bookmark from the Bookmarks toolbar of the browser. The URL is: file:///home/oracle/labs/OBE/managing_cdb_fleets/managing_cdb_fleets.html.

Practice 2-2: Managing and Using PDB Snapshots

Overview

In this practice, you enable PDBs for creating PDB snapshots and use the PDB snapshots created manually or automatically to create new PDBs, or flash back a PDB to a point in time.

Tasks

1. Before starting the practice, execute the `$HOME/labs/admin/glogin_2.sh` shell script. It sets formatting for all columns selected in queries and sets appropriate instance parameters for snapshot creation.

```
$ $HOME/labs/admin/glogin_2.sh
...
$
```

2. Create PDB_1 in ORCL and open PDB_1.

```
$ mkdir -p /u02/app/oracle/oradata/ORCL/pdb_1
$ sqlplus / AS SYSDBA

SQL> CREATE PLUGGABLE DATABASE pdb_1
      ADMIN USER pdb_1_admin IDENTIFIED BY password
      ROLES=(CONNECT)
      CREATE_FILE_DEST='/u02/app/oracle/oradata/ORCL/pdb_1';
   2       3       4

Pluggable database created.

SQL> ALTER PLUGGABLE DATABASE pdb_1 OPEN;

Pluggable database altered.

SQL>
```

Q/ Are PDBs enabled for manual snapshot creation by default?

```
SQL> SELECT pdb_name, snapshot_mode FROM cdb_pdbs;

PDB_NAME  SNAPSH
-----
PDB_1      MANUAL
PDB$SEED  MANUAL
PDB1       MANUAL

SQL>
```

A/ Yes, they are.

Q/ How many PDB snapshots can be created for a PDB?

```
SQL> SELECT property_name, property_value
      FROM database_properties
      WHERE property_name LIKE '%SNAP%';
  2      3
PROPERTY_NAME                                PROPERTY_VALUE
-----
MAX_PDB_SNAPSHOTS                           8

SQL>
```

A/ Eight, by default

Q/ Can you increase this value?

```
SQL> ALTER DATABASE SET MAX_PDB_SNAPSHOTS = 10;
ALTER DATABASE SET MAX_PDB_SNAPSHOTS = 10
*
ERROR at line 1:
ORA-65046: operation not allowed from outside a pluggable
database

SQL> CONNECT sys@PDB_1 AS SYSDBA
Enter password: password
Connected.
SQL> ALTER DATABASE SET MAX_PDB_SNAPSHOTS = 10;
ALTER DATABASE SET MAX_PDB_SNAPSHOTS = 10
*
ERROR at line 1:
ORA-65383: unable to set MAX_PDB_SNAPSHOTS property to greater
than 8

SQL>
```

A/ No. You can only decrease it. Decrease it to six.

```
SQL> ALTER DATABASE SET MAX_PDB_SNAPSHOTS = 6;

Database altered.

SQL>
```

Q/ Are there any PDB snapshots created for PDB_1?

```
SQL> SELECT con_name, snapshot_name, snapshot_scn,
           snapshot_time, full_snapshot_path
        FROM cdb_pdb_snapshots;

2      3
no rows selected

SQL>
```

A/ No. PDB_1 is enabled for manual snapshot creation.

3. Create pdb_2 in ORCL, enabled for automatic PDB snapshot creation every two minutes, and open the PDB.

```
SQL> CONNECT / AS SYSDBA
Connected.
SQL> HOST mkdir -p /u02/app/oracle/oradata/ORCL/pdb_2

SQL> CREATE PLUGGABLE DATABASE pdb_2
      ADMIN USER pdb_2_admin IDENTIFIED BY password
      ROLES=(CONNECT)
      CREATE_FILE_DEST='/u02/app/oracle/oradata/ORCL/pdb_2'
      SNAPSHOT MODE EVERY 2 MINUTES;

2      3      4      5
Pluggable database created.

SQL> ALTER PLUGGABLE DATABASE pdb_2 OPEN;

Pluggable database altered.

SQL>
```

Q/ How do you check whether the new PDB is enabled for automatic snapshot creation?

```
SQL> SELECT pdb_name, snapshot_mode FROM cdb_pdb$seeds;

PDB_NAME  SNAPSHOT_MODE
-----
PDB_1     MANUAL
PDB$SEED  MANUAL
PDB1      MANUAL
PDB_2     AUTO

SQL>
```

A/ The CDB_PDBS view contains a new column, SNAPSHOT_MODE.

4. After two minutes, you can verify that PDB snapshots are automatically created for `pdb_2`.

```
SQL> SELECT con_name, snapshot_name, snapshot_scn,
           snapshot_time, full_snapshot_path
           FROM cdb_pdb_snapshots;
 2      3
CON_NAME  SNAPSHOT_NAME          SNAPSHOT_SCN  SNAPSHOT_TIME
-----
FULL_SNAPSHOT_PATH
-----
PDB_2     SNAP_1176155135_963393302  2822136      1513934110
/u02/app/oracle/oradata/snap_1176155135_2822136.pdb

SQL>
```

Q/ What type is the PDB snapshot?

A/ The PDB snapshot is an archive file (.pdb) containing the contents of the copy of the PDB at snapshot creation.

Observe that the SCN at which the PDB snapshot was created takes part of the PDB snapshot archive file name.

5. Create PDB snapshots.
- a. Create two manual PDB snapshots for `pdb_1` with different application data. First, execute `$HOME/labs/MULTI/app1.sql` that creates a tablespace, a user, and a table and inserts rows. Then create the `pdb1_first_snap` PDB snapshot.

```
SQL> CONNECT sys@PDB_1 AS SYSDBA
```

```

Enter password: password
Connected.
SQL> @$HOME/labs/MULTI/app1.sql
...
          C LABEL
-----
          1 Label1
...
SQL>

```

```

SQL> ALTER PLUGGABLE DATABASE pdb_1 SNAPSHOT pdb1_first_snap;

Pluggable database altered.

SQL> SELECT con_name, snapshot_name, snapshot_scn,
          snapshot_time, full_snapshot_path
        FROM cdb_pdb_snapshots;

```

CON_NAME	SNAPSHOT_NAME	SNAPSHOT_SCN	SNAPSHOT_TIME	FULL_SNAPSHOT_PATH
PDB_1	PDB1_FIRST_SNAP	2823699	1513934427	/u02/app/oracle/oradata/snap_779719242_2823699.pdb

```

SQL>

```

- b. After applying, the application changes by executing
 \$HOME/labs/MULTI/app2.sql, which inserts more rows. Then create the
 pdb1_second_snap PDB snapshot.
 The errors related to the HR schema are expected to show that, in this second PDB
 snapshot, the HR schema will not be present.

```

SQL> @$HOME/labs/MULTI/app2.sql
...
SQL> INSERT INTO schema_app1.tab1 VALUES (2,'Label2');

1 row created.

SQL> CREATE TABLE schema_app1.tab2 (C number, description
VARCHAR2(40));

Table created.

```

```

SQL> INSERT INTO schema_app1.tab2 VALUES (1,'Description1');

1 row created.

SQL> COMMIT;

Commit complete.

SQL> SELECT * FROM schema_app1.tab2;

      C DESCRIPTION
-----
1 Description1

SQL> set echo off
DROP USER hr CASCADE
      *
ERROR at line 1:
ORA-01918: user 'HR' does not exist

SQL> SELECT * FROM HR.departments;
SELECT * FROM HR.departments
      *
ERROR at line 1:
ORA-00942: table or view does not exist

SQL> ALTER PLUGGABLE DATABASE pdb_1 SNAPSHOT pdb1_second_snap;

Pluggable database altered.

SQL> SELECT con_name, snapshot_name, snapshot_scn,
      snapshot_time, full_snapshot_path
      FROM cdb_pdb_snapshots;
2      3
CON_NAME  SNAPSHOT_NAME                SNAPSHOT_SCN  SNAPSHOT_TIME
-----
FULL_SNAPSHOT_PATH
-----
PDB_1      PDB1_FIRST_SNAP                2823699      1513934427
/u02/app/oracle/oradata/snap_779719242_2823699.pdb

```

```
PDB_1      PDB1_SECOND_SNAP      2824976      1513934620
/u02/app/oracle/oradata/snap_779719242_2824976.pdb
```

```
SQL>
```

- c. After creating the second snapshot, the application workload continues. Execute `$HOME/labs/MULTI/app3.sql`, which creates a new table and the HR schema.

```
SQL> @$HOME/labs/MULTI/app3.sql
```

```
SQL> CREATE TABLE schema_app1.tab3 (COL1 number, Description
VARCHAR2(10));
```

```
Table created.
```

```
SQL> INSERT INTO schema_app1.tab3 VALUES (1,'Desc1');
```

```
1 row created.
```

```
SQL> COMMIT;
```

```
Commit complete.
```

```
SQL> DROP USER hr CASCADE;
```

```
DROP USER hr CASCADE
```

```
      *
```

```
ERROR at line 1:
```

```
ORA-01918: user 'HR' does not exist
```

```
SQL> CREATE USER hr IDENTIFIED BY oracle_4U;
```

```
User created.
```

```
SQL> GRANT create session, create table, unlimited tablespace TO
hr;
```

```
Grant succeeded.
```

```
SQL> CREATE TABLE hr.departments
```

```
2      ( department_id      NUMBER(4)
```

```
3      , department_name  VARCHAR2(30)
```

```
4      CONSTRAINT dept_name_nn  NOT NULL
```

```
5      , manager_id        NUMBER(6)
```

```
6      , location_id        NUMBER(4)
```



```
7      ) ;
```

Table created.

```
SQL> INSERT INTO hr.departments VALUES
2      ( 10
3      , 'Administration'
4      , 200
5      , 1700
6      );
```

1 row created.

```
SQL>
```

```
SQL> INSERT INTO hr.departments VALUES
2      ( 20
3      , 'Marketing'
4      , 201
5      , 1800
6      );
```

1 row created.

```
SQL>
```

```
SQL> INSERT INTO hr.departments VALUES
2      ( 30
3      , 'Purchasing'
4      , 114
5      , 1700
6      );
```

1 row created.

```
SQL>
```

```
SQL> INSERT INTO hr.departments VALUES
2      ( 40
3      , 'Human Resources'
4      , 203
5      , 2400
6      );
```

1 row created.

```
SQL>
SQL> INSERT INTO hr.departments VALUES
  2      ( 50
  3      , 'Shipping'
  4      , 121
  5      , 1500
  6      );

1 row created.

SQL>
SQL> INSERT INTO hr.departments VALUES
  2      ( 60
  3      , 'IT'
  4      , 103
  5      , 1400
  6      );

1 row created.

SQL>
SQL> INSERT INTO hr.departments VALUES
  2      ( 70
  3      , 'Public Relations'
  4      , 204
  5      , 2700
  6      );

1 row created.

SQL>
SQL> INSERT INTO hr.departments VALUES
  2      ( 80
  3      , 'Sales'
  4      , 145
  5      , 2500
  6      );

1 row created.

SQL>
```

```

SQL> INSERT INTO hr.departments VALUES
2      ( 90
3      , 'Executive'
4      , 100
5      , 1700
6      );

1 row created.

SQL>
SQL> INSERT INTO hr.departments VALUES
2      ( 100
3      , 'Finance'
4      , 108
5      , 1700
6      );

1 row created.

SQL>
SQL> INSERT INTO hr.departments VALUES
2      ( 110
3      , 'Accounting'
4      , 205
5      , 1700
6      );

1 row created.

SQL>
SQL> INSERT INTO hr.departments VALUES
2      ( 120
3      , 'Treasury'
4      , NULL
5      , 1700
6      );

1 row created.

SQL> INSERT INTO hr.departments VALUES
2      ( 130
3      , 'Corporate Tax'

```

```

4          , NULL
5          , 1700
6          );

1 row created.

SQL>
SQL> INSERT INTO hr.departments VALUES
2          ( 140
3          , 'Control And Credit'
4          , NULL
5          , 1700
6          );

1 row created.

SQL>
SQL> INSERT INTO hr.departments VALUES
2          ( 150
3          , 'Shareholder Services'
4          , NULL
5          , 1700
6          );

1 row created.

SQL>
SQL> INSERT INTO hr.departments VALUES
2          ( 160
3          , 'Benefits'
4          , NULL
5          , 1700
6          );

1 row created.

SQL>
SQL> INSERT INTO hr.departments VALUES
2          ( 170
3          , 'Manufacturing'
4          , NULL
5          , 1700

```

```

6          );

1 row created.

SQL>
SQL> INSERT INTO hr.departments VALUES
2          ( 180
3            , 'Construction'
4            , NULL
5            , 1700
6            );

1 row created.

SQL>
SQL> INSERT INTO hr.departments VALUES
2          ( 270
3            , 'Payroll'
4            , NULL
5            , 1700
6            );

1 row created.

SQL>
SQL> COMMIT;

Commit complete.

SQL> SELECT * FROM HR.departments;

DEPARTMENT_ID DEPARTMENT_NAME      MANAGER_ID LOCATION_ID
-----
10 Administration      200      1700
20 Marketing           201      1800
30 Purchasing          114      1700
40 Human Resources     203      2400
50 Shipping            121      1500
60 IT                  103      1400
70 Public Relations    204      2700
80 Sales               145      2500
90 Executive           100      1700

```

```

100 Finance 108 1700
110 Accounting 205 1700
120 Treasury 1700
130 Corporate Tax 1700
140 Control And Credit 1700
150 Shareholder Services 1700
160 Benefits 1700
170 Manufacturing 1700
180 Construction 1700
270 Payroll 1700

19 rows selected.

SQL>

```

- d. Check whether PDB snapshots are still automatically created for `pdb_2`.

```

SQL> CONNECT sys@PDB_2 AS SYSDBA
Enter password: password
Connected.
SQL> SELECT con_name, snapshot_name, snapshot_scn,
           snapshot_time, full_snapshot_path
           FROM cdb_pdb_snapshots;

```

CON_NAME	SNAPSHOT_NAME	SNAPSHOT_SCN	SNAPSHOT_TIME	FULL_SNAPSHOT_PATH
PDB_2	SNAP_1176155135_963393302	2822136	1513934110	/u02/app/oracle/oradata/snap_1176155135_2822136.pdb
PDB_2	SNAP_1176155135_963393421	2822706	1513934229	/u02/app/oracle/oradata/snap_1176155135_2822706.pdb
PDB_2	SNAP_1176155135_963393541	2823167	1513934349	/u02/app/oracle/oradata/snap_1176155135_2823167.pdb
PDB_2	SNAP_1176155135_963393661	2824087	1513934473	/u02/app/oracle/oradata/snap_1176155135_2824087.pdb
PDB_2	SNAP_1176155135_963393781	2824629	1513934596	/u02/app/oracle/oradata/snap_1176155135_2824629.pdb
PDB_2	SNAP_1176155135_963393901	2825586	1513934710	

```

/u02/app/oracle/oradata/snap_1176155135_2825586.pdb

PDB_2      SNAP_1176155135_963394021      2827895      1513934827
/u02/app/oracle/oradata/snap_1176155135_2827895.pdb

7 rows selected.

SQL>

```

Q/ Why are there more PDB snapshots created than six after around 20 minutes?

A/ When you decreased the value of the maximum PDB snapshots creatable, you were logged in to PDB_1. Therefore, PDB_2 still uses the default value set at the CDB root level. You could set another value for PDB_2.

- e. After creating a new user in PDB_2, create a manual PDB snapshot for PDB_2 even if the PDB is enabled for automatic PDB snapshots.

```

SQL> CREATE USER test IDENTIFIED BY password;

```

User created.

```

SQL> ALTER PLUGGABLE DATABASE pdb_2
      SNAPSHOT pdb_2_manual_tbs_snap;

2
Pluggable database altered.

```

```

SQL> SELECT con_name, snapshot_name, snapshot_scn,
           snapshot_time, full_snapshot_path
      FROM cdb_pdb_snapshots;

2      3
CON_NAME SNAPSHOT_NAME      SNAPSHOT_SCN SNAPSHOT_TIME
-----
FULL_SNAPSHOT_PATH
-----

PDB_2      SNAP_1176155135_963393302      2822136      1513934110
/u02/app/oracle/oradata/snap_1176155135_2822136.pdb

PDB_2      SNAP_1176155135_963393421      2822706      1513934229
/u02/app/oracle/oradata/snap_1176155135_2822706.pdb

PDB_2      SNAP_1176155135_963393541      2823167      1513934349
/u02/app/oracle/oradata/snap_1176155135_2823167.pdb

```

```

PDB_2      SNAP_1176155135_963393661      2824087      1513934473
/u02/app/oracle/oradata/snap_1176155135_2824087.pdb

PDB_2      SNAP_1176155135_963393781      2824629      1513934596
/u02/app/oracle/oradata/snap_1176155135_2824629.pdb

PDB_2      SNAP_1176155135_963393901      2825586      1513934710
/u02/app/oracle/oradata/snap_1176155135_2825586.pdb

PDB_2      SNAP_1176155135_963394021      2827895      1513934827
/u02/app/oracle/oradata/snap_1176155135_2827895.pdb

PDB_2      PDB_2_MANUAL_TBS_SNAP          2828361      1513934938
/u02/app/oracle/oradata/snap_1176155135_2828361.pdb

8 rows selected.

SQL>

```

Q/ After a certain time elapses (6 minutes), can you retrieve the PDB snapshots automatically created since the beginning of the practice?

```

SQL> SELECT con_name, snapshot_name, snapshot_scn,
           snapshot_time, full_snapshot_path
        FROM cdb_pdb_snapshots;

```

CON_NAME	SNAPSHOT_NAME	SNAPSHOT_SCN	SNAPSHOT_TIME	FULL_SNAPSHOT_PATH
PDB_2	SNAP_1176155135_963393421	2822706	1513934229	/u02/app/oracle/oradata/snap_1176155135_2822706.pdb
PDB_2	SNAP_1176155135_963393541	2823167	1513934349	/u02/app/oracle/oradata/snap_1176155135_2823167.pdb
PDB_2	SNAP_1176155135_963393661	2824087	1513934473	/u02/app/oracle/oradata/snap_1176155135_2824087.pdb
PDB_2	SNAP_1176155135_963393781	2824629	1513934596	/u02/app/oracle/oradata/snap_1176155135_2824629.pdb


```

PDB_2      SNAP_1176155135_963393901      2825586      1513934710
/u02/app/oracle/oradata/snap_1176155135_2825586.pdb

PDB_2      SNAP_1176155135_963394021      2827895      1513934827
/u02/app/oracle/oradata/snap_1176155135_2827895.pdb

PDB_2      PDB_2_MANUAL_TBS_SNAP          2828361      1513934938
/u02/app/oracle/oradata/snap_1176155135_2828361.pdb

PDB_2      SNAP_1176155135_963394186      2828785      1513934995
/u02/app/oracle/oradata/snap_1176155135_2828785.pdb

8 rows selected.

SQL>

```

A/ The oldest ones were dropped to let new ones be created. It works on a FIFO basis. You can, therefore, flash back to the time or SCN of the available PDB snapshots.

6. Disable PDB snapshot creation for `pdb_2`.

```

SQL> ALTER PLUGGABLE DATABASE pdb_2 SNAPSHOT MODE NONE;

Pluggable database altered.

SQL>

```

Q/ After two minutes elapse, are there still PDB snapshots automatically created for PDB_2?

```

SQL> SELECT con_name, snapshot_name, snapshot_scn,
           snapshot_time, full_snapshot_path
        FROM cdb_pdb_snapshots;

2      3
PDB_2      SNAP_1176155135_963393541      2823167      1513934349
/u02/app/oracle/oradata/snap_1176155135_2823167.pdb

PDB_2      SNAP_1176155135_963393661      2824087      1513934473
/u02/app/oracle/oradata/snap_1176155135_2824087.pdb

PDB_2      SNAP_1176155135_963393781      2824629      1513934596
/u02/app/oracle/oradata/snap_1176155135_2824629.pdb

```

```

PDB_2      SNAP_1176155135_963393901      2825586      1513934710
/u02/app/oracle/oradata/snap_1176155135_2825586.pdb

PDB_2      SNAP_1176155135_963394021      2827895      1513934827
/u02/app/oracle/oradata/snap_1176155135_2827895.pdb

PDB_2      PDB_2_MANUAL_TBS_SNAP          2828361      1513934938
/u02/app/oracle/oradata/snap_1176155135_2828361.pdb

PDB_2      SNAP_1176155135_963394186      2828785      1513934995
/u02/app/oracle/oradata/snap_1176155135_2828785.pdb

PDB_2      SNAP_1176155135_963394261      2829474      1513935070
/u02/app/oracle/oradata/snap_1176155135_2829474.pdb

8 rows selected.

SQL>

```

A/ There could be one or two, depending on the time you took to launch the disabling command between step 5.e and step 6. You can thus flash back to the time or SCN of the available PDB snapshots.

Q/ How can you check whether the mode is now set to NONE?

```

SQL> CONNECT / AS SYSDBA
Connected.
SQL> SELECT pdb_name, snapshot_mode FROM cdb_pdb$seeds;

PDB_NAME  SNAPSHOT_MODE
-----
PDB_1     MANUAL
PDB$SEED  MANUAL
PDB1      MANUAL
PDB_2     NONE

SQL> EXIT
$

```

A/ Use the SNAPSHOT_MODE column in the CDB_PDBS.

7. A user asks you to report information about the application in PDB_1 at different times. Create PDBs from PDB snapshots.

- a. A user asks you to report information about the application in PDB_1 at pdb1_first_snap time. Create pdb_1_from_first_snap PDB from the pdb1_first_snap PDB snapshot.

```
$ mkdir -p /u02/app/oracle/oradata/ORCL/pdb_1_from_first_snap
$ sqlplus / AS SYSDBA

SQL> ALTER SESSION SET db_create_file_dest =
        '/u02/app/oracle/oradata/ORCL/pdb_1_from_first_snap';
2
Session altered.

SQL> CREATE PLUGGABLE DATABASE pdb_1_from_first_snap
        FROM pdb_1 USING SNAPSHOT pdb1_first_snap;
2
Pluggable database created.

SQL> ALTER PLUGGABLE DATABASE pdb_1_from_first_snap OPEN;

Pluggable database altered.

SQL>
```

```
SQL> CONNECT sys@pdb_1_from_first_snap AS SYSDBA
Enter password: password
Connected.
SQL> SELECT * FROM schema_app1.tab1;

          C LABEL
-----
1 Label1

SQL> SELECT * FROM schema_app1.tab2;
SELECT * FROM schema_app1.tab2
                        *

ERROR at line 1:
ORA-00942: table or view does not exist

SQL> SELECT * FROM schema_app1.tab3;
SELECT * FROM schema_app1.tab3
                        *

ERROR at line 1:
ORA-00942: table or view does not exist
```

```
SQL> SELECT * FROM hr.departments;
SELECT * FROM hr.departments
                                *

ERROR at line 1:
ORA-00942: table or view does not exist

SQL>
```

- b. The user now asks you to report information about the application in PDB_1 at pdb1_second_snap time. Create the pdb_1_from_second_snap PDB from pdb1_second_snap PDB snapshot.

```
SQL> host mkdir -p
/u02/app/oracle/oradata/ORCL/pdb_1_from_second_snap
SQL> CONNECT / AS SYSDBA
Connected.
SQL> ALTER SESSION SET db_create_file_dest =
'/u02/app/oracle/oradata/ORCL/pdb_1_from_second_snap';
2
Session altered.

SQL> CREATE PLUGGABLE DATABASE pdb_1_from_second_snap
        FROM pdb_1 USING SNAPSHOT pdb1_second_snap;
2
Pluggable database created.

SQL> ALTER PLUGGABLE DATABASE pdb_1_from_second_snap OPEN;

Pluggable database altered.

SQL> CONNECT sys@pdb_1_from_second_snap AS SYSDBA
Enter password: password
Connected.
SQL> SELECT * FROM schema_app1.tab1;

          C LABEL
-----
1 Label1
2 Label2

SQL> SELECT * FROM schema_app1.tab2;

          C DESCRIPTION
-----
```

```

1 Description1

SQL> SELECT * FROM hr.departments;
SELECT * FROM hr.departments
                        *

ERROR at line 1:
ORA-00942: table or view does not exist

SQL>
SQL> SELECT * FROM schema_app1.tab3;
SELECT * FROM schema_app1.tab3
                        *

ERROR at line 1:
ORA-00942: table or view does not exist

SQL>

```

```

SQL> CONNECT / AS SYSDBA
Connected.
SQL> SHOW PDBS

      CON_ID CON_NAME                                OPEN MODE  RESTRICTED
-----
      2 PDB$SEED                                READ ONLY  NO
      3 PDB_1                                READ WRITE NO
      4 PDB1                                READ WRITE NO
      5 PDB_2                                READ WRITE NO
      6 PDB_1_FROM_FIRST_SNAP                READ WRITE NO
      8 PDB_1_FROM_SECOND_SNAP                READ WRITE NO

SQL>

```

8. For further tests, execute `$HOME/labs/MULTI/app4.sql`. The script executes the following steps:

- Creates the `pdb1_third_snap` PDB snapshot
- Drops the `HR` schema
- Creates the `pdb1_forth_snap` PDB snapshot

```

SQL> @$HOME/labs/MULTI/app4.sql
...
SQL>

```

9. You do not need the `pdb1_first_snap` and `pdb1_second_snap` PDB snapshots. You decide to drop them.

```

SQL> CONNECT sys@PDB_1 AS SYSDBA
Enter password: password

```

```

Connected.
SQL> ALTER PLUGGABLE DATABASE pdb_1
      DROP SNAPSHOT pdb1_first_snap;

2
Pluggable database altered.

SQL> ALTER PLUGGABLE DATABASE pdb_1
      DROP SNAPSHOT pdb1_second_snap;

2
Pluggable database altered.

SQL>

```

Q/ Do the PDBs created on the dropped PDB snapshots still exist?

```

SQL> CONNECT / AS SYSDBA
Connected.
SQL> SHOW PDBS

```

CON_ID	CON_NAME	OPEN MODE	RESTRICTED
2	PDB\$SEED	READ ONLY	NO
3	PDB_1	READ WRITE	NO
4	PDB1	READ WRITE	NO
5	PDB_2	READ WRITE	NO
6	PDB_1_FROM_FIRST_SNAP	READ WRITE	NO
8	PDB_1_FROM_SECOND_SNAP	READ WRITE	NO

```

SQL>

```

```

SQL> SELECT snapshot_name, snapshot_scn,
           snapshot_time, full_snapshot_path
       FROM cdb_pdb_snapshots
       WHERE con_name = 'PDB_1';

```

SNAPSHOT_NAME	SNAPSHOT_SCN	SNAPSHOT_TIME	FULL_SNAPSHOT_PATH
PDB1_FORTH_SNAP	2835359	1513935844	/u02/app/oracle/oradata/snap_779719242_2835359.pdb
PDB1_THIRD_SNAP	2834784	1513935784	

```
/u02/app/oracle/oradata/snap_779719242_2834784.pdb
```

```
SQL>
```

A/ Yes. The PDBs created on dropped PDB snapshots are not implicitly dropped when the PDB snapshots are dropped. They are dropped with the DROP PLUGGABLE DATABASE command.

10. Imagine a situation where you are informed that a user has inadvertently dropped the HR schema in pdb_1 between the time when pdb1_third_snap and pdb1_forth_snap were created. (This is what \$HOME/labs/MULTI/app4.sql executed behind the scenes.) Recover the situation when the HR schema still existed.

- a. Close PDB_1.

```
SQL> ALTER PLUGGABLE DATABASE pdb_1 CLOSE;
```

```
Pluggable database altered.
```

```
SQL>
```

- b. Create PDB_1b from the pdb1_third_snap snapshot created before the user's error.

```
SQL> HOST mkdir /u02/app/oracle/oradata/ORCL/pdb_1b
```

```
SQL> ALTER SESSION SET  
db_create_file_dest='/u02/app/oracle/oradata/ORCL/pdb_1b';
```

```
Session altered.
```

```
SQL> CREATE PLUGGABLE DATABASE pdb_1b FROM pdb_1  
      USING SNAPSHOT pdb1_third_snap;
```

```
2      3  
Pluggable database created.
```

```
SQL> ALTER PLUGGABLE DATABASE pdb_1b OPEN;
```

```
Pluggable database altered.
```

```
SQL>
```

- c. Drop PDB_1.

```
SQL> DROP PLUGGABLE DATABASE pdb_1 INCLUDING DATAFILES;
```

```
Pluggable database dropped.
```

```
SQL>
```

- d. Rename PDB_1b to PDB_1.

```
SQL> ALTER SESSION SET CONTAINER = pdb_1b;

Session altered.

SQL> ALTER PLUGGABLE DATABASE CLOSE IMMEDIATE;

Pluggable database altered.

SQL> ALTER PLUGGABLE DATABASE OPEN RESTRICTED;

Pluggable database altered.

SQL> ALTER PLUGGABLE DATABASE RENAME GLOBAL_NAME TO pdb_1;

Pluggable database altered.

SQL> ALTER PLUGGABLE DATABASE CLOSE IMMEDIATE;

Pluggable database altered.

SQL>
```

- e. Open PDB_1.

```
SQL> ALTER PLUGGABLE DATABASE OPEN;

Pluggable database altered.

SQL>
```

- f. Verify that the HR schema now exists in PDB_1.

```
SQL> CONNECT sys@PDB_1 AS SYSDBA
Enter password: password
Connected.
SQL> SELECT * FROM hr.departments;
```

DEPARTMENT_ID	DEPARTMENT_NAME	MANAGER_ID	LOCATION_ID
10	Administration	200	1700
20	Marketing	201	1800
30	Purchasing	114	1700
40	Human Resources	203	2400
50	Shipping	121	1500
60	IT	103	1400

70	Public Relations	204	2700
80	Sales	145	2500
90	Executive	100	1700
100	Finance	108	1700
110	Accounting	205	1700
120	Treasury		1700
130	Corporate Tax		1700
140	Control And Credit		1700
150	Shareholder Services		1700
160	Benefits		1700
170	Manufacturing		1700
180	Construction		1700
270	Payroll		1700

19 rows selected.

SQL>

- g. Create a new snapshot for a potential further recovery.

```
SQL> ALTER PLUGGABLE DATABASE pdb_1 SNAPSHOT pdb1_fifth_snap;

Pluggable database altered.

SQL> EXIT
$
```

11. Execute the `$HOME/labs/MULTI/cleanup_snapshots.sh` shell script. It drops all PDB snapshots, and the `pdb_1` and `pdb_2` PDBs.

```
$ $HOME/labs/MULTI/cleanup_snapshots.sh
...
$
```

Practice 2-3: Using a Dynamic Container Map

Overview

In this practice, you will use dynamic container maps. In Oracle Database 12c, a container map enables a session connected to an application root to issue SQL statements that are routed to the appropriate PDB, depending on the value of a predicate used in the SQL statement.

Container maps enable the partitioning of data at the application PDB level when the data is not physically partitioned at the table level.

In Oracle Database 19c, a container map can be dynamically updated when new PDBs are created or dropped.

Tasks

1. Before starting the practice, execute the

`$HOME/labs/MULTI/Dynamic_container_map.sh` shell script. It creates the `HR_ROOT` application root, installs `HR_APP` in the application root, then creates the `ACCOUNTING`, `RESEARCH`, and `SALES` application PDBs. The application `HR_APP` contains two `SCOTT` tables, `DEPT` and `EMP`. There is a container map table created and set at application root `HR_ROOT` level that maps department 10 to the `ACCOUNTING` PDB, department 20 to the `RESEARCH` PDB, and department 30 to the `SALES` PDB.

```
$ $HOME/labs/MULTI/Dynamic_container_map.sh
...
$
```

2. Connect to `HR_ROOT` as `SYS` and verify the contents of the `SCOTT.DEPT` and `SCOTT.EMP` tables.

```
$ sqlplus sys@HR_ROOT AS SYSDBA

Enter password: password

SQL> SELECT deptno, dname, loc FROM scott.dept;

   DEPTNO DNAME          LOC
-----
      10 ACCOUNTING      NEW YORK
      30 SALES           CHICAGO
      20 RESEARCH        DALLAS

SQL> SELECT empno, ename, deptno FROM scott.emp;

   EMPNO ENAME          DEPTNO
-----
    7369 SMITH           20
```

```

7566 JONES                20
7788 SCOTT                20
7876 ADAMS                20
7902 FORD                 20
7782 CLARK                10
7839 KING                 10
7934 MILLER               10
7499 ALLEN                30
7521 WARD                 30
7654 MARTIN               30
7698 BLAKE                30
7844 TURNER               30
7900 JAMES                30

14 rows selected.

SQL>

```

3. Find the application PDBs associated to HR_ROOT.

```

SQL> SHOW PDBS

  CON_ID  CON_NAME                                OPEN  MODE  RESTRICTED
-----  -
          6 HR_ROOT                            READ  WRITE  NO
          7 SALES                             READ  WRITE  NO
          8 ACCOUNTING                         READ  WRITE  NO
         12 RESEARCH                           READ  WRITE  NO

SQL>

```

4. Find the container map table and the description of the container map table.

```

SQL> SELECT owner, table_name FROM dba_tables
      WHERE container_map_object='YES';

 2
OWNER          TABLE_NAME
-----
SCOTT          MAPTABLE

SQL> SELECT dbms_metadata.get_ddl('TABLE','MAPTABLE','SCOTT')
      FROM dual;

 2
DBMS_METADATA.GET_DDL('TABLE','MAPTABLE','SCOTT')
-----

CREATE TABLE "SCOTT"."MAPTABLE"

```

```

(      "DEPTNO" NUMBER,
      "NAME" VARCHAR2(30)
    ) PCTFREE 10 PCTUSED 40 INITRANS 1 MAXTRANS 255
STORAGE (
  BUFFER_POOL DEFAULT FLASH_CACHE DEFAULT CELL_FLASH_CACHE
DEFAULT)
TABLESPACE "SYSTEM"
PARTITION BY LIST ("DEPTNO")
(PARTITION "ACCOUNTING" VALUES (10) SEGMENT CREATION DEFERRED
  PCTFREE 10 PCTUSED 40 INITRANS 1 MAXTRANS 255
  NOCOMPRESS LOGGING
  STORAGE (
    BUFFER_POOL DEFAULT FLASH_CACHE DEFAULT CELL_FLASH_CACHE
DEFAULT)
    TABLESPACE "SYSTEM" ,
PARTITION "RESEARCH" VALUES (20) SEGMENT CREATION DEFERRED
  PCTFREE 10 PCTUSED 40 INITRANS 1 MAXTRANS 255
  NOCOMPRESS LOGGING
  STORAGE (
    BUFFER_POOL DEFAULT FLASH_CACHE DEFAULT CELL_FLASH_CACHE
DEFAULT)
    TABLESPACE "SYSTEM" ,
PARTITION "SALES" VALUES (30) SEGMENT CREATION DEFERRED
  PCTFREE 10 PCTUSED 40 INITRANS 1 MAXTRANS 255
  NOCOMPRESS LOGGING
  STORAGE (
    BUFFER_POOL DEFAULT FLASH_CACHE DEFAULT CELL_FLASH_CACHE
DEFAULT)
    TABLESPACE "SYSTEM" )

SQL>

```

5. Verify that the query for department 20 is routed to the appropriate PDB.

```

SQL> SET LINESIZE 180
SQL> SELECT deptno, dname, loc FROM scott.dept
      WHERE deptno = 20;

2
  DEPTNO DNAME          LOC
-----
      20 RESEARCH      DALLAS

SQL> SELECT * FROM dbms_xplan.display_cursor();

PLAN_TABLE_OUTPUT

```

```

-----
SQL_ID  fp3p8g8cxr4jm, child number 0
-----

SELECT deptno, dname, loc FROM scott.dept          WHERE deptno = 20

Plan hash value: 1918422367

-----
| Id | Operation                      | Name | Rows  | Bytes | Cost (%CPU)| Time     | Pstart |
Pstop |
-----
|  0 | SELECT STATEMENT                |      |      |      |    2 (100)|          |        |
|  1 |  PARTITION LIST SINGLE          |      |    700 | 21000 |    2 (100)| 00:00:01 |        |
2 |
|  2 |    CONTAINERS FULL              | DEPT |    700 | 21000 |    2 (100)| 00:00:01 |        |
-----

16 rows selected.

SQL>

```

The output of the execution plan is truncated to display the relevant columns only.

6. Create a new PDB for departments 60, 70, 80, 90, and 100.

```

SQL> CREATE PLUGGABLE DATABASE devt
      ADMIN USER admin IDENTIFIED BY password
      CREATE_FILE_DEST='/u02/app/oracle/oradata/ORCL/hr_root/devt'
      CONTAINER_MAP UPDATE
      (ADD PARTITION devt VALUES (60, 70, 80, 90,100));

2      3      4      5
Pluggable database created.

SQL>

```

7. Open the PDB.

```

SQL> ALTER PLUGGABLE DATABASE devt OPEN;

Pluggable database altered.

SQL>

```

8. Synchronize the application HR_APP in DEVT.

```

SQL> CONNECT sys@DEVT AS SYSDBA
Enter password: password
Connected.
SQL> ALTER PLUGGABLE DATABASE APPLICATION hr_app SYNC;

Pluggable database altered.

SQL>

```

9. Insert rows into SCOTT.DEPT and SCOTT.EMP.

```

SQL> INSERT INTO scott.DEPT VALUES (60,'DEVT','FRANCE');

1 row created.

SQL> INSERT INTO scott.DEPT VALUES (70,'DEVT','GERMANY');

1 row created.

SQL> INSERT INTO scott.DEPT VALUES (40,'Admin','Poland');
INSERT INTO scott.DEPT VALUES (40,'Admin','Poland')
      *
ERROR at line 1:
ORA-65391: violation of container map partitions

SQL> INSERT INTO scott.EMP VALUES
(7499,'ALLEN','Dev',7698,TO_date('20-2-1981','dd-mm-
yyyy'),1600,300,60);

1 row created.

SQL> INSERT INTO scott.EMP VALUES
(7521,'WARD','Dev',7698,TO_date('22-2-1981','dd-mm-
yyyy'),1250,500,70);

1 row created.

SQL> INSERT INTO scott.EMP VALUES
(7654,'MARTIN','Admin',7698,TO_date('28-9-1981','dd-mm-
yyyy'),1250,1400,40);
INSERT INTO scott.EMP VALUES
(7654,'MARTIN','Admin',7698,TO_date('28-9-1981','dd-mm-
yyyy'),1250,1400,40)
      *
ERROR at line 1:
ORA-02291: integrity constraint (SCOTT.FK_DEPTNO) violated -
parent
key not found

SQL> COMMIT;

Commit complete.

SQL>

```

Q/ Why do the third and last inserts fail?

A/ Department 40 does not refer to any PDB whose value is defined in the partitions of the LIST partitioned map table SCOTT.MAPTABLE.

10. Verify that department 60 and 70 correspond to values defined in partitions of the LIST partitioned map table SCOTT.MAPTABLE.

```
SQL> CONNECT sys@HR_ROOT AS SYSDBA
Enter password: password
Connected.
SQL> SELECT partition_name, high_value
       FROM dba_tab_partitions
       WHERE table_name='MAPTABLE' ORDER BY 1;
 2      3
PARTITION_NAME HIGH_VALUE
-----
ACCOUNTING      10
DEVT             60, 70, 80, 90, 100
RESEARCH        20
SALES           30

SQL>
```

Q/ Are the partitions automatically created in the map table?

A/ Yes. The creation of the PDB DEVT for departments 60, 70, 80, 90, and 100 defined a new partition in the LIST partitioned map table SCOTT.MAPTABLE.

11. You want to maintain departments 60 and 70 in the DEVT PDB but move the rest of the values defined in DEVT into another PDB, ADMINISTRATION.

```
SQL> CREATE PLUGGABLE DATABASE administration ADMIN USER admin
       IDENTIFIED BY password
CREATE_FILE_DEST =
       '/u02/app/oracle/oradata/ORCL/hr_root/administration'
CONTAINER_MAP UPDATE
(SPLIT PARTITION devt
 INTO (PARTITION devt values (60, 70),
      PARTITION administration));
 2  3  4  5  6  7  8
Pluggable database created.

SQL>
```

12. Open the PDB.

```
SQL> ALTER PLUGGABLE DATABASE administration OPEN;

Pluggable database altered.

SQL>
```

13. Synchronize the application HR_APP in ADMINISTRATION.

```
SQL> CONNECT sys@ADMINISTRATION AS SYSDBA
Enter password: password
Connected.
SQL> ALTER PLUGGABLE DATABASE APPLICATION hr_app SYNC;

Pluggable database altered.

SQL>
```

14. Verify that the department 60 and 70 correspond to values defined for the DEVT PDB and 80, 90, and 100 for the ADMINISTRATION PDB, and that a new partition is created in the LIST partitioned map table SCOTT.MAPTABLE.

```
SQL> CONNECT sys@HR_ROOT AS SYSDBA
Enter password: password
Connected.
SQL> SELECT partition_name, high_value
        FROM   dba_tab_partitions
        WHERE  table_name='MAPTABLE' ORDER BY 1;

  2      3
PARTITION_NAME HIGH_VALUE
-----
ACCOUNTING      10
ADMINISTRATION 80, 90, 100
DEVT            60, 70
RESEARCH        20
SALES           30

SQL> EXIT
$
```

15. Execute the \$HOME/labs/MULTI/cleanup_container_map.sh shell script. It drops the DEVT and ADMINISTRATION PDBs.

```
$ $HOME/labs/MULTI/cleanup_container_map.sh
...
$
```


Practice 2-4: Using Static and Dynamic Lockdown Profiles

Overview

In this practice, you will observe how lockdown profiles in the CDB root, in application roots, and application PDBs behave regarding inheritance. You will create and alter lockdown profiles and set the profiles at different PDB levels in a CDB. You will also discover the behavior of static and dynamic lockdown profiles created from base lockdown profiles.

Tasks

1. Before starting the practice, execute the `$HOME/labs/MULTI/HR_ROOT.sh` shell script. It creates the `HR_ROOT` application root, installs `HR_APP` in the application root, and creates the `OPERATIONS` and `SALES` application PDBs.

```
$ $HOME/labs/MULTI/HR_ROOT.sh
...
$
```

2. Create lockdown profiles at different levels in the CDB.
 - a. Create two lockdown profiles in the CDB root: `CDB_prof1` and `CDB_prof2`.

```
$ sqlplus / AS SYSDBA

SQL> CREATE LOCKDOWN PROFILE CDB_prof1;

Lockdown Profile created.

SQL> ALTER LOCKDOWN PROFILE CDB_prof1
        DISABLE STATEMENT = ('alter system');

2
Lockdown Profile altered.

SQL> ALTER LOCKDOWN PROFILE CDB_prof1
        ENABLE STATEMENT = ('alter system')
        CLAUSE = ('set');

2      3
Lockdown Profile altered.

SQL>
```

```
SQL> CREATE LOCKDOWN PROFILE CDB_prof2;

Lockdown Profile created.

SQL> ALTER LOCKDOWN PROFILE CDB_prof2
```

```

                DISABLE STATEMENT = ('alter pluggable database');

2
Lockdown Profile altered.

SQL>

```

- b. Verify the existence of lockdown profiles.

```

SQL> SELECT profile_name, rule, clause, status
        FROM   cdb_lockdown_profiles;

2
PROFILE_NAME      RULE                                CLAUSE      STATUS
-----
CDB_PROF1         ALTER SYSTEM                        DISABLE
CDB_PROF1         ALTER SYSTEM                        SET          ENABLE
CDB_PROF2         ALTER PLUGGABLE DATABASE            DISABLE
PRIVATE_DBAAS
PUBLIC_DBAAS
SAAS
                                EMPTY
                                EMPTY
                                EMPTY

6 rows selected.

SQL>

```

- c. Log in to the HR_ROOT application root as HR_LOCK_MGR to create the app_root_prof lockdown profile in HR_ROOT.

```

SQL> CONNECT hr_lock_mgr@hr_root
Enter password: password
Connected.
SQL> CREATE LOCKDOWN PROFILE app_root_prof;

Lockdown Profile created.

SQL> ALTER LOCKDOWN PROFILE app_root_prof
        DISABLE STATEMENT = ('alter system');

2
Lockdown Profile altered.

SQL> ALTER LOCKDOWN PROFILE app_root_prof
        ENABLE STATEMENT = ('alter system')
        CLAUSE = ('flush shared_pool');

2      3
Lockdown Profile altered.

SQL>

```

- d. Verify the existence of the lockdown profile in the application root.

```
SQL> SELECT profile_name, rule, clause, status
      FROM   cdb_lockdown_profiles;

2
PROFILE_NAME    RULE                CLAUSE                STATUS
-----
APP_ROOT_PROF   ALTER SYSTEM                DISABLE
APP_ROOT_PROF   ALTER SYSTEM      FLUSH SHARED_POOL  ENABLE

SQL>
```

- e. Create the `sales_prof` lockdown profile for the `SALES` application PDB in the application root.

```
SQL> CREATE LOCKDOWN PROFILE sales_prof;

Lockdown Profile created.

SQL> ALTER LOCKDOWN PROFILE sales_prof
      DISABLE STATEMENT = ('alter pluggable database');

2
Lockdown Profile altered.

SQL> SELECT profile_name, rule, clause, status
      FROM   cdb_lockdown_profiles;

2
PROFILE_NAME    RULE                CLAUSE
STATUS
-----
APP_ROOT_PROF   ALTER SYSTEM                DISABLE

APP_ROOT_PROF   ALTER SYSTEM      FLUSH SHARED_POOL  ENABLE

SALES_PROF      ALTER PLUGGABLE DATABASE  DISABLE

SQL>
```

3. Observe profile lockdown inheritance when the CDB root lockdown profile is set.
- a. Set the `CDB_prof1` lockdown profile in the CDB root.

```
SQL> CONNECT / AS SYSDBA

Connected.

SQL> SHOW PARAMETER pdb_lockdown
```

NAME	TYPE	VALUE
-----	-----	-----
pdb_lockdown	string	

SQL> **ALTER SYSTEM SET pdb_lockdown = CDB_prof1 SCOPE = both;**

System altered.

SQL> **SHOW PARAMETER pdb_lockdown**

NAME	TYPE	VALUE
-----	-----	-----
pdb_lockdown	string	CDB_PROF1

SQL>

- b. Observe how the PDB1 regular PDB inherits the restrictions of the CDB root lockdown profile.

SQL> **CONNECT sys@PDB1 as sysdba**
Enter password: *password*
Connected.

SQL> **SHOW PARAMETER pdb_lockdown**

NAME	TYPE	VALUE
-----	-----	-----
pdb_lockdown	string	CDB_PROF1

SQL> **ALTER SYSTEM SET ddl_lock_timeout=30 scope=both;**

System altered.

SQL> **ALTER SYSTEM flush shared_pool;**
ALTER SYSTEM flush shared_pool
*
ERROR at line 1:
ORA-01031: insufficient privileges

SQL> **ALTER SYSTEM CHECKPOINT;**
ALTER SYSTEM CHECKPOINT
*
ERROR at line 1:
ORA-01031: insufficient privileges

SQL>

- c. Observe how the `HR_ROOT` application root inherits the restrictions of the CDB root lockdown profile.

```
SQL> CONNECT sys@hr_root AS SYSDBA
Enter password: password
Connected.
SQL> SHOW PARAMETER pdb_lockdown
```

NAME	TYPE	VALUE
-----	-----	-----
pdb_lockdown	string	CDB_PROF1

```
SQL> ALTER SYSTEM SET ddl_lock_timeout=30 scope=both;

System altered.

SQL> ALTER SYSTEM flush shared_pool;
ALTER SYSTEM flush shared_pool
*
ERROR at line 1:
ORA-01031: insufficient privileges

SQL> ALTER SYSTEM CHECKPOINT;
ALTER SYSTEM CHECKPOINT
*
ERROR at line 1:
ORA-01031: insufficient privileges

SQL>
```

- d. Observe how the `SALES` application PDB inherits the restrictions of the CDB root lockdown profile.

```
SQL> CONNECT sys@sales AS SYSDBA
Enter password: password
Connected.
SQL> SHOW PARAMETER pdb_lockdown
```

NAME	TYPE	VALUE
-----	-----	-----
pdb_lockdown	string	CDB_PROF1

```
SQL> ALTER SYSTEM SET ddl_lock_timeout=30 scope=both;

System altered.
```

```

SQL> ALTER SYSTEM flush shared_pool;
ALTER SYSTEM flush shared_pool
*
ERROR at line 1:
ORA-01031: insufficient privileges

SQL> ALTER SYSTEM CHECKPOINT;
ALTER SYSTEM CHECKPOINT
*
ERROR at line 1:
ORA-01031: insufficient privileges

SQL>

```

4. Observe profile lockdown inheritance when the application root lockdown profile is set in the application root.
 - a. Set the `app_root_prof` lockdown profile in the application root.

```

SQL> CONNECT hr_lock_mgr@hr_root
Enter password: password
Connected.
SQL> SHOW PARAMETER pdb_lockdown

```

NAME	TYPE	VALUE
pdb_lockdown	string	CDB_PROF1

```

SQL> ALTER SYSTEM SET pdb_lockdown = app_root_prof
        SCOPE = both;
2
System altered.

SQL> SHOW PARAMETER pdb_lockdown

```

NAME	TYPE	VALUE
pdb_lockdown	string	APP_ROOT_PROF

```

SQL>

```

- b. Observe how the `HR_ROOT` application root inherits the restrictions of the CDB root lockdown profile.

```

SQL> CONNECT system@hr_root
Enter password: password
Connected.
SQL> SHOW PARAMETER pdb_lockdown

```

NAME	TYPE	VALUE
-----	-----	-----
pdb_lockdown	string	APP_ROOT_PROF

```

SQL> ALTER SYSTEM SET ddl_lock_timeout=30 scope=both;

System altered.

SQL> ALTER SYSTEM flush shared_pool;
ALTER SYSTEM flush shared_pool
*
ERROR at line 1:
ORA-01031: insufficient privileges

SQL> ALTER SYSTEM CHECKPOINT;
ALTER SYSTEM CHECKPOINT
*
ERROR at line 1:
ORA-01031: insufficient privileges

SQL>

```

Observe that `app_root_prof` affects all application PDBs in the application container, but the application root still inherits the rules of its nearest ancestor, `CDB_prof1`.

- c. Observe how the `SALES` application PDB inherits the restrictions of the application root lockdown profile.

```

SQL> CONNECT sys@sales AS SYSDBA
Enter password: password
Connected.

SQL> ALTER SYSTEM SET ddl_lock_timeout=30 scope=both;
ALTER SYSTEM SET ddl_lock_timeout=30 scope=both
*
ERROR at line 1:
ORA-01031: insufficient privileges

SQL> ALTER SYSTEM flush shared_pool;
ALTER SYSTEM flush shared_pool
*

```

```

ERROR at line 1:
ORA-01031: insufficient privileges

SQL> ALTER SYSTEM CHECKPOINT;
ALTER SYSTEM CHECKPOINT
*
ERROR at line 1:
ORA-01031: insufficient privileges

SQL>

```

The `SALES` application PDB inherits the restrictions of the application root lockdown profile in addition to those set in its nearest ancestor, `CDB_prof1`.

5. Observe profile lockdown inheritance when the application root lockdown profile is set in the application PDB.
 - a. Set the `sales_prof` lockdown profile in the `SALES` application PDB.

```

SQL> CONNECT hr_lock_mgr@sales
Enter password: password
Connected.
SQL> SHOW PARAMETER pdb_lockdown

NAME                                TYPE        VALUE
-----                                -
pdb_lockdown                        string      APP_ROOT_PROF
SQL> ALTER SYSTEM SET pdb_lockdown = '' SCOPE = both;
ALTER SYSTEM SET pdb_lockdown = '' SCOPE = both
*
ERROR at line 1:
ORA-01031: insufficient privileges

SQL>

```

- b. Temporarily unset `app_root_prof` so as to be able to set the `sales_prof` lockdown profile in the `SALES` application PDB. Then reset `app_root_prof`.

```

SQL> CONNECT hr_lock_mgr@hr_root
Enter password: password
Connected.
SQL> ALTER SYSTEM SET pdb_lockdown = '' SCOPE = both;

```



```
System altered.
```

```
SQL>
```

- c. Set the `sales_prof` lockdown profile in the `SALES` application PDB.

```
SQL> CONNECT hr_lock_mgr@sales
```

```
Enter password: password
```

```
Connected.
```

```
SQL> ALTER SYSTEM SET pdb_lockdown = sales_prof SCOPE = both;
```

```
System altered.
```

```
SQL> SHOW PARAMETER pdb_lockdown
```

NAME	TYPE	VALUE
-----	-----	-----
pdb_lockdown	string	SALES_PROF

```
SQL>
```

- d. Reset `app_root_prof` in the application root.

```
SQL> CONNECT hr_lock_mgr@hr_root
```

```
Enter password: password
```

```
Connected.
```

```
SQL> ALTER SYSTEM SET pdb_lockdown = app_root_prof  
SCOPE = both;
```

```
2
```

```
System altered.
```

```
SQL>
```

- e. Observe how the `SALES` application PDB still inherits the restrictions of the application root lockdown profile and the CDB root lockdown profile.

```
SQL> CONNECT sys@sales AS SYSDBA
```

```
Enter password: password
```

```
Connected.
```

```
SQL> ALTER PLUGGABLE DATABASE sales CLOSE;
```

```
ALTER PLUGGABLE DATABASE sales CLOSE
```

```
*
```

```
ERROR at line 1:
```

```
ORA-01031: insufficient privileges
```

```
SQL> ALTER SYSTEM SET ddl_lock_timeout=30 scope=both;
```

```

System altered.

SQL> ALTER SYSTEM flush shared_pool;
ALTER SYSTEM flush shared_pool
*
ERROR at line 1:
ORA-01031: insufficient privileges

SQL> ALTER SYSTEM CHECKPOINT;
ALTER SYSTEM CHECKPOINT
*
ERROR at line 1:
ORA-01031: insufficient privileges

SQL>

```

Because a lockdown profile is set in the application PDB, the application PDB inherits the restrictions of the application PDB lockdown profile and those of the CDB root lockdown profile and not from the application root lockdown profile.

6. Drop lockdown profiles.

- a. Drop the CDB_prof1 and CDB_prof2 lockdown profiles in the CDB root.

```

SQL> CONNECT / AS SYSDBA
Connected.
SQL> ALTER SYSTEM SET pdb_lockdown = '' SCOPE = both;

System altered.

SQL> DROP LOCKDOWN PROFILE CDB_prof1;

Lockdown Profile dropped.

SQL> DROP LOCKDOWN PROFILE CDB_prof2;

Lockdown Profile dropped.

SQL>

```

- b. Drop the `sales_prof` lockdown profile in the `SALES` application PDB.

```
SQL> CONNECT sys@sales AS SYSDBA
Enter password: password
Connected.
SQL> ALTER SYSTEM SET pdb_lockdown = '' SCOPE = both;

System altered.

SQL> CONNECT hr_lock_mgr@hr_root
Enter password: password
Connected.
SQL> DROP LOCKDOWN PROFILE sales_prof;

Lockdown Profile dropped.

SQL>
```

- c. Update the existing `app_root_prof` to remove `ALTER SYSTEM SET` restriction.

```
SQL> ALTER LOCKDOWN PROFILE app_root_prof
      ENABLE STATEMENT = ('alter system')
      CLAUSE = ('set');

2      3
Lockdown Profile altered.

SQL>
```

7. In the second part of this practice, you will create static and dynamic lockdown profiles. You create and set, respectively, a static lockdown profile in the `SALES` application PDB from a basic application root lockdown profile, and a dynamic lockdown profile in the `OPERATIONS` application PDB from the same basic application root lockdown profile.

- a. Log in to the application root to create the static application root lockdown profile `static_sales_prof` from the `app_root_prof` lockdown profile for the `SALES` application PDB.

```
SQL> CREATE LOCKDOWN PROFILE static_sales_prof
      FROM app_root_prof;

2
Lockdown Profile created.

SQL> ALTER LOCKDOWN PROFILE static_sales_prof
      DISABLE STATEMENT = ('alter pluggable database');

2
Lockdown Profile altered.
```

```

SQL> SELECT profile_name, rule, clause, status
      FROM   cdb_lockdown_profiles;

  2
PROFILE_NAME          RULE                                CLAUSE
-----
STATUS
-----
APP_ROOT_PROF        ALTER SYSTEM
DISABLE

APP_ROOT_PROF        ALTER SYSTEM                                FLUSH SHARED_POOL
ENABLE

APP_ROOT_PROF        ALTER SYSTEM                                SET
ENABLE

STATIC_SALES_PROF    ALTER PLUGGABLE DATABASE
DISABLE

STATIC_SALES_PROF    ALTER SYSTEM
DISABLE

STATIC_SALES_PROF    ALTER SYSTEM                                FLUSH SHARED_POOL
ENABLE

STATIC_SALES_PROF    ALTER SYSTEM                                SET
ENABLE

7 rows selected.

SQL>

```

Q/ What do you observe about `static_sales_prof` in the `CDB_LOCKDOWN_PROFILES` view?

A/ The rules from the basic `app_root_prof` lockdown profile are all copied into the `static_sales_prof` lockdown profile.

- b. Log in to the `SALES` application PDB to set the static `static_sales_prof` lockdown profile.

```

SQL> CONNECT hr_lock_mgr@sales
Enter password: password

```

```

Connected.
SQL> ALTER SYSTEM SET pdb_lockdown = '' SCOPE = both;

System altered.

SQL> ALTER SYSTEM SET pdb_lockdown = static_sales_prof
        SCOPE = both;

2
System altered.

SQL> SHOW PARAMETER pdb_lockdown

```

NAME	TYPE	VALUE
pdb_lockdown	string	STATIC_SALES_PROF

```

SQL>

```

- c. Before creating the dynamic lockdown profile in the `OPERATIONS` application PDB, log in to the application PDB to verify that you can still create a partitioned table because the restriction rule of the dynamic lockdown profile will be the partitioning feature.

```

SQL> CONNECT hr_lock_mgr@operations
Enter password: password
Connected.
SQL> SHOW PARAMETER pdb_lockdown

```

NAME	TYPE	VALUE
pdb_lockdown	string	APP_ROOT_PROF

```

SQL>
SQL> CREATE TABLE sales ( SALESMAN_ID NUMBER(5),
        SALESMAN_NAME VARCHAR2(30), SALES_STATE VARCHAR2(20))
        PARTITION BY LIST (SALES_STATE) AUTOMATIC
        (PARTITION P_CAL VALUES ('CALIFORNIA'));

2      3      4
Table created.

SQL> DROP TABLE sales;

Table dropped.

SQL>

```

- d. Log in to the application root to create the dynamic application root lockdown profile `dynamic_op_prof` from `app_root_prof` for the `OPERATIONS` application PDB.

```
SQL> CONNECT hr_lock_mgr@hr_root
Enter password: password
Connected.
SQL> CREATE LOCKDOWN PROFILE dynamic_op_prof
      INCLUDING app_root_prof;
2
Lockdown Profile created.

SQL> ALTER LOCKDOWN PROFILE dynamic_op_prof
      DISABLE OPTION = ('PARTITIONING');
2
Lockdown Profile altered.

SQL> SELECT profile_name, rule, clause, status
      FROM cdb_lockdown_profiles;
2
```

PROFILE_NAME	RULE	CLAUSE

STATUS		

APP_ROOT_PROF	ALTER SYSTEM	
DISABLE		
APP_ROOT_PROF	ALTER SYSTEM	FLUSH SHARED_POOL
ENABLE		
APP_ROOT_PROF	ALTER SYSTEM	SET
ENABLE		
DYNAMIC_OP_PROF	PARTITIONING	
DISABLE		
STATIC_SALES_PROF	ALTER PLUGGABLE DATABASE	
DISABLE		
STATIC_SALES_PROF	ALTER SYSTEM	
DISABLE		
STATIC_SALES_PROF	ALTER SYSTEM	FLUSH SHARED_POOL
ENABLE		

```

STATIC_SALES_PROF ALTER SYSTEM SET
ENABLE

8 rows selected.

SQL>

```

Q/ What do you observe about `dynamic_op_prof` in the `CDB_LOCKDOWN_PROFILES` view?

A/ The rules from the basic `app_root_prof` lockdown profile are not copied into the `dynamic_op_prof` lockdown profile. You will check whether they are in effect.

- e. Log in to the `OPERATIONS` application PDB to set the dynamic `dynamic_op_prof` lockdown profile.

```

SQL> CONNECT hr_lock_mgr@operations
Enter password: password
Connected.
SQL> ALTER SYSTEM SET pdb_lockdown = '' SCOPE = both;

System altered.

SQL> ALTER SYSTEM SET pdb_lockdown = dynamic_op_prof
SCOPE = both;

2
System altered.

SQL> SHOW PARAMETER pdb_lockdown

```

NAME	TYPE	VALUE
pdb_lockdown	string	DYNAMIC_OP_PROF

```

SQL>

```

8. Test the behavior of static and dynamic lockdown profiles set in the `SALES` and `OPERATIONS` application PDBs, respectively.

```

SQL> CONNECT sys@sales AS SYSDBA
Enter password: password
Connected.
SQL> ALTER SYSTEM SET ddl_lock_timeout=30 scope=both;

System altered.

```

```

SQL> ALTER SYSTEM flush shared_pool;

System altered.

SQL> ALTER SYSTEM CHECKPOINT;
ALTER SYSTEM CHECKPOINT
*
ERROR at line 1:
ORA-01031: insufficient privileges

SQL> ALTER PLUGGABLE DATABASE CLOSE;
ALTER PLUGGABLE DATABASE CLOSE
*
ERROR at line 1:
ORA-01031: insufficient privileges

SQL> CREATE TABLE sales
      ( SALESMAN_ID  NUMBER(5), SALESMAN_NAME VARCHAR2(30),
        SALES_STATE VARCHAR2(20))
      PARTITION BY LIST (SALES_STATE) AUTOMATIC
        (PARTITION P_CAL VALUES ('CALIFORNIA'));
      2      3      4      5
Table created.

SQL>

```

The first, second, and third commands respect rules inherited from `app_root_prof`.
The fourth command respects rules inherited from `static_sales_prof`.
There is no restriction in `static_sales_prof` for the fifth command to execute.

9. Test in the `OPERATIONS` application PDB if the restriction rules of the dynamic `dynamic_op_prof` lockdown profile using the values from the existing basic `app_root_prof` lockdown profile apply.

```

SQL> CONNECT sys@operations AS SYSDBA
Enter password: password
Connected.
SQL> ALTER SYSTEM SET ddl_lock_timeout=30 scope=both;

System altered.

```



```

SQL> ALTER SYSTEM flush shared_pool;

System altered.

SQL> ALTER SYSTEM CHECKPOINT;
ALTER SYSTEM CHECKPOINT
*
ERROR at line 1:
ORA-01031: insufficient privileges

SQL> ALTER PLUGGABLE DATABASE CLOSE;

Pluggable database altered.

SQL> ALTER PLUGGABLE DATABASE OPEN;

Pluggable database altered.

SQL> CREATE TABLE sales
      ( SALESMAN_ID NUMBER(5), SALESMAN_NAME VARCHAR2(30),
        SALES_STATE VARCHAR2(20))
      PARTITION BY LIST (SALES_STATE) AUTOMATIC
      (PARTITION P_CAL VALUES ('CALIFORNIA'));
      2      3      4      5
CREATE TABLE sales
*
ERROR at line 1:
ORA-00439: feature not enabled: Partitioning

SQL>

```

The first, second, and third commands respect rules inherited from `app_root_prof`. There is no restriction in `dynamic_op_prof` for the fourth command to execute. The fifth command respects rules inherited from `dynamic_op_prof`.

10. Test how any subsequent changes to the existing basic `app_root_prof` lockdown profile impact the derived `static_sales_prof` and `dynamic_op_prof` lockdown profiles. Add a rule to `app_root_prof`.

```

SQL> CONNECT hr_lock_mgr@hr_root
Enter password: password

```

Connected.

```
SQL> ALTER LOCKDOWN PROFILE App_root_prof
      ENABLE STATEMENT = ('alter system')
      CLAUSE = ('checkpoint');
```

2 3

Lockdown Profile altered.

```
SQL> SELECT profile_name, rule, clause, status
      FROM cdb_lockdown_profiles;
```

2

PROFILE_NAME	RULE	CLAUSE
APP_ROOT_PROF	ALTER SYSTEM	
APP_ROOT_PROF	ALTER SYSTEM	FLUSH SHARED_POOL
APP_ROOT_PROF	ALTER SYSTEM	SET
APP_ROOT_PROF	ALTER SYSTEM	CHECKPOINT
DYNAMIC_OP_PROF	PARTITIONING	
STATIC_SALES_PROF	ALTER PLUGGABLE DATABASE	
STATIC_SALES_PROF	ALTER SYSTEM	
STATIC_SALES_PROF	ALTER SYSTEM	FLUSH SHARED_POOL
STATIC_SALES_PROF	ALTER SYSTEM	SET

```
9 rows selected.
```

```
SQL>
```

11. Check whether the rule added to the existing basic `app_root_prof` lockdown profile does not impact the derived `static_sales_prof`.

```
SQL> CONNECT sys@sales AS SYSDBA
Enter password: password
Connected.
SQL> ALTER SYSTEM CHECKPOINT;
ALTER SYSTEM CHECKPOINT
*
ERROR at line 1:
ORA-01031: insufficient privileges

SQL>
```

12. Check whether the rule added to the existing basic `app_root_prof` lockdown profile impacts the derived `dynamic_op_prof`.

```
SQL> CONNECT sys@operations AS SYSDBA
Enter password: password
Connected.
SQL> ALTER SYSTEM CHECKPOINT;

System altered.

SQL> CREATE TABLE sales
      ( SALESMAN_ID NUMBER(5), SALESMAN_NAME VARCHAR2(30),
        SALES_STATE VARCHAR2(20))
      PARTITION BY LIST (SALES_STATE) AUTOMATIC
      (PARTITION P_CAL VALUES ('CALIFORNIA'));
2      3      4      5
CREATE TABLE sales
*
ERROR at line 1:
ORA-00439: feature not enabled: Partitioning

SQL> EXIT
$
```

13. Execute the `$HOME/labs/MULTI/cleanup_profiles.sh` shell script. It drops the `HR_ROOT` application root PDB and its associated application PDBs—`SALES` and `OPERATIONS`—and therefore, all lockdown profiles created in these PDBs.

```
$ $HOME/labs/MULTI/cleanup_profiles.sh
...
$
```

Practice 2-5: Switching Over Refreshable Cloned PDBs

Overview

In this practice, you reverse the roles of a PDB and its refreshable cloned PDB. The refreshable cloned PDB can be made the primary PDB while the primary PDB becomes the refreshable cloned PDB.

Tasks

1. Before starting the practice, execute the `$HOME/labs/admin/glogin_2b.sh` shell script. It sets formatting for all columns selected in queries.

```
$ $HOME/labs/admin/glogin_2b.sh
...
$
```

2. Configure the primary CDB and PDB. The session in `ORCL` is maintained and called *Primary Session*.
 - a. Execute the `$HOME/labs/MULTI/create_PDB1.sql` SQL script. It re-creates `PDB1` using Oracle Managed Files (OMF) and drops `PDB_REFRESHED` if it exists. Then grant the `SYSOPER` privilege required to create refreshable cloned PDBs. Then create the database link to connect to the future refreshable PDB in `ORCL`.

```
$ sqlplus / AS SYSDBA

SQL> @$HOME/labs/MULTI/create_PDB1.sql
...
SQL>
```

- b. Create a common user in `ORCL` and grant the common user the `SYSOPER` system privilege and others.

```
SQL> DROP USER c##u1 CASCADE;
DROP USER c##u1 CASCADE
          *
ERROR at line 1:
ORA-01918: user 'C##U1' does not exist

SQL> CREATE USER c##u1 IDENTIFIED BY password;

User created.

SQL> GRANT create session, resource, create any table,
      unlimited tablespace, sysoper, create pluggable database
      TO c##u1 CONTAINER = ALL;

 2      3
Grant succeeded.

SQL>
```

- c. Create in ORCL a public fixed-user database link relying on the common user created in the previous step.

```
SQL> DROP PUBLIC DATABASE LINK link_ORCL;
DROP PUBLIC DATABASE LINK link_ORCL
                                *
ERROR at line 1:
ORA-02024: database link not found

SQL> CREATE PUBLIC DATABASE LINK link_ORCL
      CONNECT TO c##u1 IDENTIFIED BY password
      USING 'ORCL';
      2      3
Database link created.

SQL>
```

- d. In PDB1, execute the \$HOME/labs/MULTI/create_userswitch_tab.sql SQL script that creates the USERSWITCH.BIGTAB table with 10000 rows.

```
SQL> CONNECT system@PDB1
Enter password: password
Connected.
SQL> SET sqlprompt "SQL_prim> "
SQL_prim> @$HOME/labs/MULTI/create_userswitch_tab.sql
...
SQL_prim>
```

- e. Count the rows and display the data in the rows.

```
SQL_prim> SELECT count(*) FROM userswitch.bigtab;

COUNT (*)
-----
10000

SQL_prim> SELECT DISTINCT label FROM userswitch.bigtab;

LABEL
-----
DATA FROM PDB1

SQL_prim>
```

3. Open another terminal window called *Switch Session* where you log in to ORCL and create the refreshable cloned PDB `PDB_REFRESHED`. Create `PDB_REFRESHED` automatically refreshed from `PDB1` every two minutes.

```
$ mkdir -p /u02/app/oracle/oradata/ORCL/pdb_refreshed
$ sqlplus / AS SYSDBA

SQL> SET sqlprompt "SQL_switch> "
SQL_switch> CREATE PLUGGABLE DATABASE pdb_refreshed
                FROM pdb1@link_ORCL
                CREATE_FILE_DEST =
                '/u02/app/oracle/oradata/ORCL/pdb_refreshed'
                REFRESH MODE every 2 minutes;

      2      3      4      5      6
Pluggable database created.

SQL_switch>
```

- a. Check the status of the refreshable PDB.

```
SQL_switch> SELECT pdb_name, status FROM cdb_pdb$seeds;

PDB_NAME          STATUS
-----
PDB1              NORMAL
PDB$SEED          NORMAL
PDB_REFRESHED     REFRESHING

SQL_switch>
```

4. In *Primary Session*, start and commit a transaction in the production `PDB1`.

```
SQL_prim> INSERT INTO userswitch.bigtab
                VALUES ('New DATA FROM PDB1');

      2
1 row created.

SQL_prim> COMMIT;

Commit complete.

SQL_prim> SELECT DISTINCT label FROM userswitch.bigtab;

LABEL
-----
DATA FROM PDB1
New DATA FROM PDB1

SQL_prim>
```

5. In *Switch Session*, wait two minutes before you can observe that the data in PDB_REFRESHED is refreshed.

```
SQL_switch> ALTER SESSION SET container = pdb_refreshed;

Session altered.

SQL_switch> ALTER PLUGGABLE DATABASE OPEN READ ONLY;

Pluggable database altered.

SQL_switch> SELECT DISTINCT label FROM userswitch.bigtab;

LABEL
-----
DATA FROM PDB1
New DATA FROM PDB1

SQL_switch> SELECT count(*) FROM userswitch.bigtab;

COUNT(*)
-----
10001

SQL_switch>
```

6. Switch over the refreshable cloned PDB.

- a. In *Primary Session*, switch the primary PDB1 PDB over the refreshable clone PDB_REFRESHED PDB.

```
SQL_prim> CONNECT sys@ORCL AS SYSDBA
Enter password: password
Connected.
SQL_prim> ALTER SESSION SET CONTAINER = pdb1;

Session altered.

SQL_prim> ALTER PLUGGABLE DATABASE
            REFRESH MODE EVERY 2 MINUTES
            FROM pdb_refreshed@link_ORCL SWITCHOVER;

2      3
Pluggable database altered.

SQL_prim>
```


- b. Verify the status of the refreshable PDB1.

```
SQL_prim> ALTER SESSION SET CONTAINER = CDB$ROOT;

Session altered.

SQL_prim> SELECT pdb_name, status FROM cdb_pdbs;

PDB_NAME          STATUS
-----
PDB$SEED          NORMAL
PDB_REFRESHED     NORMAL
PDB1             REFRESHING

SQL_prim>
```

- c. In *Switch Session*, start and commit a transaction in the primary PDB_REFRESHED.

```
SQL_switch> ALTER SESSION SET CONTAINER = pdb_refreshed;

Session altered.

SQL_switch> INSERT INTO userswitch.bigtab
              VALUES ('New DATA FROM pdb_refreshed');
           2
1 row created.

SQL_switch> COMMIT;

Commit complete.

SQL_switch> SELECT DISTINCT label FROM userswitch.bigtab;

LABEL
-----
New DATA FROM pdb_refreshed
DATA FROM PDB1
New DATA FROM PDB1

SQL_switch>
```

- d. Back in *Primary Session*, wait two minutes before you can observe that the data in PDB1 is refreshed.

```
SQL_prim> ALTER SESSION SET CONTAINER = pdb1;

Session altered.

SQL_prim> ALTER PLUGGABLE DATABASE OPEN READ ONLY;
```

```
Pluggable database altered.
```

```
SQL_prim> SELECT DISTINCT label FROM userswitch.bigtab;
```

```
LABEL
```

```
-----
```

```
New DATA FROM pdb_refreshed
```

```
DATA FROM PDB1
```

```
New DATA FROM PDB1
```

```
SQL_prim> SELECT count(*) FROM userswitch.bigtab;
```

```
COUNT(*)
```

```
-----
```

```
10002
```

```
SQL_prim>
```

- e. Do not forget to close the refreshable PDB to be refreshed.

```
SQL_prim> ALTER PLUGGABLE DATABASE CLOSE;
```

```
Pluggable database altered.
```

```
SQL_prim>
```

7. Re-switch. You now observe that you can switch over back and forth.

- a. Before you switch the primary PDB_REFRESHED PDB over the refreshable cloned PDB1, update data in the USERSWITCH.BIGTAB table in the primary PDB. Proceed in the *Switch Session*.

```
SQL_switch> UPDATE userswitch.bigtab
```

```
SET label = 'DATA updated IN pdb_refreshed';
```

```
2
```

```
10002 rows updated.
```

```
SQL_switch> COMMIT;
```

```
Commit complete.
```

```
SQL_switch> SELECT DISTINCT label FROM userswitch.bigtab;
```

```
LABEL
```

```
-----
```

```
DATA updated IN pdb_refreshed
```

```
SQL_switch>
```

- b. In the *Primary Session*, wait two minutes before you can observe that the data in PDB1 is refreshed.

```
SQL_prim> ALTER PLUGGABLE DATABASE OPEN READ ONLY;

Pluggable database altered.

SQL_prim> SELECT DISTINCT label FROM userswitch.bigtab;

LABEL
-----
DATA updated IN pdb_refreshed

SQL_prim> EXIT
$
```

- c. From the *Switch Session*, switch the primary PDB_REFRESHED PDB over the refreshable cloned PDB1 PDB.

```
SQL_switch> ALTER SESSION SET CONTAINER = pdb_refreshed;

Session altered.

SQL_switch> ALTER PLUGGABLE DATABASE
              REFRESH MODE EVERY 2 MINUTES
              FROM pdb1@link_ORCL SWITCHOVER;

  2      3
Pluggable database altered.

SQL_switch>
```

- d. Verify the status of the refreshable PDB_REFRESHED.

```
SQL_switch> ALTER SESSION SET CONTAINER = CDB$ROOT;

Session altered.

SQL_switch> SELECT pdb_name, status FROM cdb_pdbs;

PDB_NAME          STATUS
-----
PDB$SEED          NORMAL
PDB_REFRESHED    REFRESHING
PDB1              NORMAL

SQL_switch> EXIT
$
```

8. Execute the `$HOME/labs/MULTI/cleanup_PDB_REFRESHED.sh` shell script to drop PDB_REFRESHED in ORCL.

```
$ $HOME/labs/MULTI/cleanup_PDB_REFRESHED.sh
...
$
```

Practices for Lesson 3: Managing Security

Practices for Lesson 3: Overview

Overview

In these practices, you will discover how to create schema-only accounts, manage PDB keystores, create user-defined TDE master keys, export and import fixed-user database links without the database links password, and finally encrypt sensitive data in Database Replay files.

Practice 3-1: Creating Schema-Only Accounts

Overview

In this practice, you will create a user with no authentication. Application designers may want to create a user account that contains the application data dictionary, but should not be allowed to log in. This can be used to enforce data access through the application, separation of duties at the application level, and other security mechanisms.

Tasks

1. Before starting the practice, execute the `$HOME/labs/admin/glogin_3.sh` shell script. It sets formatting for all columns selected in queries.

```
$ $HOME/labs/admin/glogin_3.sh
...
$
```

2. Log in to PDB1 as SYSTEM to create the SCHEMA_NOAUTH schema-only account. If PDB1 does not exist, execute the `$HOME/labs/admin/PDB1.sh` shell script.

```
$ sqlplus system@PDB1

Enter password: password

SQL> CREATE USER schema_noauth NO AUTHENTICATION;

User created.

SQL>
```

3. Grant CREATE TABLE and UNLIMITED TABLESPACE to SCHEMA_NOAUTH.

```
SQL> GRANT create table, unlimited tablespace
      TO schema_noauth;

2
Grant succeeded.

SQL>
```

4. Verify that a schema-only account with sufficient privileges can create objects in the schema but cannot connect under the schema-only account directly.

```
SQL> CREATE TABLE schema_noauth.tab1 (c NUMBER);

Table created.

SQL>
```

```
SQL> CONNECT schema_noauth@PDB1

Enter password:
```



```
ERROR:
ORA-01005: null password given; logon denied

Warning: You are no longer connected to ORACLE.
SQL>
```

The schema-only account does not have any authentication, neither basic by password, nor strong by global authentication and, therefore, cannot connect.

5. List the schema-only accounts.

```
SQL> CONNECT system@PDB1
Enter password: password
Connected.
SQL> SELECT username, authentication_type FROM dba_users
        ORDER BY 2;

 2
-----
USERNAME                                AUTHENTI
-----
SCHEMA_NOAUTH                           NONE
OJVMSYS                                 NONE
DVSYS                                   NONE
LBACSYS                                 NONE
DVF                                     NONE
AUDSYS                                 NONE
REMOTE_SCHEDULER_AGENT                  PASSWORD
GSMUSER                                 PASSWORD
...
ORDSYS                                 PASSWORD
SYS                                     PASSWORD
SYSTEM                                 PASSWORD

39 rows selected.

SQL>
```

You can observe that some Oracle-supplied schema accounts are also schema-only accounts like Database Vault `DVSYS` and `DVF`, Oracle Label Security `LBACSYS` and Unified Auditing `AUDSYS`.

Q/ Can you change the authentication type of `SCHEMA_NOAUTH`?

```
SQL> ALTER USER schema_noauth IDENTIFIED BY password;

User altered.

SQL>
```

A/ Yes.

6. Verify that the user can now log in to the database.

```
SQL> CONNECT schema_noauth@PDB1
Enter password: password
ERROR:
ORA-01045: user SCHEMA_NOAUTH lacks CREATE SESSION privilege;
logon denied

Warning: You are no longer connected to ORACLE.
SQL>
```

7. Grant the user the CREATE SESSION privilege.

```
SQL> CONNECT system@PDB1
Enter password: password
Connected.
SQL> GRANT create session TO schema_noauth;

Grant succeeded.

SQL> CONNECT schema_noauth@PDB1
Enter password: password
Connected.
SQL>
```

8. Drop the SCHEMA_NOAUTH schema-only account.

```
SQL> CONNECT system@PDB1
Enter password: password
Connected.
SQL> DROP USER schema_noauth CASCADE;

User dropped.

SQL> EXIT
$
```

Practice 3-2: Managing PDB Keystores

Overview

In this practice, you will learn about the keystore isolated and united modes, and the reason for introducing these modes, and set up isolated mode PDB keystores.

Tasks

1. Execute the `$HOME/labs/admin/PDB1.sh` shell script to recreate PDB1.

```
$ $HOME/labs/admin/PDB1.sh
...
$
```

2. You plan to isolate PDB keystores. To use isolated mode, the `WALLET_ROOT` initialization parameter must be set. Configure the root directory location of the CDB keystore and the keystore type.
 - a. Configure the wallet root location and restart the database instance.

```
$ sqlplus / AS SYSDBA

SQL> SHOW PARAMETER wallet

NAME                                TYPE        VALUE
-----                                -
wallet_root                          string

SQL> ALTER SYSTEM SET
      wallet_root = '/u01/app/oracle/admin/ORCL/tde_keystore'
      SCOPE=SPFILE;

      2      3
System altered.

SQL>
```

```
SQL> SHUTDOWN IMMEDIATE
Database closed.
Database dismounted.
ORACLE instance shut down.
SQL> STARTUP
ORACLE instance started.

Total System Global Area 1426060816 bytes
Fixed Size                  8893968 bytes
Variable Size               570425344 bytes
```

```

Database Buffers          838860800 bytes
Redo Buffers              7880704 bytes
Database mounted.
Database opened.
SQL> ALTER PLUGGABLE DATABASE all OPEN;

Pluggable database altered.

SQL> SHOW PARAMETER wallet_root

NAME                      TYPE      VALUE
-----
wallet_root              string    /u01/app/oracle/admin/ORCL/tde_keystore
SQL>

```

- b. Configure the default keystore type for any future isolated PDB. In this case, the isolated keystores will be stored as OS files.

```

SQL> ALTER SYSTEM SET tde_configuration =
                        'KEYSTORE_CONFIGURATION=FILE'
                        SCOPE=BOTH;

      2      3
System altered.

SQL>

```

KEYSTORE_CONFIGURATION can take the values FILE, OKV, HSM, FILE|OKV, FILE|HSM, OKV|FILE, and HSM|FILE to configure the type of isolated keystore for each PDB.

```

SQL> SELECT wrl_parameter, status, keystore_mode
        FROM v$encryption_wallet;

      2
WRL_PARAMETER                                STATUS
KEYSTORE
-----
/u01/app/oracle/admin/ORCL/tde_keystore/tde/ NOT_AVAILABLE  NONE
UNITED                                         NOT_AVAILABLE
UNITED                                         NOT_AVAILABLE
SQL>

```

- c. Create the CDB root keystore with its own password.

```
SQL> ADMINISTER KEY MANAGEMENT CREATE KEYSTORE IDENTIFIED BY
password_CDB;

keystore altered.

SQL>
```

Observe that you do not have to provide the path of the keystore. It is already defined by using the `WALLET_ROOT` initialization parameter replacing the `ENCRYPTION_WALLET_LOCATION` parameter in the `$ORACLE_HOME/network/admin/sqlnet.ora` file. After you create the keystore, it is closed by default.

```
SQL> SELECT wrl_parameter, status, keystore_mode
        FROM v$encryption_wallet;

 2
WRL_PARAMETER                                STATUS
KEYSTORE
-----
/u01/app/oracle/admin/ORCL/tde_keystore/tde/ CLOSED      NONE
                                           CLOSED
UNITED
                                           CLOSED
UNITED

SQL> !ls -l /u01/app/oracle/admin/ORCL/tde_keystore/tde/*
-rw----- 1 oracle oinstall 2555 Nov 24 02:50
/u01/app/oracle/admin/ORCL/tde_keystore/tde/ewallet.p12

SQL>
```

- d. Open the CDB root keystore.

```
SQL> ADMINISTER KEY MANAGEMENT
      SET KEYSTORE OPEN IDENTIFIED BY password_CDB
      CONTAINER=ALL;

 2      3
keystore altered.

SQL> SELECT con_id, wrl_parameter, status, keystore_mode
        FROM v$encryption_wallet;

 2
CON_ID WRL_PARAMETER
```

```

-----
STATUS                KEYSTORE
-----
1 /u01/app/oracle/admin/ORCL/tde_keystore/tde/
OPEN_NO_MASTER_KEY NONE

2

CLOSED                UNITED

3

OPEN_NO_MASTER_KEY UNITED

SQL>

```

The CDB root keystore is opened without a master key set.

- e. Set the master key in the CDB root keystore.

```

SQL> ADMINISTER KEY MANAGEMENT
      SET KEY IDENTIFIED BY password_CDB WITH BACKUP
      CONTAINER=ALL;

2      3
keystore altered.

SQL>

```

An ORA-46665 error message may appear because some PDBs did not open the keystore in their container. Therefore, the TDE master encryption key for those PDBs cannot be activated until they are opened. Nevertheless, the CDB root keystore is opened with a TDE master encryption key set.

```

SQL> SELECT con_id, wrl_parameter, status, wallet_type,
           keystore_mode
      FROM v$encryption_wallet;

2      3
CON_ID WRL_PARAMETER
-----
STATUS                WALLET_TYPE                KEYSTORE
-----
1 /u01/app/oracle/admin/ORCL/tde_keystore/tde/
OPEN                PASSWORD                NONE

2

CLOSED                UNKNOWN                UNITED

```

	3		
OPEN		PASSWORD	UNITED
SQL>			

The new `KEYSTORE_MODE` column displays `NONE` for the CDB root. This value is seen when `V$ENCRYPTION_WALLET` is queried from the CDB root, or when the system is a non-CDB. In the latter case, the keystore mode does not apply.

The new `KEYSTORE_MODE` column displays `UNITED` for the PDBs. This is the default mode.

3. For further tasks in the practice, set the key in `PDB1`.

```
SQL> CONNECT sys@PDB1 AS SYSDBA
Enter password: password
Connected.
SQL> ADMINISTER KEY MANAGEMENT
      SET KEY IDENTIFIED BY password_CDB WITH BACKUP;
      2
keystore altered.
SQL>
```

4. Create a PDB to run in isolated mode.
 - a. Create the new `PDB_ISOLATED` PDB and open it. By default, the `PDB_ISOLATED` PDB uses the default mode of the CDB, which is the united mode.

```
SQL> CONNECT / AS SYSDBA
Connected.
SQL> !mkdir /u02/app/oracle/oradata/ORCL/PDB_ISOLATED

SQL> CREATE PLUGGABLE DATABASE pdb_isolated
      ADMIN USER pdb2_admin IDENTIFIED BY password
      ROLES=(CONNECT)
CREATE_FILE_DEST='/u02/app/oracle/oradata/ORCL/PDB_ISOLATED';
      2      3
Pluggable database created.

SQL> ALTER PLUGGABLE DATABASE pdb_isolated OPEN;

Pluggable database altered.

SQL>
```

- b. Isolate the PDB keystore for the newly created PDB.

```
SQL> CONNECT sys@pdb_isolated AS SYSDBA
Enter password: password
Connected.
SQL> ALTER SYSTEM SET tde_configuration =
        'KEYSTORE_CONFIGURATION=FILE' SCOPE=BOTH;
2
System altered.

SQL> SELECT wrl_parameter, status, wallet_type, keystore_mode
        FROM v$encryption_wallet;
2
WRL_PARAMETER                                STATUS
-----
WALLET_TYPE                                KEYSTORE
-----
/u01/app/oracle/admin/ORCL/tde_keystore/5EB2 NOT_AVAILABLE
7B9A049874D8E0532133960A3FA3/tde/
UNKNOWN                                ISOLATED

SQL>
```

The new `KEYSTORE_MODE` column displays `ISOLATED`. This value is seen when `V$ENCRYPTION_WALLET` is queried from the PDB. The PDB is configured to use its own keystore and to open the keystore using its own isolated keystore password.

- c. Create the security officer in the PDB responsible for PDB keystore management.

```
SQL> CREATE USER sec_tde IDENTIFIED BY password;

User created.

SQL> GRANT create session, administer key management,
        select any dictionary, create any table TO sec_tde;
2
Grant succeeded.

SQL>
```


- d. Now create the PDB keystore with its own password, open the PDB keystore, and set the PDB TDE master encryption key.
- e. Create the PDB keystore with its own password.

```
SQL> CONNECT sec_tde@pdb_isolated
Enter password: password
Connected.
SQL> ADMINISTER KEY MANAGEMENT
        CREATE KEYSTORE IDENTIFIED BY password_PDB;

2
keystore altered.

SQL>
```

- f. Open the PDB keystore.

```
SQL> ADMINISTER KEY MANAGEMENT
        SET KEYSTORE OPEN IDENTIFIED BY password_PDB;

2
keystore altered.

SQL>
```

- g. Set the PDB TDE master encryption key.

```
SQL> ADMINISTER KEY MANAGEMENT
        SET KEY IDENTIFIED BY password_PDB WITH BACKUP;

2
keystore altered.

SQL>
```

- h. Verify that the PDB holds its own key.

```
SQL> SELECT wrl_parameter, status, wallet_type, keystore_mode
        FROM v$encryption_wallet;

2
WRL_PARAMETER                                STATUS
-----
WALLET_TYPE                                KEYSTORE
-----
/u01/app/oracle/admin/ORCL/tde_keystore/5EB2 OPEN
7B9A049874D8E0532133960A3FA3/tde/
PASSWORD                                ISOLATED

SQL>
```

- i. Check the existence of the `ewallet.p12` file in the subdirectory created under the `WALLET_ROOT` directory.

```
SQL> !ls -l
/u01/app/oracle/admin/ORCL/tde_keystore/5EB27B9A049874D8E0532133
960A3FA3/tde/
total 4
-rw----- 1 oracle oinstall 2555 Nov 24 04:10 ewallet.p12
-rw----- 1 oracle oinstall 2555 Jan 5 15:57
ewallet_2018010515572043.p12

SQL>
```

- j. Create an encrypted table.

```
SQL> CREATE TABLE system.test (c NUMBER ENCRYPT);

Table created.

SQL>
```

5. Convert a united mode PDB to run in isolated mode.

- a. Find an existing PDB that holds its TDE master encryption key in the CDB root keystore, working in united mode.

```
SQL> CONNECT / AS SYSDBA
Connected.
SQL> SELECT pdb_name, keystore_mode
      FROM v$encryption_wallet w, cdb_pdb$seeds p
      WHERE w.con_id = p.pdb_id;

 2      3
PDB_NAME      KEYSTORE
-----
PDB1          UNITED
PDB$SEED      UNITED
PDB_ISOLATED  ISOLATED

SQL>
```

```
SQL> CONNECT sys@PDB1 AS SYSDBA
Enter password: password
Connected.
SQL> SELECT wrl_parameter, status, wallet_type, keystore_mode
      FROM v$encryption_wallet;

 2
WRL_PARAMETER      STATUS
-----
```

WALLET_TYPE	KEYSTORE	
-----	-----	
		OPEN
PASSWORD	UNITED	
SQL>		

Observe that `WRL_PARAMETER` in the `V$ENCRYPTION_WALLET` view displays a null value. There is no value because the PDB TDE master encryption keys for united PDBs are stored in the CDB root keystore.

The new `KEYSTORE_MODE` column displays `UNITED`. This value is seen when `V$ENCRYPTION_WALLET` is queried from the PDB. The PDB is configured to use the CDB root keystore to store its PDB TDE master encryption key.

If you wanted the PDB keystore to be of another keystore type than the default value set at the CDB root level, then you would set the `TDE_CONFIGURATION` parameter to another value.

- b. Create and open the PDB keystore with its own password.

```
SQL> ADMINISTER KEY MANAGEMENT ISOLATE KEYSTORE
                                IDENTIFIED BY password_PDB
                                FROM ROOT KEYSTORE IDENTIFIED BY password_CDB
                                WITH BACKUP;

  2      3      4
keystore altered.

SQL>
```

All the previously active (historical) TDE master encryption keys associated with the PDB are copied to the isolated keystore.

- c. Find the directory where the PDB keystore is created.

```
SQL> SELECT wrl_parameter, status, wallet_type, keystore_mode
       FROM v$encryption_wallet;

  2
WRL_PARAMETER                                STATUS
-----
WALLET_TYPE                                KEYSTORE
-----
/u01/app/oracle/admin/ORCL/tde_keystore/6413 OPEN
979102FD7535E053BA05C40AAAD/tde/
```

```
PASSWORD ISOLATED
```

```
SQL>
```

- d. Check the existence of the `ewallet.p12` file in the subdirectory created under the `WALLET_ROOT` directory.

```
SQL> !ls -l  
/u01/app/oracle/admin/ORCL/tde_keystore/6413979102FD7535E053BA05  
C40AAAAD/tde
```

```
total 4  
-rw----- 1 oracle oinstall 2059 Nov 27 02:07 ewallet.p12
```

```
SQL>
```

- e. Create an encrypted table.

```
SQL> CREATE TABLE system.test (c NUMBER ENCRYPT);
```

```
Table created.
```

```
SQL>
```

6. Convert an isolated mode PDB to run in united mode.

- a. Find an existing PDB that runs in isolated mode.

```
SQL> CONNECT / AS SYSDBA
```

```
Connected.
```

```
SQL> SELECT pdb_name, keystore_mode  
FROM v$encryption_wallet w, cdb_pdb$seeds p  
WHERE w.con_id = p.pdb_id;
```

2	3
PDB_NAME	KEYSTORE
-----	-----
PDB1	ISOLATED
PDB\$SEED	UNITED
PDB_ISOLATED	ISOLATED

```
SQL>
```

- b. Log in to `PDB_ISOLATED` as `SEC_TDE` to convert `PDB_ISOLATED` to run in united mode.

```
SQL> CONNECT sec_tde@PDB_ISOLATED
```

```
Enter password: password
```

```
Connected.
```

```
SQL> ADMINISTER KEY MANAGEMENT UNITE KEYSTORE
```

```
IDENTIFIED BY password_PDB  
WITH ROOT KEYSTORE IDENTIFIED BY password_CDB  
WITH BACKUP;
```

```

2 ADMINISTER KEY MANAGEMENT UNITE KEYSTORE IDENTIFIED BY
password_PDB
*
ERROR at line 1:
ORA-01031: insufficient privileges

SQL>

```

Observe that even if the `SEC_TDE` local user has the `ADMINISTER KEY MANAGEMENT` privilege, it is not sufficient to convert the PDB keystore to united mode.

- c. Create a common user in the CDB root and grant the user the `ADMINISTER KEY MANAGEMENT` privilege commonly.

```

SQL> CONNECT / AS SYSDBA
Connected.
SQL> CREATE USER c##sec_off IDENTIFIED BY password
CONTAINER=ALL;

2
User created.

SQL> GRANT create session, administer key management,
select any dictionary, create any table
TO c##sec_off CONTAINER=ALL;

2 3
Grant succeeded.

SQL>

```

- d. Log in to `PDB_ISOLATED` as the common user to convert the PDB to run in united mode.

```

SQL> CONNECT c##sec_off@PDB_ISOLATED
Enter password: password
Connected.
SQL> ADMINISTER KEY MANAGEMENT UNITE KEYSTORE
IDENTIFIED BY password_CDB
WITH ROOT KEYSTORE IDENTIFIED BY password_CDB
WITH BACKUP;

2 3 4
ADMINISTER KEY MANAGEMENT UNITE KEYSTORE
*
ERROR at line 1:
ORA-28357: password required to open the wallet

SQL>

```

Q/ Which keystore password does the error refer to?

A/ It refers to the PDB isolated keystore.

```
SQL> ADMINISTER KEY MANAGEMENT UNITE KEYSTORE
                                IDENTIFIED BY password_PDB
                                WITH ROOT KEYSTORE IDENTIFIED BY password_CDB
                                WITH BACKUP;

  2      3      4
keystore altered.

SQL>
```

- e. Verify that the PDB is now running in united mode.

```
SQL> SELECT wrl_parameter, status, wallet_type, keystore_mode
      FROM v$encryption_wallet;

  2
WRL_PARAMETER          STATUS      WALLET_TYPE          KEYSTORE
-----
                                OPEN          PASSWORD          UNITED

SQL>
```

All the previously active (historical) TDE master encryption keys associated with the PDB are copied to the keystore of the CDB root.

- f. Create an encrypted table.

```
SQL> CREATE TABLE system.test2 (c NUMBER ENCRYPT);

Table created.

SQL> EXIT
$
```

7. Execute the `$HOME/labs/admin/cleanup_PDBs.sh` shell script to drop PDB1 and PDB_ISOLATED in ORCL.

If you used other passwords for the keystores of PDB1 and PDB_ISOLATED in ORCL, update the passwords in the shell script and SQL scripts.

```
$ $HOME/labs/admin/cleanup_PDBs.sh
...
$
```

Practice 3-3: Creating User-Defined TDE Master Encryption Keys

Overview

In this practice, you will set your own PDB TDE master encryption key and create and later use your own PDB TDE master encryption key.

Tasks

1. Re-create PDB1 and create PDB2 by executing the `$HOME/labs/SEC/PDBs.sh` shell script.

```
$ $HOME/labs/SEC/PDBs.sh
...
$
```

2. Set your own TDE master encryption key.
 - a. Log in to the CDB root and verify that PDB1 PDB exists and is open.

```
$ sqlplus / AS SYSDBA

SQL> SHOW PDBS

      CON_ID CON_NAME                                OPEN MODE  RESTRICTED
-----
          2 PDB$SEED                                READ ONLY  NO
          5 PDB1                                    READ WRITE NO
          7 PDB2                                    READ WRITE NO

SQL>
```

- b. Check whether the CDB root keystore exists and is opened with a TDE master encryption key set.

```
SQL> SELECT wrl_parameter, status, keystore_mode
       FROM v$encryption_wallet
       WHERE con_id = 1;

 2      3
WRL_PARAMETER                                STATUS
-----
KEYSTORE
-----
/u01/app/oracle/admin/ORCL/tde_keystore/tde/ OPEN
NONE

SQL>
```

- c. Log in to PDB1 and verify that the PDB keystore is not created yet.

```
SQL> CONNECT sys@PDB1 AS SYSDBA
Enter password: password
Connected.

SQL> SELECT wrl_parameter, status, keystore_mode
        FROM v$encryption_wallet;

2
WRL_PARAMETER          STATUS          KEYSTORE
-----
                      CLOSED          UNITED

SQL>SQL>
```

- d. Create the security officer who will define and activate a TDE master encryption for PDB1 by using a defined raw binary value.

```
SQL> CREATE USER sec_tde IDENTIFIED BY password;

User created.

SQL> GRANT create session, administer key management,
        unlimited tablespace, create any table,
        select any dictionary TO sec_tde;

2      3
Grant succeeded.

SQL>
```

- e. Connect as the security officer and create the PDB keystore.

```
SQL> CONNECT sec_tde@PDB1
Enter password: password
Connected.

SQL> ADMINISTER KEY MANAGEMENT CREATE KEYSTORE
        IDENTIFIED BY password_PDB;

2
keystore altered.

SQL>
```

- f. Check whether the PDB keystore is created.

```
SQL> SELECT wrl_parameter, status, keystore_mode
        FROM v$encryption_wallet;

2
WRL_PARAMETER          STATUS
-----
KEYSTORE
-----
```



```
/u01/app/oracle/admin/ORCL/tde_keystore/624 CLOSED
B9740A1760A7E053CE4CC40A6F4F/tde/
ISOLATED

SQL>
```

- g. Open the keystore of the PDB.

```
SQL> ADMINISTER KEY MANAGEMENT SET KEYSTORE OPEN
        IDENTIFIED BY password_PDB;

      2
keystore altered.

SQL>
```

- h. Define and activate your own TDE master encryption key for PDB1.

```
SQL> ADMINISTER KEY MANAGEMENT CREATE KEY '17BA048707B402493C'
        IDENTIFIED BY password_PDB WITH BACKUP;
ADMINISTER KEY MANAGEMENT CREATE KEY '17BA048707B402493C'
*
ERROR at line 1:
ORA-46685: invalid master key identifier or master key value

SQL>
```

```
SQL> !oerr ora 46685
46685, 00000, "invalid master key identifier or master key
value"
// *Cause: The master key identifier or master key value was
invalid.
//           To be valid, the values supplied for both the master
key identifier and the master key value must be hex-encoded and
must have the following properties:
//           - For a master key value to be valid, its length
must be a valid key length for the specified encryption
algorithm.
//           - For a master key identifier to be valid, it must
have a length of 16 bytes.
//           - For a master key identifier to be valid, it must
have a non-zero value.
// *Action: Specify a correct value for the master key
identifier or master key in the master key string.

SQL>
```

- i. Define a proper value.

```
SQL> ADMINISTER KEY MANAGEMENT CREATE KEY
      '10203040506070801112131415161718:3D432109DF88967A541967062A6F4E
      460E892318E307F017BA048707B402493C'
      IDENTIFIED BY password_PDB WITH BACKUP;
      2      3
keystore altered.

SQL>
```

Q/ You could have received the following error message:

ORA-46658: keystore not open in the container

What would have been the reason of the failure?

A/ The PDB keystore must be opened. The keystore can be temporarily opened with the `FORCE KEYSTORE` clause. In this case, you would have used the following command:

```
SQL> ADMINISTER KEY MANAGEMENT SET KEY '17BA048707B402493C'
      FORCE KEYSTORE IDENTIFIED BY password_PDB WITH BACKUP;
```

- j. Verify that the key is set in the PDB keystore.

```
SQL> SELECT key_id, activating_pdbname FROM v$encryption_keys;

KEY_ID
-----
ACTIVATING_PDBNAME
-----
ARAgMEBQYHCAERITFBUWFxgAAAAAAAAAAAAAAAAAAAAAAAAAAAAA

SQL>
```

Q/ Is the key activated in the PDB?

A/ It seems not. The `ACTIVATING_PDBNAME` column does not show any value.

- k. Test if you can create a table with an encrypted column.

```
SQL> CREATE TABLE test (c NUMBER ENCRYPT);
CREATE TABLE test (c NUMBER ENCRYPT)
*
ERROR at line 1:
ORA-28361: master key not yet set

SQL>
```

This confirms our assumption. By creating the key with the `ADMINISTER KEY MANAGEMENT CREATE KEY` command, you can activate the key whenever you want.

- l. Use the `key_id` found in the `v$encryption_keys` view in Step 2-j.

```
SQL> ADMINISTER KEY MANAGEMENT USE KEY
      'ARAgMEBQYHCAERITFBUWFxgAAAAAAAAAAAAAAAAAAAAAAAAAAAA'
      FORCE KEYSTORE IDENTIFIED BY password_PDB WITH BACKUP;
  2      3
keystore altered.

SQL>
```

- m. Test if you can create a table with an encrypted column.

```
SQL> CREATE TABLE test (c NUMBER ENCRYPT);

Table created.

SQL> DESC test
      Name                                     Null?      Type
-----
C                                     NUMBER ENCRYPT

SQL>
```

- n. Verify that the key is now set in the PDB keystore.

```
SQL> SELECT key_id, activating_pdbname FROM v$encryption_keys;

KEY_ID
-----
ACTIVATING_PDBNAME
-----
ARAgMEBQYHCAERITFBUWFxgAAAAAAAAAAAAAAAAAAAAAAAAAAAA
PDB1

SQL>
```

3. Define and use your own TDE master encryption key in one single statement.

- a. Log in to `PDB2` and verify that the PDB keystore does not exist yet.

```
SQL> CONNECT sys@PDB2 AS SYSDBA
Enter password: password
Connected.
SQL> SELECT wrl_parameter, status, keystore_mode
      FROM v$encryption_wallet;
  2
```

WRL_PARAMETER	STATUS
-----	-----
KEystore	

	CLOSED
UNITED	
SQL>	

- b. Connect as the security officer and create the PDB keystore.

```
SQL> CREATE USER sec_tde IDENTIFIED BY password;

User created.

SQL> GRANT create session, ADMINISTER KEY MANAGEMENT,
      select any dictionary, unlimited tablespace,
      create any table
      TO sec_tde;
  2      3      4
Grant succeeded.

SQL>
```

```
SQL> CONNECT sec_tde@PDB2
Enter password: password
Connected.
SQL> ADMINISTER KEY MANAGEMENT CREATE KEYSTORE
      IDENTIFIED BY password_PDB;
  2
keystore altered.

SQL>
```

- c. Check whether the PDB keystore is created.

```
SQL> SELECT wrl_parameter, status, keystore_mode
      FROM v$encryption_wallet;
  2
WRL_PARAMETER                                STATUS
-----
KEystore
-----
/u01/app/oracle/admin/ORCL/tde_keystore/6249 CLOSED
3C4D0CFD420EE053CE4CC40A86D4/tde/
```

```
ISOLATED
```

```
SQL>
```

- d. Define and use your own TDE master encryption key for PDB2.

```
SQL> ADMINISTER KEY MANAGEMENT SET KEY  
'10203040506070801112131415161718:3D432109DF88967A541967062A6F4E  
460E892318E307F017BA048707B402493C'
```

```
FORCE KEYSTORE IDENTIFIED BY password_PDB WITH BACKUP;
```

```
2 3
```

```
keystore altered.
```

```
SQL>
```

- e. Verify that the key is set in the PDB keystore.

```
SQL> SELECT key_id, activating_pdbname FROM v$encryption_keys;
```

```
KEY_ID
```

```
-----  
ACTIVATING_PDBNAME
```

```
-----  
ARAgMEBQYHCAERITFBUWFxgAAAAAAAAAAAAAAAAAAAAAAAAAAAAA
```

```
PDB2
```

```
SQL>
```

Observe that the key is activated. Compare with the result of Step 2-j.

- f. Test if you can create a table with an encrypted column.

```
SQL> CREATE TABLE test (c NUMBER ENCRYPT);
```

```
Table created.
```

```
SQL> DESC test
```

```
Name
```

```
Null?
```

```
Type
```

```
-----  
C
```

```
-----  
NUMBER ENCRYPT
```

```
SQL> EXIT
```

```
$
```

4. Execute the `$HOME/labs/admin/cleanup_PDBs.sh` shell script to drop PDB1 and PDB2 in ORCL.

```
$ $HOME/labs/admin/cleanup_PDBs.sh
...
$
```

Practice 3-4: Exporting and Importing Fixed-User Database Links

Overview

In this practice, you will prevent an intruder from de-obfuscating the obfuscated password of fixed-user database links by not exporting the password, but the metadata only. After exporting database links with invalid password values, you must reset the passwords after reimporting them.

Tasks

1. Execute the `$HOME/labs/SEC/PDBs_TEST_TAB.sh` shell script to re-create PDB1 and PDB2, and create `SYSTEM.TEST` in PDB2.

```
$ $HOME/labs/SEC/PDBs_TEST_TAB.sh
...
$
```

2. Export and import database links **with obfuscated passwords**.
 - a. Log in to the PDB1 PDB and create a fixed-user database link that points to PDB2. (This could be another CDB if there was another one on the VM.)

```
$ sqlplus system@PDB1

Enter password: password

SQL> CREATE DATABASE LINK test CONNECT TO system
        IDENTIFIED BY password
        USING 'PDB2';

      2      3
Database link created.

SQL>
```

- b. Test that the database link allows you to read the `SYSTEM.TEST` table in PDB2.

```
SQL> SELECT * FROM system.test@test;

          C
-----
          1

SQL>
```

- c. Check the enforcement status of credentials encryption in the dictionary.

```
SQL> SELECT enforcement FROM dictionary_credentials_encrypt;

ENFORCEM
-----
DISABLED

SQL>
```

- d. Before exporting the database link, create a logical directory for the export dump file.

```
SQL> CREATE OR REPLACE DIRECTORY dp_dump AS '/tmp';

Directory created.

SQL> EXIT
$
```

- e. Export the fixed-user database link.

```
$ rm /tmp/dp_test.dmp
rm: cannot remove `/tmp/dp_test.dmp': No such file or directory
$ expdp system@PDB1 directory=dp_dump dumpfile=dp_test
Password: password
...
Processing object type SCHEMA_EXPORT/DB_LINK
Processing object type SCHEMA_EXPORT/TABLE/TABLE
Processing object type SCHEMA_EXPORT/TABLE/COMMENT
Processing object type SCHEMA_EXPORT/TABLE/INDEX/INDEX
Processing object type SCHEMA_EXPORT/TABLE/CONSTRAINT/CONSTRAINT
Master table "SYSTEM"."SYS_EXPORT_SCHEMA_01" successfully
loaded/unloaded
*****
Dump file set for SYSTEM.SYS_EXPORT_SCHEMA_01 is:
  /tmp/dp_test.dmp
Job "SYSTEM"."SYS_EXPORT_SCHEMA_01" successfully completed at
Tue Nov 28 09:12:10 2017 elapsed 0 00:00:54

$
```

The obfuscated password of the exported database link is not protected in the dump file.

- f. Drop the fixed-user database link.

```
$ sqlplus system@PDB1

Enter password: password

SQL> DROP DATABASE LINK test;

Database link dropped.

SQL> EXIT

$
```

- g. Import the fixed-user database link.

```
$ impdp system@PDB1 directory=dp_dump dumpfile=dp_test
Password: password
...
Processing object type SCHEMA_EXPORT/DB_LINK
Job "SYSTEM"."SYS_IMPORT_FULL_01" successfully completed at Tue
Nov 28 09:14:50 2017 elapsed 0 00:00:05
$
```

- h. Test that the database link works to read the `TEST` table in `PDB1` of `cdb1`.

```
$ sqlplus system@PDB1

Enter password: password

SQL> SELECT * FROM system.test@test;

          C
-----
          1

SQL> EXIT

$
```

- i. To avoid any security risk, execute the export operation with the recommended Data Pump parameter that encrypts the data in the dump file. By default, all data including metadata is encrypted in the export dump file.

```
$ rm /tmp/dp_test.dmp
$ expdp system@PDB1 directory=dp_dump dumpfile=dp_test
ENCRYPTION_PWD_PROMPT=YES
Enter password: password
...
Encryption Password: enc_password
...
Processing object type SCHEMA_EXPORT/DB_LINK
```

```

Processing object type SCHEMA_EXPORT/TABLE/TABLE
Processing object type SCHEMA_EXPORT/TABLE/COMMENT
Processing object type SCHEMA_EXPORT/TABLE/INDEX/INDEX
Processing object type SCHEMA_EXPORT/TABLE/CONSTRAINT/CONSTRAINT
Master table "SYSTEM"."SYS_EXPORT_SCHEMA_01" successfully
loaded/unloaded
*****
Dump file set for SYSTEM.SYS_EXPORT_SCHEMA_01 is:
  /tmp/dp_test.dmp
Job "SYSTEM"."SYS_EXPORT_SCHEMA_01" successfully completed at
Tue Nov 28 09:18:52 2017 elapsed 0 00:00:52
$

```

- j. After dropping the database link, import it back.

```

$ sqlplus system@PDB1

Enter password: password

SQL> DROP DATABASE LINK test;

Database link dropped.

SQL> EXIT
$

```

```

$ impdp system@PDB1 directory=dp_dump dumpfile=dp_test
ENCRYPTION_PWD_PROMPT=YES
Password: password
...
Encryption Password: enc_password
...
Processing object type SCHEMA_EXPORT/DB_LINK
Job "SYSTEM"."SYS_IMPORT_FULL_01" successfully completed at Tue
Nov 28 09:21:25 2017 elapsed 0 00:00:09

$

```

Remark: Then during the import operation, whether the keystore is opened or closed affects whether or not an encryption password must be provided. If the keystore is opened during the export operation and you provided an encryption password, then you do not need to provide the password during the import operation. If the keystore is closed during the export operation, then you must provide the password during the import operation.

- k. Test that the database link works to read the `TEST` table in `PDB1`.

```
$ sqlplus system@PDB1

Enter password: password

SQL> SELECT * FROM system.test@test;

          C
-----
          1

SQL>
```

If you want to protect only the passwords of database links because there is no other sensitive data in the data to be exported, then use the method explained in the next task.

3. Enable credentials encryption.
- a. Enable the credentials encryption to allow the fixed-user database links to be exported with an invalid password value in the dump file during export. Log in to the PDB and enable the enforcement of credentials encryption in the dictionary.

```
SQL> ALTER DATABASE dictionary encrypt credentials;
ALTER DATABASE dictionary encrypt credentials
*
ERROR at line 1:
ORA-28447: insufficient privilege to execute ALTER DATABASE
DICTIONARY statement

SQL> CONNECT sys@PDB1 AS SYSDBA
Enter password: password
Connected.

SQL> ALTER DATABASE dictionary encrypt credentials;
ALTER DATABASE dictionary encrypt credentials
*
ERROR at line 1:
ORA-28447: insufficient privilege to execute ALTER DATABASE
DICTIONARY statement

SQL>
```

- b. Create a user with the `SYSKM` privilege.

```
SQL> CREATE USER enccreds IDENTIFIED BY password;

User created.

SQL> GRANT create session, syskm TO enccreds;

Grant succeeded.

SQL>
```

- c. Connect as the `SYSKM` privileged user in `PDB1` and enable the enforcement of credentials encryption in the dictionary.

```
SQL> CONNECT enccreds@PDB1 AS SYSKM
Enter password: password
Connected.
SQL> ALTER DATABASE dictionary encrypt credentials;

Database dictionary altered.

SQL> SELECT enforcement FROM dictionary_credentials_encrypt;

ENFORCEM
-----
ENABLED

SQL> EXIT
$
```

4. Export and import database links with invalid passwords.

- a. Export the fixed-user database link.

```
$ rm /tmp/dp_test.dmp
$ expdp system@PDB1 directory=dp_dump dumpfile=dp_test
Enter password: password
...
Processing object type SCHEMA_EXPORT/DB_LINK
ORA-39395: Warning: object SYSTEM.TEST requires password reset
after import

Processing object type SCHEMA_EXPORT/TABLE/TABLE
Processing object type SCHEMA_EXPORT/TABLE/COMMENT
Processing object type SCHEMA_EXPORT/TABLE/INDEX/INDEX
Processing object type SCHEMA_EXPORT/TABLE/CONSTRAINT/CONSTRAINT
```

```

Master table "SYSTEM"."SYS_EXPORT_SCHEMA_01" successfully
loaded/unloaded
*****

Dump file set for SYSTEM.SYS_EXPORT_SCHEMA_01 is:
  /tmp/dp_test.dmp
Job "SYSTEM"."SYS_EXPORT_SCHEMA_01" completed with 1 error(s) at
Tue Nov 28 09:50:14 2017 elapsed 0 00:00:44

$

```

The database link password is stored as an invalid value in the dump file. This is the reason for the warning message that tells you that after importing the database link, you must reset the password of the database link.

- b. Drop the fixed-user database link.

```

$ sqlplus system@PDB1

Enter password: password

SQL> DROP DATABASE LINK test;

Database link dropped.

SQL> EXIT

$

```

- c. Import the fixed-user database link. The database link password is still stored as an invalid value in the dump file. This is the reason for the warning message.

```

$ impdp system@PDB1 directory=dp_dump dumpfile=dp_test
Password: password
...
Processing object type SCHEMA_EXPORT/DB_LINK
ORA-39395: Warning: object SYSTEM.TEST requires password reset
after import

Job "SYSTEM"."SYS_IMPORT_FULL_01" completed with 1 error(s) at
Tue Nov 28 09:53:08 2017 elapsed 0 00:00:05

$

```

- d. Test that the database link works to read the `TEST` table in `PDB1` of `cdb1`. If the `SYSTEM` user is not granted the `ALTER DATABASE LINK system` privilege, ask the user with the `SYSDBA` privilege to grant the privilege to `SYSTEM`.

```
$ sqlplus system@PDB1

Enter password: password

SQL> SELECT * FROM system.test@test;
SELECT * FROM system.test@test
                        *
ERROR at line 1:
ORA-28449: cannot use an invalidated database link

SQL> ALTER DATABASE LINK test CONNECT TO system
                        IDENTIFIED BY password;
ALTER DATABASE LINK test CONNECT TO system IDENTIFIED BY
password
*
ERROR at line 1:
ORA-01031: insufficient privileges

SQL> CONNECT sys@PDB1 AS SYSDBA
Enter password: password
Connected.
SQL> GRANT alter database link TO system;

Granted succeeded.

SQL> CONNECT system@PDB1
Enter password: password
Connected.
SQL> ALTER DATABASE LINK test CONNECT TO system
                        IDENTIFIED BY password;
Database link altered.

SQL> SELECT * FROM system.test@test;

      C
-----
      1
```

```
SQL> EXIT  
$
```

5. Execute the `$HOME/labs/SEC/cleanup_DBlinks.sh` shell script to drop the database link created in PDB1 and then the `$HOME/labs/admin/cleanup_PDBs.sh` shell script.

```
$ $HOME/labs/SEC/cleanup_DBlinks.sh  
...  
$ $HOME/labs/admin/cleanup_PDBs.sh  
...  
$
```

Practice 3-5: Encrypting Sensitive Data in Database Replay Files

Overview

In this practice, you will protect sensitive data when the data is captured by Database Replay in capture files. You will encrypt the data captured by Database Replay in capture files.

Tasks

1. Before starting the practice, execute the `$HOME/labs/SEC/DBReplay.sh` shell script. The script re-creates PDB1 and sets the master key set in PDB1.

```
$ $HOME/labs/SEC/DBReplay.sh
...
$
```

2. Set the password for the Database Replay user in the keystore.
 - a. Log in to ORCL and verify the master key set.

```
$ sqlplus / AS SYSDBA

SQL> SELECT wrl_parameter, status, keystore_mode, wallet_type
      FROM v$encryption_wallet WHERE con_id = 1;

  2
WRL_PARAMETER                                STATUS
-----
KEYSTORE WALLET_TYPE
-----
/u01/app/oracle/admin/ORCL/tde_keystore/tde/ OPEN
NONE      PASSWORD

SQL>
```

- b. Set the password for the `oracle.rat.database_replay.encryption` user in the keystore.

```
SQL> ADMINISTER KEY MANAGEMENT ADD SECRET 'replay_password'
      FOR CLIENT 'oracle.rat.database_replay.encryption'
      IDENTIFIED BY password WITH BACKUP;

  2      3
keystore altered.

SQL>
```


3. Capture workload data by using Database Replay.

- a. The Database Replay capture creates files in a directory. Create the logical directory for the capture files.

```
SQL> CONNECT system
Enter password: password
Connected.
SQL> HOST mkdir /home/oracle/replay

SQL> CREATE OR REPLACE DIRECTORY oltp
                AS '/home/oracle/replay';

2
Directory created.

SQL>
```

- b. Start capturing data with the Database Replay procedure. Specify the encryption algorithm used to encrypt the data in the workload capture files.

```
SQL> exec DBMS_WORKLOAD_CAPTURE.START_CAPTURE (name =>
'OLTP_peak', dir => 'OLTP', ENCRYPTION => 'AES256')

PL/SQL> procedure successfully completed.

SQL>
```

- c. During the capture, in another terminal window session called *Workload Session*, execute the workload on the CDB and PDBs.

```
$ $HOME/labs/SEC/workload.sh
...
$
```

- d. When you consider that the workload is sufficient for replay testing, stop the capture.

```
SQL> exec DBMS_WORKLOAD_CAPTURE.FINISH_CAPTURE ()

PL/SQL> procedure successfully completed.

SQL>
```

The data in capture files is encrypted with an encryption key stored in the header of each file. The encryption key was generated with the encryption key mapped to the `oracle.rat.database_replay.encryption` user retrieved from the keystore.

4. Process the capture files and replay the workload.
 - a. As in the normal whole process of Database Replay, after capturing the workload into files, you process the capture files.

```
SQL> exec DBMS_WORKLOAD_REPLAY.PROCESS_CAPTURE (capture_dir =>
'OLTP')

PL/SQL> procedure successfully completed.

SQL>
```

- b. Then you initialize and prepare the replay.

```
SQL> exec DBMS_WORKLOAD_REPLAY.INITIALIZE_REPLAY(replay_name =>
'R', replay_dir => 'OLTP')

PL/SQL> procedure successfully completed.

SQL> exec DBMS_WORKLOAD_REPLAY.PREPARE_REPLAY ( )

PL/SQL> procedure successfully completed.

SQL>
```

During the last three steps, the keystore must be open. The password for the Database Replay user `oracle.rat.database_replay.encryption` is retrieved and verified in the keystore.

- c. You are ready to start workload clients to replay the captured workload with `wrc` clients. In *Workload Session*, if the workload is not finished, interrupt the `$HOME/labs/SEC/workload.sh` shell script and start the `wrc` process.

```
$ wrc REPLAYDIR=/home/oracle/replay USERID=system
...
Password: password
(wrc_main_7555.trc) ORA-15526: failed to validate user password
for capture
$
```

Observe that the client fails with an error message.

The client user password needs to be validated against the `oracle.rat.database_replay.encryption` user password stored in the keystore. On the client side, the `oracle.rat.database_replay.encryption` client password is set in an SSO wallet.

- d. Create the SSO wallet to store the `oracle.rat.database_replay.encryption` client password that must match the password set in step 2-b. First, create the directory for the `cwallet.sso` file. Then create the `cwallet.sso` file.

```
$ mkdir /tmp/replay_encrypt_cwallet
$ mkstore -wrl /tmp/replay_encrypt_cwallet -create
...
Enter password: cwallet_password
Enter password again: cwallet_password
$
```

- e. Add the credentials for the `oracle.rat.database_replay.encryption` user.

```
$ mkstore -wrl /tmp/replay_encrypt_cwallet -createEntry
'oracle.rat.database_replay.encryption' "replay_password"
...
Enter wallet password: cwallet_password
$
```

- f. Start the workload clients to replay the captured workload with `wrc` clients.

```
$ wrc REPLAYDIR=/home/oracle/replay USERID=system
WALLETDIR=/tmp/replay_encrypt_cwallet
...
Password: password
Wait for the replay to start (17:20:04)
```

The password required is the `SYSTEM` user password.

- g. The `wrc` client is waiting for DB Replay to start in the database. Back to the initial session in the CDB, execute the `START_REPLAY` procedure.

```
SQL> exec DBMS_WORKLOAD_REPLAY.START_REPLAY ()

PL/SQL> procedure successfully completed.

SQL> EXIT
$
```

- h. As soon as the DB Replay procedure is started in the database, the clients start replaying and finally completes.

```
Replay client 1 started (17:22:33)
Replay client 1 finished (17:24:24)
$
```

5. Execute the `$HOME/labs/SEC/cleanup_DBReplay.sh` shell script to drop the DB Replay capture files.

```
$ $HOME/labs/SEC/cleanup_DBReplay.sh
...
$
```


Practices for Lesson 4: Using RMAN Enhancements

Practices for Lesson 4: Overview

Overview

In these practices, you recover a plugged non-CDB and PDB using preplugin backups, duplicate a PDB into an existing CDB without instantiating a new instance, and finally duplicate an on-premises CDB for migration to cloud.

Practice 4-1: Recovering a Plugged Non-CDB Using Preplugin Backups

Overview

In this practice, you recover a plugged non-CDB using preplugin backups. Preplugin backups are all backups of a non-CDB taken before unplugging/plugging operations, usable after plugging into a new CDB.

Tasks

1. You use the Oracle By Example “Recovering Plugged Non-CDB Using Preplugin Backups” to see how to proceed. Launch a browser and click the bookmark from the Bookmarks toolbar of the browser. The URL is
`file:///home/oracle/labs/OBE/recovering_plugged_noncdb_using_preplugin_backups/recovering_plugged_noncdb_using_preplugin_backups.html`

Practice 4-2: Recovering a Plugged PDB Using Preplugin Backups

Overview

In this practice, you recover a plugged PDB using preplugin backups.

Tasks

1. Before starting the practice, execute the `$HOME/labs/admin/glogin_4.sh` shell script. It sets formatting for all columns selected in queries.

```
$ $HOME/labs/admin/glogin_4.sh
...
$
```

2. Re-create PDB1 for the practice by executing the `$HOME/labs/RMAN/PDB1_backup.sh` shell script. The shell script also sets the PDB key in the keystore (because PDBs use encryption on cloud VMs) and creates the `HR` schema in PDB1. In the case you run out of space, clean up the Fast Recovery Area with commands such as:

```
rman target /
DELETE OBSOLETE;
DELETE ARCHIVELOG ALL;
ALTER SYSTEM SET DB_RECOVERY_FILE_DEST_SIZE=30G;
EXIT
```

```
$ $HOME/labs/RMAN/PDB1_backup.sh
...
$
```

3. Prepare the PDB before unplugging to use preplugin backups.
 - a. In your current terminal window (called *Session1*), display the content of the `HR.DEPARTMENTS` table in PDB1.

```
$ sqlplus system@PDB1

Enter password: password

SQL> SELECT * FROM hr.departments;

DEPARTMENT_ID DEPARTMENT_NAME      MANAGER_ID LOCATION_ID
-----
          10 Administration          200         1700
          20 Marketing              201         1800
          30 Purchasing              114         1700
          40 Human Resources          203         2400
          50 Shipping                121         1500
          60 IT                      103         1400
```

70	Public Relations	204	2700
80	Sales	145	2500
90	Executive	100	1700
100	Finance	108	1700
110	Accounting	205	1700
120	Treasury		1700
130	Corporate Tax		1700
140	Control And Credit		1700
150	Shareholder Services		1700
160	Benefits		1700
170	Manufacturing		1700
180	Construction		1700
190	Contracting		1700
200	Operations		1700
210	IT Support		1700
220	NOC		1700
230	IT Helpdesk		1700
240	Government Sales		1700
250	Retail Sales		1700
260	Recruiting		1700
270	Payroll		1700

27 rows selected.

SQL>

- b. In another terminal window (called *Session2*), launch `rman` to create backups for `PDB1` of `ORCL` before unplugging. It is recommended to back up all the archive logs before the unplug operation so that they can be restored from preplugin backups when a recovery might be required in the destination CDB.

```
$ rman target /

connected to target database: ORCL (DBID=1488874136)

RMAN> BACKUP PLUGGABLE DATABASE pdb1 PLUS ARCHIVELOG;

Starting backup at 30-NOV-17
current log archived
using target database control file instead of recovery catalog
allocated channel: ORA_DISK_1
channel ORA_DISK_1: SID=43 device type=DISK
skipping archived logs of thread 1 from sequence 121 to 175;
already backed up
channel ORA_DISK_1: starting archived log backup set
channel ORA_DISK_1: specifying archived log(s) in backup set
```

```

input archived log thread=1 sequence=176 RECID=176
STAMP=961428612
input archived log thread=1 sequence=177 RECID=177
STAMP=961428715
channel ORA_DISK_1: starting piece 1 at 30-NOV-17
channel ORA_DISK_1: finished piece 1 at 30-NOV-17
piece
handle=/u03/app/oracle/fast_recovery_area/ORCL/backupset/2017_11
_30/o1_mf_annnn_TAG20171130T153156_f2093d7p_.bkp
tag=TAG20171130T153156 comment=NONE
channel ORA_DISK_1: backup set complete, elapsed time: 00:00:01
Finished backup at 30-NOV-17

Starting backup at 30-NOV-17
using channel ORA_DISK_1
channel ORA_DISK_1: starting full datafile backup set
channel ORA_DISK_1: specifying datafile(s) in backup set
input datafile file number=00110
name=/u02/app/oracle/oradata/ORCL/pdb1/ORCL/5F35EC627C3606EEE053
0241C40A1875/datafile/o1_mf_sysaux_f208zk5x_.dbf
input datafile file number=00109
name=/u02/app/oracle/oradata/ORCL/pdb1/ORCL/5F35EC627C3606EEE053
0241C40A1875/datafile/o1_mf_system_f208zk5t_.dbf
input datafile file number=00111
name=/u02/app/oracle/oradata/ORCL/pdb1/ORCL/5F35EC627C3606EEE053
0241C40A1875/datafile/o1_mf_undotbs1_f208zk60_.dbf
channel ORA_DISK_1: starting piece 1 at 30-NOV-17
channel ORA_DISK_1: finished piece 1 at 30-NOV-17
piece
handle=/u03/app/oracle/fast_recovery_area/ORCL/5F35EC627C3606EEE
0530241C40A1875/backupset/2017_11_30/o1_mf_nnndf_TAG20171130T153
157_f2093ghl_.bkp tag=TAG20171130T153157 comment=NONE
channel ORA_DISK_1: backup set complete, elapsed time: 00:00:15
Finished backup at 30-NOV-17

Starting backup at 30-NOV-17
current log archived
using channel ORA_DISK_1
channel ORA_DISK_1: starting archived log backup set
channel ORA_DISK_1: specifying archived log(s) in backup set
input archived log thread=1 sequence=178 RECID=178
STAMP=961428733
channel ORA_DISK_1: starting piece 1 at 30-NOV-17
channel ORA_DISK_1: finished piece 1 at 30-NOV-17
piece
handle=/u03/app/oracle/fast_recovery_area/ORCL/backupset/2017_11

```

```

_30/o1_mf_annnn_TAG20171130T153214_f2093y4m_.bkp
tag=TAG20171130T153214 comment=NONE
channel ORA_DISK_1: backup set complete, elapsed time: 00:00:01
Finished backup at 30-NOV-17

Starting Control File and SPFILE Autobackup at 30-NOV-17
piece
handle=/u03/app/oracle/fast_recovery_area/ORCL/autobackup/2017_1
1_30/o1_mf_s_961428735_f2093zo7_.bkp comment=NONE
Finished Control File and SPFILE Autobackup at 30-NOV-17

RMAN>

```

- c. Back in *Session1*, export RMAN backup information that belongs to PDB1 to its dictionary before unplugging so that preplugin backups can be used in the target CDB even if no new backup is performed after plugging. The metadata is transported along with the PDB during the migration.

```

SQL> exec DBMS_PDB.EXPORTRMANBACKUP('PDB1')

PL/SQL procedure successfully completed.

SQL>

```

4. Unplug and plug the PDB.

- a. Unplug PDB1 from ORCL and export the PDB key. (This is required on cloud VMs because the PDBs use encryption.)

```

SQL> CONNECT / AS SYSDBA
Connected.
SQL> ALTER PLUGGABLE DATABASE pdb1 CLOSE;

Pluggable database altered.

SQL> HOST rm /tmp/pdb1.xml
rm: Cannot remove '/tmp/pdb1.xml': No such file or directory

SQL> ALTER PLUGGABLE DATABASE pdb1
        UNPLUG INTO '/tmp/pdb1.xml'
        ENCRYPT USING "unplug_password";

  2      3
Pluggable database altered.

SQL>

```

- b. Archive the current redo log and quit the session.

```
SQL> ALTER SYSTEM SWITCH LOGFILE;

System altered.

SQL> SELECT name, next_change#, sequence#
       FROM v$archived_log ORDER BY 3;

NAME
-----
-----
NEXT_CHANGE#  SEQUENCE#
-----
...
/u03/app/oracle/fast_recovery_area/ORCL/archivelog/2017_11_30/o1
_mf_1_178_f2093xs8_.arc
          3500964          178

/u03/app/oracle/fast_recovery_area/ORCL/archivelog/2017_11_30/o1
_mf_1_179_f209rf05_.arc
          3502788          179

179 rows selected.

SQL>
```

- c. Plug PDB1 as ORCL_PDB1 into ORCL and open ORCL_PDB1. (The target CDB could be another CDB. There is no other CDB on the VM.)

```
SQL> HOST mkdir /u02/app/oracle/oradata/ORCL/orcl_pdb1

SQL> CREATE PLUGGABLE DATABASE orcl_pdb1 AS CLONE
       USING '/tmp/pdb1.xml'
       KEYSTORE IDENTIFIED BY CDB_keystore_password
       DECRYPT USING "unplug_password"
       CREATE_FILE_DEST = '/u02/app/oracle/oradata/ORCL/orcl_pdb1';
2      3      4      5
Pluggable database created.

SQL> ALTER PLUGGABLE DATABASE orcl_pdb1 OPEN;

Pluggable database altered.

SQL>
```

- d. In *Session2*, check whether the preplugin backups for ORCL_PDB1 are cataloged in ORCL.

```
RMAN> SET PREPLUGIN CONTAINER=orcl_pdb1;

executing command: SET PREPLUGIN CONTAINER

RMAN>
```

```
RMAN> LIST PREPLUGIN BACKUP;
```

List of **Backup Sets**

=====

BS Key	Size	Device Type	Elapsed Time	Completion Time
--------	------	-------------	--------------	-----------------

3	768.00K	SBT_TAPE	00:00:01	27-NOV-17
BP Key: 3 Status: AVAILABLE Compressed: YES Tag: TAG20171127T080017				
Handle: 05skjnki_1_1 Media: storage.oraclecorp.com/v1/Storage-jpizana/dbaaslrg				

List of Archived Logs in backup set 3

Thrd	Seq	Low SCN	Low Time	Next SCN	Next Time
1	3	2510936	27-NOV-17	2512455	27-NOV-17
1	4	2512455	27-NOV-17	2612473	27-NOV-17

...

BS Key	Size	Device Type	Elapsed Time	Completion Time
--------	------	-------------	--------------	-----------------

201	28.50K	DISK	00:00:00	30-NOV-17
BP Key: 210 Status: AVAILABLE Compressed: NO Tag: TAG20171130T153214				

Piece Name:

/u03/app/oracle/fast_recovery_area/ORCL/backupset/2017_11_30/o1_mf_annnn_TAG20171130T153214_f2093y4m_.bkp

List of Archived Logs in backup set 201

Thrd	Seq	Low SCN	Low Time	Next SCN	Next Time
1	178	3500932	30-NOV-17	3500964	30-NOV-17

```
RMAN>
```

- e. Check whether the preplugin archive log files for ORCL_PDB1 are cataloged in ORCL.

```

RMAN> LIST PREPLUGIN ARCHIVELOG ALL;

List of Archived Log Copies for database with db_unique_name
ORCL
=====
Key          Thrd Seq          S Low Time
-----
121          1      121          A 29-NOV-17
          Name:
/u03/app/oracle/fast_recovery_area/ORCL/archivelog/2017_11_29/o1
_mf_1_121_f1xpddp4_.arc

...

178          1      178          A 30-NOV-17
          Name:
/u03/app/oracle/fast_recovery_area/ORCL/archivelog/2017_11_30/o1
_mf_1_178_f2093xs8_.arc

RMAN>

```

Q/ Has the last redo log file archived been cataloged?

A/ No, because it was created after the PDB was unplugged.

- f. Catalog the last archived log file missing, generated after the unplug operation.

```

RMAN> CATALOG PREPLUGIN ARCHIVELOG
'/u03/app/oracle/fast_recovery_area/ORCL/archivelog/2017_11_30/o
1_mf_1_179_f209rf05_.arc';

cataloged archived log
archived log file
name=/u03/app/oracle/fast_recovery_area/ORCL/archivelog/2017_11_
30/o1_mf_1_179_f209rf05_.arc RECID=179 STAMP=0

RMAN>

```

- g. Verify that the cataloged preplugin backups and archive log files are available on disk.

```

RMAN> CROSSCHECK PREPLUGIN BACKUP;

allocated channel: ORA_SBT_TAPE_1
channel ORA_SBT_TAPE_1: SID=38 device type=SBT_TAPE

```



```

channel ORA_SBT_TAPE_1: Oracle Database Backup Service Library
VER=3.17.9.5
...
backup piece
handle=/u03/app/oracle/fast_recovery_area/ORCL/5F35EC627C3606EEE
0530241C40A1875/backupset/2017_11_30/o1_mf_nnndf_TAG20171130T153
157_f2093ghl_.bkp RECID=209 STAMP=961428718
crosschecked backup piece: found to be 'AVAILABLE'
backup piece
handle=/u03/app/oracle/fast_recovery_area/ORCL/backupset/2017_11
_30/o1_mf_annnn_TAG20171130T153214_f2093y4m_.bkp RECID=210
STAMP=961428734
Crosschecked 94 objects

RMAN>

```

```

RMAN> CROSSCHECK PREPLUGIN ARCHIVELOG ALL;

released channel: ORA_SBT_TAPE_1
released channel: ORA_SBT_TAPE_2
released channel: ORA_SBT_TAPE_3
released channel: ORA_SBT_TAPE_4
released channel: ORA_SBT_TAPE_5
released channel: ORA_DISK_1
allocated channel: ORA_DISK_1
channel ORA_DISK_1: SID=38 device type=DISK
validation succeeded for archived log
archived log file
name=/u03/app/oracle/fast_recovery_area/ORCL/archivelog/2017_11_
29/o1_mf_1_122_flxr4mvc_.arc RECID=122 STAMP=961345811
...
validation succeeded for archived log
archived log file
name=/u03/app/oracle/fast_recovery_area/ORCL/archivelog/2017_1
1_30/o1_mf_1_178_f2093xs8_.arc RECID=178 STAMP=961428733
validation succeeded for archived log
archived log file
name=/u03/app/oracle/fast_recovery_area/ORCL/archivelog/2017_11_
30/o1_mf_1_179_f209rf05_.arc RECID=179 STAMP=961413125
Crosschecked 59 objects

RMAN>

```

5. In *Session1*, a data file is unintentionally removed from ORCL_PDB1.
- a. Remove the system data file of ORCL_PDB1.

```
SQL> CONNECT sys@ORCL_PDB1 AS SYSDBA
Enter password: password
Connected.
SQL> SELECT name FROM v$datafile;

NAME
-----
/u02/app/oracle/oradata/ORCL/orcl_pdb1/ORCL/5F35FE374FC20D36E053
0241C40AA7AB/datafile/o1_mf_system_f20bbbnk_.dbf

/u02/app/oracle/oradata/ORCL/orcl_pdb1/ORCL/5F35FE374FC20D36E053
0241C40AA7AB/datafile/o1_mf_sysaux_f20bbbnk_.dbf

/u02/app/oracle/oradata/ORCL/orcl_pdb1/ORCL/5F35FE374FC20D36E053
0241C40AA7AB/datafile/o1_mf_undotbs1_f20bbbnk_.dbf

SQL> HOST rm
/u02/app/oracle/oradata/ORCL/orcl_pdb1/ORCL/5F35FE374FC20D36E053
0241C40AA7AB/datafile/o1_mf_system_f20bbbnk_.dbf

SQL>
```

- b. Try to connect to the PDB. You may be able to connect but then get an error, or maybe not. In the latter case, you would get the “You are no longer connected to ORACLE.” error message.

```
SQL> CONNECT system@ORCL_PDB1
Enter password: password
ERROR:
ORA-02002: error while writing to audit trail
ORA-01116: error in opening database file 112
ORA-01110: data file 112:
'/u02/app/oracle/oradata/ORCL/orcl_pdb1/ORCL/5F35FE374FC20D36E05
3024
1C40AA7AB/datafile/o1_mf_system_f20bbbnk_.dbf'
ORA-27041: unable to open file
Linux-x86_64 Error: 2: No such file or directory
Additional information: 3

Error accessing package DBMS_APPLICATION_INFO
Connected.
SQL>
```

- c. Close ORCL_PDB1 if it is not already closed.

```
SQL> CONNECT / AS SYSDBA
Connected.
SQL> SHOW PDBS
```

CON_ID	CON_NAME	OPEN MODE	RESTRICTED
2	PDB\$SEED	READ ONLY	NO
4	PDB1	MOUNTED	
5	ORCL_PDB1	READ WRITE	NO

```
SQL> ALTER PLUGGABLE DATABASE orcl_pdb1 CLOSE;

Pluggable database altered.

SQL> SHOW PDBS
```

CON_ID	CON_NAME	OPEN MODE	RESTRICTED
2	PDB\$SEED	READ ONLY	NO
4	PDB1	MOUNTED	
5	ORCL_PDB1	MOUNTED	

```
SQL>
```

6. Recover the PDB from preplugin backups.

- a. To recover the situation, because there is no new backup completed after the plug-in operation, you use the preplugin backups. The data files are restored from backups taken before the PDB was plugged in.

In *Session2*, restore and recover ORCL_PDB1 from the preplugin backups. Because you are working on a cloud VM, encryption is set and, therefore, RMAN requires decryption from backups.

```
RMAN> SET DECRYPTION IDENTIFIED BY CDB_keystore_password;

executing command: SET decryption

RMAN>
```

- b. Run a normal recovery after preplugin recovery in order to synchronize the PDB with the CDB. If the channel used is different, ignore this potential discrepancy.

```
RMAN> RUN
      { RESTORE PLUGGABLE DATABASE orcl_pdb1 FROM PREPLUGIN;
        RECOVER PLUGGABLE DATABASE orcl_pdb1 FROM PREPLUGIN;
      }
2> 3> 4>
```

```

Starting restore at 30-NOV-17
using channel ORA_SBT_TAPE_1
using channel ORA_SBT_TAPE_2
using channel ORA_SBT_TAPE_3
using channel ORA_SBT_TAPE_4
using channel ORA_SBT_TAPE_5
using channel ORA_DISK_1

channel ORA_DISK_1: starting datafile backup set restore
channel ORA_DISK_1: specifying datafile(s) to restore from
backup set
channel ORA_DISK_1: restoring datafile 00112 to
/u02/app/oracle/oradata/ORCL/orcl_pdb1/ORCL/5F35FE374FC20D36E053
0241C40AA7AB/datafile/ol_mf_system_f20bbbnk_.dbf
channel ORA_DISK_1: restoring datafile 00113 to
/u02/app/oracle/oradata/ORCL/orcl_pdb1/ORCL/5F35FE374FC20D36E053
0241C40AA7AB/datafile/ol_mf_sysaux_f20bbbns_.dbf
channel ORA_DISK_1: restoring datafile 00114 to
/u02/app/oracle/oradata/ORCL/orcl_pdb1/ORCL/5F35FE374FC20D36E053
0241C40AA7AB/datafile/ol_mf_undotbs1_f20bbbnz_.dbf
channel ORA_DISK_1: reading from backup piece
/u03/app/oracle/fast_recovery_area/ORCL/5F35EC627C3606EEE0530241
C40A1875/backupset/2017_11_30/ol_mf_nnndf_TAG20171130T153157_f20
93ghl_.bkp
channel ORA_DISK_1: piece
handle=/u03/app/oracle/fast_recovery_area/ORCL/5F35EC627C3606EEE
0530241C40A1875/backupset/2017_11_30/ol_mf_nnndf_TAG20171130T153
157_f2093ghl_.bkp tag=TAG20171130T153157
channel ORA_DISK_1: restored backup piece 1
channel ORA_DISK_1: restore complete, elapsed time: 00:00:15
Finished restore at 30-NOV-17

Starting recover at 30-NOV-17
using channel ORA_SBT_TAPE_1
using channel ORA_SBT_TAPE_2
using channel ORA_SBT_TAPE_3
using channel ORA_SBT_TAPE_4
using channel ORA_SBT_TAPE_5
using channel ORA_DISK_1

starting media recovery

archived log for thread 1 with sequence 178 is already on disk
as file
/u03/app/oracle/fast_recovery_area/ORCL/archivelog/2017_11_30/ol
_mf_1_178_f2093xs8_.arc

```

```
archived log for thread 1 with sequence 179 is already on disk
as file
/u03/app/oracle/fast_recovery_area/ORCL/archivelog/2017_11_30/o1
_mf_1_179_f209rf05_.arc
media recovery complete, elapsed time: 00:00:00
Finished recover at 30-NOV-17

RMAN>
```

- c. In *Session1*, open ORCL_PDB1.

```
SQL> ALTER PLUGGABLE DATABASE orcl_pdb1 OPEN;
ALTER PLUGGABLE DATABASE ORCL_PDB1 open
*
ERROR at line 1:
ORA-01122: database file 114 failed verification check
ORA-01110: data file 114:
'/u02/app/oracle/oradata/ORCL/orcl_pdb1/ORCL/5F35FE374FC20D36E05
3024
1C40AA7AB/datafile/o1_mf_undotbs1_f20bbbnz_.dbf'
ORA-01204: file number is 111 rather than 114 - wrong file

SQL>
```

- d. Because the PDB is in LOCAL UNDO mode, you have to recover the UNDO tablespace. In *Session2*, perform the recovery operation for the PDB.

```
RMAN> RECOVER PLUGGABLE DATABASE orcl_pdb1;

Starting recover at 30-NOV-17
using channel ORA_SBT_TAPE_1
using channel ORA_SBT_TAPE_2
using channel ORA_SBT_TAPE_3
using channel ORA_SBT_TAPE_4
using channel ORA_SBT_TAPE_5
using channel ORA_DISK_1

starting media recovery
media recovery complete, elapsed time: 00:00:01

Finished recover at 30-NOV-17

RMAN> EXIT
$
```

- e. Back in *Session1*, open ORCL_PDB1.

```
SQL> ALTER PLUGGABLE DATABASE orcl_pdb1 OPEN;
```

```
Pluggable database altered.
```

```
SQL>
```

- f. Verify the content of the HR.DEPARTMENTS table in ORCL_PDB1.

```
SQL> CONNECT hr@ORCL_PDB1
```

```
Enter password: password
```

```
Connected.
```

```
SQL> SELECT * FROM hr.departments;
```

DEPARTMENT_ID	DEPARTMENT_NAME	MANAGER_ID	LOCATION_ID
10	Administration	200	1700
20	Marketing	201	1800
30	Purchasing	114	1700
40	Human Resources	203	2400
50	Shipping	121	1500
60	IT	103	1400
70	Public Relations	204	2700
80	Sales	145	2500
90	Executive	100	1700
100	Finance	108	1700
110	Accounting	205	1700
120	Treasury		1700
130	Corporate Tax		1700
140	Control And Credit		1700
150	Shareholder Services		1700
160	Benefits		1700
170	Manufacturing		1700
180	Construction		1700
190	Contracting		1700
200	Operations		1700
210	IT Support		1700
220	NOC		1700
230	IT Helpdesk		1700
240	Government Sales		1700
250	Retail Sales		1700
260	Recruiting		1700
270	Payroll		1700

```
27 rows selected.
```

```
SQL> EXIT
```

```
$
```

7. Execute the `$HOME/labs/RMAN/cleanup_preplugin_PDBs.sh` shell script to clean up preplugin backups and drop the PDBs.

```
$ $HOME/labs/RMAN/cleanup_preplugin_PDBs.sh
```

```
...
```

```
$
```

Practice 4-3: Duplicating a PDB into an Existing CDB

Overview

In this practice, you clone a PDB into an existing open read/write CDB by using the RMAN `Duplicate` command without the use of a fresh auxiliary instance.

Tasks

1. You use the Oracle By Example “Duplicating Active PDBs” to see how to proceed. Launch a browser and click the bookmark from the Bookmarks toolbar of the browser. The URL is `file:///home/oracle/labs/OBE/duplicating_active_pdb/duplicating_active_pdb.html`

Practice 4-4: Duplicating an On-Premises CDB for Cloud

Overview

In this practice, you duplicate an on-premises CDB with encrypted tablespaces to be compatible on cloud. Duplication as encrypted or decrypted aids with the transition from on-premises databases to cloud and vice versa.

Tasks

1. You use the Oracle By Example “Duplicating CDB as Encrypted” to see how to proceed. Launch a browser and click the bookmark from the Bookmarks toolbar of the browser. The URL is
`file:///home/oracle/labs/OBE/duplicating_cdb_as_encrypted/duplicating_cdb_as_encrypted.html`

Practices for Lesson 5:
Using General Database Enhancements

Practices for Lesson 5: Overview

Overview

In these practices, you use private temporary tables, use Data Pump import with a new option to load corrupted files, convert partitioned tables online, and use batched DDL.

Practice 5-1: Managing Private Temporary Tables

Overview

In this practice, you create and use private temporary tables in your session.

Tasks

1. Before starting the practice, execute the `$HOME/labs/admin/glogin_5.sh` shell script. It sets formatting for all columns selected in queries.

```
$ $HOME/labs/admin/glogin_5.sh
...
$
```

2. Re-create PDB1 for the practice by executing the `$HOME/labs/DB/create_PDB1_HR.sh` shell script. The shell script also creates the HR schema in PDB1.

```
$ $HOME/labs/DB/create_PDB1_HR.sh
...
$
```

3. Create a private temporary table (PTT).
 - a. Log in to PDB1 and start a session as HR. This is *Session1*.

```
$ sqlplus hr@PDB1

Enter password: password

SQL>
```

- b. Create a simple table with some data values. The TEST table is available only for transactions committed besides the PTT.

```
SQL> CREATE TABLE test (x NUMBER, y VARCHAR2(10));

Table created.

SQL> INSERT INTO test VALUES (1,'A');

1 row created.

SQL> COMMIT;

Commit complete.

SQL>
```

- c. Create a PTT.

```
SQL> CREATE PRIVATE TEMPORARY TABLE mine
      (x NUMBER, y VARCHAR2(10));
CREATE PRIVATE TEMPORARY TABLE mine
*
ERROR at line 1:
ORA-00903: invalid table name

SQL>
```

Q/ Why does the attempt fail?

A/ Check the prefix defined by default.

```
SQL> CONNECT / AS SYSDBA
Connected.
SQL> SHOW PARAMETER PRIVATE_TEMP_TABLE_PREFIX

NAME                                TYPE        VALUE
-----
private_temp_table_prefix           string      ORA$PTT_

SQL>
```

- d. Create a PTT with the appropriate prefix.

```
SQL> CONNECT hr@PDB1
Enter password: password
Connected.
SQL>
SQL> CREATE PRIVATE TEMPORARY TABLE ora$ptt_mine
      (x NUMBER, y VARCHAR2(10));
2
Table created.

SQL>
```

- e. Insert rows in the PTT.

```
SQL> INSERT INTO ora$ptt_mine VALUES (1,'Work1');

1 row created.

SQL>
```

- f. Display data from the PTT.

```
SQL> SELECT * FROM ora$ptt_mine;

          X Y
-----
          1 Work1

SQL>
```

4. Identify PTTs between sessions.

- a. In *Session1*, describe the `ORA$PTT_mine` PTT.

```
SQL> DESC hr.ORA$PTT_mine

Name                               Null?    Type
-----
X                                   NUMBER
Y                                   VARCHAR2(10)

SQL>
```

- b. Find all information related to the PTT.

```
SQL> SELECT sid, serial#, table_name, tablespace_name,
           duration
        FROM user_private_temp_tables;

 2      3
SID SERIAL# TABLE_NAME      TABLESPACE_NAME DURATION
-----
271   58468 ORA$PTT_MINE      TEMP              TRANSACTION

SQL>
```

Observe that the PTT is of `TRANSACTION` type. This is the default duration type. This means that the PTT is automatically dropped at the end of the transaction in which the PTT has been created.

- c. Log in to `PDB1` in another terminal window and start a session as `HR`. This is *Session2*. Verify that the PTT created by *Session1* is not visible to *Session2*.

```
$ sqlplus hr@PDB1

Enter password: password

SQL> DESC hr.ORA$PTT_mine
ERROR:
```



```
ORA-04043: object hr.ORA$PTT_mine does not exist

SQL> EXIT
$
```

5. Use PTTs.

- a. In *Session1*, insert a row into `TEST` and roll back the transaction.

```
SQL> INSERT INTO test VALUES (2,'B');

1 row created.

SQL> ROLLBACK;

Rollback complete.

SQL>
```

Q/ What happens to the PTT?

```
SQL> SELECT sid, serial#, table_name, tablespace_name,
           duration
        FROM user_private_temp_tables;

 2      3
no rows selected

SQL>
```

A/ Observe that the TRANSACTION duration type PTT is automatically dropped with the ROLLBACK statement. A COMMIT would have also dropped the PTT.

- b. Create a new PTT of `SESSION` duration type that will last until your session ends.

```
SQL> CREATE PRIVATE TEMPORARY TABLE ora$ptt_mine2
           (x NUMBER, y VARCHAR2(10))
           ON COMMIT PRESERVE DEFINITION;

 2      3
Table created.

SQL>
```

- c. Find all information related to the PTT.

```
SQL> SELECT sid, serial#, table_name, tablespace_name,
           duration
        FROM user_private_temp_tables;

 2      3
```

SID	SERIAL#	TABLE_NAME	TABLESPACE_NAME	DURATION
271	58468	ORA\$PTT_MINE2	TEMP	SESSION

SQL>

- d. Insert rows in the PTT.

```
SQL> INSERT INTO ora$ptt_mine2 VALUES (2,'Work2');

1 row created.

SQL>
```

- e. Display data from the PTT.

```
SQL> SELECT * FROM ora$ptt_mine2;

      X Y
-----
      2 Work2

SQL>
```

- f. Insert a row into TEST and commit the transaction.

```
SQL> INSERT INTO test VALUES (3,'C');

1 row created.

SQL> COMMIT;

Commit complete.

SQL>
```

Q/ What happens to the PTT?

```
SQL> SELECT sid, serial#, table_name, tablespace_name,
           duration
        FROM user_private_temp_tables;

 2      3
SID SERIAL# TABLE_NAME     TABLESPACE_NAME DURATION
-----
271   58468 ORA$PTT_MINE2    TEMP              SESSION

SQL>
```

A/ The PTT still exists because it will be dropped at the end of the session.

- g. Display data from the PTT.

```
SQL> SELECT * FROM ora$ptt_mine2;

          X Y
-----
          2 Work2

SQL>
```

- h. Quit Session1.

```
SQL> EXIT
$
```

- i. Verify that the PTT is automatically dropped.

```
$ sqlplus hr@PDB1

Enter password: password

SQL> SELECT sid, serial#, table_name, tablespace_name,
          duration
          FROM user_private_temp_tables;
 2      3
no rows selected

SQL>
```

6. Use PTTs with save points.

- a. Create the EMP table with EMPLOYEES rows.

```
SQL> DROP TABLE emp PURGE;
DROP TABLE emp PURGE
      *
ERROR at line 1:
ORA-00942: table or view does not exist

SQL> CREATE TABLE emp AS SELECT * FROM employees;

Table created.

SQL>
```

- b. Create a PTT with the 107 rows of EMP.

```
SQL> CREATE PRIVATE TEMPORARY TABLE ora$ptt_emp
      AS SELECT * FROM emp;

2
Table created.

SQL>
```

- c. Insert another set of rows.

```
SQL> INSERT INTO ora$ptt_emp SELECT * FROM emp;

107 rows created.

SQL>
```

- d. Create a first save point.

```
SQL> SAVEPOINT point1;

Savepoint created.

SQL>
```

- e. Count the number of rows in the PTT.

```
SQL> SELECT count(*) FROM ora$ptt_emp;

COUNT (*)
-----
214

SQL>
```

- f. Find all information related to the PTT.

```
SQL> SELECT sid, serial#, table_name, tablespace_name,
      duration, num_rows
      FROM user_private_temp_tables;

2    3
SID SERIAL# TABLE_NAME      TABLESPACE_NAME DURATION
NUM_ROWS
-----
-
271    28457 ORA$PTT_EMP      TEMP              TRANSACTION
107

SQL>
```

The number of rows corresponds to the number of rows in the PTT at the time of the PTT's creation.

- g. Insert another set of rows.

```
SQL> INSERT INTO ora$ptt_emp SELECT * FROM emp;

107 rows created.

SQL>
```

- h. Count the number of rows in the PTT.

```
SQL> SELECT count(*) FROM ora$ptt_emp;

COUNT (*)
-----
          321

SQL>
```

- i. Create the second save point.

```
SQL> SAVEPOINT point2;

Savepoint created.

SQL>
```

- j. Insert another set of rows.

```
SQL> INSERT INTO ora$ptt_emp SELECT * FROM emp;

107 rows created.

SQL>
```

- k. Count the number of rows in the PTT.

```
SQL> SELECT count(*) FROM ora$ptt_emp;

COUNT (*)
-----
          428

SQL>
```

- l. You discover that you should keep only the first set of rows inserted in step 3. Roll back to `point1`.

```
SQL> ROLLBACK TO point1;
```

```
Rollback complete.
```

```
SQL>
```

Q/ What would have happened if you had used the `ROLLBACK` command?

A/ The PTT would have been dropped.

- m. Count the number of rows in the PTT.

```
SQL> SELECT count(*) FROM ora$ptt_emp;
```

```
  COUNT (*)
-----
         214
```

```
SQL>
```

- n. You can insert temporary rows into EMP.

```
SQL> INSERT INTO emp SELECT * FROM ORA$PTT_emp;
```

```
214 rows created.
```

```
SQL>
```

- o. Commit and count the number of rows in EMP.

```
SQL> COMMIT;
```

```
Commit complete.
```

```
SQL> SELECT count(*) FROM emp;
```

```
  COUNT (*)
-----
         321
```

```
SQL>
```

Q/ Does the PTT exist?

```
SQL> SELECT sid, serial#, table_name, tablespace_name,
           duration, num_rows
           FROM user_private_temp_tables;
```

```
2      3
```

```
no rows selected
```

```
SQL> EXIT
```

```
$
```

A/ No. The `COMMIT` statement has dropped the `TRANSACTION` duration `PTT`.

Practice 5-2: Using Data Pump Import with CONTINUE_LOAD_ON_FORMAT_ERROR

Overview

In this practice, you use the Data Pump import option to load data while detecting a data stream error. Oracle Database 19c Data Pump can resume loading at the next granule boundary instead of aborting a Data Pump import operation.

Tasks

1. Detect rows with errors and continue importing.
 - a. Before importing data with Data Pump import, execute the `$HOME/labs/DB/SCOTT.sh` shell script. The script creates the `SCOTT` schema in `PDB1`, then creates a logical directory where the dump file resides, and grants the appropriate privilege on the directory to `SCOTT`. This is completed in a terminal window that you will call *Session_Sqlplus*.

```
$ $HOME/labs/DB/SCOTT.sh
...
$
```

- b. In another terminal window that you call *Session_Import*, import as `SYSTEM` into `PDB1` corrupted data by using a corrupted dump file that loads rows into the `SCOTT.L_FACT_PAGE` table.

```
$ impdp system@PDB1 FULL=y DUMPFILE=exp_corrupt1.dmp DIRECTORY=dp_dir
TABLE_EXISTS_ACTION=replace REMAP_TABLESPACE=TBS_6:SYSTEM
Password: password
...
Master table "SYSTEM"."SYS_IMPORT_FULL_01" successfully
loaded/unloaded
import done in AL32UTF8 character set and AL16UTF16 NCHAR
character set
export done in WE8DEC character set and AL16UTF16 NCHAR
character set
Warning: possible data loss in character set conversions
Starting "SYSTEM"."SYS_IMPORT_FULL_01": system/*****@PDB1
FULL=y dumpfile=exp_corrupt1.dmp DIRECTORY=dp_dir
TABLE_EXISTS_ACTION=replace REMAP_TABLESPACE=TBS_6:SYSTEM
Processing object type TABLE_EXPORT/TABLE/TABLE
Processing object type TABLE_EXPORT/TABLE/TABLE_DATA
ORA-02374: conversion error loading table "SCOTT"."L_FACT_PAGE"
ORA-26093: input data column size (63) exceeds the maximum input
size (40)
ORA-02372: data for row: REFER_PAGE_NAME :
0X'5B5D0B4C6974657261747572652F02C102018002C102018009'
```



```

ORA-31693: Table data object "SCOTT"."L_FACT_PAGE" failed to
load/unload and is being skipped due to error:
ORA-02354: error in exporting/importing data
ORA-39840: A data load operation has detected data stream format
error 3000.

Job "SYSTEM"."SYS_IMPORT_FULL_01" completed with 1 error(s) at
Fri Dec 1 13:34:35 2017 elapsed 0 00:00:05

$

```

- c. In *Session_Sqlplus*, check if the SCOTT.L_FACT_PAGE table is partially loaded.

```

$ sqlplus scott@PDB1

Enter password: password

SQL> SELECT * FROM scott.l_fact_page;

no rows selected

SQL>

```

- d. In *Session_Import*, re-import the data into the SCOTT.L_FACT_PAGE table, skipping the corrupted records. Use the CONTINUE_LOAD_ON_FORMAT_ERROR value of the DATA_OPTIONS parameter.

```

$ impdp system@PDB1 FULL=y DUMPFILE=exp_corrupt1.dmp DIRECTORY=dp_dir
TABLE_EXISTS_ACTION=replace REMAP_TABLESPACE=TBS_6:SYSTEM
DATA_OPTIONS=CONTINUE_LOAD_ON_FORMAT_ERROR
Password: password
...
Master table "SYSTEM"."SYS_IMPORT_FULL_01" successfully
loaded/unloaded
import done in AL32UTF8 character set and AL16UTF16 NCHAR
character set
export done in WE8DEC character set and AL16UTF16 NCHAR
character set
Warning: possible data loss in character set conversions
Starting "SYSTEM"."SYS_IMPORT_FULL_01": system/*****@PDB1
FULL=y DUMPFILE=exp_corrupt1.dmp DIRECTORY=dp_dir
TABLE_EXISTS_ACTION=replace REMAP_TABLESPACE=TBS_6:SYSTEM
DATA_OPTIONS=CONTINUE_LOAD_ON_FORMAT_ERROR
Processing object type TABLE_EXPORT/TABLE/TABLE
Processing object type TABLE_EXPORT/TABLE/TABLE_DATA
ORA-02374: conversion error loading table "SCOTT"."L_FACT_PAGE"

```

```

ORA-26093: input data column size (63) exceeds the maximum input
size (40)
ORA-02372: data for row: REFER_PAGE_NAME :
0X'5B5D0B4C6974657261747572652F02C102018002C102018009'

ORA-02374: conversion error loading table "SCOTT"."L_FACT_PAGE"
ORA-39840: A data load operation has detected data stream format
error 3000.
ORA-02372: data for row: REFER_PAGE_NAME :
0X'5B5D0B4C6974657261747572652F02C102018002C102018009'

ORA-39840: A data load operation has detected data stream format
error 3000.
. . imported "SCOTT"."L_FACT_PAGE"                      18.06
MB    91538 out of 200000 rows
Job "SYSTEM"."SYS_IMPORT_FULL_01" successfully completed at Fri
Dec 1 13:38:58 2017 elapsed 0 00:00:03
$

```

Before Oracle Database 19c, a detected corruption would cause the table load to fail. With this new parameter, the Data Pump skips over the bad row and then starts loading at a known good spot in the data stream. This feature helps you in cases when the Data Pump dump file is corrupt and you have no way to recreate the dump file with the correct data. This allows you to recover at least some of the data.

- e. In *Session_Sqlplus*, check that the `SCOTT.L_FACT_PAGE` table is partially loaded with 91538 rows as the import log mentions.

```

SQL> SELECT COUNT(*) FROM scott.l_fact_page;

COUNT (*)
-----
          91538

SQL>

```

- f. Drop the table before completing another import.

```

SQL> DROP TABLE scott.l_fact_page;

Table dropped.

SQL>

```

- g. In *Session_Import*, import as SYSTEM into PDB1 by using another corrupted dump file that loads rows into the SCOTT.L_FACT_PAGE table.

```
$ impdp system@PDB1 FULL=y DUMPFILE=exp_corrupt2.dmp DIRECTORY=dp_dir
TABLE_EXISTS_ACTION=replace REMAP_TABLESPACE=TBS_6:SYSTEM
Password: password
...
ORA-39002: invalid operation
ORA-31694: master table "SYSTEM"."SYS_IMPORT_FULL_01" failed to
load/unload
ORA-02354: error in exporting/importing data
ORA-02357: header in file /home/oracle/labs/DB/exp_corrupt2.dmp
may not contain correct information
$
```

- h. In *Session_Sqlplus*, check if the SCOTT.L_FACT_PAGE table is not loaded at all, or created.

```
SQL> SELECT COUNT(*) FROM scott.l_fact_page;
SELECT COUNT(*) FROM scott.l_fact_page
                        *
ERROR at line 1:
ORA-00942: table or view does not exist

SQL>
```

Q/ Why does the import fail to load rows?

A/ The error does not rely on corruption in rows but in the header of the dump file.

Q/ Will the CONTINUE_LOAD_ON_FORMAT_ERROR value of the DATA_OPTIONS parameter be sufficient to load data?

- i. In *Session_Import*, re-import the data into the SCOTT.L_FACT_PAGE table, skipping the corrupted records. Use the CONTINUE_LOAD_ON_FORMAT_ERROR value of the DATA_OPTIONS parameter.

```
$ impdp system@PDB1 FULL=y DUMPFILE=exp_corrupt2.dmp DIRECTORY=dp_dir
TABLE_EXISTS_ACTION=replace REMAP_TABLESPACE=TBS_6:SYSTEM
DATA_OPTIONS=CONTINUE_LOAD_ON_FORMAT_ERROR
Password: password
...
ORA-39002: invalid operation
```

```
ORA-31694: master table "SYSTEM"."SYS_IMPORT_FULL_01" failed to
load/unload
ORA-02354: error in exporting/importing data
ORA-02357: header in file /home/oracle/labs/DB/exp_corrupt2.dmp
may not contain correct information

$
```

A/ The CONTINUE_LOAD_ON_FORMAT_ERROR value of the DATA_OPTIONS parameter is not sufficient to load data because the error resides in the header of the dump file, which is not readable. A corrupt header for a data pump dump file will prevent Data Pump from doing anything. The CONTINUE_LOAD_ON_FORMAT_ERROR parameter does not change the behavior. This parameter only applies when corruptions are found in the data stream for table data.

- j. In *Session_Sqlplus*, check that the SCOTT.L_FACT_PAGE table is not created.

```
SQL> SELECT COUNT(*) FROM scott.l_fact_page;
SELECT COUNT(*) FROM scott.l_fact_page
                        *
ERROR at line 1:
ORA-00942: table or view does not exist
```

2. Import using the DBMS_DATAPUMP API.

- a. Log in as SYSTEM in PDB1.

```
SQL> CONNECT system@PDB1
Enter password: password
Connected.
SQL> SET SERVEROUTPUT ON
SQL>
```

- b. Use API without the new DBMS_DATAPUMP.KU\$_DATAOPT_CONT_LD_ON_FMT_ERR argument.

```
SQL> DECLARE
    dp_handle NUMBER;
    job_status VARCHAR2(30);
    dir_name varchar2(30);
BEGIN
    dir_name := 'DP_DIR';
    dp_handle := dbms_datapump.open (operation => 'IMPORT',
    job_mode => 'FULL');
    dbms_datapump.add_file (handle => dp_handle, filename =>
    'exp_corrupt1.dmp', directory => dir_name, filetype => 1);
```

```

        dbms_datapump.add_file (handle => dp_handle, filename =>
'EXP.LOG', directory => dir_name, filetype => 3);
        dbms_datapump.set_parallel (dp_handle, 1);
        dbms_datapump.set_parameter (handle => dp_handle, name =>
'TABLE_EXISTS_ACTION', value => 'REPLACE');
        dbms_datapump.metadata_remap (handle => dp_handle, name =>
'REMAP_TABLESPACE', old_value => 'TBS_6', value => 'SYSTEM');
        dbms_datapump.start_job (dp_handle);
        dbms_datapump.wait_for_job (handle => dp_handle, job_state
=> job_status);
        dbms_output.put_line ('DataPump Export - '|| to_char
(sysdate, 'DD/MM/YYYY HH24:MI:SS')|| ' Status '||job_status);
        dbms_datapump.detach (handle => dp_handle);
END;
/
DataPump Export - 01/12/2017 14:12:36 Status COMPLETED

PL/SQL procedure successfully completed.

SQL>

```

- c. Check if the SCOTT.L_FACT_PAGE table is partially loaded.

```

SQL> SELECT COUNT(*) FROM scott.l_fact_page;

COUNT(*)
-----
0

SQL>

```

- d. Reuse API with the new DBMS_DATAPUMP.KU\$_DATAOPT_CONT_LD_ON_FMT_ERR argument.

```

SQL> DECLARE
    dp_handle NUMBER;
    job_status VARCHAR2(30);
    dir_name varchar2(30);
BEGIN
    dir_name := 'DP_DIR';
    dp_handle := dbms_datapump.open (operation => 'IMPORT',
job_mode => 'FULL');
    dbms_datapump.add_file (handle => dp_handle, filename =>
'exp_corrupt1.dmp', directory => dir_name, filetype =>
DBMS_DATAPUMP.KU$_FILE_TYPE_DUMP_FILE);
    dbms_datapump.add_file (handle => dp_handle, filename =>
'EXP.LOG', directory => dir_name, filetype =>
DBMS_DATAPUMP.KU$_FILE_TYPE_LOG_FILE);

```

```

dbms_datapump.set_parallel (dp_handle, 1);
dbms_datapump.set_parameter (handle => dp_handle, name =>
'TABLE_EXISTS_ACTION', value => 'REPLACE');
dbms_datapump.metadata_remap (handle => dp_handle, name =>
'REMAP_TABLESPACE', old_value => 'TBS_6', value => 'SYSTEM');
dbms_datapump.set_parameter (handle => dp_handle, name =>
'DATA_OPTIONS', value =>
DBMS_DATAPUMP.KU$_DATAOPT_CONT_LD_ON_FMT_ERR);
dbms_datapump.start_job (dp_handle);
dbms_datapump.wait_for_job (handle => dp_handle, job_state =>
job_status);
dbms_output.put_line ('DataPump Export - '|| to_char (sysdate,
'DD/MM/YYYY HH24:MI:SS')|| ' Status '||job_status);
dbms_datapump.detach (handle => dp_handle);
end;
/
DataPump Export - 01/12/2017 14:14:53 Status COMPLETED

PL/SQL procedure successfully completed.

SQL>

```

- e. Check that the SCOTT.L_FACT_PAGE table is partially loaded.

```

SQL> SELECT COUNT(*) FROM scott.l_fact_page;

COUNT (*)
-----
          91538

SQL> EXIT
$

```

3. Execute the \$HOME/labs/DB/cleanup_DP.sh shell script to drop the SCOTT.L_FACT_PAGE table and the dump file.

```

$ $HOME/labs/DB/cleanup_DP.sh
...
$

```

Practice 5-3: Converting HASH Partitioned Tables to RANGE Partitioned Tables Online

Overview

In this practice, you convert a `HASH` partitioned table to a `RANGE` partitioned table online.

Tasks

1. Before starting the practice, execute the `$HOME/labs/DB/part_hash.sh` shell script. The SQL script creates the `SH` and `HR` users, and grants the users the `DBA` role and `READ` and `WRITE` privileges on the `DP_PDB_ORCL` logical directory.

```
$ $HOME/labs/DB/part_hash.sh
...
$
```

2. Log in to the `PDB1` PDB (Session1) and create an `HR.EMP` `HASH` partitioned table.

```
$ sqlplus system@PDB1

Enter password: password

SQL> DROP TABLE hr.emp CASCADE CONSTRAINTS PURGE;
DROP TABLE hr.emp CASCADE CONSTRAINTS PURGE
          *
ERROR at line 1:
ORA-00942: table or view does not exist

SQL> CREATE TABLE hr.emp (empno NUMBER NOT NULL,
                           ename VARCHAR2(10), job VARCHAR2(9),
                           mgr NUMBER(4), sal NUMBER(7,2))
                           PARTITION BY HASH (sal) PARTITIONS 4;
      2      3      4
Table created.

SQL>
```

3. Insert rows into the `HR.EMP` table. Execute the `$HOME/labs/DB/insert1.sql` SQL script.

```
SQL> @$HOME/labs/DB/insert1.sql
...
SQL>
```

4. Create three indexes on the partitioned table:

- Local I1_SAL index on the SAL column
- Global unique I2_EMPNO index on the EMPNO column
- Global I3_MGR index on the MGR column

```
SQL> CREATE INDEX hr.i1_sal ON hr.emp (sal) LOCAL;

Index created.

SQL> CREATE UNIQUE INDEX hr.i2_empno ON hr.emp (empno);

Index created.

SQL> CREATE INDEX hr.i3_mgr ON hr.emp (mgr);

Index created.

SQL>
```

5. Display the partitioning method of the table.

```
SQL> SELECT owner, table_name, partitioning_type AS type,
          autolist, interval, interval_subpartition,
          autolist_subpartition
        FROM dba_part_tables WHERE owner = 'HR';
 2      3      4
OWNER  TABLE_NAME TYPE  AUTOLIST INTERVAL          INTERVAL_SUB
-----
AUTOLIST_SUB
-----
HR      EMP          HASH  NO
NO
```

```
SQL>
```

6. Display the partitions of the table.

```
SQL> SELECT composite, partition_name, high_value
        FROM dba_tab_partitions
        WHERE table_name = 'EMP' and table_owner='HR';
 2      3
COM PARTITION_NAME HIGH_VALUE
-----
NO SYS_P228
NO SYS_P229
```



```

NO SYS_P230
NO SYS_P231

SQL>

```

7. Display the list of indexes and whether they are partitioned or not.

```

SQL> SELECT index_name, partitioned
      FROM dba_indexes
      WHERE index_name LIKE 'I%' and owner='HR';

 2      3
INDEX_NAME  PAR
-----
I1_SAL      YES
I3_MGR      NO
I2_EMPNO    NO

SQL>

```

8. Display the type of partitioning of the partitioned indexes.

```

SQL> SELECT index_name, locality, partitioning_type AS type,
      autolist, interval, interval_subpartition,
      autolist_subpartition
      FROM dba_part_indexes
      WHERE owner='HR';

 2      3      4      5
INDEX_NAME  LOCALI  TYPE  AUTOLIST  INTERVAL  INTERVAL_SUB
-----
AUTOLIST_SUB
-----
I1_SAL      LOCAL  HASH  NO
NO

SQL>

```

9. Display the partitions of the index.

```

SQL> SELECT index_name, partition_name, high_value
      FROM dba_ind_partitions
      WHERE index_name IN ('I1_SAL', 'I2_EMPNO', 'I3_MGR')
      AND index_owner='HR';

 2      3      4
INDEX_NAME  PARTITION_NAME  HIGH_VALUE
-----
I1_SAL      SYS_P228

```

```

I1_SAL      SYS_P229
I1_SAL      SYS_P230
I1_SAL      SYS_P231

SQL>

```

10. In another terminal window, log in to PDB1 as HR and update the employees' salary (Session2).

```

$ sqlplus hr@PDB1

Enter password: password

SQL> UPDATE hr.emp SET mgr=208 WHERE empno=100;

1 row updated.

SQL>

```

11. In Session1, attempt to convert the HASH partitioned table to a RANGE partitioned table.

```

SQL> ALTER TABLE hr.emp MODIFY
      PARTITION BY RANGE (empno) INTERVAL (100)
      (PARTITION p1 VALUES LESS THAN (200),
      PARTITION p2 VALUES LESS THAN (500)
      )
      UPDATE INDEXES
      (hr.i1_sal LOCAL, hr.i2_empno GLOBAL PARTITION BY RANGE
      (empno)
      (PARTITION ip1 VALUES LESS THAN (MAXVALUE)));
ALTER TABLE hr.emp MODIFY
      *
ERROR at line 1:
ORA-00054: resource busy and acquire with NOWAIT specified or
timeout expired

SQL>

```

Because the `ONLINE` keyword is not added to the statement, the operation cannot get the exclusive lock.

12. Re-execute the operation with the `ONLINE` keyword.

```

SQL> ALTER TABLE hr.emp MODIFY
      PARTITION BY RANGE (empno) INTERVAL (100)
      (PARTITION p1 VALUES LESS THAN (200),
      PARTITION p2 VALUES LESS THAN (500)
      )

```

```

ONLINE
UPDATE INDEXES
  (hr.i1_sal LOCAL,
   hr.i2_empno GLOBAL PARTITION BY RANGE (empno)
   (PARTITION ip1 VALUES LESS THAN (MAXVALUE)));

```

The statement waits for the `UPDATE` in *Session2* to complete.

13. In *Session2*, commit the `UPDATE` statement.

```

SQL> COMMIT;

Commit complete.

SQL> EXIT
$

```

14. In *Session1*, the `ALTER TABLE` statement completes.

```

Table altered.

SQL>

```

15. Display the new partitioning method of the table.

```

SQL> SELECT owner, table_name, partitioning_type AS type,
          autolist, interval, interval_subpartition,
          autolist_subpartition
        FROM dba_part_tables WHERE owner = 'HR';

```

2	3	4			
OWNER	TABLE_NAME	TYPE	AUTOLIST	INTERVAL	INTERVAL_SUB
HR	EMP	RANGE	NO	100	
NO					

```

SQL>

```

16. Display the new partitions of the table.

```

SQL> SELECT composite, partition_name, high_value
        FROM dba_tab_partitions
        WHERE table_name = 'EMP' and table_owner='HR';

```

2	3	
COM	PARTITION_NAME	HIGH_VALUE

```

-----
NO   P1                200
NO   P2                500

SQL>

```

17. Display the list of indexes and whether they are partitioned or not.

```

SQL> SELECT index_name, partitioned
      FROM   dba_indexes
      WHERE  index_name LIKE 'I%' and owner='HR';

 2      3
INDEX_NAME  PAR
-----
I1_SAL      YES
I2_EMPNO    YES
I3_MGR      NO

SQL>

```

18. Display the partitions of the indexes.

```

SQL> SELECT index_name, partition_name, high_value
      FROM   dba_ind_partitions
      WHERE  index_name IN ('I1_SAL', 'I2_EMPNO', 'I3_MGR')
      AND    index_owner='HR';

 2      3      4
INDEX_NAME  PARTITION_NAME  HIGH_VALUE
-----
I1_SAL      P1                200
I1_SAL      P2                500
I2_EMPNO    IP1              MAXVALUE

SQL>

```

19. Display the type of partitioning of the partitioned indexes.

```

SQL> SELECT index_name, locality, partitioning_type AS type,
      autolist, interval, interval_subpartition,
      autolist_subpartition
      FROM   dba_part_indexes
      WHERE  owner='HR';

 2      3      4      5
INDEX_NAME  LOCALI  TYPE  AUTOLIST  INTERVAL  INTERVAL_SUB
-----
AUTOLIST_SUB
-----

```

```
I2_EMPNO    GLOBAL RANGE NO
NO

I1_SAL      LOCAL  RANGE NO      100
NO

SQL> EXIT
$
```

20. Execute the `$HOME/labs/DB/cleanup_hash.sh` shell script to drop the `HR.EMP` partitioned table and indexes.

```
$ $HOME/labs/DB/cleanup_hash.sh
...
$
```

Practice 5-4: Converting LIST Partitioned Tables to LIST AUTOMATIC Partitioned Tables Online

Overview

In this practice, you convert a LIST partitioned table, partitioned on two keys, to a LIST AUTOMATIC partitioned table, online.

Tasks

1. Create the SH.SALES_BY_REGION_AND_CHANNEL LIST partitioned table. The table is partitioned by LIST on two keys, STATE and CHANNEL.

```
$ sqlplus system@PDB1

Enter password: password

SQL> DROP TABLE sh.sales_by_region_and_channel CASCADE
CONSTRAINTS PURGE;
DROP TABLE sh.sales_by_region_and_channel CASCADE CONSTRAINTS
PURGE
          *
ERROR at line 1:
ORA-00942: table or view does not exist

SQL> CREATE TABLE sh.sales_by_region_and_channel
      (deptno number, deptname varchar2(20),
       quarterly_sales number(10, 2),
       state varchar2(2), channel varchar2(1))
      PARTITION BY LIST (state, channel)
      (PARTITION q1_northwest_direct VALUES (('OR', 'D'),
('WA', 'D')),
       PARTITION q1_northwest_indirect VALUES (('OR', 'I'),
('WA', 'I')),
       PARTITION q1_southwest_direct VALUES (('AZ', 'D'),
('UT', 'D'), ('NM', 'D')),
       PARTITION q1_ca_direct VALUES ('CA', 'D'),
       PARTITION rest VALUES (DEFAULT));
 2      3      4      5      6      7      8      9
Table created.

SQL>
```

2. Insert rows into the SH.SALES_BY_REGION_AND_CHANNEL table. Execute the \$HOME/labs/DB/insert2.sql SQL script.

```
SQL> @$HOME/labs/DB/insert2.sql
...
SQL>
```

3. Display the partitioning method of the table.

```
SQL> SELECT owner, table_name, partitioning_type AS type,
          autolist, interval, interval_subpartition,
          autolist_subpartition
        FROM dba_part_tables WHERE owner = 'SH';
```

2	3	4	
OWNER	TABLE_NAME	TYPE	AUTOLIST INTERVAL

-			
	INTERVAL_SUB	AUTOLIST_SUB	

SH	SALES_BY_REGION_AND_CHANNEL	LIST	NO
		NO	

```
SQL>
```

4. Display the partitions of the table and the high values in each partition.

```
SQL> SELECT partition_name, high_value
        FROM dba_tab_partitions
        WHERE table_owner = 'SH';
```

2	3	
PARTITION_NAME	HIGH_VALUE	

-		
Q1_CA_DIRECT	('CA', 'D')	
Q1_NORTHWEST_DIRECT	('OR', 'D'), ('WA', 'D')	
Q1_NORTHWEST_INDIRECT	('OR', 'I'), ('WA', 'I')	
Q1_SOUTHWEST_DIRECT	('AZ', 'D'), ('UT', 'D'), ('NM', 'D')	
REST	DEFAULT	

```
SQL>
```

5. In Session2, log in to PDB1 as SH and increase the quarterly sales.

```
$ sqlplus sh@PDB1
```

```
Enter password: password
```

```
SQL> UPDATE sh.sales_by_region_and_channel
      SET quarterly_sales=quarterly_sales*10;

2
7 rows updated.

SQL>
```

6. In *Session1*, attempt to convert the LIST partitioned table on two keys to a LIST AUTOMATIC partitioned table on one key.

```
SQL> ALTER TABLE sh.sales_by_region_and_channel MODIFY
      PARTITION BY LIST (state) AUTOMATIC
      (PARTITION northwest VALUES ('OR', 'WA'),
       PARTITION southwest VALUES ('AZ', 'UT', 'NM'),
       PARTITION california VALUES ('CA'),
       PARTITION rest VALUES (DEFAULT))
      ONLINE;

ALTER TABLE sh.sales_by_region_and_channel MODIFY
              *

ERROR at line 1:
ORA-14851: DEFAULT [sub]partition cannot be specified for
AUTOLIST [sub]partitioned objects.

SQL>
```

7. Re-execute the operation without the DEFAULT partition.

```
SQL> ALTER TABLE sh.sales_by_region_and_channel MODIFY
      PARTITION BY LIST (state) AUTOMATIC
      (PARTITION northwest VALUES ('OR', 'WA'),
       PARTITION southwest VALUES ('AZ', 'UT', 'NM'),
       PARTITION california VALUES ('CA'))
      ONLINE UPDATE INDEXES;
```

The statement waits for the UPDATE in *Session2* to complete.

8. In *Session2*, commit the UPDATE statement.

```
SQL> COMMIT;

Commit complete.

SQL> EXIT
$
```


9. In *Session1*, the ALTER TABLE statement completes.

```
Table altered.
```

```
SQL>
```

10. Display the new partitioning method of the table.

```
SQL> SELECT owner, table_name, partitioning_type AS type,  
          autolist, interval, interval_subpartition,  
          autolist_subpartition  
        FROM dba_part_tables WHERE owner = 'SH';
```

```
2      3      4  
OWNER  TABLE_NAME                                TYPE  AUTOLIST  INTERVAL  
-----  
-  
INTERVAL_SUB  AUTOLIST_SUB  
-----  
SH           SALES_BY_REGION_AND_CHANNEL LIST  YES  
                                         NO
```

```
SQL>
```

11. Display the partitions of the table and the high values in each partition.

```
SQL> SELECT partition_name, high_value FROM dba_tab_partitions  
       WHERE table_name = 'SALES_BY_REGION_AND_CHANNEL';
```

```
2  
PARTITION_NAME          HIGH_VALUE  
-----  
CALIFORNIA              'CA'  
NORTHWEST               'OR', 'WA'  
SOUTHWEST               'AZ', 'UT', 'NM'  
SYS_P248                'FL'  
SYS_P249                'TX'  
SYS_P250                ''
```

```
6 rows selected.
```

```
SQL> EXIT
```

```
$
```

12. Drop the `SH.SALES_BY_REGION_AND_CHANNEL` table. Execute the `$HOME/labs/DB/cleanup_list.sh` shell script.

```
$ $HOME/labs/DB/cleanup_list.sh
...
$
```

Practice 5-5: Converting LIST AUTOMATIC Partitioned Tables to Subpartitioned Tables Online

Overview

In this practice, you convert a LIST AUTOMATIC partitioned table to a LIST AUTOMATIC partitioned table with SUBPARTITIONING, online.

Tasks

1. Create the HR.T LIST AUTOMATIC partitioned table. The table is partitioned by LIST on one key, C1.

```
$ sqlplus system@PDB1

Enter password: password

SQL> DROP TABLE hr.t CASCADE CONSTRAINTS PURGE;
DROP TABLE hr.t CASCADE CONSTRAINTS PURGE
          *
ERROR at line 1:
ORA-00942: table or view does not exist

SQL> CREATE TABLE hr.t (c1 number, c2 number)
      PARTITION BY LIST (c1) AUTOMATIC
      ( PARTITION p1 values (1),
        PARTITION p2 values (2),
        PARTITION p3 values (3));
  2      3      4      5
Table created.

SQL>
```

2. Insert rows into the HR.T table. Execute the \$HOME/labs/DB/insert3.sql SQL script.

```
SQL> @$HOME/labs/DB/insert3.sql
...
SQL>
```

3. Display the partitioning method of the table.

```
SQL> SELECT owner, table_name, partitioning_type AS type,
          autolist, subpartitioning_type
      FROM dba_part_tables WHERE table_name = 'T';
  2      3
OWNER  TABLE_NAME TYPE  AUTOLIST SUBPARTIT
-----

```

```

HR      T      LIST  YES      NONE
SQL>

```

4. Display the partitions of the table and the high values in each partition.

```

SQL> SELECT partition_name, high_value FROM dba_tab_partitions
      WHERE table_name = 'T';

  2
PARTITION_NAME HIGH_VALUE
-----
P1              1
P2              2
P3              3
SQL>

```

5. In *Session2*, log in to PDB1 as HR and increase the values in the C2 column.

```

$ sqlplus hr@PDB1

Enter password: password

SQL> UPDATE hr.t SET c2=c2*10;

11 rows updated.

SQL>

```

6. In *Session1*, convert the LIST AUTOMATIC partitioned table to a LIST AUTOMATIC subpartitioned table.

```

SQL> ALTER TABLE hr.t MODIFY PARTITION BY LIST (c1) AUTOMATIC
      SUBPARTITION BY list (c2)
      SUBPARTITION TEMPLATE
      ( SUBPARTITION sp1 VALUES (1),
        SUBPARTITION sp2 VALUES (2),
        SUBPARTITION sp3 VALUES (3),
        SUBPARTITION sp_unknown VALUES (DEFAULT)
      )
      ( PARTITION p1 VALUES (1),
        PARTITION p2 VALUES (2),
        PARTITION p3 VALUES (3)
      )
      ONLINE;

```

The statement waits for the `UPDATE` in the second session to complete.

7. In *Session2*, commit the `UPDATE` statement.

```
SQL> COMMIT;

Commit complete.

SQL> EXIT
$
```

8. Back in the conversion session, the `ALTER TABLE` statement completes.

```
Table altered.

SQL>
```

9. Insert rows into `HR.T`. Execute the `$HOME/labs/DB/insert4.sql` SQL script.

```
SQL> @$HOME/labs/DB/insert4.sql
...
SQL>
```

10. Display the new subpartitioning method of the table.

```
SQL> SELECT owner, table_name, partitioning_type AS type,
          autolist, subpartitioning_type
        FROM dba_part_tables WHERE table_name = 'T';

 2      3
OWNER TABLE_NAME TYPE  AUTOLIST SUBPARTIT
-----
HR      T          LIST  YES      LIST

SQL>
```

11. Display the partitions of the table and the high values in each partition.

```
SQL> SELECT partition_name, subpartition_name, high_value
        FROM dba_tab_subpartitions
        WHERE table_name = 'T';

 2      3
PARTITION_NAME SUBPARTITION_ HIGH_VALUE
-----
P1              P1_SP1          1
P1              P1_SP2          2
P1              P1_SP3          3
P1              P1_SP_UNKNOWN DEFAULT
P2              P2_SP1          1
P2              P2_SP2          2
P2              P2_SP3          3
```

```

P2          P2_SP_UNKNOWN DEFAULT
P3          P3_SP1          1
P3          P3_SP2          2
P3          P3_SP3          3
P3          P3_SP_UNKNOWN DEFAULT

```

12 rows selected.

SQL>

12. Display values from the subpartitions.

```
SQL> SELECT * FROM hr.t SUBPARTITION (P1_SP2);
```

C1	C2
1	2

```
SQL> SELECT * FROM hr.t SUBPARTITION (P2_SP3);
```

C1	C2
2	3

```
SQL> SELECT * FROM hr.t SUBPARTITION (P2_SP_UNKNOWN);
```

C1	C2
2	10
2	20
2	20
2	30

```
SQL> SELECT * FROM hr.t SUBPARTITION (P3_SP_UNKNOWN);
```

C1	C2
3	10
3	20
3	20
3	4
3	5

```
SQL> EXIT
```

\$

13. Drop the `HR.T` table. Execute the `$HOME/labs/DB/cleanup_list_auto.sh` shell script.

```
$ $HOME/labs/DB/cleanup_list_auto.sh
...
$
```

Practice 5-6: Merging Partitions of a Partitioned Table Online

Overview

In this practice, you merge partitions of a `RANGE` partitioned table online.

Tasks

1. Before starting the practice, execute the `$HOME/labs/DB/part_merge.sh` shell script. The shell script creates a directory for the future import and the `RANGE` partitioned table `SH.SALES`, imports rows with a dump file, and creates the `SH.I2_PROMO_ID` on the `PROMO_ID` column of the `SH.SALES` table. Errors will occur during the import operation. Ignore them.

```
$ $HOME/labs/DB/part_merge.sh
...
$
```

2. Log in to `PDB1` again as `SYSTEM` and display the partitioning method of the table.

```
$ sqlplus system@PDB1

Enter password: password

SQL> SELECT owner, table_name, partitioning_type AS type,
          autolist, interval
        FROM dba_part_tables WHERE owner='SH';

 2
OWNER TABLE_NAME TYPE  AUTOLIST INTERVAL
-----
SH      SALES      RANGE NO

SQL>
```

3. Display the partitions of the table and the high values in each partition.

```
SQL> COL high_value FORMAT A44
SQL> SELECT partition_name, high_value FROM dba_tab_partitions
        WHERE table_name = 'SALES' and table_owner='SH';

 2
PARTITION_NAME HIGH_VALUE
-----
SALES_Q1_1998 TO_DATE(' 1998-04-01 00:00:00', 'SYYYY-MM-DD
                HH24:MI:SS', 'NLS_CALENDAR=GREGORIA
SALES_Q2_1998 TO_DATE(' 1998-07-01 00:00:00', 'SYYYY-MM-DD
                HH24:MI:SS', 'NLS_CALENDAR=GREGORIA
```



```

SALES_Q3_1998 TO_DATE(' 1998-10-01 00:00:00', 'SYYYY-MM-DD
                    HH24:MI:SS', 'NLS_CALEDAR=GREGORIA

SALES_Q4_1998 TO_DATE(' 1999-01-01 00:00:00', 'SYYYY-MM-DD
                    HH24:MI:SS', 'NLS_CALEDAR=GREGORIA

SALES_Q1_1999 TO_DATE(' 1999-04-01 00:00:00', 'SYYYY-MM-DD
                    HH24:MI:SS', 'NLS_CALEDAR=GREGORIA

SALES_Q2_1999 TO_DATE(' 1999-07-01 00:00:00', 'SYYYY-MM-DD
                    HH24:MI:SS', 'NLS_CALEDAR=GREGORIA

SALES_Q3_1999 TO_DATE(' 1999-10-01 00:00:00', 'SYYYY-MM-DD
                    HH24:MI:SS', 'NLS_CALEDAR=GREGORIA

SALES_Q4_1999 TO_DATE(' 2000-01-01 00:00:00', 'SYYYY-MM-DD
                    HH24:MI:SS', 'NLS_CALEDAR=GREGORIA

SALES_Q1_2000 TO_DATE(' 2000-04-01 00:00:00', 'SYYYY-MM-DD
                    HH24:MI:SS', 'NLS_CALEDAR=GREGORIA

SALES_Q2_2000 TO_DATE(' 2000-07-01 00:00:00', 'SYYYY-MM-DD
                    HH24:MI:SS', 'NLS_CALEDAR=GREGORIA

SALES_Q3_2000 TO_DATE(' 2000-10-01 00:00:00', 'SYYYY-MM-DD
                    HH24:MI:SS', 'NLS_CALEDAR=GREGORIA

SALES_Q4_2000 TO_DATE(' 2001-01-01 00:00:00', 'SYYYY-MM-DD
                    HH24:MI:SS', 'NLS_CALEDAR=GREGORIA

SALES_Q1_2001 TO_DATE(' 2001-04-01 00:00:00', 'SYYYY-MM-DD
                    HH24:MI:SS', 'NLS_CALEDAR=GREGORIA

SALES_Q2_2001 TO_DATE(' 2001-07-01 00:00:00', 'SYYYY-MM-DD
                    HH24:MI:SS', 'NLS_CALEDAR=GREGORIA

SALES_Q3_2001 TO_DATE(' 2001-10-01 00:00:00', 'SYYYY-MM-DD
                    HH24:MI:SS', 'NLS_CALEDAR=GREGORIA

SALES_Q4_2001 TO_DATE(' 2002-01-01 00:00:00', 'SYYYY-MM-DD
                    HH24:MI:SS', 'NLS_CALEDAR=GREGORIA

```

```

SALES_1995      TO_DATE(' 1996-01-01 00:00:00', 'SYYYY-MM-DD
                HH24:MI:SS', 'NLS_CALENDAR=GREGORIA

SALES_1996      TO_DATE(' 1997-01-01 00:00:00', 'SYYYY-MM-DD
                HH24:MI:SS', 'NLS_CALENDAR=GREGORIA

SALES_H1_1997   TO_DATE(' 1997-07-01 00:00:00', 'SYYYY-MM-DD
                HH24:MI:SS', 'NLS_CALENDAR=GREGORIA

SALES_H2_1997   TO_DATE(' 1998-01-01 00:00:00', 'SYYYY-MM-DD
                HH24:MI:SS', 'NLS_CALENDAR=GREGORIA

SALES_Q1_2002   TO_DATE(' 2002-04-01 00:00:00', 'SYYYY-MM-DD
                HH24:MI:SS', 'NLS_CALENDAR=GREGORIA

SALES_Q1_2003   TO_DATE(' 2003-04-01 00:00:00', 'SYYYY-MM-DD
                HH24:MI:SS', 'NLS_CALENDAR=GREGORIA

SALES_Q2_2002   TO_DATE(' 2002-07-01 00:00:00', 'SYYYY-MM-DD
                HH24:MI:SS', 'NLS_CALENDAR=GREGORIA

SALES_Q2_2003   TO_DATE(' 2003-07-01 00:00:00', 'SYYYY-MM-DD
                HH24:MI:SS', 'NLS_CALENDAR=GREGORIA

SALES_Q3_2002   TO_DATE(' 2002-10-01 00:00:00', 'SYYYY-MM-DD
                HH24:MI:SS', 'NLS_CALENDAR=GREGORIA

SALES_Q3_2003   TO_DATE(' 2003-10-01 00:00:00', 'SYYYY-MM-DD
                HH24:MI:SS', 'NLS_CALENDAR=GREGORIA

SALES_Q4_2002   TO_DATE(' 2003-01-01 00:00:00', 'SYYYY-MM-DD
                HH24:MI:SS', 'NLS_CALENDAR=GREGORIA

SALES_Q4_2003   TO_DATE(' 2004-01-01 00:00:00', 'SYYYY-MM-DD
                HH24:MI:SS', 'NLS_CALENDAR=GREGORIA

28 rows selected.

SQL>

```

4. Display the indexes of the table and check whether the indexes are partitioned or not.

```
SQL> SELECT index_name, index_type, partitioned
      FROM dba_indexes
      WHERE table_name='SALES';

2
INDEX_NAME                                INDEX_TYPE                                PAR
-----
SALES_PROD_BIX                           BITMAP                                    YES
SALES_CUST_BIX                           BITMAP                                    YES
SALES_TIME_BIX                           BITMAP                                    YES
SALES_CHANNEL_BIX                        BITMAP                                    YES
I2_PROMO_ID                              NORMAL                                   NO

SQL>
```

5. Display the type of partitioning of the partitioned indexes.

```
SQL> SELECT index_name, locality, partitioning_type AS type,
      autolist, interval
      FROM dba_part_indexes WHERE owner='SH';

2      3
INDEX_NAME                                LOCALI TYPE  AUTOLIST  INTERVAL
-----
SALES_CHANNEL_BIX  LOCAL  RANGE  NO
SALES_CUST_BIX    LOCAL  RANGE  NO
SALES_PROD_BIX    LOCAL  RANGE  NO
SALES_TIME_BIX    LOCAL  RANGE  NO

SQL>
```

6. In *Session2*, log in to PDB1 as SH and increase the quantity sold in the SALES_Q1_2000 partition.

```
$ sqlplus sh@PDB1

Enter password: password

SQL> UPDATE sh.sales PARTITION (sales_q1_2000)
      SET QUANTITY_SOLD=QUANTITY_SOLD*10;

2
62197 rows updated.

SQL>
```

7. In *Session1*, merge the partition of 2000 year to a single partition. This operation should be possible concurrently with the DML operation.

```
SQL> ALTER TABLE sh.sales
      MERGE PARTITIONS sales_q1_2000, sales_q2_2000,
```

```

                                sales_q3_2000,sales_q4_2000
                                INTO PARTITION sales_2000
                                COMPRESS UPDATE INDEXES ONLINE;
2      3      4      5

```

The statement waits for the `UPDATE` in the second session to complete.

8. In the `UPDATE` session, commit the `UPDATE` statement.

```

SQL> COMMIT;

Commit complete.

SQL> EXIT
$

```

9. In the merge session, the `ALTER TABLE` statement completes.

```

Table altered.

SQL>

```

10. Verify that the four partitions are merged into one single partition. The query of step 5 reported 28 partitions. The current query reports 25 partitions.

```

SQL> SELECT partition_name, high_value
       FROM dba_tab_partitions
       WHERE table_name = 'SALES' AND table_owner = 'SH';
2      3

```

PARTITION_NAME	HIGH_VALUE
SALES_1995	TO_DATE(' 1996-01-01 00:00:00', 'SYYYY-MM-DD HH24:MI:SS', 'NLS_CALENDAR=GREGORIA
SALES_1996	TO_DATE(' 1997-01-01 00:00:00', 'SYYYY-MM-DD HH24:MI:SS', 'NLS_CALENDAR=GREGORIA
SALES_2000	TO_DATE(' 2001-01-01 00:00:00', 'SYYYY-MM-DD HH24:MI:SS', 'NLS_CALENDAR=GREGORIA
SALES_H1_1997	TO_DATE(' 1997-07-01 00:00:00', 'SYYYY-MM-DD HH24:MI:SS', 'NLS_CALENDAR=GREGORIA
SALES_H2_1997	TO_DATE(' 1998-01-01 00:00:00', 'SYYYY-MM-DD HH24:MI:SS', 'NLS_CALENDAR=GREGORIA

```

SALES_Q1_1998 TO_DATE(' 1998-04-01 00:00:00', 'SYYYY-MM-DD
                    HH24:MI:SS', 'NLS_CALEDAR=GREGORIA

SALES_Q1_1999 TO_DATE(' 1999-04-01 00:00:00', 'SYYYY-MM-DD
                    HH24:MI:SS', 'NLS_CALEDAR=GREGORIA

SALES_Q1_2001 TO_DATE(' 2001-04-01 00:00:00', 'SYYYY-MM-DD
                    HH24:MI:SS', 'NLS_CALEDAR=GREGORIA

SALES_Q1_2002 TO_DATE(' 2002-04-01 00:00:00', 'SYYYY-MM-DD
                    HH24:MI:SS', 'NLS_CALEDAR=GREGORIA

SALES_Q1_2003 TO_DATE(' 2003-04-01 00:00:00', 'SYYYY-MM-DD
                    HH24:MI:SS', 'NLS_CALEDAR=GREGORIA

SALES_Q2_1998 TO_DATE(' 1998-07-01 00:00:00', 'SYYYY-MM-DD
                    HH24:MI:SS', 'NLS_CALEDAR=GREGORIA

SALES_Q2_1999 TO_DATE(' 1999-07-01 00:00:00', 'SYYYY-MM-DD
                    HH24:MI:SS', 'NLS_CALEDAR=GREGORIA

SALES_Q2_2001 TO_DATE(' 2001-07-01 00:00:00', 'SYYYY-MM-DD
                    HH24:MI:SS', 'NLS_CALEDAR=GREGORIA

SALES_Q2_2002 TO_DATE(' 2002-07-01 00:00:00', 'SYYYY-MM-DD
                    HH24:MI:SS', 'NLS_CALEDAR=GREGORIA

SALES_Q2_2003 TO_DATE(' 2003-07-01 00:00:00', 'SYYYY-MM-DD
                    HH24:MI:SS', 'NLS_CALEDAR=GREGORIA

SALES_Q3_1998 TO_DATE(' 1998-10-01 00:00:00', 'SYYYY-MM-DD
                    HH24:MI:SS', 'NLS_CALEDAR=GREGORIA

SALES_Q3_1999 TO_DATE(' 1999-10-01 00:00:00', 'SYYYY-MM-DD
                    HH24:MI:SS', 'NLS_CALEDAR=GREGORIA

SALES_Q3_2001 TO_DATE(' 2001-10-01 00:00:00', 'SYYYY-MM-DD
                    HH24:MI:SS', 'NLS_CALEDAR=GREGORIA

SALES_Q3_2002 TO_DATE(' 2002-10-01 00:00:00', 'SYYYY-MM-DD
                    HH24:MI:SS', 'NLS_CALEDAR=GREGORIA

```

```
SALES_Q3_2003 TO_DATE(' 2003-10-01 00:00:00', 'SYYYY-MM-DD
HH24:MI:SS', 'NLS_CALEDAR=GREGORIA

SALES_Q4_1998 TO_DATE(' 1999-01-01 00:00:00', 'SYYYY-MM-DD
HH24:MI:SS', 'NLS_CALEDAR=GREGORIA

SALES_Q4_1999 TO_DATE(' 2000-01-01 00:00:00', 'SYYYY-MM-DD
HH24:MI:SS', 'NLS_CALEDAR=GREGORIA

SALES_Q4_2001 TO_DATE(' 2002-01-01 00:00:00', 'SYYYY-MM-DD
HH24:MI:SS', 'NLS_CALEDAR=GREGORIA

SALES_Q4_2002 TO_DATE(' 2003-01-01 00:00:00', 'SYYYY-MM-DD
HH24:MI:SS', 'NLS_CALEDAR=GREGORIA

SALES_Q4_2003 TO_DATE(' 2004-01-01 00:00:00', 'SYYYY-MM-DD
HH24:MI:SS', 'NLS_CALEDAR=GREGORIA

25 rows selected.

SQL> EXIT
$
```

11. Drop the SH.SALES table. Execute the \$HOME/labs/DB/cleanup_merge.sh shell script.

```
$ $HOME/labs/DB/cleanup_merge.sh
...
$
```

Practice 5-7: Using Batched DDL

Overview

In this practice, you identify differences between two table definitions and generate a single `ALTER TABLE` statement that can be executed for one of the two tables to make it look like the other table.

Tasks

1. Before starting the practice, execute the `$HOME/labs/DB/DDL.sh` shell script. The SQL script recreates `PDB1` and the `SCOTT` schema with two tables, `TEST1` and `TEST2`.

```
$ $HOME/labs/DB/DDL.sh
...
$
```

2. Use the existing `DBMS_METADATA.SET_TRANSFORM_PARAM` procedure to compare the two tables without the new `BATCHED_ALTER_DDL` name of the transform parameter.

```
$ sqlplus SCOTT@PDB1

Enter password: password

SQL>
SQL> SET SERVEROUTPUT ON
```

```
DECLARE
h          NUMBER;    -- Open handle
och        NUMBER;    -- openc handle
th         NUMBER;    -- transform handle
clob1      CLOB;      -- object1 sxml doc
clob2      CLOB;      -- object2 sxml doc
difdoc     CLOB;      -- difference doc
altxml     CLOB;      -- alter xmldoc
altddl     CLOB;      -- alter ddl doc

BEGIN
-- Fetch SXML metadata for the first doc
h      := dbms_metadata.open('TABLE');
dbms_metadata.set_filter(h, 'NAME', 'TEST1');
th     := dbms_metadata.add_transform(h, 'SXML');
clob1  := dbms_metadata.fetch_clob(h);
dbms_metadata.close(h);

-- Fetch SXML metadata for the second doc
h      := dbms_metadata.open('TABLE');
dbms_metadata.set_filter(h, 'NAME', 'TEST2');
```

```

th      := dbms_metadata.add_transform(h, 'SXML');
clob2 := dbms_metadata.fetch_clob(h);
dbms_metadata.close(h);

-- Compare the two objects and generate an SXML difference document
och := dbms_metadata_diff.openc('TABLE');
dbms_metadata_diff.add_document(och, clob1);
dbms_metadata_diff.add_document(och, clob2);
difdoc := dbms_metadata_diff.fetch_clob(och);
dbms_metadata_diff.close(och);

-- Using the difdoc get the ALTER xmldoc
h := dbms_metadata.openw('TABLE');
th := dbms_metadata.add_transform(h, 'ALTERXML');
dbms_metadata.set_parse_item(h, 'XPATH');
dbms_metadata.set_parse_item(h, 'ALTERABLE');
dbms_metadata.set_parse_item(h, 'CLAUSE_TYPE');
dbms_metadata.set_parse_item(h, 'NAME');
dbms_metadata.set_parse_item(h, 'COLUMN_ATTRIBUTE');
dbms_metadata.set_parse_item(h, 'CONSTRAINT_TYPE');
dbms_metadata.set_parse_item(h, 'CONSTRAINT_STATE');
dbms_metadata.set_parse_item(h, 'PARTITION_TYPE');

-- Get the alter xmldoc
dbms_lob.createtemporary(altxml, TRUE);
dbms_metadata.convert (h, difdoc, altxml);
dbms_metadata.close(h);

-- convert the alter xmldoc to ddl
h := dbms_metadata.openw('TABLE');
th := dbms_metadata.add_transform(h, 'ALTERDDL');
dbms_lob.createtemporary(altdddl, TRUE);
dbms_metadata.convert (h, altxml, altdddl);

-- display the ddl resulting from this compare
dbms_output.put_line(altdddl);

END;
/

```

```

SQL> DECLARE
  2  h          NUMBER;    -- Open handle
  3  och        NUMBER;    -- openc handle
  4  th         NUMBER;    -- transform handle
  5  clob1      CLOB;      -- object1 sxml doc

```



```

6  clob2      CLOB;      -- object2 sxml doc
7  difdoc     CLOB;      -- difference doc
8  altxml     CLOB;      -- alter xmldoc
9  altddl     CLOB;      -- alter ddl doc
10
11 BEGIN
12 -- Fetch SXML metadata for the first doc
13 h      := dbms_metadata.open('TABLE');
14 dbms_metadata.set_filter(h, 'NAME', 'TEST1');
15 th     := dbms_metadata.add_transform(h, 'SXML');
16 clob1  := dbms_metadata.fetch_clob(h);
17 dbms_metadata.close(h);
18
19 -- Fetch SXML metadata for the second doc
20 h      := dbms_metadata.open('TABLE');
21 dbms_metadata.set_filter(h, 'NAME', 'TEST2');
22 th     := dbms_metadata.add_transform(h, 'SXML');
23 clob2  := dbms_metadata.fetch_clob(h);
24 dbms_metadata.close(h);
25
26 -- Compare the two objects and generate an SXML difference
document
27 och := dbms_metadata_diff.openc('TABLE');
28 dbms_metadata_diff.add_document(och, clob1);
29 dbms_metadata_diff.add_document(och, clob2);
30 difdoc := dbms_metadata_diff.fetch_clob(och);
31 dbms_metadata_diff.close(och);
32
33 -- Using the difdoc get the ALTER xmldoc
34 h := dbms_metadata.openw('TABLE');
35 th := dbms_metadata.add_transform(h, 'ALTERXML');
36 dbms_metadata.set_parse_item(h, 'XPATH');
37 dbms_metadata.set_parse_item(h, 'ALTERABLE');
38 dbms_metadata.set_parse_item(h, 'CLAUSE_TYPE');
39 dbms_metadata.set_parse_item(h, 'NAME');
40 dbms_metadata.set_parse_item(h, 'COLUMN_ATTRIBUTE');
41 dbms_metadata.set_parse_item(h, 'CONSTRAINT_TYPE');
42 dbms_metadata.set_parse_item(h, 'CONSTRAINT_STATE');
43 dbms_metadata.set_parse_item(h, 'PARTITION_TYPE');
44
45 -- Get the alter xmldoc
46 dbms_lob.createtemporary(altxml, TRUE);
47 dbms_metadata.convert (h, difdoc, altxml);

```

```

48  dbms_metadata.close(h) ;
49
50  -- convert the alter xmldoc to ddl
51  h := dbms_metadata.openw('TABLE') ;
52  th := dbms_metadata.add_transform(h, 'ALTERDDL') ;
53  dbms_lob.createtemporary(altdddl, TRUE) ;
54  dbms_metadata.convert (h, altxml, altdddl) ;
55
56  -- display the ddl resulting from this compare
57  dbms_output.put_line(altdddl) ;
58
59  END;
60  /
ALTER TABLE "SCOTT"."TEST1" ADD ("T" VARCHAR2(30))
      ALTER TABLE "SCOTT"."TEST1" MODIFY ("Y" VARCHAR2(40))
      ALTER TABLE "SCOTT"."TEST1" RENAME TO "TEST2"

PL/SQL procedure successfully completed.

SQL>

```

Observe that there are three `ALTER TABLE` commands for the three differences detected. The reason is that the three clauses required to make the first `SCOTT.TEST1` table look like the second `SCOTT.TEST2` table—`ADD`, `MODIFY`, and `RENAME`—cannot be combined within a single `ALTER TABLE` statement. In this case, Oracle Database 12c would anyway display the same result.

3. Re-create the first `SCOTT.TEST2` table and compare.

```

SQL> DROP TABLE scott.test2;

Table dropped.

SQL> CREATE TABLE scott.test2 (x NUMBER) ;

Table created.

SQL>

```

4. Use the same `DBMS_METADATA.SET_TRANSFORM_PARAM` procedure to compare the two tables without the new `BATCHED_ALTER_DDL` name of the transform parameter.

```
DECLARE
h          NUMBER;    -- Open handle
och        NUMBER;    -- openc handle
th         NUMBER;    -- transform handle
clob1      CLOB;      -- object1 sxml doc
clob2      CLOB;      -- object2 sxml doc
difdoc     CLOB;      -- difference doc
altxml     CLOB;      -- alter xmldoc
altddl     CLOB;      -- alter ddl doc

BEGIN
-- Fetch SXML metadata for the first doc
h      := dbms_metadata.open('TABLE');
dbms_metadata.set_filter(h, 'NAME', 'TEST1');
th     := dbms_metadata.add_transform(h, 'SXML');
clob1  := dbms_metadata.fetch_clob(h);
dbms_metadata.close(h);

-- Fetch SXML metadata for the second doc
h      := dbms_metadata.open('TABLE');
dbms_metadata.set_filter(h, 'NAME', 'TEST2');
th     := dbms_metadata.add_transform(h, 'SXML');
clob2  := dbms_metadata.fetch_clob(h);
dbms_metadata.close(h);

-- Compare the two objects and generate an SXML difference document
och := dbms_metadata_diff.openc('TABLE');
dbms_metadata_diff.add_document(och, clob1);
dbms_metadata_diff.add_document(och, clob2);
difdoc := dbms_metadata_diff.fetch_clob(och);
dbms_metadata_diff.close(och);

-- Using the difdoc get the ALTER xmldoc
h := dbms_metadata.openw('TABLE');
th := dbms_metadata.add_transform(h, 'ALTERXML');
dbms_metadata.set_parse_item(h, 'XPATH');
dbms_metadata.set_parse_item(h, 'ALTERABLE');
dbms_metadata.set_parse_item(h, 'CLAUSE_TYPE');
dbms_metadata.set_parse_item(h, 'NAME');
dbms_metadata.set_parse_item(h, 'COLUMN_ATTRIBUTE');
dbms_metadata.set_parse_item(h, 'CONSTRAINT_TYPE');
dbms_metadata.set_parse_item(h, 'CONSTRAINT_STATE');
dbms_metadata.set_parse_item(h, 'PARTITION_TYPE');
```

```

-- Get the alter xmldoc
dbms_lob.createtemporary(altxml, TRUE);
dbms_metadata.convert (h, difdoc, altxml);
dbms_metadata.close(h);

-- convert the alter xmldoc to ddl
h := dbms_metadata.openw('TABLE');
th := dbms_metadata.add_transform(h, 'ALTERDDL');
dbms_lob.createtemporary(altdll, TRUE);
dbms_metadata.convert (h, altxml, altdll);

-- display the ddl resulting from this compare
dbms_output.put_line(altdll);

END;
/

```

```

SQL> DECLARE
  2  h          NUMBER;    -- Open handle
  3  och        NUMBER;    -- openc handle
  4  th         NUMBER;    -- transform handle
  5  clob1      CLOB;      -- object1 sxml doc
  6  clob2      CLOB;      -- object2 sxml doc
  7  difdoc     CLOB;      -- difference doc
  8  altxml     CLOB;      -- alter xmldoc
  9  altdll     CLOB;      -- alter ddl doc
 10
 11 BEGIN
 12  -- Fetch SXML metadata for the first doc
 13  h      := dbms_metadata.open('TABLE');
 14  dbms_metadata.set_filter(h, 'NAME', 'TEST1');
 15  th     := dbms_metadata.add_transform(h, 'SXML');
 16  clob1 := dbms_metadata.fetch_clob(h);
 17  dbms_metadata.close(h);
 18
 19  -- Fetch SXML metadata for the second doc
 20  h      := dbms_metadata.open('TABLE');
 21  dbms_metadata.set_filter(h, 'NAME', 'TEST2');
 22  th     := dbms_metadata.add_transform(h, 'SXML');
 23  clob2 := dbms_metadata.fetch_clob(h);
 24  dbms_metadata.close(h);
 25

```

```

26 -- Compare the two objects and generate an SXML difference
document
27 och := dbms_metadata_diff.openc('TABLE');
28 dbms_metadata_diff.add_document(och, clob1);
29 dbms_metadata_diff.add_document(och, clob2);
30 difdoc := dbms_metadata_diff.fetch_clob(och);
31 dbms_metadata_diff.close(och);
32
33 -- Using the difdoc get the ALTER xmldoc
h := dbms_metadata.openw('TABLE');
34 35 th := dbms_metadata.add_transform(h, 'ALTERXML');
36 dbms_metadata.set_parse_item(h, 'XPATH');
37 dbms_metadata.set_parse_item(h, 'ALTERABLE');
38 dbms_metadata.set_parse_item(h, 'CLAUSE_TYPE');
39 dbms_metadata.set_parse_item(h, 'NAME');
40 dbms_metadata.set_parse_item(h, 'COLUMN_ATTRIBUTE');
41 dbms_metadata.set_parse_item(h, 'CONSTRAINT_TYPE');
42 dbms_metadata.set_parse_item(h, 'CONSTRAINT_STATE');
43 dbms_metadata.set_parse_item(h, 'PARTITION_TYPE');
44
45 -- Get the alter xmldoc
46 dbms_lob.createtemporary(altxml, TRUE);
47 dbms_metadata.convert (h, difdoc, altxml);
48 dbms_metadata.close(h);
49
50 -- convert the alter xmldoc to ddl
51 h := dbms_metadata.openw('TABLE');
52 th := dbms_metadata.add_transform(h, 'ALTERDDL');
53 dbms_lob.createtemporary(altdddl, TRUE);
54 dbms_metadata.convert (h, altxml, altdddl);
55
56 -- display the ddl resulting from this compare
57 dbms_output.put_line(altdddl);
58 END;
59 /
ALTER TABLE "SCOTT"."TEST1" DROP ("Y")
ALTER TABLE "SCOTT"."TEST1" DROP ("Z")
ALTER TABLE "SCOTT"."TEST1" RENAME TO "TEST2"

PL/SQL procedure successfully completed.

SQL>

```

Observe that there is now one single `ALTER TABLE` command for the first two differences detected. The reason is that these two differences correspond to the same `DROP` clause required to make the first `SCOTT.TEST1` table look like the second `SCOTT.TEST2` table.

In this case, Oracle Database 12c would display two `ALTER TABLE` commands for the first two differences detected:

```
ALTER TABLE "SCOTT"."TEST1" DROP ("Y")
ALTER TABLE "SCOTT"."TEST1" DROP ("Z")
```

5. Use the same `DBMS_METADATA.SET_TRANSFORM_PARAM` procedure to compare the two tables with the new `BATCHED_ALTER_DDL` name of the transform parameter.

```
DECLARE
h          NUMBER;    -- Open handle
och        NUMBER;    -- openc handle
th         NUMBER;    -- transform handle
clob1      CLOB;      -- object1 sxml doc
clob2      CLOB;      -- object2 sxml doc
difdoc     CLOB;      -- difference doc
altxml     CLOB;      -- alter xmldoc
altddl     CLOB;      -- alter ddl doc

BEGIN
-- Fetch SXML metadata for the first doc
h      := dbms_metadata.open('TABLE');
dbms_metadata.set_filter(h, 'NAME', 'TEST1');
th     := dbms_metadata.add_transform(h, 'SXML');
clob1  := dbms_metadata.fetch_clob(h);
dbms_metadata.close(h);

-- Fetch SXML metadata for the second doc
h      := dbms_metadata.open('TABLE');
dbms_metadata.set_filter(h, 'NAME', 'TEST2');
th     := dbms_metadata.add_transform(h, 'SXML');
clob2  := dbms_metadata.fetch_clob(h);
dbms_metadata.close(h);

-- Compare the two objects and generate an SXML difference document
och := dbms_metadata_diff.openc('TABLE');
dbms_metadata_diff.add_document(och, clob1);
dbms_metadata_diff.add_document(och, clob2);
difdoc := dbms_metadata_diff.fetch_clob(och);
dbms_metadata_diff.close(och);

-- Using the difdoc get the ALTER xmldoc
h := dbms_metadata.openw('TABLE');
```

```

th := dbms_metadata.add_transform(h, 'ALTERXML');
dbms_metadata.set_parse_item(h, 'XPATH');
dbms_metadata.set_parse_item(h, 'ALTERABLE');
dbms_metadata.set_parse_item(h, 'CLAUSE_TYPE');
dbms_metadata.set_parse_item(h, 'NAME');
dbms_metadata.set_parse_item(h, 'COLUMN_ATTRIBUTE');
dbms_metadata.set_parse_item(h, 'CONSTRAINT_TYPE');
dbms_metadata.set_parse_item(h, 'CONSTRAINT_STATE');
dbms_metadata.set_parse_item(h, 'PARTITION_TYPE');

-- Get the alter xmldoc
dbms_metadata.set_transform_param(th, 'BATCH_ALTER_DDL', true);
dbms_lob.createtemporary(altxml, TRUE);
dbms_metadata.convert (h, difdoc, altxml);
dbms_metadata.close(h);

-- convert the alter xmldoc to ddl
h := dbms_metadata.openw('TABLE');
th := dbms_metadata.add_transform(h, 'ALTERDDL');
dbms_lob.createtemporary(altdddl, TRUE);
dbms_metadata.convert (h, altxml, altdddl);

-- display the ddl resulting from this compare
dbms_output.put_line(altdddl);

END;
/

```

```

SQL> DECLARE
  2  h          NUMBER;    -- Open handle
  3  och        NUMBER;    -- openc handle
  4  th         NUMBER;    -- transform handle
  5  clob1      CLOB;      -- object1 sxml doc
  6  clob2      CLOB;      -- object2 sxml doc
  7  difdoc     CLOB;      -- difference doc
  8  altxml     CLOB;      -- alter xmldoc
  9  altdddl    CLOB;      -- alter ddl doc
 10
 11 BEGIN
 12  -- Fetch SXML metadata for the first doc
 13  h      := dbms_metadata.open('TABLE');
 14  dbms_metadata.set_filter(h, 'NAME', 'TEST1');
 15  th     := dbms_metadata.add_transform(h, 'SXML');
 16  clob1  := dbms_metadata.fetch_clob(h);
 17  dbms_metadata.close(h);

```

```

18
19 -- Fetch SXML metadata for the second doc
20 h      := dbms_metadata.open('TABLE');
21 dbms_metadata.set_filter(h, 'NAME', 'TEST2');
22 th     := dbms_metadata.add_transform(h, 'SXML');
23 clob2 := dbms_metadata.fetch_clob(h);
24 dbms_metadata.close(h);
25
26 -- Compare the two objects and generate an SXML difference
document
27 och := dbms_metadata_diff.openc('TABLE');
28 dbms_metadata_diff.add_document(och, clob1);
29 dbms_metadata_diff.add_document(och, clob2);
30 difdoc := dbms_metadata_diff.fetch_clob(och);
31 dbms_metadata_diff.close(och);
32
33 -- Using the difdoc get the ALTER xmldoc
34 h := dbms_metadata.openw('TABLE');
35 th := dbms_metadata.add_transform(h, 'ALTERXML');
36 dbms_metadata.set_parse_item(h, 'XPATH');
37 dbms_metadata.set_parse_item(h, 'ALTERABLE');
38 dbms_metadata.set_parse_item(h, 'CLAUSE_TYPE');
39 dbms_metadata.set_parse_item(h, 'NAME');
40 dbms_metadata.set_parse_item(h, 'COLUMN_ATTRIBUTE');
41 dbms_metadata.set_parse_item(h, 'CONSTRAINT_TYPE');
42 dbms_metadata.set_parse_item(h, 'CONSTRAINT_STATE');
43 dbms_metadata.set_parse_item(h, 'PARTITION_TYPE');
44
45 -- Get the alter xmldoc
46 dbms_metadata.set_transform_param(th, 'BATCH_ALTER_DDL',
true);
47 dbms_lob.createtemporary(altxml, TRUE);
48 dbms_metadata.convert (h, difdoc, altxml);
49 dbms_metadata.close(h);
50
51 -- convert the alter xmldoc to ddl
52 h := dbms_metadata.openw('TABLE');
53 th := dbms_metadata.add_transform(h, 'ALTERDDL');
54 dbms_lob.createtemporary(altdddl, TRUE);
55 dbms_metadata.convert (h, altxml, altdddl);
56
57 -- display the ddl resulting from this compare
58 dbms_output.put_line(altdddl);

```



```

59
60 END;
61 /
ALTER TABLE "SCOTT"."TEST1" DROP ("Y", "Z")
    ALTER TABLE "SCOTT"."TEST1" RENAME TO "TEST2"

PL/SQL procedure successfully completed.

SQL>

```

Observe that there is now one single `ALTER TABLE` command for the first two differences detected. The reason is that these two differences correspond to the same `DROP` clause required to make the first `SCOTT.TEST1` table look like the second `SCOTT.TEST2` table.

6. Re-create the tables and compare.

```

SQL> DROP TABLE scott.test1;

Table dropped.

SQL> CREATE TABLE scott.test1 (x NUMBER);

Table created.

SQL> DROP TABLE scott.test2;

Table dropped.

SQL> CREATE TABLE scott.test2 (x NUMBER, y VARCHAR2(40), t
VARCHAR2(30), z DATE);

Table created.

SQL>

```

7. Use the `DBMS_METADATA.SET_TRANSFORM_PARAM` procedure to compare the two tables with the new `BATCHED_ALTER_DDL` name of the transform parameter.

```

DECLARE
h          NUMBER;    -- Open handle
och        NUMBER;    -- openc handle
th         NUMBER;    -- transform handle
clob1      CLOB;       -- object1 sxml doc

```

```

clob2      CLOB;      -- object2 sxml doc
difdoc     CLOB;      -- difference doc
altxml     CLOB;      -- alter xmldoc
altddl     CLOB;      -- alter ddl doc

BEGIN
-- Fetch SXML metadata for the first doc
h      := dbms_metadata.open('TABLE');
dbms_metadata.set_filter(h, 'NAME', 'TEST1');
th     := dbms_metadata.add_transform(h, 'SXML');
clob1  := dbms_metadata.fetch_clob(h);
dbms_metadata.close(h);

-- Fetch SXML metadata for the second doc
h      := dbms_metadata.open('TABLE');
dbms_metadata.set_filter(h, 'NAME', 'TEST2');
th     := dbms_metadata.add_transform(h, 'SXML');
clob2  := dbms_metadata.fetch_clob(h);
dbms_metadata.close(h);

-- Compare the two objects and generate an SXML difference document
och := dbms_metadata_diff.openc('TABLE');
dbms_metadata_diff.add_document(och, clob1);
dbms_metadata_diff.add_document(och, clob2);
difdoc := dbms_metadata_diff.fetch_clob(och);
dbms_metadata_diff.close(och);

-- Using the difdoc get the ALTER xmldoc
h := dbms_metadata.openw('TABLE');
th := dbms_metadata.add_transform(h, 'ALTERXML');
dbms_metadata.set_parse_item(h, 'XPATH');
dbms_metadata.set_parse_item(h, 'ALTERABLE');
dbms_metadata.set_parse_item(h, 'CLAUSE_TYPE');
dbms_metadata.set_parse_item(h, 'NAME');
dbms_metadata.set_parse_item(h, 'COLUMN_ATTRIBUTE');
dbms_metadata.set_parse_item(h, 'CONSTRAINT_TYPE');
dbms_metadata.set_parse_item(h, 'CONSTRAINT_STATE');
dbms_metadata.set_parse_item(h, 'PARTITION_TYPE');

-- Get the alter xmldoc
dbms_metadata.set_transform_param(th, 'BATCH_ALTER_DDL', true);
dbms_lob.createtemporary(altxml, TRUE);
dbms_metadata.convert(h, difdoc, altxml);
dbms_metadata.close(h);

-- convert the alter xmldoc to ddl

```

```

h := dbms_metadata.openw('TABLE');
th := dbms_metadata.add_transform(h,'ALTERDDL');
dbms_lob.createtemporary(altdddl, TRUE);
dbms_metadata.convert (h, altxml, altdddl);

-- display the ddl resulting from this compare
dbms_output.put_line(altdddl);

END;
/

```

```

SQL> DECLARE
2  h          NUMBER;    -- Open handle
3  och        NUMBER;    -- openc handle
4  th         NUMBER;    -- transform handle
5  clob1      CLOB;      -- object1 sxml doc
6  clob2      CLOB;      -- object2 sxml doc
7  difdoc     CLOB;      -- difference doc
8  altxml     CLOB;      -- alter xmldoc
9  altdddl    CLOB;      -- alter ddl doc
10
11 BEGIN
12 -- Fetch SXML metadata for the first doc
13 h      := dbms_metadata.open('TABLE');
14 dbms_metadata.set_filter(h,'NAME', 'TEST1');
15 th     := dbms_metadata.add_transform(h,'SXML');
16 clob1 := dbms_metadata.fetch_clob(h);
17 dbms_metadata.close(h);
18
19 -- Fetch SXML metadata for the second doc
20 h      := dbms_metadata.open('TABLE');
21 dbms_metadata.set_filter(h, 'NAME', 'TEST2');
22 th     := dbms_metadata.add_transform(h, 'SXML');
23 clob2 := dbms_metadata.fetch_clob(h);
24 dbms_metadata.close(h);
25
26 -- Compare the two objects and generate an SXML difference
document
27 och := dbms_metadata_diff.openc('TABLE');
28 dbms_metadata_diff.add_document(och, clob1);
29 dbms_metadata_diff.add_document(och, clob2);
30 difdoc := dbms_metadata_diff.fetch_clob(och);
31 dbms_metadata_diff.close(och);

```

```

32
33  -- Using the difdoc get the ALTER xmldoc
34  h := dbms_metadata.openw('TABLE');
35  th := dbms_metadata.add_transform(h,'ALTERXML');
36  dbms_metadata.set_parse_item(h,'XPATH');
37  dbms_metadata.set_parse_item(h,'ALTERABLE');
38  dbms_metadata.set_parse_item(h,'CLAUSE_TYPE');
39  dbms_metadata.set_parse_item(h,'NAME');
40  dbms_metadata.set_parse_item(h,'COLUMN_ATTRIBUTE');
41  dbms_metadata.set_parse_item(h,'CONSTRAINT_TYPE');
42  dbms_metadata.set_parse_item(h,'CONSTRAINT_STATE');
43  dbms_metadata.set_parse_item(h,'PARTITION_TYPE');
44
45  -- Get the alter xmldoc
46  dbms_metadata.set_transform_param(th, 'BATCH_ALTER_DDL',
true);
47  dbms_lob.createtemporary(altxml, TRUE);
48  dbms_metadata.convert (h, difdoc, altxml);
49  dbms_metadata.close(h);
50
51  -- convert the alter xmldoc to ddl
52  h := dbms_metadata.openw('TABLE');
53  th := dbms_metadata.add_transform(h,'ALTERDDL');
54  dbms_lob.createtemporary(altdddl, TRUE);
55  dbms_metadata.convert (h, altxml, altdddl);
56
57  -- display the ddl resulting from this compare
58  dbms_output.put_line(altdddl);
59
60  END;
61  /
ALTER TABLE "SCOTT"."TEST1" ADD ("Y" VARCHAR2(40),
"T" VARCHAR2(30), "Z" DATE)
    ALTER TABLE "SCOTT"."TEST1" RENAME TO "TEST2"

PL/SQL procedure successfully completed.

SQL>

```

Observe that there is now one single `ALTER TABLE` command for the first three differences detected. The reason is that these three differences correspond to the same `ADD` clause required to make the first `SCOTT.TEST1` table look like the second

SCOTT.TEST2 table.

In this case, Oracle Database 12c would display three ALTER TABLE commands for the first three differences detected:

```
ALTER TABLE "SCOTT"."TEST1" ADD ("Y" VARCHAR2(40))
ALTER TABLE "SCOTT"."TEST1" ADD ("T" VARCHAR2(30))
ALTER TABLE "SCOTT"."TEST1" ADD ("Z" DATE)
```

Q/ What happens if the BATCHED_ALTER_DDL name of the transform parameter is set to FALSE or is simply not defined?

```
DECLARE
h          NUMBER;    -- Open handle
och        NUMBER;    -- openc handle
th         NUMBER;    -- transform handle
clob1      CLOB;      -- object1 sxml doc
clob2      CLOB;      -- object2 sxml doc
difdoc     CLOB;      -- difference doc
altxml     CLOB;      -- alter xmldoc
altddl     CLOB;      -- alter ddl doc

BEGIN
-- Fetch SXML metadata for the first doc
h      := dbms_metadata.open('TABLE');
dbms_metadata.set_filter(h, 'NAME', 'TEST1');
th     := dbms_metadata.add_transform(h, 'SXML');
clob1  := dbms_metadata.fetch_clob(h);
dbms_metadata.close(h);

-- Fetch SXML metadata for the second doc
h      := dbms_metadata.open('TABLE');
dbms_metadata.set_filter(h, 'NAME', 'TEST2');
th     := dbms_metadata.add_transform(h, 'SXML');
clob2  := dbms_metadata.fetch_clob(h);
dbms_metadata.close(h);

-- Compare the two objects and generate an SXML difference document
och := dbms_metadata_diff.openc('TABLE');
dbms_metadata_diff.add_document(och, clob1);
dbms_metadata_diff.add_document(och, clob2);
difdoc := dbms_metadata_diff.fetch_clob(och);
dbms_metadata_diff.close(och);

-- Using the difdoc get the ALTER xmldoc
h := dbms_metadata.openw('TABLE');
th := dbms_metadata.add_transform(h, 'ALTERXML');
```

```

dbms_metadata.set_parse_item(h,'XPATH');
dbms_metadata.set_parse_item(h,'ALTERABLE');
dbms_metadata.set_parse_item(h,'CLAUSE_TYPE');
dbms_metadata.set_parse_item(h,'NAME');
dbms_metadata.set_parse_item(h,'COLUMN_ATTRIBUTE');
dbms_metadata.set_parse_item(h,'CONSTRAINT_TYPE');
dbms_metadata.set_parse_item(h,'CONSTRAINT_STATE');
dbms_metadata.set_parse_item(h,'PARTITION_TYPE');

-- Get the alter xmldoc
dbms_metadata.set_transform_param(th, 'BATCH_ALTER_DDL', false);
dbms_lob.createtemporary(altxml, TRUE);
dbms_metadata.convert (h, difdoc, altxml);
dbms_metadata.close(h);

-- convert the alter xmldoc to ddl
h := dbms_metadata.openw('TABLE');
th := dbms_metadata.add_transform(h,'ALTERDDL');
dbms_lob.createtemporary(altddl, TRUE);
dbms_metadata.convert (h, altxml, altddl);

-- display the ddl resulting from this compare
dbms_output.put_line(altddl);

END;
/

```

```

SQL> DECLARE
  2  h          NUMBER;    -- Open handle
  3  och        NUMBER;    -- openc handle
  4  th         NUMBER;    -- transform handle
  5  clob1      CLOB;      -- object1 sxml doc
  6  clob2      CLOB;      -- object2 sxml doc
  7  difdoc     CLOB;      -- difference doc
  8  altxml     CLOB;      -- alter xmldoc
  9  altddl     CLOB;      -- alter ddl doc
 10
 11 BEGIN
 12  -- Fetch SXML metadata for the first doc
h    := dbms_metadata.open('TABLE');
 13  14 dbms_metadata.set_filter(h,'NAME', 'TEST1');
 15  th    := dbms_metadata.add_transform(h,'SXML');
 16  clob1 := dbms_metadata.fetch_clob(h);
 17  dbms_metadata.close(h);

```

```

18
19 -- Fetch SXML metadata for the second doc
20 h      := dbms_metadata.open('TABLE');
21 dbms_metadata.set_filter(h, 'NAME', 'TEST2');
22 th     := dbms_metadata.add_transform(h, 'SXML');
23 clob2 := dbms_metadata.fetch_clob(h);
24 dbms_metadata.close(h);
25
26 -- Compare the two objects and generate an SXML difference
document
27 och := dbms_metadata_diff.openc('TABLE');
28 dbms_metadata_diff.add_document(och, clob1);
29 dbms_metadata_diff.add_document(och, clob2);
30 difdoc := dbms_metadata_diff.fetch_clob(och);
31 dbms_metadata_diff.close(och);
32
33 -- Using the difdoc get the ALTER xmldoc
34 h := dbms_metadata.openw('TABLE');
35 th := dbms_metadata.add_transform(h, 'ALTERXML');
36 dbms_metadata.set_parse_item(h, 'XPATH');
37 dbms_metadata.set_parse_item(h, 'ALTERABLE');
38 dbms_metadata.set_parse_item(h, 'CLAUSE_TYPE');
39 dbms_metadata.set_parse_item(h, 'NAME');
40 dbms_metadata.set_parse_item(h, 'COLUMN_ATTRIBUTE');
41 dbms_metadata.set_parse_item(h, 'CONSTRAINT_TYPE');
42 dbms_metadata.set_parse_item(h, 'CONSTRAINT_STATE');
43 dbms_metadata.set_parse_item(h, 'PARTITION_TYPE');
44
45 -- Get the alter xmldoc
46 dbms_metadata.set_transform_param(th, 'BATCH_ALTER_DDL',
false);
47 dbms_lob.createtemporary(altxml, TRUE);
48 dbms_metadata.convert (h, difdoc, altxml);
49 dbms_metadata.close(h);
50
51 -- convert the alter xmldoc to ddl
52 h := dbms_metadata.openw('TABLE');
53 th := dbms_metadata.add_transform(h, 'ALTERDDL');
54 dbms_lob.createtemporary(altdddl, TRUE);
55 dbms_metadata.convert (h, altxml, altdddl);
56
57 -- display the ddl resulting from this compare
58 dbms_output.put_line(altdddl);

```

```
59
60 END;
61 /
ALTER TABLE "SCOTT"."TEST1" ADD ("Y" VARCHAR2(40))
ALTER TABLE "SCOTT"."TEST1" ADD ("T" VARCHAR2(30))
ALTER TABLE "SCOTT"."TEST1" ADD ("Z" DATE)
ALTER TABLE "SCOTT"."TEST1" RENAME TO "TEST2"

PL/SQL procedure successfully completed.

SQL> EXIT
$
```

A/ This is Oracle Database 12c behavior.

8. Drop the SCOTT schema account. Execute the \$HOME/labs/DB/cleanup_DDL.sh shell script.

```
$ $HOME/labs/DB/cleanup_DDL.sh
...
$
```


Practices for Lesson 6: Improving Performance

Practices for Lesson 6: Overview

Overview

In these practices, you configure and use Automatic In-Memory Management (AIM) and track in-memory expressions (IMEs) within a window capture.

Practice 6-1: Configuring and Using AIM

Overview

In this practice, you configure and use Automatic In-Memory (AIM) to avoid setting in-memory Automatic Data Optimization (ADO) policies.

Tasks

1. Before starting the practice, execute the `$HOME/labs/admin/glogin_6.sh` shell script. It sets formatting for all columns selected in queries.

```
$ $HOME/labs/admin/glogin_6.sh
...
$
```

2. Configure the in-memory column store size to 450M, and create and load in-memory tables `OE.PART`, `OE.SUPPLIER`, `OE.DATE_DIM`, `OE.CUSTOMER`, and `OE.LINEORDER`. Execute the `$HOME/labs/PERF/IM_tables.sh` shell script to complete these tasks.

```
$ $HOME/labs/PERF/IM_tables.sh
...
$
```

3. Verify that the in-memory column store size is set to 464M.

```
$ sqlplus / AS SYSDBA

SQL> SHOW PARAMETER inmemory_size

NAME                                TYPE                                VALUE
-----                                -                                -
inmemory_size                       big integer                        464M
SQL>
```

4. Verify that Automatic In-Memory (AIM) is not yet activated.

```
SQL> SHOW PARAMETER inmemory_automatic_level

NAME                                TYPE                                VALUE
-----                                -                                -
inmemory_automatic_level            string                             OFF
SQL>
```

5. Activate heat map statistics as you would when you define Automatic Data Optimization (ADO) policies that are able to evict in-memory segments from the IM column store when under pressure.

```
SQL> ALTER SYSTEM SET heat_map=ON;

System altered.
```

```
SQL>
```

6. Check whether in-memory tables are populated into the IM column store.

```
SQL> SELECT segment_name, bytes, inmemory_size,
           bytes_not_populated, populate_status
        FROM v$im_segments;

2      3
no rows selected

SQL>
```

Q/ Why are the in-memory tables not populated into the IM column store?

A/ The in-memory tables are not populated into the IM column store because AIM is not enabled and on-demand population has not been requested.

7. Perform an on-demand population of the OE.PART, OE.SUPPLIER, OE.DATE_DIM, OE.CUSTOMER, and OE.LINEORDER tables.

```
SQL> CONNECT system@PDB1
Enter password: password
Connected.
SQL> SELECT count(*) from oe.LINEORDER;

COUNT (*)
-----
13619440

SQL> SELECT count(*) from oe.PART;

COUNT (*)
-----
1600000

SQL> SELECT count(*) from oe.SUPPLIER;

COUNT (*)
-----
16000

SQL> SELECT count(*) from oe.CUSTOMER;

COUNT (*)
```

```

-----
240000

SQL> SELECT count(*) from oe.DATE_DIM;

COUNT (*)
-----
2554

SQL>

```

Q/ Are all in-memory tables populated into the IM column store?

```

SQL> SELECT segment_name, bytes, inmemory_size,
           bytes_not_populated, populate_status
        FROM   v$sql_segments;

2      3
SEGMENT_NAME          BYTES INMEMORY_SIZE BYTES_NOT_POPULATED
-----
POPULATE_STATUS
-----
LINEORDER      1583767552      174653440      1180688384
STARTED

SQL>

```

After several minutes of loading, the populating status changes.

```

SQL> SELECT segment_name, bytes, inmemory_size,
           bytes_not_populated, populate_status
        FROM   v$sql_segments;

2      3
SEGMENT_NAME          BYTES INMEMORY_SIZE BYTES_NOT_POPULATED
-----
POPULATE_STATUS
-----
LINEORDER      1583767552      337838080      828366848
OUT OF MEMORY

SQL>

```

A/ You discover that the OE.SUPPLIER table is not populated into the IM column store because there is not enough IM column store space for all in-memory segments.

8. Set the INMEMORY_AUTOMATIC_LEVEL initialization parameter to ensure that any hot segment that is not populated because of memory pressure is populated first.
Note that the AIM action will happen when the IMCO background process wakes up to act. The IMCO cycle is 2 minutes.

```
SQL> ALTER SYSTEM SET inmemory_automatic_level = MEDIUM;

System altered.

SQL> EXIT
$
```

9. Now execute the \$HOME/labs/PERF/loop.sh shell script. The script executes many SELECT statements on the OE.SUPPLIER table. This generates heat map statistics that will trigger an eviction AIM action of the cold segment, OE.LINEORDER, and populate action of other hotter segments.

Note: Because there is no ADO policy on tables allowing eviction and therefore overriding AIM, AIM can submit tasks to evict the cold segments and populate hotter segments.

```
$ $HOME/labs/PERF/loop.sh
...
$
```

10. Check that the task in charge of evicting the in-memory segment, OE.LINEORDER, and populating hotter in-memory segments is running. You will have to reiterate the query until the IMCO background process wakes up to act.

```
$ sqlplus system@PDB1

Enter password: password

SQL> SELECT object_name, action, im.status
        FROM   v$im_adotaskdetails im, dba_objects o
        WHERE  im.obj# = o.object_id;

 2      3
OBJECT_NAME      ACTION      STATUS
-----
DATE_DIM         POPULATE    DONE
CUSTOMER         POPULATE    DONE
LINEORDER        POPULATE    PROCESSING
SUPPLIER         POPULATE    DONE
PART             POPULATE    DONE

SQL>
```



```
SQL> SELECT object_name, action, im.status
      FROM v$im_adotaskdetails im, dba_objects o
      WHERE im.obj# = o.object_id;
```

```

2      3
OBJECT_NAME      ACTION              STATUS
-----
LINEORDER      PARTIAL  POPULATE  DONE
LINEORDER      PARTIAL  POPULATE  DONE
PART            POPULATE  DONE
PART            POPULATE  DONE
SUPPLIER        POPULATE  DONE
SUPPLIER        POPULATE  DONE
CUSTOMER        POPULATE  DONE
CUSTOMER        POPULATE  DONE
DATE_DIM        POPULATE  DONE
DATE_DIM        POPULATE  DONE
```

10 rows selected.

```
SQL>
```

11. Find the evicted in-memory segment and verify that the OE.LINEORDER table is now fully populated into the IM column store.

```
SQL> SELECT segment_name, bytes, inmemory_size,
      bytes_not_populated, populate_status
      FROM v$im_segments;
```

```

2      3
SEGMENT_NAME      BYTES INMEMORY_SIZE BYTES_NOT_POPULATED
-----
POPULATE_STATUS
-----
DATE_DIM          319488      1310720      0
COMPLETED

PART              166526976    35258368      0
COMPLETED

CUSTOMER          30670848     7602176      0
COMPLETED

LINEORDER         1583767552    271908864    945807360
COMPLETED
```

SUPPLIER	1974272	1310720	0
COMPLETED			

SQL>

Q/ Why are there more and more actions from the IMCO background process whereas the hottest in-memory segments have been populated?

```
SQL> SELECT im.obj#, object_name, action, im.status
        FROM v$im_adotaskdetails im, dba_objects o
        WHERE im.obj# = o.object_id;
```

2	3			
	OBJ#	OBJECT_NAME	ACTION	STATUS
	77666	LINEORDER	PARTIAL POPULATE	DONE
	77666	LINEORDER	PARTIAL POPULATE	DONE
	77666	LINEORDER	PARTIAL POPULATE	DONE
	77667	PART	POPULATE	DONE
	77667	PART	POPULATE	DONE
	77667	PART	POPULATE	DONE
	77668	SUPPLIER	POPULATE	DONE
	77668	SUPPLIER	POPULATE	DONE
	77668	SUPPLIER	POPULATE	DONE
	77669	CUSTOMER	POPULATE	DONE
	77669	CUSTOMER	POPULATE	DONE
	77669	CUSTOMER	POPULATE	DONE
	77670	DATE_DIM	POPULATE	DONE
	77670	DATE_DIM	POPULATE	DONE
	77670	DATE_DIM	POPULATE	DONE

15 rows selected.

```
SQL> SELECT object_name, action, im.status
        FROM v$im_adotaskdetails im, dba_objects o
        WHERE im.obj# = o.object_id;
```

2	3			
	OBJECT_NAME	ACTION	STATUS	
	LINEORDER	PARTIAL POPULATE	DONE	
	LINEORDER	PARTIAL POPULATE	DONE	
	LINEORDER	PARTIAL POPULATE	DONE	
	LINEORDER	PARTIAL POPULATE	DONE	

PART	POPULATE	DONE
PART	POPULATE	DONE
PART	POPULATE	DONE
PART	POPULATE	DONE
SUPPLIER	POPULATE	DONE
SUPPLIER	POPULATE	DONE
SUPPLIER	POPULATE	DONE
SUPPLIER	POPULATE	DONE
CUSTOMER	POPULATE	DONE
CUSTOMER	POPULATE	DONE
CUSTOMER	POPULATE	DONE
CUSTOMER	POPULATE	DONE
DATE_DIM	POPULATE	DONE
DATE_DIM	POPULATE	DONE
DATE_DIM	POPULATE	DONE
DATE_DIM	POPULATE	DONE

20 rows selected.

SQL> **EXIT**

\$

A/ The IMCO background process regularly checks whether there are no hotter in-memory segments to populate.

Practice 6-2: Tracking IM Expressions Within a Window Capture

Overview

In this practice, you track expressions populated into the in-memory column store within a defined capture window.

Tasks

1. Execute the `$HOME/labs/PERF/IM_tables.sh` shell script. It configures the in-memory column store size, and creates and loads in-memory tables. Even if the script was already executed in Practice 6-1, re-execute it.

```
$ $HOME/labs/PERF/IM_tables.sh
...
$
```

In case you encounter disk space issues with an Archiver error, proceed with the following commands answering YES to both:

```
$ rman target /

RMAN> DELETE ARCHIVELOG ALL;
...
RMAN> DELETE OBSOLETE;
...
RMAN> EXIT;
$
```

2. In the current session called *Session1*, create another in-memory small table.

```
$ sqlplus system@PDB1

Enter password: password

SQL> CREATE TABLE OE.test ( c1 NUMBER, c2 NUMBER,
                             vc1 AS (c1+c2), vc2 AS (c1*2))
                             INMEMORY;

      2      3
Table created.

SQL> INSERT INTO OE.test (c1, c2) VALUES (1,2);

1 row created.

SQL> INSERT INTO OE.test (c1, c2) VALUES (2,3);
```

```
1 row created.
```

```
SQL> COMMIT;
```

```
Commit complete.
```

```
SQL>
```

3. Execute the `$HOME/labs/PERF/loop2.sql` SQL script that contains queries with expressions on the `OE.LINEORDER`, `OE.DATE_DIM`, and `OE.TEST` in-memory tables.

```
SQL> @$HOME/labs/PERF/loop2.sql
```

```
...
```

```
SQL>
```

4. In another terminal window, called *Session2*, while the loop is still executing, track expressions within a capture window. Log in to `PDB1` as `OE`.

```
$ sqlplus oe@PDB1
```

```
Enter password: password
```

```
Connected.
```

```
SQL>
```

- a. In *Session2*, open a capture window to signal the beginning of an expression capture window.

```
SQL> exec DBMS_INMEMORY_ADMIN.IME_OPEN_CAPTURE_WINDOW()
```

```
PL/SQL procedure successfully completed.
```

```
SQL>
```

- b. In *Session2*, get the current capture state of the expression capture window and the time stamp of the most recent modification.

```
SQL> SET SERVEROUTPUT ON
```

```
SQL> DECLARE
```

```
    P_CAPTURE_STATE          varchar2(40);
```

```
    P_LAST_MODIFIED          timestamp;
```

```
    BEGIN
```

```
        DBMS_INMEMORY_ADMIN.IME_GET_CAPTURE_STATE(P_CAPTURE_STATE,  
P_LAST_MODIFIED);
```

```
        dbms_output.put_line('-----');
```

```
        dbms_output.put_line('State = '||P_CAPTURE_STATE);
```

```
        dbms_output.put_line('Date = '||P_LAST_MODIFIED);
```

```
        dbms_output.put_line('-----');
```

```
    END;
```

```
    /
```

```
-----  
State = OPEN
```

```
Date = 05-DEC-17 09.05.09.162892 AM
```

```
-----
```

```
PL/SQL procedure successfully completed.
```

```
SQL>
```

- c. In *Session2*, close the window to signal the end of the expression capture window.

```
SQL> exec DBMS_INMEMORY_ADMIN.IME_CLOSE_CAPTURE_WINDOW()
```

```
PL/SQL procedure successfully completed.
```

```
SQL>
```

Q/ What is the state of the capture?

```
SQL> DECLARE
```

```
    P_CAPTURE_STATE          varchar2(40);
```

```
    P_LAST_MODIFIED          timestamp;
```

```
    BEGIN
```

```
    DBMS_INMEMORY_ADMIN.IME_GET_CAPTURE_STATE(P_CAPTURE_STATE,  
P_LAST_MODIFIED);
```

```
    dbms_output.put_line('-----');
```

```
    dbms_output.put_line('State = '||P_CAPTURE_STATE);
```

```
    dbms_output.put_line('Date = '||P_LAST_MODIFIED);
```

```
    dbms_output.put_line('-----');
```

```
    END;
```

```
    /
```

```
-----
```

```
State = CLOSE
```

```
Date = 05-DEC-17 10.20.06.273889 AM
```

```
-----
```

```
PL/SQL procedure successfully completed.
```

```
SQL>
```

A/ The state displays CLOSE.

- d. Still in *Session2*, add all the hot expressions captured in the previous window into the IM column store. The database considers statistics for expressions tracked in the most recent capture window.

```
SQL> exec DBMS_INMEMORY_ADMIN.IME_CAPTURE_EXPRESSIONS('WINDOW')

PL/SQL procedure successfully completed.

SQL>
```

5. Display expression statistics.

- a. Still in *Session2*, display the statistics of expressions that have been tracked in the current window.

```
SQL> SELECT owner, table_name, evaluation_count AS COUNT,
          created, expression_text
        FROM dba_expression_statistics
        WHERE snapshot = 'WINDOW'
        AND owner='OE';
```

2	3	4	5
OWNER	TABLE_NAM	COUNT	CREATED
OE	LINEORDER	108592	05-DEC-17
	"LO_ORDERDATE"		
OE	LINEORDER	450209408	05-DEC-17
	"LO_DISCOUNT"		
OE	LINEORDER	450209408	05-DEC-17
	"LO_QUANTITY"		
OE	LINEORDER	450209408	05-DEC-17
	"LO_SHIPPRIORITY"		
OE	LINEORDER	450209408	05-DEC-17
	"LO_SHIPMODE"		
OE	LINEORDER	450209408	05-DEC-17
	"LO_EXTENDEDPRICE"		
OE	LINEORDER	450209408	05-DEC-17
	SYS_OP_BLOOM_FILTER(:BF0000,"LO_ORDERDATE")		

```

OE      LINEORDER 450209408 05-DEC-17
"LO_TAX"

OE      LINEORDER 443278176 05-DEC-17
"LO_EXTENDEDPRICE"*(1-"LO_DISCOUNT")

OE      LINEORDER 443278176 05-DEC-17
1+"LO_TAX"

OE      DATE_DIM      217184 05-DEC-17
"D_DATEKEY"

OE      DATE_DIM      108592 05-DEC-17
TO_DATE(TO_CHAR("D_DATEKEY"), 'YYYY-MM-DD')

OE      TEST          88 05-DEC-17
"C1"*2

OE      TEST          176 05-DEC-17
"C1"

OE      TEST          176 05-DEC-17
"C2"

OE      TEST          88 05-DEC-17
"C1"+"C2"

16 rows selected.

SQL>

```

Expression Statistics Store (ESS) information is stored in the data dictionary. The `DBA_EXPRESSION_STATISTICS` view shows the metadata that the optimizer has collected in the ESS.

- b. In-Memory expressions are exposed as system-generated virtual columns, prefixed by the string `SYS_IME`, in the `USER_IM_EXPRESSIONS` view. The result may differ from the result below. It depends on the frequency and statistics collected.

```

SQL> SELECT table_name, column_name, sql_expression
       FROM   user_im_expressions;

2

```



```

TABLE_NAM COLUMN_NAME
-----
SQL_EXPRESSION
-----
DATE_DIM SYS_IME00010000000B474F
TO_DATE(TO_CHAR("D_DATEKEY"), 'YYYY-MM-DD')

LINEORDER SYS_IME00010000000B4753
1+"LO_TAX"

LINEORDER SYS_IME00010000000B4754
"LO_EXTENDEDPRICE"*(1-"LO_DISCOUNT")

LINEORDER SYS_IME00010000000B475C
"LO_EXTENDEDPRICE"*(1-"LO_DISCOUNT")*(1+"LO_TAX")

LINEORDER SYS_IME0001000000099170
"LO_EXTENDEDPRICE"*(1-"LO_DISCOUNT")

SQL>

```

6. Still in *Session2*, clear the expression virtual columns. Before that, interrupt the loop in *Session1* in case it is still executing.
 - a. Clear all `SYS_IME` expression virtual columns in the database, including those tracked in the last window.

```

SQL> exec DBMS_INMEMORY_ADMIN.IME_DROP_ALL_EXPRESSIONS()

PL/SQL procedure successfully completed.

SQL>

```

- b. Check that all `SYS_IME` expression virtual columns in the database are cleared.

Q/ Are the expressions statistics that have been tracked in the current window cleared?

```

SQL> SELECT table_name, column_name, sql_expression
       FROM   user_im_expressions;

2
no rows selected

SQL>

```

A/ Yes.

Q/ Are the expressions statistics that have been tracked in the previous window cleared?

```
SQL> SELECT owner, table_name, evaluation_count AS COUNT,
          created, expression_text
        FROM dba_expression_statistics
        WHERE snapshot = 'WINDOW'
        AND    owner='OE';
```

2	3	4	5
OWNER	TABLE_NAM	COUNT	CREATED

EXPRESSION_TEXT			

OE	LINEORDER	108592	05-DEC-17
	"LO_ORDERDATE"		
OE	LINEORDER	450209408	05-DEC-17
	"LO_DISCOUNT"		
OE	LINEORDER	450209408	05-DEC-17
	"LO_QUANTITY"		
OE	LINEORDER	450209408	05-DEC-17
	"LO_SHIPRIORITY"		
OE	LINEORDER	450209408	05-DEC-17
	"LO_SHIPMODE"		
OE	LINEORDER	450209408	05-DEC-17
	"LO_EXTENDEDPRICE"		
OE	LINEORDER	450209408	05-DEC-17
	SYS_OP_BLOOM_FILTER(:BF0000,"LO_ORDERDATE")		
OE	LINEORDER	450209408	05-DEC-17
	"LO_TAX"		
OE	LINEORDER	443278176	05-DEC-17
	"LO_EXTENDEDPRICE"*(1-"LO_DISCOUNT")		
OE	LINEORDER	443278176	05-DEC-17
	1+"LO_TAX"		

```

OE      DATE_DIM      217184 05-DEC-17
"D_DATEKEY"

OE      DATE_DIM      108592 05-DEC-17
TO_DATE(TO_CHAR("D_DATEKEY"), 'YYYY-MM-DD')

OE      TEST          88 05-DEC-17
"C1"*2

OE      TEST          176 05-DEC-17
"C1"

OE      TEST          176 05-DEC-17
"C2"

OE      TEST          88 05-DEC-17
"C1"+"C2"

16 rows selected.

SQL>

```

A/ No.

Q/ What will happen when you open a new capture window?

```

SQL> exec DBMS_INMEMORY_ADMIN.IME_OPEN_CAPTURE_WINDOW()

PL/SQL procedure successfully completed.

SQL>

```

```

SQL> SELECT owner, table_name, evaluation_count AS COUNT,
          created, expression_text
        FROM dba_expression_statistics
        WHERE snapshot = 'WINDOW'
        AND owner='OE';
 2      3      4      5
no rows selected

```

```
SQL> EXIT  
$
```

A/ New expressions will be tracked and collected within the new capture window.

7. Quit Session1.

```
SQL> EXIT  
$
```

8. Execute the `$HOME/labs/PERF/cleanup_IM_tables.sh` shell script to drop in-memory tables in PDB1 and disable the IM column store usage.

```
$ $HOME/labs/PERF/cleanup_IM_tables.sh  
...  
$
```

Practices for Lesson 7:

Handling Enhancements in Big Data and Data Warehousing

Practices for Lesson 7: Overview

Overview

In these practices, you query inlined external tables, populate external tables into the IM column store, use hierarchy-based predicates and calculated measures on analytic views, and use polymorphic table functions (PTFs).

Practice 7-1: Querying Inlined External Tables

Overview

In this practice, you query inlined external tables without creating a persistent object in the data dictionary.

Tasks

1. Before starting the practice, execute the `$HOME/labs/admin/glogin_7.sh` shell script. It sets formatting for all columns selected in queries.

```
$ $HOME/labs/admin/glogin_7.sh
...
$
```

2. Execute the `$HOME/labs/DW/ext_tables.sh` shell script. The script completes the following operations:

- Creates the `TEST` user
- Grants the `TEST` user the required privileges
- Creates the `ext_dir` directory AS `'/home/oracle/labs'`
- Grants `READ` and `WRITE` privileges to `TEST`

```
$ $HOME/labs/DW/ext_tables.sh
...
$
```

3. Display the content of the `/home/oracle/labs/DW/empext1.dat` external data file.

```
$ cat /home/oracle/labs/DW/empext1.dat
114,Den,Raphaely,DRAPHEAL,515.127.4561,07-DEC-
02,PU_MAN,11000,,100,30
115,Alexander,Khoo,AKHOO,515.127.4562,18-MAY-
03,PU_CLERK,3100,,114,30
116,Shelli,Baida,SBAIDA,515.127.4563,24-DEC-
05,PU_CLERK,2900,,114,30
117,Sigal,Tobias,STOBIAS,515.127.4564,24-JUL-
05,PU_CLERK,2800,,114,30
...
199,Douglas,Grant,DGRANT,650.507.9844,13-JAN-
08,SH_CLERK,2600,,124,50
200,Jennifer,Whalen,JWHALEN,515.123.4444,17-SEP-
03,AD_ASST,4400,,101,10
201,Michael,Hartstein,MHARTSTE,515.123.5555,17-FEB-
04,MK_MAN,13000,,100,20
202,Pat,Fay,PFAY,603.123.6666,17-AUG-05,MK_REP,6000,,201,20
203,Susan,Mavris,SMAVRIS,515.123.7777,07-JUN-
02,HR_REP,6500,,101,40
$
```


4. Connect as `TEST` and query the inlined external table.

```
$ sqlplus test@PDB1

Enter password: password

SQL> SELECT ext_emp.employee_id, ext_emp.first_name,
           ext_emp.last_name
FROM EXTERNAL
  (( EMPLOYEE_ID NUMBER(6,0), FIRST_NAME VARCHAR2(20),
     LAST_NAME VARCHAR2(25), EMAIL VARCHAR2(25),
     PHONE_NUMBER VARCHAR2(20), HIRE_DATE DATE,
     JOB_ID VARCHAR2(10), SALARY NUMBER(8,2),
     COMMISSION_PCT NUMBER(2,2), MANAGER_ID NUMBER(6,0),
     DEPARTMENT_ID NUMBER(4,0))
TYPE ORACLE_LOADER
DEFAULT DIRECTORY ext_dir
ACCESS PARAMETERS
  ( records delimited by newline
    badfile ext_dir:'empxt%a_%p.bad'
    logfile ext_dir:'empxt%a_%p.log'
    fields terminated by ','
      missing field values are null
    ( employee_id, first_name, last_name, email,
      phone_number, hire_date, job_id, salary,
      commission_pct, manager_id, department_id)
  )
LOCATION ('empext1.dat')
REJECT LIMIT UNLIMITED
) ext_emp;
```

2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17
18	19	20	21	22	23	24	25								
EMPLOYEE_ID	FIRST_NAME		LAST_NAME												

	114	Den	Raphaely												
	115	Alexander	Khoo												
...															
	201	Michael	Hartstein												
...															
	202	Pat	Fay												
	203	Susan	Mavris												

2078 rows selected.

```
SQL>
```

5. Query the inlined external table and display the number of employees for each department. Order the result by department number.

```
SQL> SELECT ext_emp.department_id, count(*) AS dept_count
FROM EXTERNAL
  (( EMPLOYEE_ID NUMBER(6,0), FIRST_NAME VARCHAR2(20),
    LAST_NAME VARCHAR2(25), EMAIL VARCHAR2(25),
    PHONE_NUMBER VARCHAR2(20), HIRE_DATE DATE,
    JOB_ID VARCHAR2(10), SALARY NUMBER(8,2),
    COMMISSION_PCT NUMBER(2,2), MANAGER_ID NUMBER(6,0),
    DEPARTMENT_ID NUMBER(4,0))
TYPE ORACLE_LOADER
DEFAULT DIRECTORY ext_dir
ACCESS PARAMETERS
  ( records delimited by newline
    badfile ext_dir:'empxt%a_%p.bad'
    logfile ext_dir:'empxt%a_%p.log'
    fields terminated by ','
      missing field values are null
    ( employee_id, first_name, last_name, email,
      phone_number, hire_date, job_id, salary,
      commission_pct, manager_id, department_id)
  )
LOCATION ('empext1.dat')
REJECT LIMIT UNLIMITED
) ext_emp
GROUP BY ext_emp.department_id
ORDER BY ext_emp.department_id;
```

DEPARTMENT_ID	DEPT_COUNT
10	244
20	507
30	123
40	257
50	945
4114	1
	1

7 rows selected.

```
SQL> EXIT  
$
```

Q/ Should you drop the inlined external table?

A/ There is no table to drop in the database because there is no persistent object in the data dictionary associated to the inlined external table.

Practice 7-2: Populating External Tables in the In-Memory Column Store

Overview

In this practice, you create, populate, and then query in-memory external tables.

Tasks

1. Because the in-memory external table loads its data from the `/home/oracle/labs/DW/empext1.dat` external file and requires space in the IM column store, execute the `$HOME/labs/DW/IM_ext_tables.sh` shell script that completes the following operations:

- Configures the in-memory column store size to 750M
- Creates the `TEST` user
- Grants `TEST` the required privileges
- Creates the `ext_dir` directory
- Grants `READ` and `WRITE` privileges to `TEST`

```
$ $HOME/labs/DW/IM_ext_tables.sh
...
$
```

2. Log in to `PDB1` as `TEST` and create the in-memory external table.

```
$ sqlplus test@PDB1

Enter password: password

SQL> CREATE TABLE test.ext_emp (
    id NUMBER(6), first_name VARCHAR2(20),
    last_name VARCHAR2(25), email VARCHAR2(25),
    phone_number VARCHAR2(20), hire_date DATE,
    job_id VARCHAR2(10), salary NUMBER(8,2),
    commission_pct NUMBER(2,2), manager_id NUMBER(6),
    department_id NUMBER(4))
    ORGANIZATION EXTERNAL
    (TYPE ORACLE_LOADER DEFAULT DIRECTORY ext_dir
    ACCESS PARAMETERS
        ( records delimited by newline
          badfile ext_dir:'empxt%a_%p.bad'
          logfile ext_dir:'empxt%a_%p.log'
          fields terminated by ','
              missing field values are null
        (ID, FIRST_NAME, LAST_NAME, EMAIL, PHONE_NUMBER,
          HIRE_DATE, JOB_ID, SALARY, COMMISSION_PCT,
```

```

MANAGER_ID, DEPARTMENT_ID))
LOCATION ('empext1.dat')
)
REJECT LIMIT UNLIMITED
INMEMORY;
2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17
18 19 20 21 22
Table created.

SQL>

```

3. Verify that the external table is an in-memory object.

```

SQL> SELECT table_name, inmemory, inmemory_compression
      FROM user_external_tables;

2
TABLE_NAME          INMEMORY INMEMORY_COMPRESS
-----
EXT_EMP             ENABLED FOR QUERY LOW

SQL>

```

4. Populate in-memory external tables by counting the number of rows in the table.

```

SQL> SELECT count(*) FROM test.ext_emp;

COUNT (*)
-----
2078

SQL>

```

Q/ Is the table populated into IMCS after the query? (Connect as SYSTEM to verify)

```

SQL> CONNECT system@PDB1
Enter password: password
Connected.
SQL> SELECT owner, segment_name, populate_status, con_id
      FROM v$im_segments;

2
no rows selected

SQL>

```

A/ No.

- a. Populate the external table into the IM column store.

```
SQL> EXEC dbms_inmemory.populate ('TEST','EXT_EMP')

PL/SQL procedure successfully completed.

SQL>
```

- b. Verify that the segment is populated into the IM column store.

```
SQL> SELECT owner, segment_name, populate_status, con_id
       FROM v$im_segments;

 2
OWNER          SEGMENT_NAME POPULATE_STAT      CON_ID
-----
TEST           EXT_EMP      COMPLETED              4

SQL>
```

5. Query in-memory external tables.

- a. Log in to PDB1 as user TEST and query the in-memory external table.

```
SQL> CONNECT test@PDB1
Enter password: password
Connected.
SQL> SELECT count(*) FROM test.ext_emp;

COUNT (*)
-----
        2078

SQL>
```

- b. Check whether the in-memory external table data is queried from the IM column store or buffer cache.

```
SQL> SELECT * FROM dbms_xplan.display_cursor();

PLAN_TABLE_OUTPUT
-----
SQL_ID 7382qx3bj1x4n, child number 0
-----
SELECT count(*) FROM test.ext_emp

Plan hash value: 546827939
-----
```

Id	Operation	Name	Rows	Cost
(%CPU)	Time			
0	SELECT STATEMENT			341
(100)				
1	SORT AGGREGATE		1	
2	EXTERNAL TABLE ACCESS FULL	EXT_EMP	102K	341
(1)	00:00:01			

14 rows selected.

SQL>

The execution plan shows that the table was accessed from the buffer cache.

- c. To access the in-memory table data from the IM column store, set the `QUERY_REWRITE_INTEGRITY` parameter to `STALE_TOLERATED` in your session.

```
SQL> ALTER SESSION SET
          QUERY_REWRITE_INTEGRITY = STALE_TOLERATED;
2
Session altered.

SQL>
```

- d. Re-query the table.

```
SQL> SELECT count(*) FROM test.ext_emp;

COUNT (*)
-----
        2078

SQL>
```

Q/ Does the execution plan show that the table was still accessed from the buffer cache?

```
SQL> SELECT * FROM dbms_xplan.display_cursor();

PLAN_TABLE_OUTPUT
-----
SQL_ID 7382qx3bj1x4n, child number 0
```

```
-----  
SELECT count(*) FROM test.ext_emp
```

```
Plan hash value: 546827939
```

```
-----  
| Id  | Operation                                | Name      | Rows  
| Cost (%CPU)| Time              |           |  
-----  
|  0  | SELECT STATEMENT                        |           |  
|    341 (100)|                   |           |  
|  1  | SORT AGGREGATE                          |           |  
|    1  |                                           |           |  
|  2  | EXTERNAL TABLE ACCESS INMEMORY FULL | EXT_EMP |  
102K|    341 (1)| 00:00:01 |
```

```
-----  
14 rows selected.
```

```
SQL>
```

A/ No. The execution plan shows that the table is now accessed from the IM column store.

6. Update the external table to observe how queries on in-memory external tables behave.
 - a. Add records to the external file.

```
SQL> host echo "202,Pat,Fay,PFAY,603.123.6666,17-AUG-  
05,MK_REP,6000,,201,20" >> /home/oracle/labs/DW/empext1.dat
```

```
SQL> host echo "203,Susan,Mavris,SMAVRIS,515.123.7777,07-JUN-  
02,HR_REP,6500,,101,40" >> /home/oracle/labs/DW/empext1.dat
```

```
SQL>
```

Q/ Is the in-memory external table automatically repopulated into the IM column store?

```
SQL> SELECT count(*) FROM test.ext_emp;
```

```
COUNT (*)
```

```
-----  
2078
```

```
SQL>
```


A/ No. The query still displays the same number of employees.

- b. Repopulate the in-memory external table.

```
SQL> EXEC dbms_inmemory.repopulate ('TEST','EXT_EMP')

PL/SQL procedure successfully completed.

SQL>
```

- c. Verify that the in-memory external table is repopulated into the IM column store.

```
SQL> SELECT count(*) FROM test.ext_emp;

COUNT (*)
-----
        2080

SQL>
```

7. Look at the system statistics related to external in-memory segments.

```
SQL> CONNECT system@PDB1
Enter password: password
Connected.
SQL> SELECT display_name, value FROM v$sysstat s, v$statname n
       WHERE s.statistic# = n.statistic#
       AND (display_name LIKE 'IM XT%'
           OR display_name LIKE 'IM%external%')
       AND value <> 0;

 2   3   4   5
DISPLAY_NAME                                VALUE
-----
IM populate external table read time (ms)      886
IM XT populate rows                          4158
IM XT populate CUs                           2
IM XT populate segments                      2

SQL> EXIT
$
```

8. Execute the \$HOME/labs/DW/cleanup_IM_tables.sh shell script to drop the in-memory external table.

```
$ $HOME/labs/DW/cleanup_IM_tables.sh
...
$
```

Practice 7-3: Using Hierarchy-Based Predicates and Calculated Measures on Analytic Views

Overview

In this practice, you query analytic views (AVs) with SQL by using the same capabilities as Microsoft's Multidimensional Expression (MDX) language, or by using the `DBMS_MDX_ODBO` package, the MDX interface provided by PL/SQL.

Tasks

1. Execute the `$HOME/labs/DW/AV.sh` shell script. The script creates an analytic view (AV) based on fact and dimension tables, all loaded with data.

```
$ $HOME/labs/DW/AV.sh
...
$
```

2. Query the AV.
 - a. Look at the `TIME_HIER` hierarchy.

```
$ sqlplus av@PDB1

Enter password: password

SQL> SELECT year_name, quarter_name, month_name, depth,
           parent_level_name
       FROM   time_hier;
 2      3
YEAR_NAM QUARTER_ MONTH_NA      DEPTH PARENT_
-----
CY2011    Q1CY2011  Feb-11          3 QUARTER
CY2011    Q1CY2011  Jan-11          3 QUARTER
CY2011    Q1CY2011  Mar-11          3 QUARTER
CY2012    Q1CY2012  Mar-12          3 QUARTER
CY2012    Q1CY2012  Jan-12          3 QUARTER
CY2012    Q1CY2012  Feb-12          3 QUARTER
CY2013    Q1CY2013  Jan-13          3 QUARTER
CY2013    Q1CY2013  Feb-13          3 QUARTER
CY2013    Q1CY2013  Mar-13          3 QUARTER
CY2014    Q1CY2014  Feb-14          3 QUARTER
CY2014    Q1CY2014  Mar-14          3 QUARTER
CY2014    Q1CY2014  Jan-14          3 QUARTER
CY2015    Q1CY2015  Feb-15          3 QUARTER
CY2015    Q1CY2015  Jan-15          3 QUARTER
CY2015    Q1CY2015  Mar-15          3 QUARTER
```

CY2011	Q2CY2011	Jun-11	3	QUARTER
CY2011	Q2CY2011	May-11	3	QUARTER
CY2011	Q2CY2011	Apr-11	3	QUARTER
CY2012	Q2CY2012	Jun-12	3	QUARTER
CY2012	Q2CY2012	Apr-12	3	QUARTER
CY2012	Q2CY2012	May-12	3	QUARTER
CY2013	Q2CY2013	May-13	3	QUARTER
CY2013	Q2CY2013	Jun-13	3	QUARTER
CY2013	Q2CY2013	Apr-13	3	QUARTER
CY2014	Q2CY2014	May-14	3	QUARTER
CY2014	Q2CY2014	Jun-14	3	QUARTER
CY2014	Q2CY2014	Apr-14	3	QUARTER
CY2015	Q2CY2015	Jun-15	3	QUARTER
CY2015	Q2CY2015	May-15	3	QUARTER
CY2015	Q2CY2015	Apr-15	3	QUARTER
CY2011	Q3CY2011	Jul-11	3	QUARTER
CY2011	Q3CY2011	Sep-11	3	QUARTER
CY2011	Q3CY2011	Aug-11	3	QUARTER
CY2012	Q3CY2012	Aug-12	3	QUARTER
CY2012	Q3CY2012	Sep-12	3	QUARTER
CY2012	Q3CY2012	Jul-12	3	QUARTER
CY2013	Q3CY2013	Sep-13	3	QUARTER
CY2013	Q3CY2013	Aug-13	3	QUARTER
CY2013	Q3CY2013	Jul-13	3	QUARTER
CY2014	Q3CY2014	Jul-14	3	QUARTER
CY2014	Q3CY2014	Sep-14	3	QUARTER
CY2014	Q3CY2014	Aug-14	3	QUARTER
CY2015	Q3CY2015	Aug-15	3	QUARTER
CY2015	Q3CY2015	Jul-15	3	QUARTER
CY2015	Q3CY2015	Sep-15	3	QUARTER
CY2011	Q4CY2011	Dec-11	3	QUARTER
CY2011	Q4CY2011	Oct-11	3	QUARTER
CY2011	Q4CY2011	Nov-11	3	QUARTER
CY2012	Q4CY2012	Nov-12	3	QUARTER
CY2012	Q4CY2012	Oct-12	3	QUARTER
CY2012	Q4CY2012	Dec-12	3	QUARTER
CY2013	Q4CY2013	Dec-13	3	QUARTER
CY2013	Q4CY2013	Nov-13	3	QUARTER
CY2013	Q4CY2013	Oct-13	3	QUARTER
CY2014	Q4CY2014	Dec-14	3	QUARTER
CY2014	Q4CY2014	Oct-14	3	QUARTER
CY2014	Q4CY2014	Nov-14	3	QUARTER

CY2015	Q4CY2015	Oct-15	3	QUARTER
CY2015	Q4CY2015	Dec-15	3	QUARTER
CY2015	Q4CY2015	Nov-15	3	QUARTER
CY2011	Q1CY2011		2	YEAR
CY2012	Q1CY2012		2	YEAR
CY2013	Q1CY2013		2	YEAR
CY2014	Q1CY2014		2	YEAR
CY2015	Q1CY2015		2	YEAR
CY2011	Q2CY2011		2	YEAR
CY2012	Q2CY2012		2	YEAR
CY2013	Q2CY2013		2	YEAR
CY2014	Q2CY2014		2	YEAR
CY2015	Q2CY2015		2	YEAR
CY2011	Q3CY2011		2	YEAR
CY2012	Q3CY2012		2	YEAR
CY2013	Q3CY2013		2	YEAR
CY2014	Q3CY2014		2	YEAR
CY2015	Q3CY2015		2	YEAR
CY2011	Q4CY2011		2	YEAR
CY2012	Q4CY2012		2	YEAR
CY2013	Q4CY2013		2	YEAR
CY2014	Q4CY2014		2	YEAR
CY2015	Q4CY2015		2	YEAR
CY2011			1	ALL
CY2012			1	ALL
CY2013			1	ALL
CY2014			1	ALL
CY2015			1	ALL
			0	

86 rows selected.

SQL>

- b. Query the AV to find the values of sales for all years, per year and per quarter.

```
SQL> SELECT time_hier.member_name, sales
        FROM   sales_av HIERARCHIES(time_hier)
        WHERE  time_hier.level_name IN  ('ALL','YEAR','QUARTER')
        ORDER BY time_hier.hier_order;

  2      3      4
MEMBER_N      SALES
-----
ALL          1970258105
```

```

CY2011      593507775
Q1CY2011    7068888
Q2CY2011    556371578
Q4CY2011    30067309
CY2012      568622304
Q2CY2012    564675917
Q4CY2012    3946387
CY2013      611827904
Q1CY2013    5327980
Q2CY2013    592275258
Q3CY2013    10287982
Q4CY2013    3936684
CY2014      196300122
Q1CY2014    5327980
Q2CY2014    183612568
Q3CY2014    7359574

```

17 rows selected.

SQL>

- c. Query the AV to find the values of sales for all months for the first and second quarters of all years.

```

SQL> SELECT time_hier.member_name, sales
      FROM sales_av HIERARCHIES(time_hier)
      WHERE time_hier.level_name = 'MONTH'
      AND    TO_CHAR(time_hier.month_end_date,'Q') IN (1,2)
      ORDER BY time_hier.hier_order;

```

2	3	4	5
MEMBER_N		SALES	
Jan-11		1740908	
Feb-11		5327980	
Apr-11		556371578	
Apr-12		564675917	
Mar-13		5327980	
Apr-13		590711509	
May-13		1563749	
Feb-14		5327980	
Apr-14		183612568	

9 rows selected.

SQL>

- Use filter-before aggregate predicates to query the AV. The `USING` and `FILTER FACT` clauses allow a query to filter fact rows before being aggregated by the analytic view. Use the `FILTER FACT` clause to change the values of aggregate-level hierarchy members. The `USING` clause can be thought of as an inner query, similar to a `WITH` clause. Report sales at the year and quarter levels, but only for the first half of each year. This example filters months where the quarter of the year of the `MONTH_END_DATE` attribute is 1 or 2. The hierarchy used is `TIME_HIER`.

```
SQL> SELECT time_hier.member_name, sales
      FROM ANALYTIC VIEW (
        USING sales_av HIERARCHIES(time_hier)
        FILTER FACT (time_hier TO level_name = 'MONTH'
                     AND TO_CHAR(month_end_date,'Q') IN (1,2)))
      WHERE time_hier.level_name IN ('ALL','YEAR','QUARTER')
      ORDER BY time_hier.hier_order;
```

2	3	4	5	6	7
MEMBER_N				SALES	
ALL				1914660169	
CY2011				563440466	
Q1CY2011				7068888	
Q2CY2011				556371578	
CY2012				564675917	
Q2CY2012				564675917	
CY2013				597603238	
Q1CY2013				5327980	
Q2CY2013				592275258	
CY2014				188940548	
Q1CY2014				5327980	
Q2CY2014				183612568	

12 rows selected.

SQL>

Q/ Do you get visual or nonvisual totals?

A/ The filter-before aggregate predicate for the 2011 year filtered out the fourth quarter (see Q4CY2011 in the previous query). The resulting total of sales for year 2011 aggregates only the first two quarters. This is the same behavior for all years. The aggregated sum of sales for all years is smaller in the second query than in the previous query.

SQL execution with filter-before aggregate predicates manages to report visual totals.

- a. Now filter on two hierarchies, `TIME_HIER` and `GEOGRAPHY_HIER`. Report sales for the first half of each year (first and second quarters) in Mexico and Canada.

```
SQL> SELECT time_hier.member_name AS time,
           geography_hier.member_name AS geography, sales
FROM ANALYTIC VIEW (
  USING sales_av HIERARCHIES(time_hier, geography_hier)
  FILTER FACT (time_hier TO level_name = 'QUARTER' AND
               (quarter_name like 'Q1%'
                OR quarter_name like 'Q2%'),
               geography_hier TO level_name = 'COUNTRY'
               AND country_name in ('Mexico','Canada'))
WHERE time_hier.level_name = 'YEAR'
AND geography_hier.level_name = 'REGION'
ORDER BY time_hier.hier_order;
```

2	3	4	5	6	7	8	9	10	11	12
TIME	GEOGRAPHY						SALES			

CY2011	North America						38931636			
CY2012	North America						39129154			
CY2013	North America						41586079			
CY2014	North America						13119131			

```
SQL>
```

- b. Re-execute the same query, adding another country. Report sales for the first half of each year (first and second quarters) in Mexico, Canada, and Chile.

```
SQL> SELECT time_hier.member_name AS time,
           geography_hier.member_name AS geography, sales
FROM ANALYTIC VIEW (
  USING sales_av HIERARCHIES(time_hier, geography_hier)
  FILTER FACT
    (time_hier TO level_name = 'QUARTER' AND
     (quarter_name like 'Q1%' OR quarter_name like 'Q2%'),
     geography_hier TO level_name = 'COUNTRY'
     AND country_name in ('Mexico','Canada' , 'Chile'))
WHERE time_hier.level_name = 'YEAR'
AND geography_hier.level_name = 'REGION'
ORDER BY time_hier.hier_order;
```

2	3	4	5	6	7	8	9	10	11	12
TIME	GEOGRAPHY						SALES			

CY2011	South America						4880427			
CY2011	North America						38931636			

```

CY2012 South America      4764838
CY2012 North America      39129154
CY2013 North America      41586079
CY2013 South America      5423944
CY2014 North America      13119131
CY2014 South America      1618375

```

8 rows selected.

SQL>

- c. Re-execute the same query and report the global sales total for each year of each region.

```

SQL> SELECT time_hier.member_name AS time,
           geography_hier.member_name AS geography, sales
FROM ANALYTIC VIEW (
  USING sales_av HIERARCHIES(time_hier, geography_hier)
  FILTER FACT
    (time_hier TO level_name = 'QUARTER' AND
     (quarter_name like 'Q1%' OR quarter_name like 'Q2%'),
     geography_hier TO level_name = 'COUNTRY'
     AND country_name in ('Mexico','Canada' ,'Chile'))
WHERE time_hier.level_name IN ('ALL','YEAR')
AND geography_hier.level_name = 'REGION'
ORDER BY time_hier.hier_order;

```

```

  2  3  4  5  6  7  8  9 10 11 12
TIME  GEOGRAPHY              SALES
-----

```

```

ALL      South America      16687584
ALL      North America      132766000
CY2011 South America      4880427
CY2011 North America      38931636
CY2012 South America      4764838
CY2012 North America      39129154
CY2013 South America      5423944
CY2013 North America      41586079
CY2014 South America      1618375
CY2014 North America      13119131

```

10 rows selected.

SQL>

Q/ Does the sales total for North America include all years, all quarters, and all countries?

A/ No.

Each hierarchy specifies a filter-before aggregate predicate, which serves to filter the leaves of that hierarchy before aggregating the measures. The fact rows are filtered to include only the leaf descendants of those members.

4. Use calculated measures to query the AV. The `USING` and `ADD MEASURES` clauses can be used to define a calculated measure within a `SELECT` statement. Define two calculated measures by adding the total of sales for the period prior to each period and the percent change of sales against the prior period.

```
SQL> SELECT time_hier.member_name, sales, sales_prior_period,
          ROUND(sales_prior_period_pct_change,3) AS
          percent_change_sales
FROM ANALYTIC VIEW (
  USING sales_av HIERARCHIES(time_hier)
  ADD MEASURES (sales_prior_period AS (LAG(sales)
          OVER (HIERARCHY time_hier OFFSET 1)),
  sales_prior_period_pct_change AS
          (LAG_DIFF_PERCENT(sales)
          OVER (HIERARCHY time_hier OFFSET 1))
  )
  )
WHERE time_hier.level_name = 'YEAR'
ORDER BY time_hier.hier_order;
```

2	3	4	5	6	7	8	9	10	11	12	13
MEMBER_N		SALES	SALES_PRIOR_PERIOD							PERCENT_CHANGE_SALES	
CY2011		593507775									
CY2012		568622304						593507775		-.042	
CY2013		611827904						568622304		.076	
CY2014		196300122						611827904		-.679	

```
SQL>
```

5. Use both filter-before aggregate predicates and calculated measures to query the AV.
 - a. Report sales, the total of sales for the period prior to each period, and the percent change of sales against the prior period for the first half of years in Mexico and Canada.

```
SQL> SELECT time_hier.member_name AS time,
          geography_hier.member_name AS geography, sales,
          sales_prior_period,
          ROUND(sales_prior_period_pct_change,3)
```

```

        AS percent_change_sales
FROM ANALYTIC VIEW (
  USING sales_av HIERARCHIES(time_hier, geography_hier)
  FILTER FACT ( time_hier TO level_name = 'QUARTER' AND
                (quarter_name like 'Q1%'
                 OR quarter_name like 'Q2%'),
                geography_hier TO level_name = 'COUNTRY'
                AND country_name in ('Mexico','Canada'))
  ADD MEASURES ( sales_prior_period AS (LAG(sales) OVER
                (HIERARCHY time_hier OFFSET 1)),
                sales_prior_period_pct_change AS
                (LAG_DIFF_PERCENT(sales) OVER
                (HIERARCHY time_hier OFFSET 1)))
        )
WHERE time_hier.level_name = 'YEAR'
AND   geography_hier.level_name = 'REGION'
ORDER BY time_hier.hier_order;

```

TIME	GEOGRAPHY	SALES	SALES_PRIOR_PERIOD
PERCENT_CHANGE_SALES			
CY2011	North America	38931636	
CY2012	North America	39129154	38931636
		.005	
CY2013	North America	41586079	39129154
		.063	
CY2014	North America	13119131	41586079
		-.685	

SQL>

- b. Add a WHERE clause and check whether it impacts the aggregated measure data.

```

SQL> SELECT time_hier.member_name AS time,
           geography_hier.member_name AS geography, sales,
           sales_prior_period,
           ROUND(sales_prior_period_pct_change,3)
           AS percent_change_sales
FROM ANALYTIC VIEW (
  USING sales_av HIERARCHIES(time_hier, geography_hier)
  FILTER FACT ( time_hier TO level_name = 'QUARTER' AND
                (quarter_name like 'Q1%'
                 OR quarter_name like 'Q2%'),

```

```

        geography_hier TO level_name = 'COUNTRY'
        AND country_name in ('Mexico','Canada'))
ADD MEASURES ( sales_prior_period AS (LAG(sales) OVER
        (HIERARCHY time_hier OFFSET 1)),
        sales_prior_period_pct_change AS
        (LAG_DIFF_PERCENT(sales) OVER
        (HIERARCHY time_hier OFFSET 1)))
    )
WHERE time_hier.level_name = 'YEAR'
AND geography_hier.level_name = 'REGION'
AND sales > 30000000
ORDER BY time_hier.hier_order;
2   3   4   5   6   7   8   9  10  11  12  13  14  15  16  17  18  19
20  21  22
TIME      GEOGRAPHY              SALES  SALES_PRIOR_PERIOD
-----
PERCENT_CHANGE_SALES
-----
CY2011 North America      38931636
CY2012 North America      39129154          38931636
                        .005
CY2013 North America      41586079          39129154
                        .063

SQL> EXIT
$

```

An appended `WHERE` clause does not impact the aggregated measure data. It simply reduces the rows returned after the filter-before aggregation and the measures calculation are applied.

6. Execute the `$HOME/labs/DW/cleanup_AV.sh` shell script to drop the AV schema, including fact and dimension tables.

```

$ $HOME/labs/DW/cleanup_AV.sh
...
$

```

Practice 7-4: Using Polymorphic Table Functions

Overview

In this practice, you create and use polymorphic table functions (PTFs).

Tasks

1. Execute the `$HOME/labs/DW/PTF.sh` shell script. The script creates the `SCOTT.DEPT` and `SCOTT.EMP` tables.

```
$ $HOME/labs/DW/PTF.sh
...
$
```

2. In the first example, you create the `GET_COL` PTF that returns the columns of a given data type from a given table.

- a. Log in to the `PDB1` PDB as `SCOTT` to create a PTF that returns the columns of a given data type from a given table.

```
$ sqlplus scott@PDB1

Enter password: password

SQL>
```

- b. Create the PL/SQL package that contains the functions/procedures that the PTF implementation must use:
 - The supplied `DESCRIBE` function, which returns the new table “shape”
 - The supplied `FETCH_ROWS` procedure is not required because it does not need to produce associated new column values for a given subset of rows.
 - In the package, the parameters are:
 - `tab`: Input table
 - `type_name`: A string representing the type of columns to return
 - `flip`: “False” (default value). Only matched columns with the given type name are returned; otherwise, the PTF ignores the columns with the given type name and returns the other columns.

```
CREATE OR REPLACE PACKAGE get_col_p AS
FUNCTION describe (tab IN OUT DBMS_TF.table_t,
                  type_name VARCHAR2,
                  flip VARCHAR2 DEFAULT 'False')
RETURN DBMS_TF.describe_t;
END get_col_p;
/
```

```

SQL> CREATE OR REPLACE PACKAGE get_col_p AS
FUNCTION describe (tab IN OUT DBMS_TF.table_t,
                  type_name VARCHAR2,
                  flip VARCHAR2 DEFAULT 'False')
    RETURN DBMS_TF.describe_t;
END get_col_p;
/
  2   3   4   5   6   7
Package created.

SQL>

```

- c. Create the package body containing the DESCRIBE function.

```

CREATE OR REPLACE PACKAGE BODY GET_COL_P as
FUNCTION describe (tab IN OUT DBMS_TF.table_t,
                  type_name VARCHAR2,
                  flip VARCHAR2 DEFAULT 'False')
    RETURN DBMS_TF.describe_t
AS
    typ constant varchar2(1024) := upper(trim(type_name));
BEGIN
    FOR i IN 1 .. tab.column.count() LOOP
        tab.column(i).pass_through := CASE upper(substr(flip,1,1))
            WHEN 'F' THEN
                DBMS_TF.Column_Type_Name(tab.column(i).description)=typ
            ELSE
                DBMS_TF.Column_Type_Name(tab.column(i).description)!=typ
            END /* case */;
    END LOOP;
RETURN null;
END;
END GET_COL_P;
/

```

```

SQL> CREATE OR REPLACE PACKAGE BODY GET_COL_P as
FUNCTION describe (tab IN OUT DBMS_TF.table_t,
                  type_name VARCHAR2,
                  flip VARCHAR2 DEFAULT 'False')
    RETURN DBMS_TF.describe_t
AS
    typ constant varchar2(1024) := upper(trim(type_name));
BEGIN
    FOR i IN 1 .. tab.column.count() LOOP

```

```

        tab.column(i).pass_through := CASE upper(substr(flip,1,1))
            WHEN 'F' THEN
                DBMS_TF.Column_Type_Name(tab.column(i).description)=typ
            ELSE
                DBMS_TF.Column_Type_Name(tab.column(i).description)!=typ
            END /* case */;
    END LOOP;
RETURN null;
END;
END GET_COL_P;
/
2   3   4   5   6   7   8   9  10  11  12  13  14  15  16  17
18  19  20  21
Package body created.

SQL>

```

- d. Create the PTF. Specify exactly one formal argument of type `TABLE`, specify the return type of the PTF as `TABLE`, specify the type of the PTF function (`ROW` or `TABLE POLYMORPHIC`), and indicate which package contains the actual PTF implementation.

```

SQL> CREATE OR REPLACE FUNCTION GET_COL (
            tab          TABLE,
            type_name VARCHAR2,
            flip          VARCHAR2 DEFAULT 'False')
RETURN TABLE
    PIPELINED ROW POLYMORPHIC using GET_COL_P;
/
2   3   4   5   6   7
Function created.

SQL>

```

3. Query tables by using the `GET_COL` PTF.

- a. Query the `SCOTT.DEPT` table.

```

SQL> SELECT * FROM scott.dept;

```

DEPTNO	DNAME	LOC
10	ACCOUNTING	NEW YORK
20	RESEARCH	DALLAS
30	SALES	CHICAGO
40	OPERATIONS	BOSTON

```

SQL>

```

- b. Use the GET_COL PTF to report from the SCOTT.DEPT table only columns whose type is VARCHAR2.

```
SQL> SELECT * FROM GET_COL(scott.dept, 'varchar2');
```

DNAME	LOC
ACCOUNTING	NEW YORK
RESEARCH	DALLAS
SALES	CHICAGO
OPERATIONS	BOSTON

```
SQL>
```

- c. Use the same GET_COL PTF to report from another table, SCOTT.EMP table, only columns whose type is not VARCHAR2.

```
SQL> SELECT *  
      FROM GET_COL(scott.emp, 'varchar2', flip => 'True');
```

```
2  
EMPNO      MGR HIREDATE      SAL      COMM      DEPTNO  
-----  
7369      7902 17-DEC-80      800  
7499      7698 20-FEB-81     1600      300      30  
7521      7698 22-FEB-81     1250      500      30  
7566      7839 02-APR-81     2975  
7654      7698 28-SEP-81     1250     1400      30  
7698      7839 01-MAY-81     2850  
7782      7839 09-JUN-81     2450      10  
7788      7566 19-APR-87     3000      20  
7839              17-NOV-81     5000      10  
7844      7698 08-SEP-81     1500      0      30  
7876      7788 23-MAY-87     1100      20  
7900      7698 03-DEC-81      950      30  
7902      7566 03-DEC-81     3000      20  
7934      7782 23-JAN-82     1300      10
```

```
14 rows selected.
```

```
SQL>
```

- d. Use the GET_COL PTF to report for ANALYST and PRESIDENT employees only columns whose type is VARCHAR2.

```
SQL> SELECT * FROM GET_COL(scott.emp, 'varchar2')  
      WHERE job IN ('ANALYST' , 'PRESIDENT');
```

```
2
```

ENAME	JOB
-----	-----
SCOTT	ANALYST
KING	PRESIDENT
FORD	ANALYST
SQL>	

Only `ENAME` and `JOB` columns of data type `VARCHAR2` are displayed.

- e. Use the `GET_COL` PTF to report about employees whose commission is greater than 1000 and only columns whose type is not `VARCHAR2`.

```
SQL> SELECT *
      FROM GET_COL(scott.emp, 'varchar2', flip => 'True')
      WHERE comm > 1000;
```

2	3				
EMPNO	MGR	HIREDATE	SAL	COMM	DEPTNO
-----	-----	-----	-----	-----	-----
7654	7698	28-SEP-81	1250	1400	30

```
SQL>
```

All columns except the `ENAME` and `JOB` columns are displayed.

- f. Use the `GET_COL` PTF to report about employees whose commission is greater than 1000 and only columns whose type is `VARCHAR2`.

```
SQL> SELECT * FROM GET_COL(scott.emp, 'varchar2')
      WHERE comm > 1000;
```

```
WHERE comm > 1000
      *
```

```
ERROR at line 1:
ORA-00904: "COMM": invalid identifier
```

```
SQL>
```

Because `COMM` is not a `VARCHAR2` type of column, the selection cannot apply.

- g. Use the `GET_COL` PTF to report about employees whose job is either `ANALYST` or `PRESIDENT` and only columns whose type is not `VARCHAR2`.

```
SQL> SELECT *
      FROM GET_COL(scott.emp, 'varchar2', flip => 'True')
```



```

WHERE job IN ('ANALYST' , 'PRESIDENT');
WHERE job IN ('ANALYST' , 'PRESIDENT')
      *
ERROR at line 2:
ORA-00904: "JOB": invalid identifier

SQL>

```

Because `JOB` is a `VARCHAR2` type of column, the selection cannot apply.

4. In a second example, you create the `CHANGE_CASE_P` PTF that changes the case of all the `VARCHAR2` columns in the input table to the specified case.
 - a. Create the PL/SQL package body. The PL/SQL package that contains the functions/procedures for the PTF implementation must use:
 - The supplied `DESCRIBE` function that returns the new table “shape”
 - The supplied `FETCH_ROWS` procedure required to produce the associated new column values for a given subset of rows
 - In the package, the parameters are:
 - `tab`: Input table
 - `type_name`: The case-change function ('upper', 'lower', or 'initcap')

```

SQL> CREATE OR REPLACE PACKAGE CHANGE_CASE_P AS
      FUNCTION Describe(tab IN OUT DBMS_TF.Table_t,
                        new_case VARCHAR2)
                        RETURN DBMS_TF.describe_t;
      PROCEDURE Fetch_Rows(new_case varchar2);
END CHANGE_CASE_P;
/
2    3    4    5    6    7
Package created.

SQL>

```

- b. Create the package body containing the `DESCRIBE` function and `FETCH_ROWS` procedure.

```

SQL> CREATE OR REPLACE PACKAGE BODY CHANGE_CASE_P AS
      FUNCTION Describe(tab IN OUT DBMS_TF.table_t,
                        new_case VARCHAR2)
                        RETURN DBMS_TF.describe_t
      AS
        new_cols DBMS_TF.columns_new_t;
        col_id   pls_integer := 1;

```

```

BEGIN
    FOR i IN 1 .. tab.column.count() LOOP
        IF tab.column(i).description.type =
            DBMS_TF.TYPE_VARCHAR2 then
            tab.column(i).pass_through := FALSE;
            tab.column(i).for_read      := TRUE;
            new_cols(col_id) := tab.column(i).description;
            col_id := col_id + 1;
        END IF;
    END LOOP;
    RETURN DBMS_TF.describe_t(new_columns => new_cols);
END;

PROCEDURE Fetch_Rows(new_case varchar2)
AS
    case_function constant varchar2(100) :=
        upper(trim(new_case));
    rowset          DBMS_TF.Row_Set_t;
BEGIN
    DBMS_TF.Get_Row_Set(rowset);
    IF case_function in ('UPPER', 'LOWER', 'INITCAP')
    THEN
        FOR c IN 1 .. rowset.count() LOOP
            FOR r IN 1 .. rowset(c).tab_varchar2.count() LOOP
                rowset(c).tab_varchar2(r) :=
                    CASE case_function WHEN 'UPPER'
                        THEN upper(rowset(c).tab_varchar2(r))
                        WHEN 'LOWER'
                        THEN lower(rowset(c).tab_varchar2(r))
                        ELSE initcap(rowset(c).tab_varchar2(r))
                    END;
            END LOOP;
        END LOOP;
    END IF;
    DBMS_TF.Put_Row_Set(rowset);
END;

END CHANGE_CASE_P;
/
  2   3   4   ...
Package body created.

SQL>

```

- c. Create the PTF. Specify exactly one formal argument of type `TABLE`, the return type of the PTF as `TABLE`, the type of the PTF function (`ROW` or `TABLE POLYMORPHIC`), and indicate which package contains the actual PTF implementation.

```
SQL> CREATE OR REPLACE FUNCTION CHANGE_CASE
      (tab table, new_case varchar2)
      RETURN TABLE
      PIPELINED ROW POLYMORPHIC using CHANGE_CASE_P;
/
2      3      4      5
Function created.

SQL>
```

5. Query tables by using the `CHANGE_CASE` PTF.

- a. Display the content of the `SCOTT.DEPT` table after changing the case of all the `VARCHAR2` columns to the specified case.

```
SQL> SELECT * FROM CHANGE_CASE(scott.dept, 'initcap');

      DEPTNO DNAME          LOC
-----
10 Accounting      New York
20 Research        Dallas
30 Sales            Chicago
40 Operations      Boston

SQL>
```

- b. Display the content of another table, the `SCOTT.EMP` table, after changing the case of all the `VARCHAR2` columns to the specified case.

```
SQL> SELECT * FROM CHANGE_CASE(scott.emp, 'lower')
      WHERE comm > 1000;

2
EMPNO      MGR HIREDATE      SAL      COMM      DEPTNO ENAME
-----
JOB
-----
7654      7698 28-SEP-81    1250     1400           30 martin
salesman

SQL>
```

- c. Display the content of the `SCOTT.EMP` table after changing the case of all the `VARCHAR2` columns to the specified case and selecting only rows whose `JOB` is `SALESMAN`.

```
SQL> SELECT * FROM CHANGE_CASE(scott.emp, 'lower')
      WHERE job = 'SALESMAN';

2
no rows selected

SQL>
```

Q/ Why are there no rows selected?

A/ Because `JOB` is a `VARCHAR2` type of column, the PTF being applied on this `VARCHAR2` column, `SALESMAN`, becomes `salesman`.

```
SQL> SELECT * FROM CHANGE_CASE(scott.emp, 'lower')
      WHERE job = 'salesman';

2

EMPNO      MGR HIREDATE          SAL      COMM      DEPTNO  ENAME
-----
JOB
-----
  7499    7698 20-FEB-81      1600      300         30  allen
salesman

  7521    7698 22-FEB-81      1250      500         30  ward
salesman

  7654    7698 28-SEP-81      1250     1400         30  martin
salesman

  7844    7698 08-SEP-81      1500         0         30  turner
salesman

SQL>
```

- d. View the execution plan. The `POLYMORPHIC TABLE FUNCTION` operation indicates the row source for a polymorphic table function, which is a table function whose return type is determined by its arguments.

```
SQL> SELECT * FROM dbms_xplan.display_cursor();

PLAN_TABLE_OUTPUT
```

```
SQL_ID 090hgbkyy2jy0, child number 0
```

```
SELECT * FROM CHANGE_CASE(scott.emp, 'lower') WHERE job  
='salesman'
```

```
Plan hash value: 3956160932
```

Id	Operation	Name	Rows
0	SELECT STATEMENT		
2 (100)			
1	POLYMORPHIC TABLE FUNCTION	CHANGE_CASE	14
2	TABLE ACCESS FULL	EMP	14

Cost (%CPU) | Time |

Note

- dynamic statistics used: dynamic sampling (level=2)

18 rows selected.

SQL>

Q/ What happens if you use both PTFs on the SCOTT.EMP table?

```
SQL> SELECT *  
      FROM CHANGE_CASE(GET_COL(scott.emp, 'varchar2'));  
FROM CHANGE_CASE(GET_COL(scott.emp, 'varchar2'))  
      *  
ERROR at line 1:  
ORA-62569: nested polymorphic table function is disallowed  
  
SQL>
```

A/ A nested polymorphic table function is disallowed.

6. To use both PTFs, consider using a with-clause.

```
SQL> WITH T AS (  
          SELECT * FROM GET_COL(scott.emp, 'varchar2')  
        )  
    SELECT * FROM CHANGE_CASE(T, 'initcap');  
  2    3    4  
ENAME      JOB  
-----  
Smith      Clerk  
Allen      Salesman  
Ward       Salesman  
Jones      Manager  
Martin     Salesman  
Blake      Manager  
Clark      Manager  
Scott      Analyst  
King       President  
Turner     Salesman  
Adams      Clerk  
James      Clerk  
Ford       Analyst  
Miller     Clerk  
  
14 rows selected.  
  
SQL> EXIT  
$
```

7. Execute the `$HOME/labs/DW/cleanup_PTF.sh` shell script to drop the PTFs.

```
$ $HOME/labs/DW/cleanup_PTF.sh  
...  
$
```


Practices for Lesson 8:
Describing Sharding Enhancements

Practices for Lesson 8

There are no practices for this lesson.